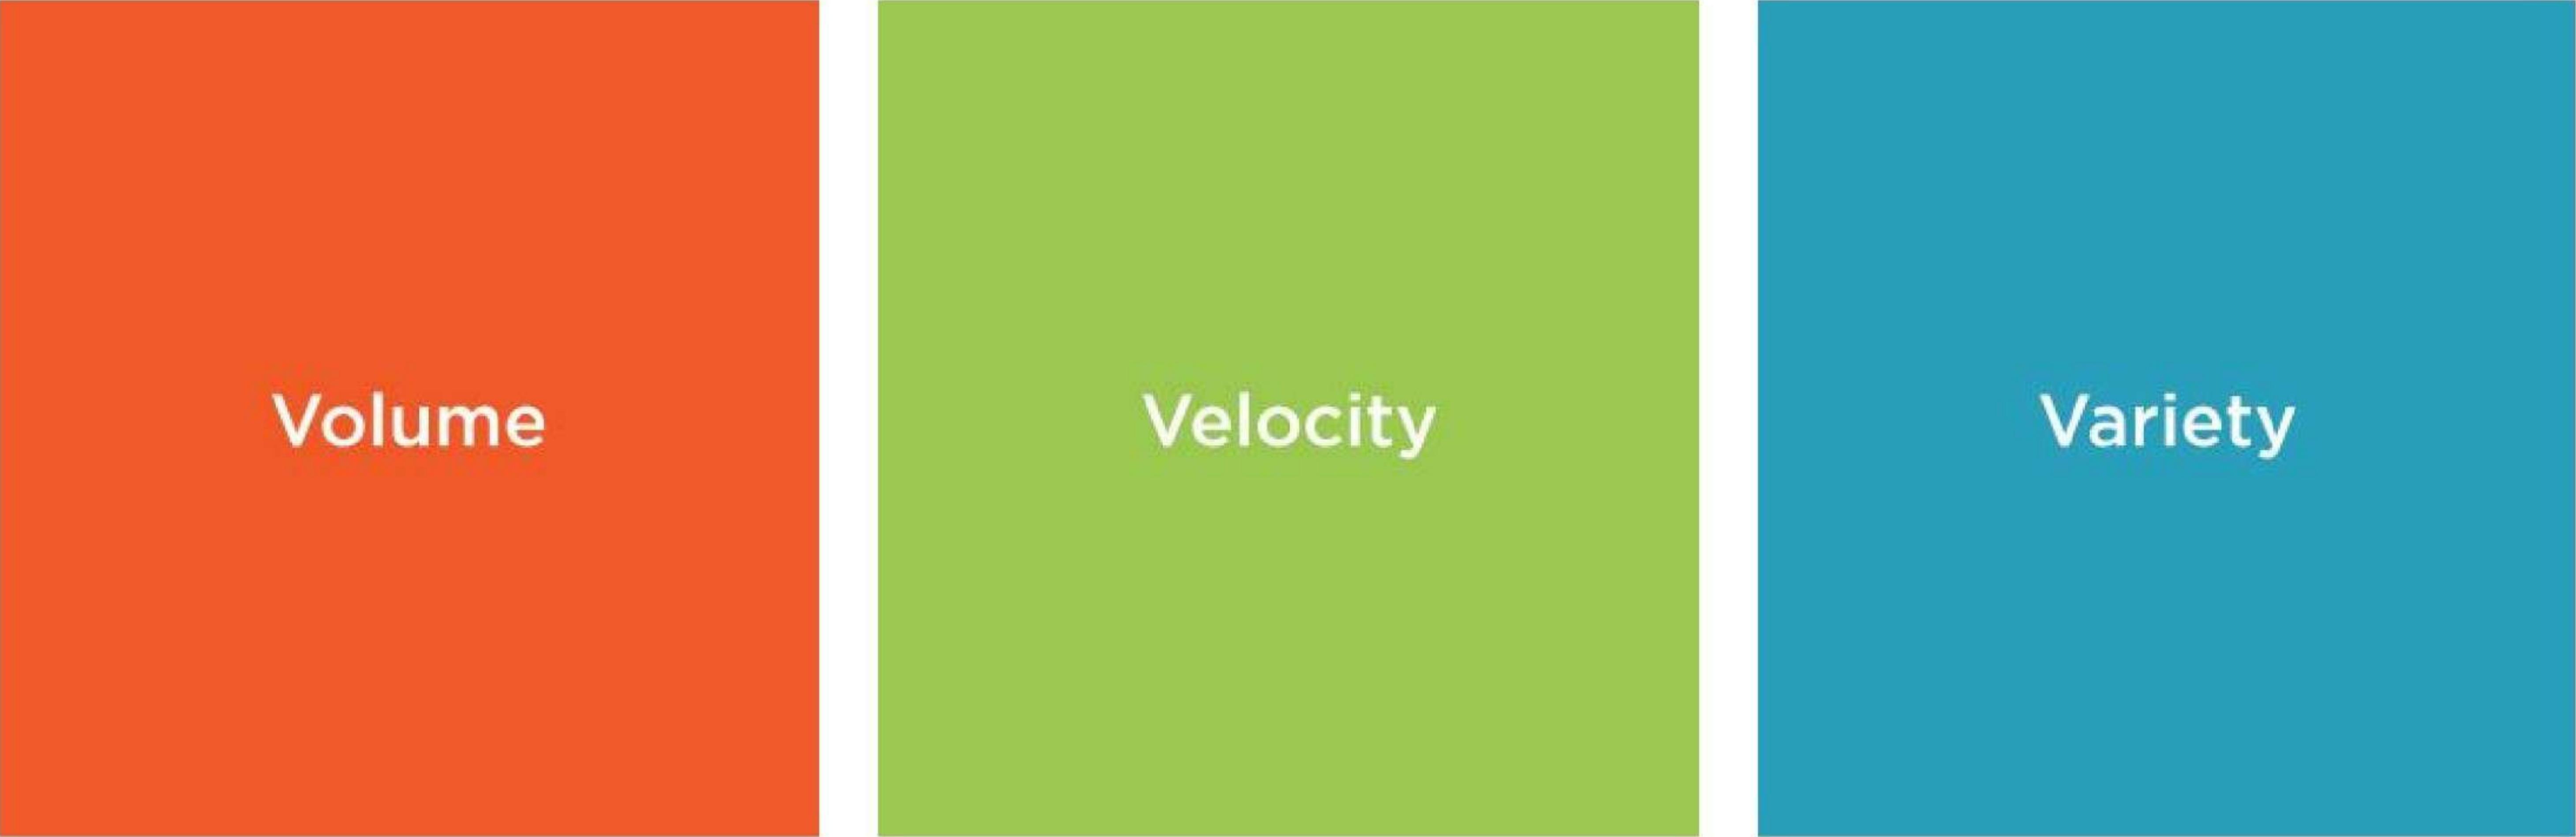


What Is Big Data?



Big Data Characteristics

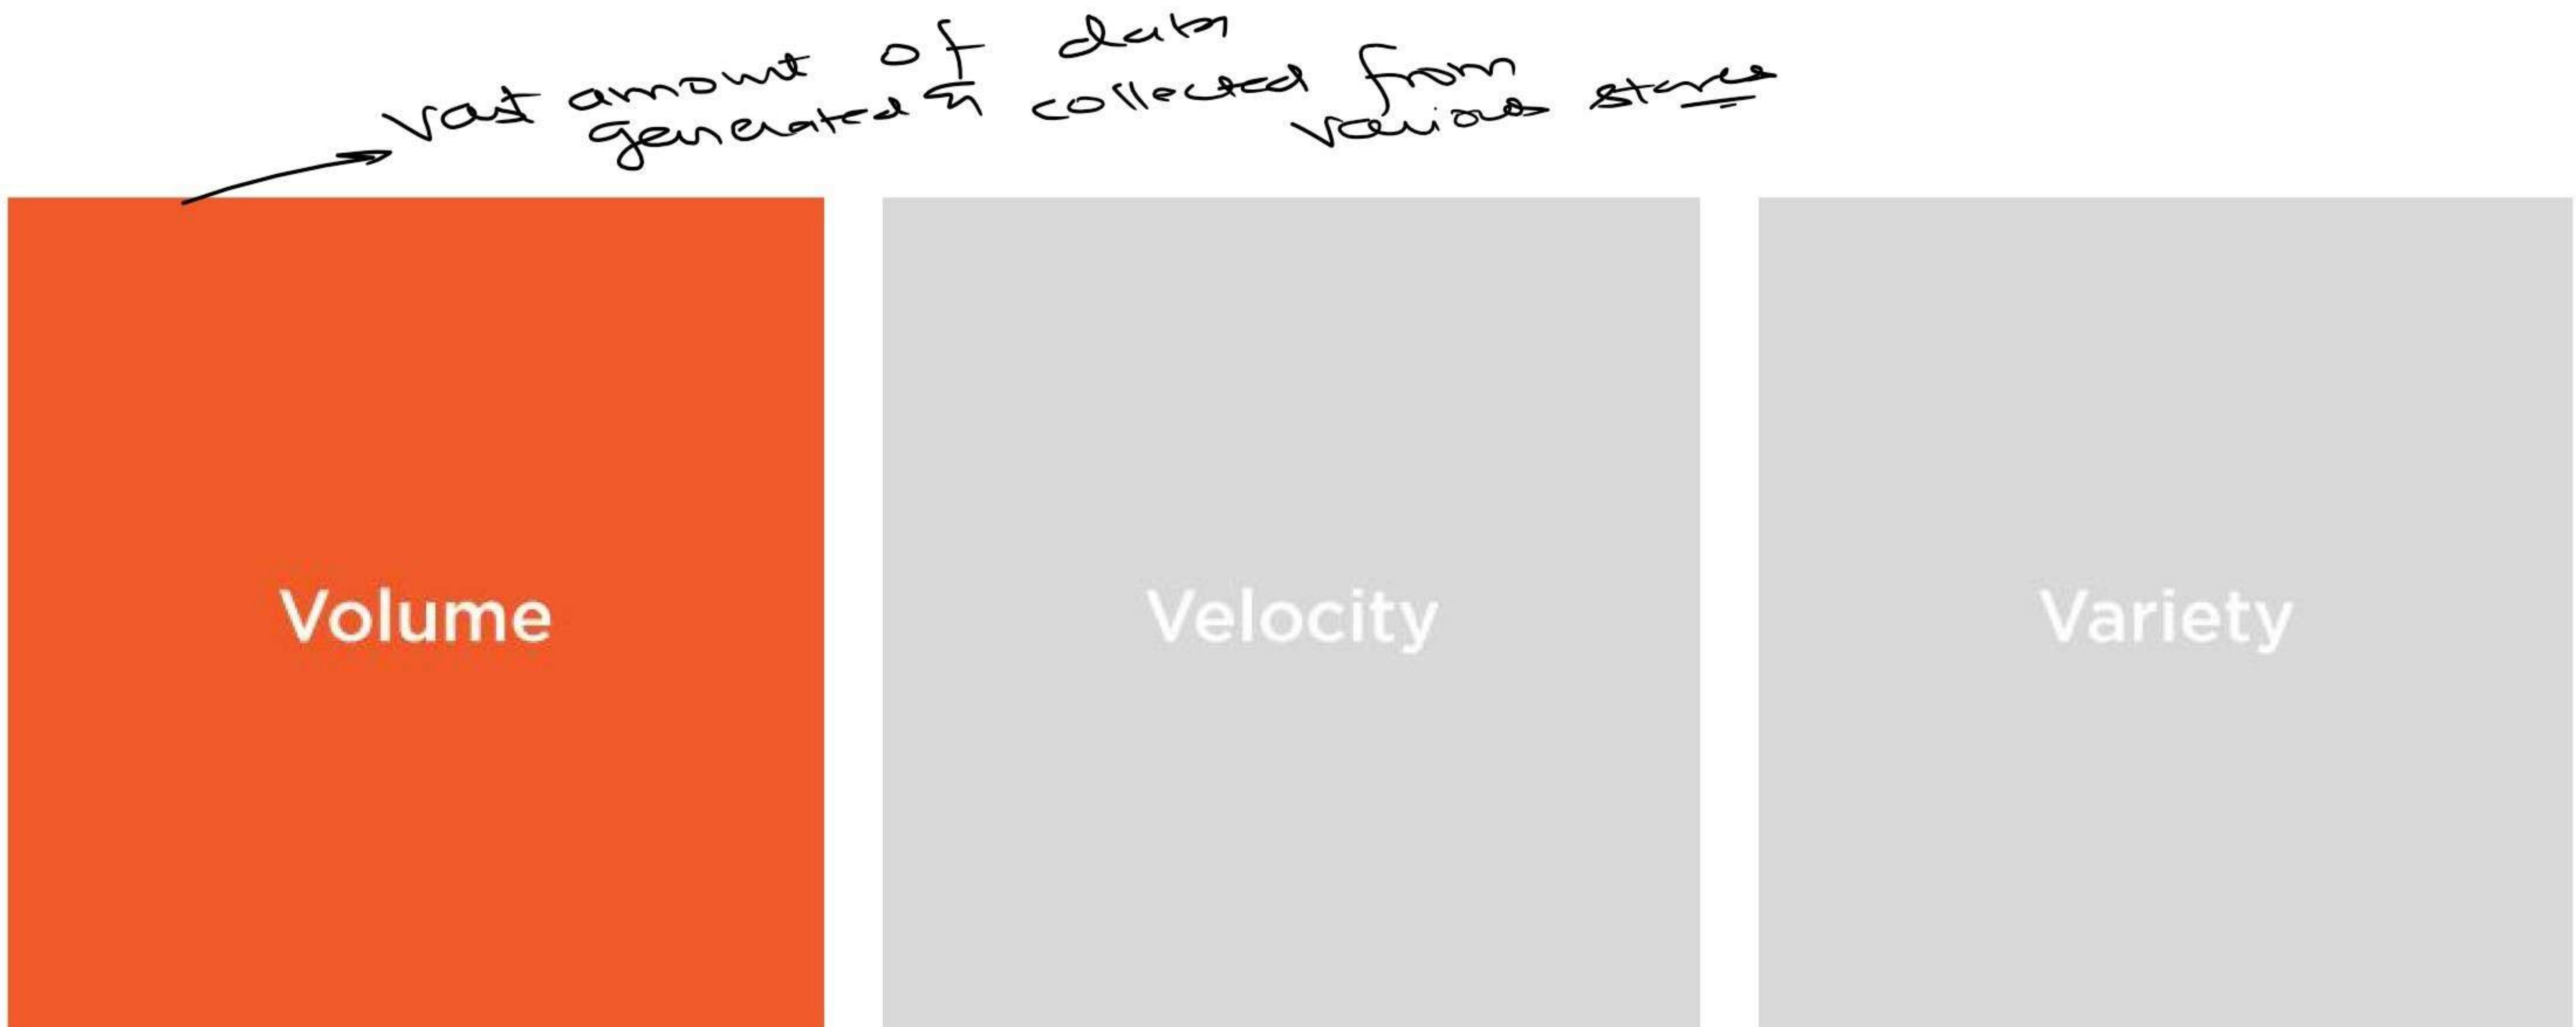


Volume

Velocity

Variety

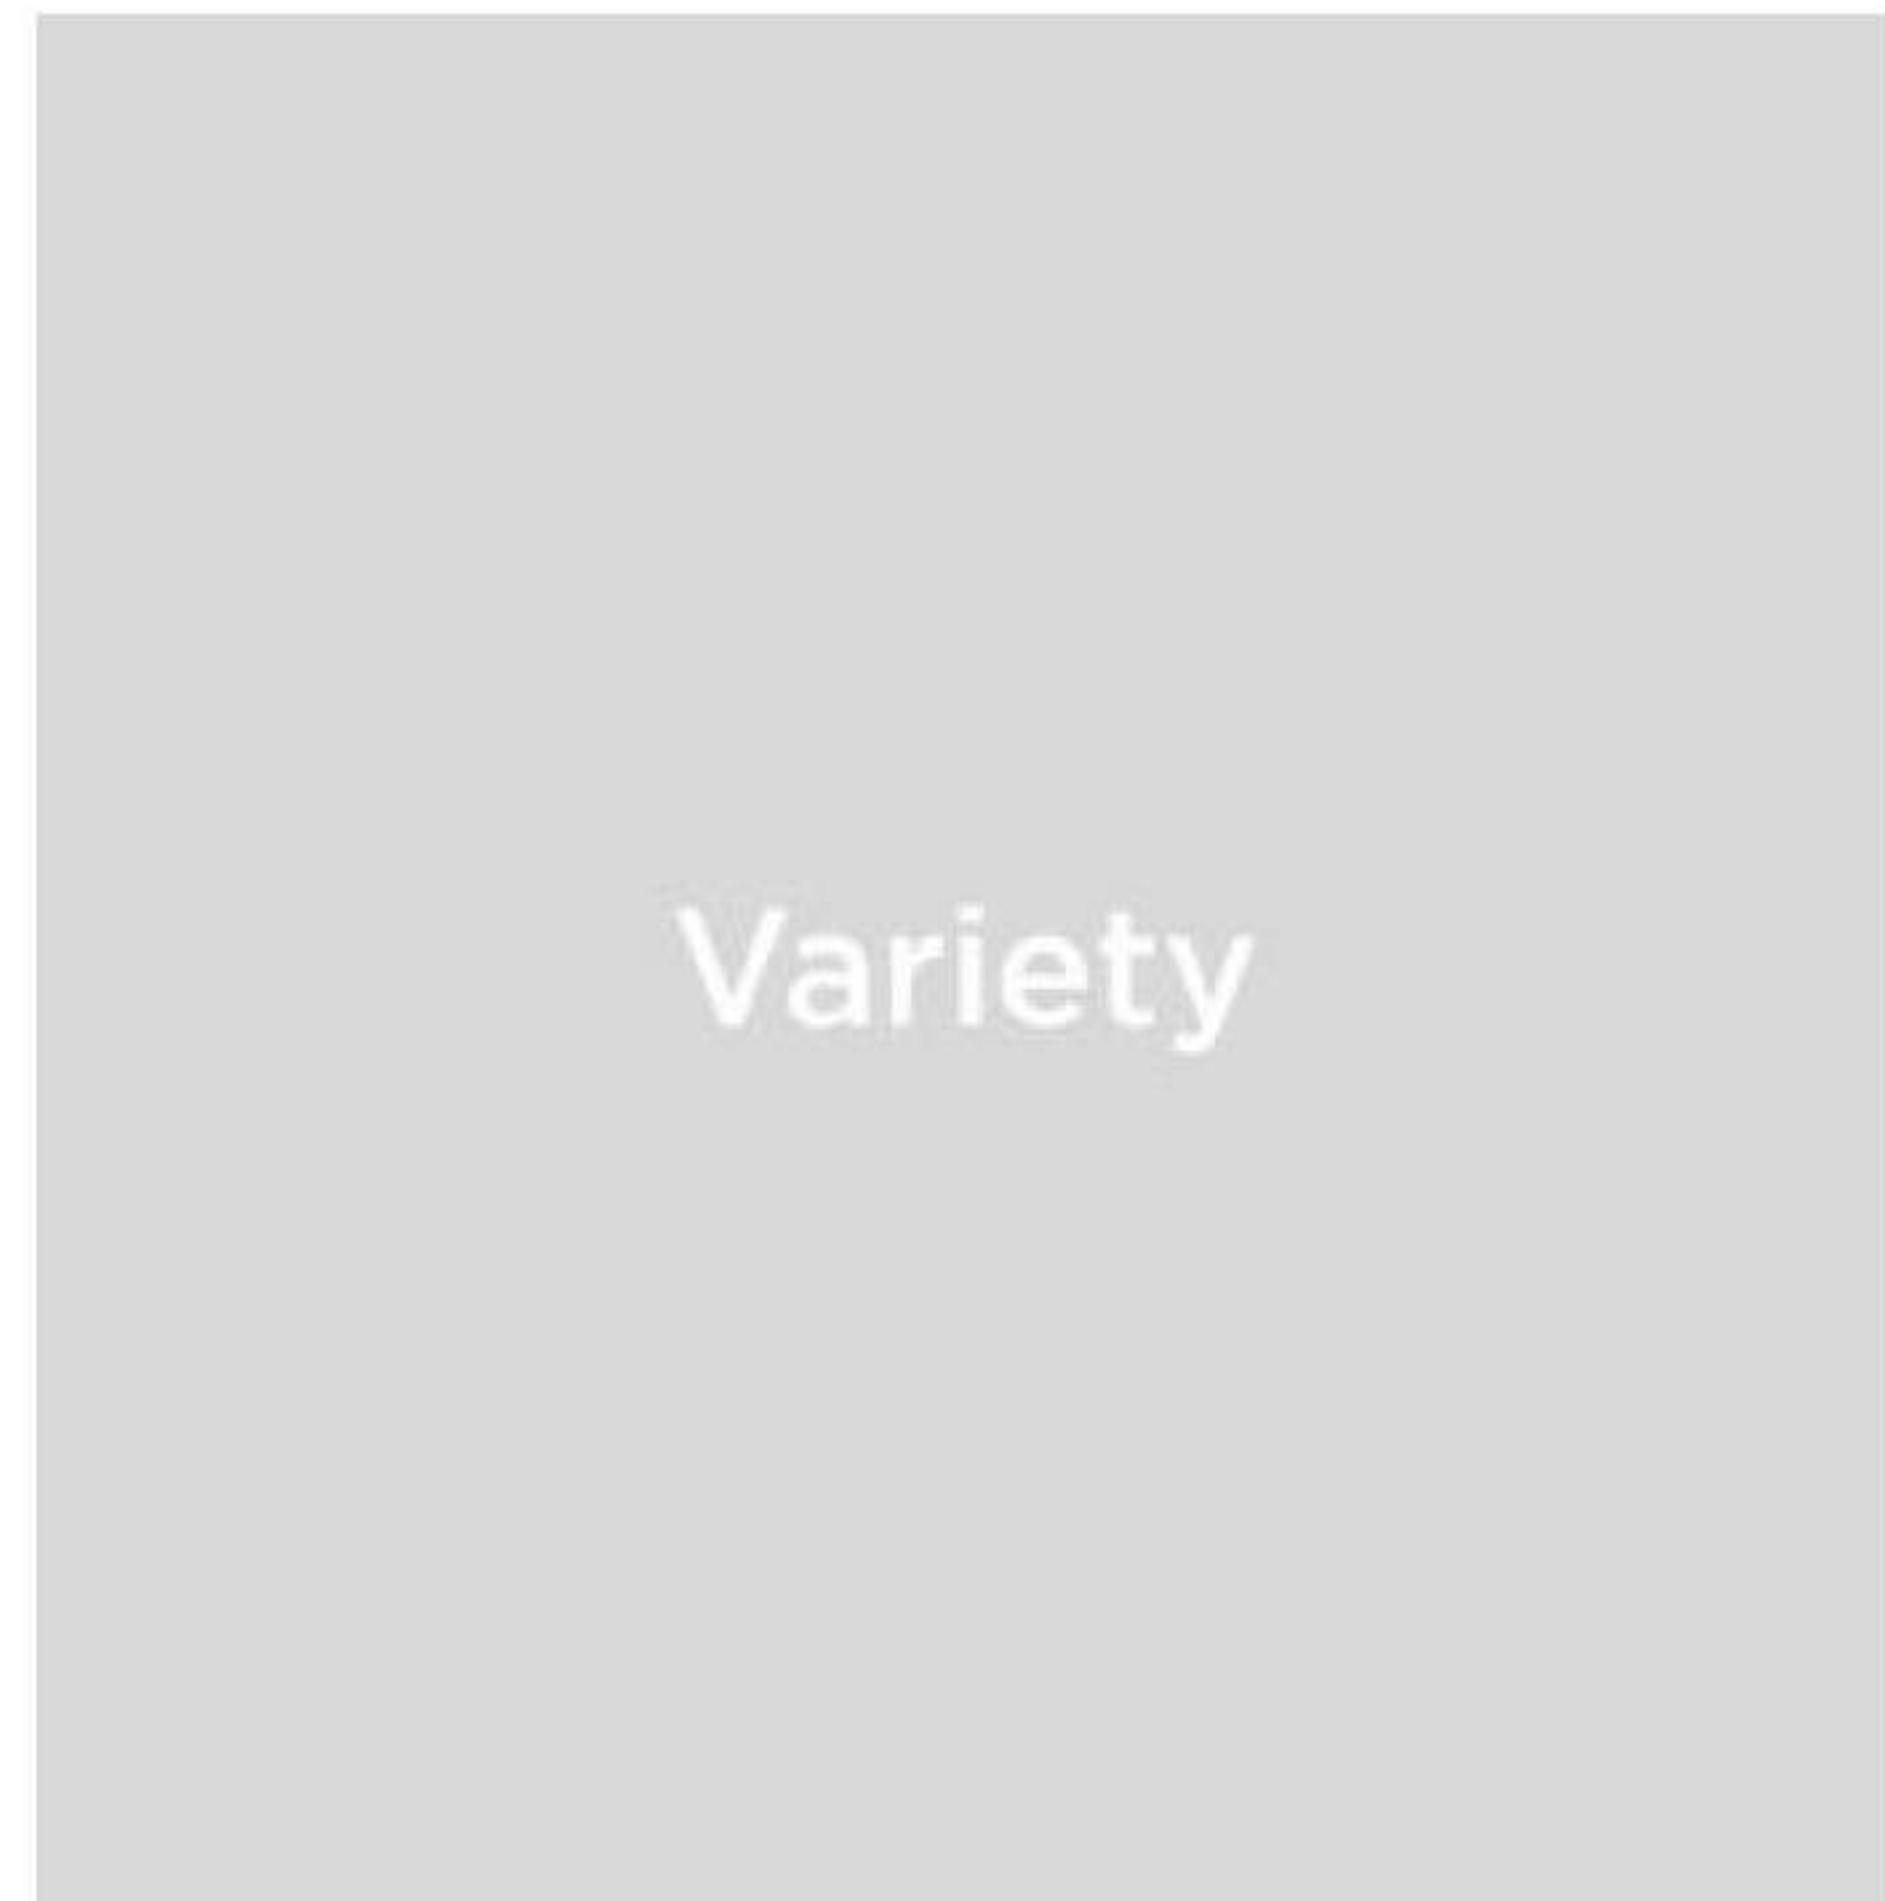
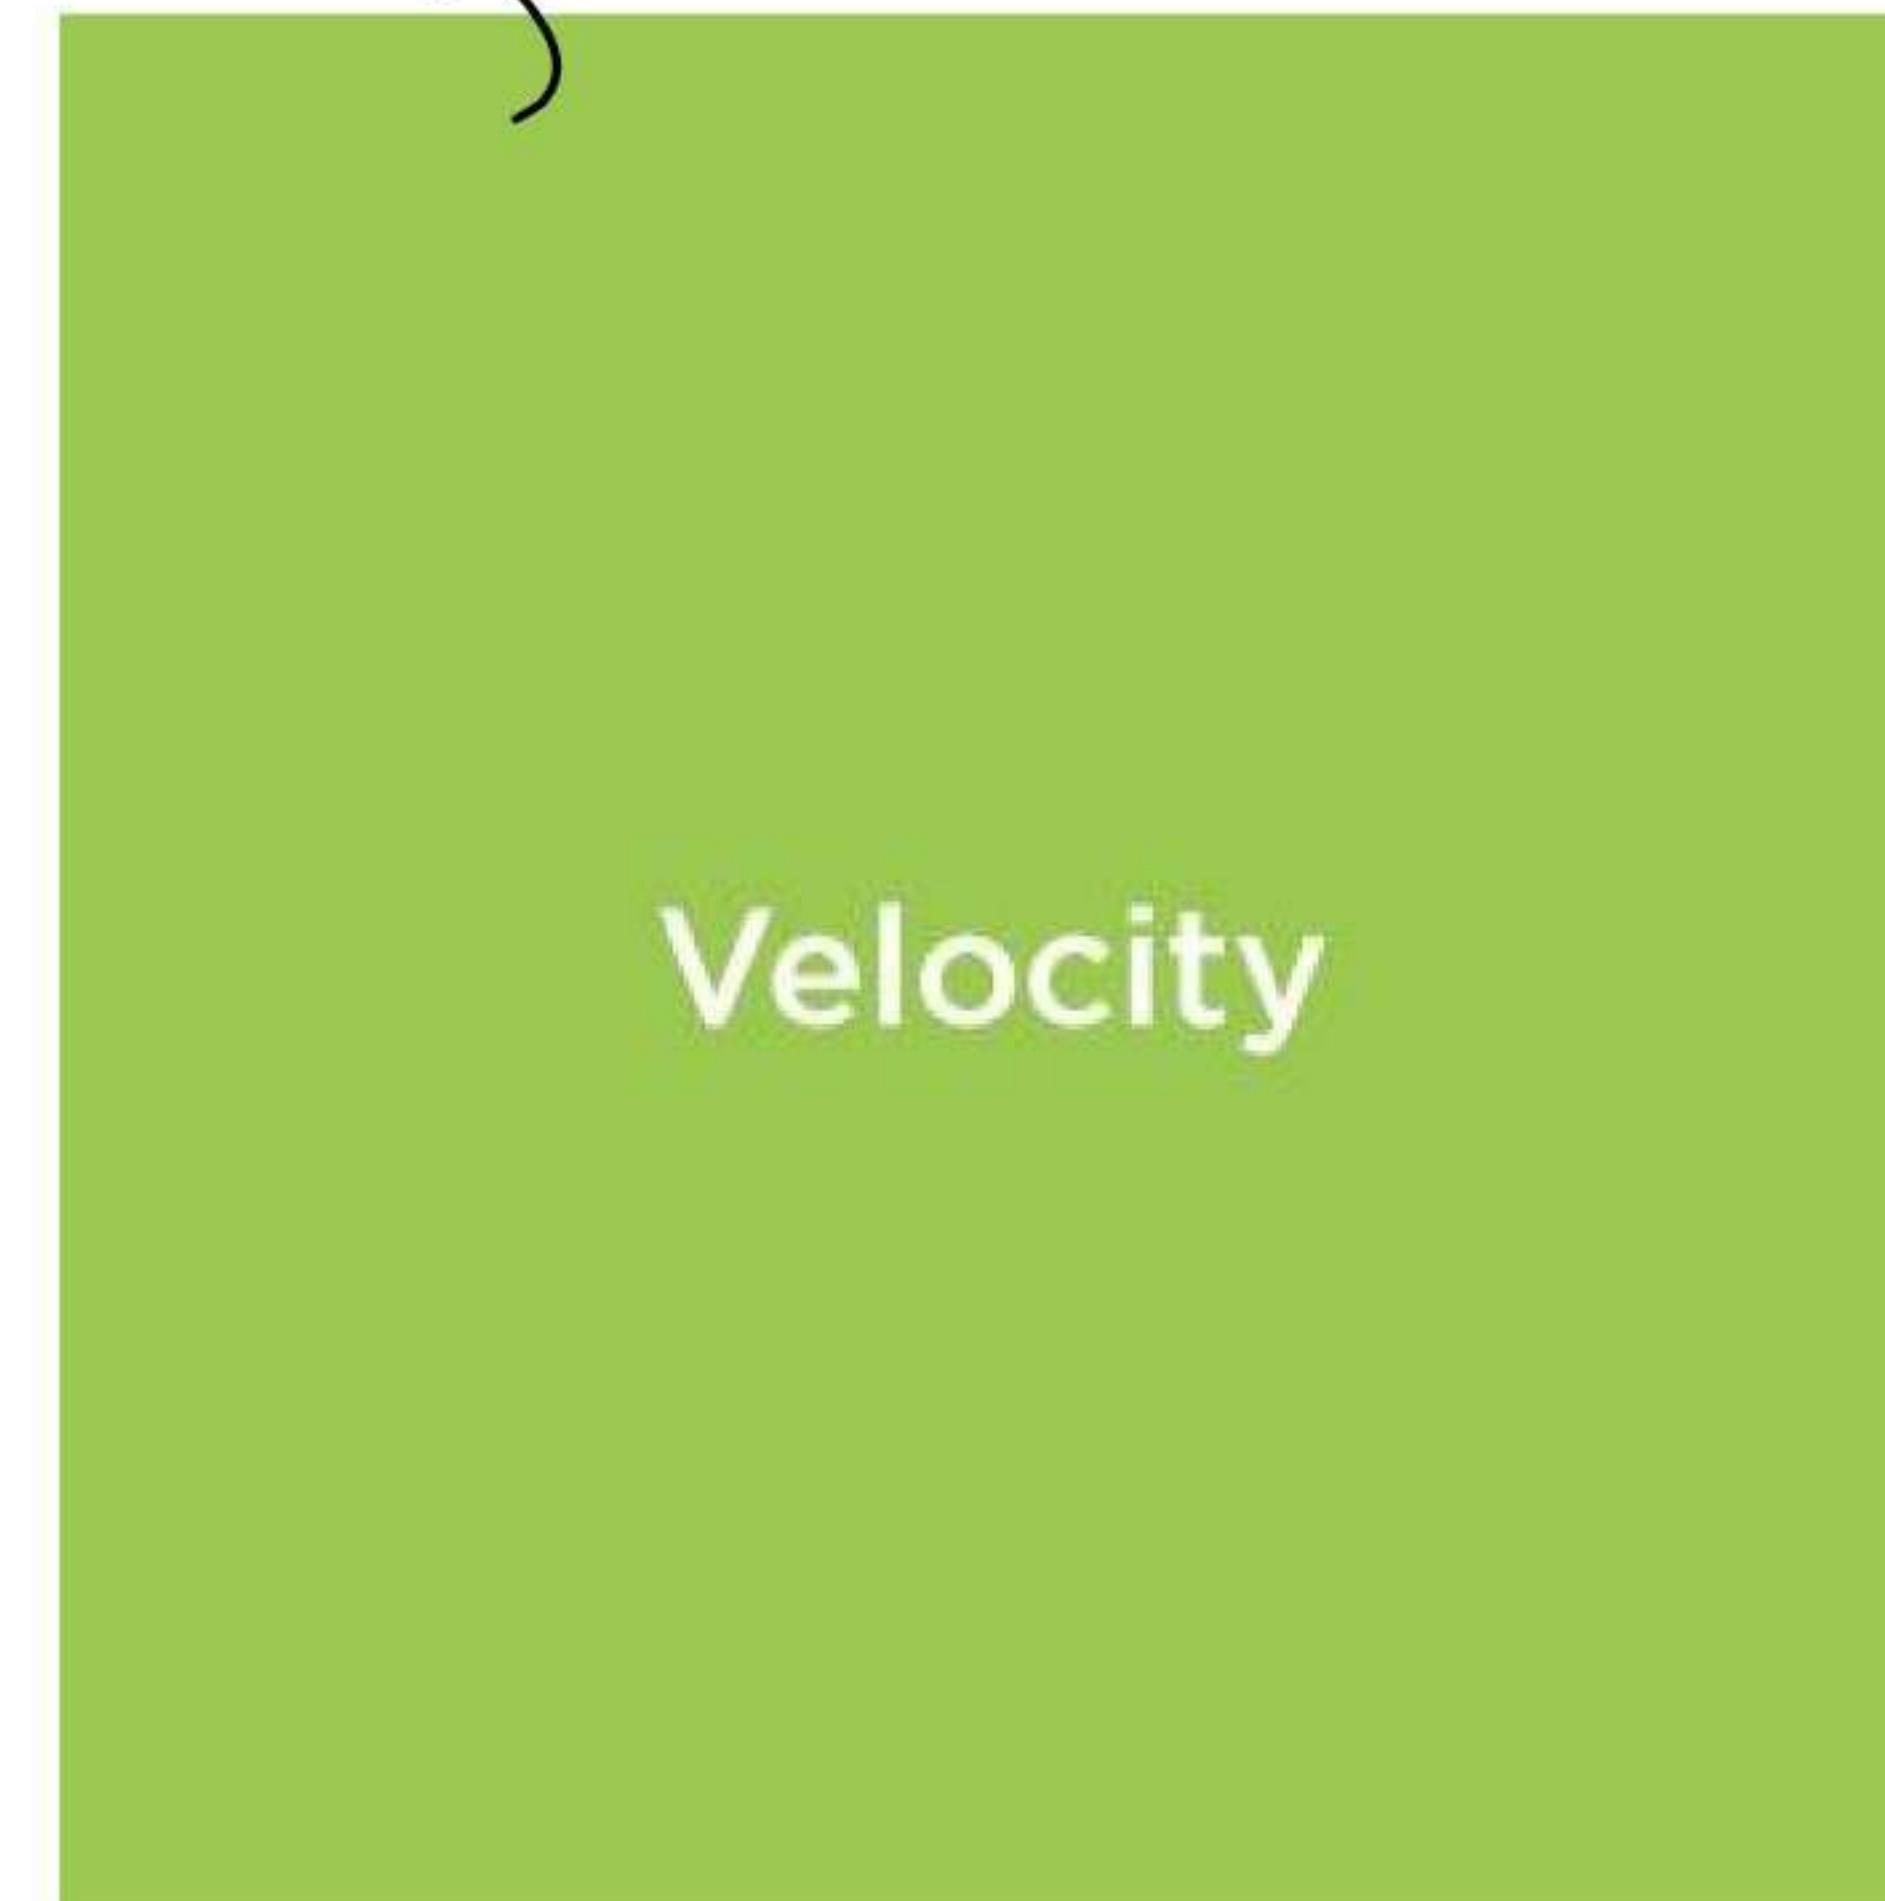
Big Data Characteristics

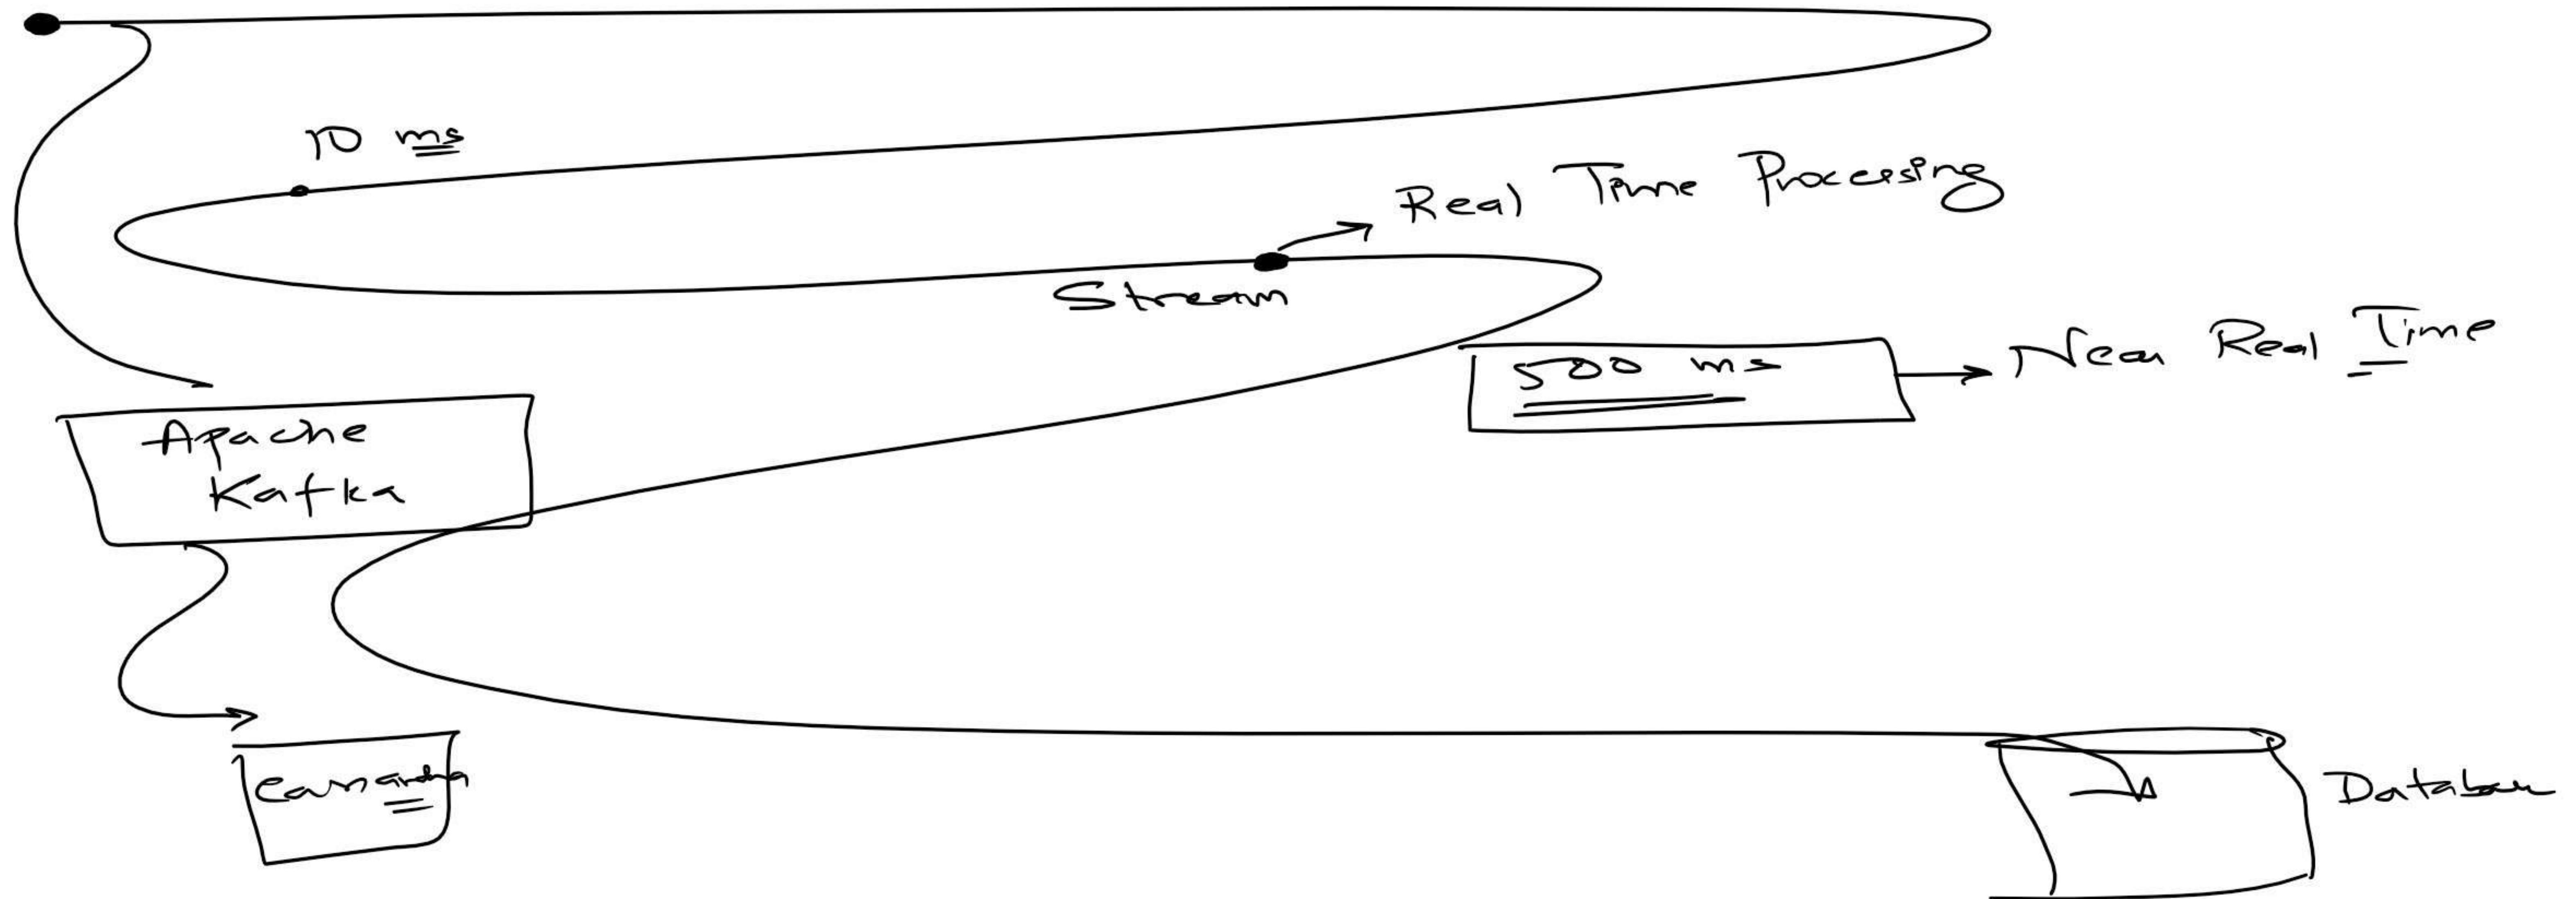


Big Data Characteristics

speed at which data is generated

$$\frac{RT}{\underline{NRT}}$$





Big Data Characteristics

Volume

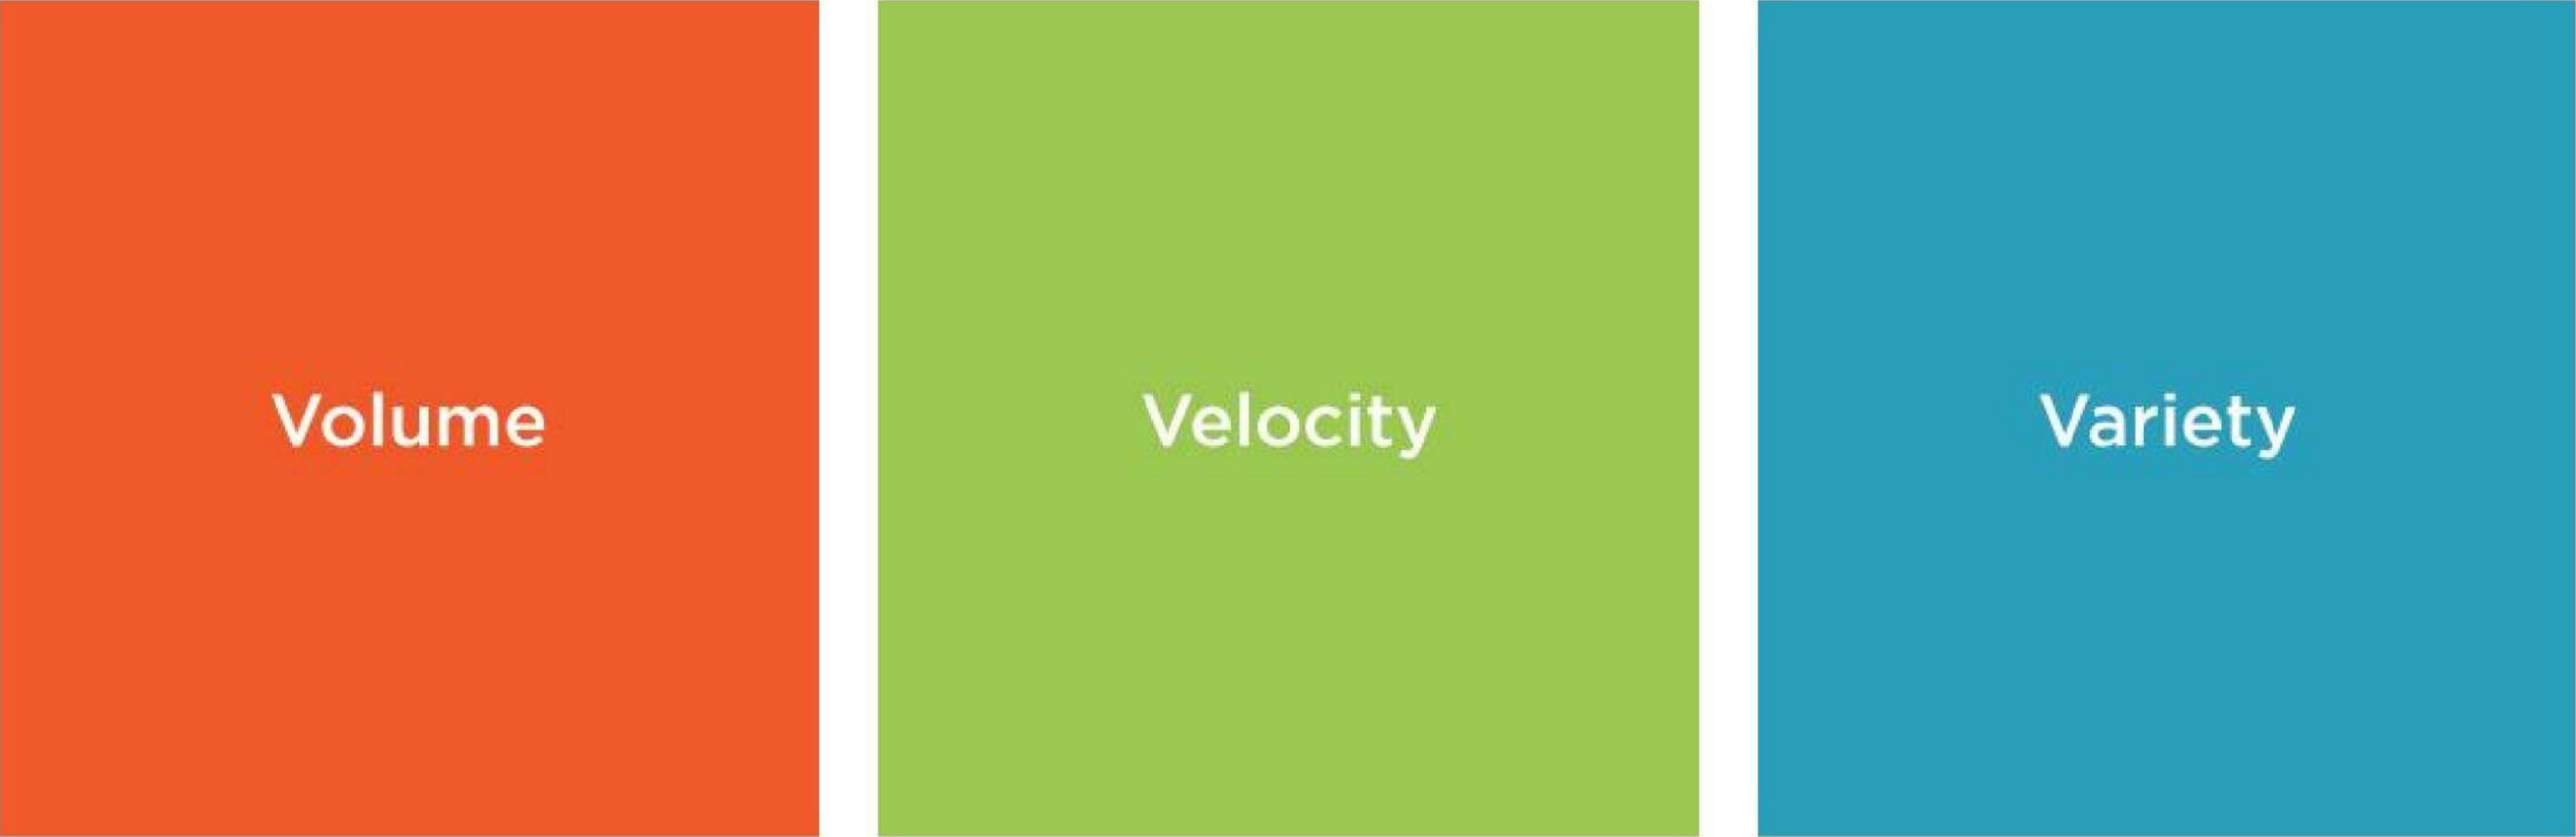
Velocity

Variety

different type of
data

Structured
Semi - Structured
Unstructured

Big Data Characteristics

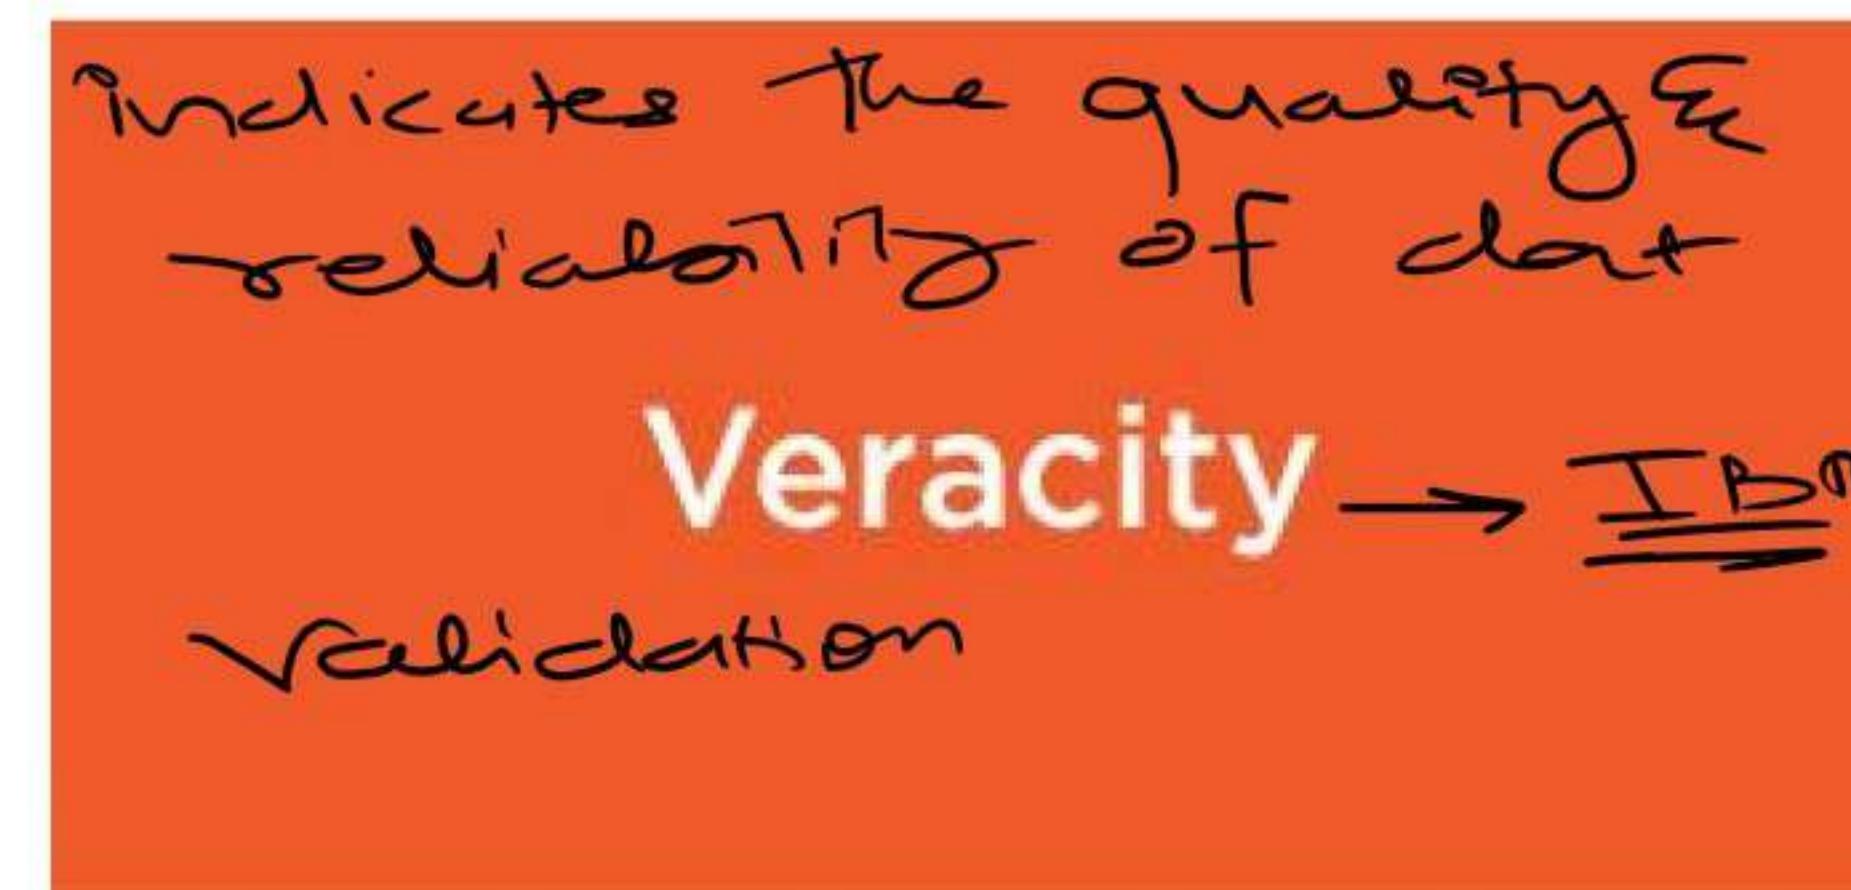


Volume

Velocity

Variety

Big Data Characteristics



Big Data Characteristics

Volume

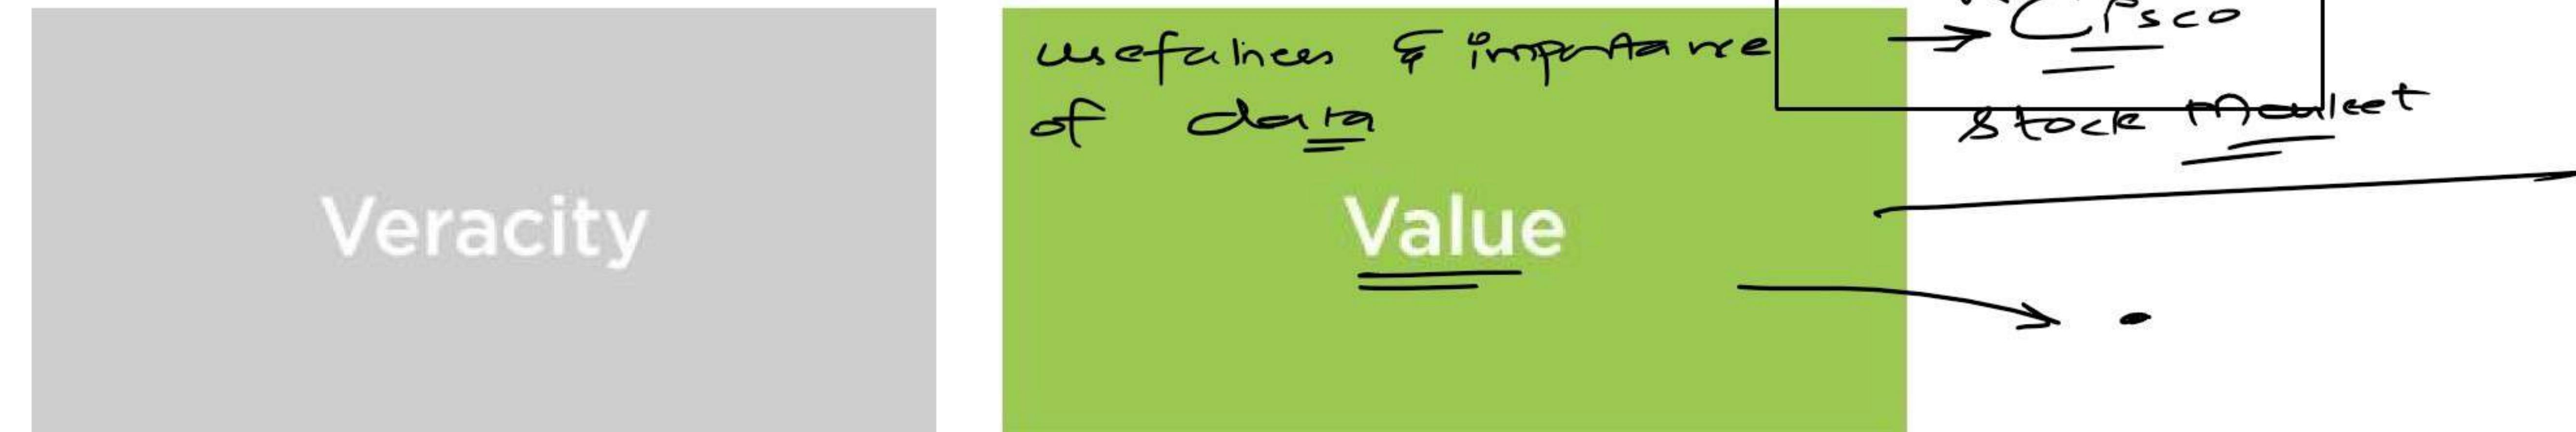
Velocity

Variety

Veracity

Value

Big Data Characteristics



{ Batch Processing { Stock Trading
{ Real Time Processing { Future Trading

Types of Big Data

Structured

Semi-structured

Unstructured

WHERE [Delimiters]

WHAT [what that part
exactly means]

UNSTRUCTURED

X

text extract, images, videos

audio, pdf

SEMI-STRUCTURED

✓

csv, tsv

X

STRUCTURED

✓

json, db-tables

✓

hello how are you?

1, Tom, IT, 25000, 2014-01-09

i am fine, how

{ id: 1,

how (unstructured?)

name: 'Tomcat',

no

age: "unknown"

}

id	name	city	age
-	-	-	-
-	-	-	-

Types of Big Data

Structured

Semi-structured

Unstructured

Types of Big Data

Structured

Semi-structured

Unstructured

Types of Big Data

Structured

Semi-structured

Unstructured

(A short) History of Big Data

Big Data Projects





Projects Directory

[Home](#)[Committees](#)[Projects](#)[Releases](#)[Statistics](#)[Timelines](#)

Search...

[About](#)**Project listings:** [By Name](#) [By Committee](#) [By Category](#) [By Programming Language](#) [By Number of Committers](#)

Projects by category:

- **big-data (49):**

- [Apache Accumulo](#)
- [Apache Airavata](#)
- [Apache Ambari](#)
- [Apache Apex \(in the Attic\)](#)
- [Apache Avro](#)
- [Apache Beam](#)
- [Apache Bigtop](#)
- [Apache BookKeeper](#)
- [Apache Calcite](#)
- [Apache Camel](#)
- [Apache CarbonData](#)
- [Apache CouchDB](#)
- [Apache Crunch \(in the Attic\)](#)
- [Apache Daffodil](#)
- [Apache DataFu \(Incubating\)](#)
- [Apache DirectMemory \(in the Attic\)](#)
- [Apache Drill](#)
- [Apache Edgent \(Incubating\)](#)
- [Apache Falcon \(in the Attic\)](#)
- [Apache Flink](#)
- [Apache Flume](#)
- [Apache Fluo](#)
- [Apache Fluo Recipes](#)

Big Data Goes Mainstream

What Do These Companies Have In Common?

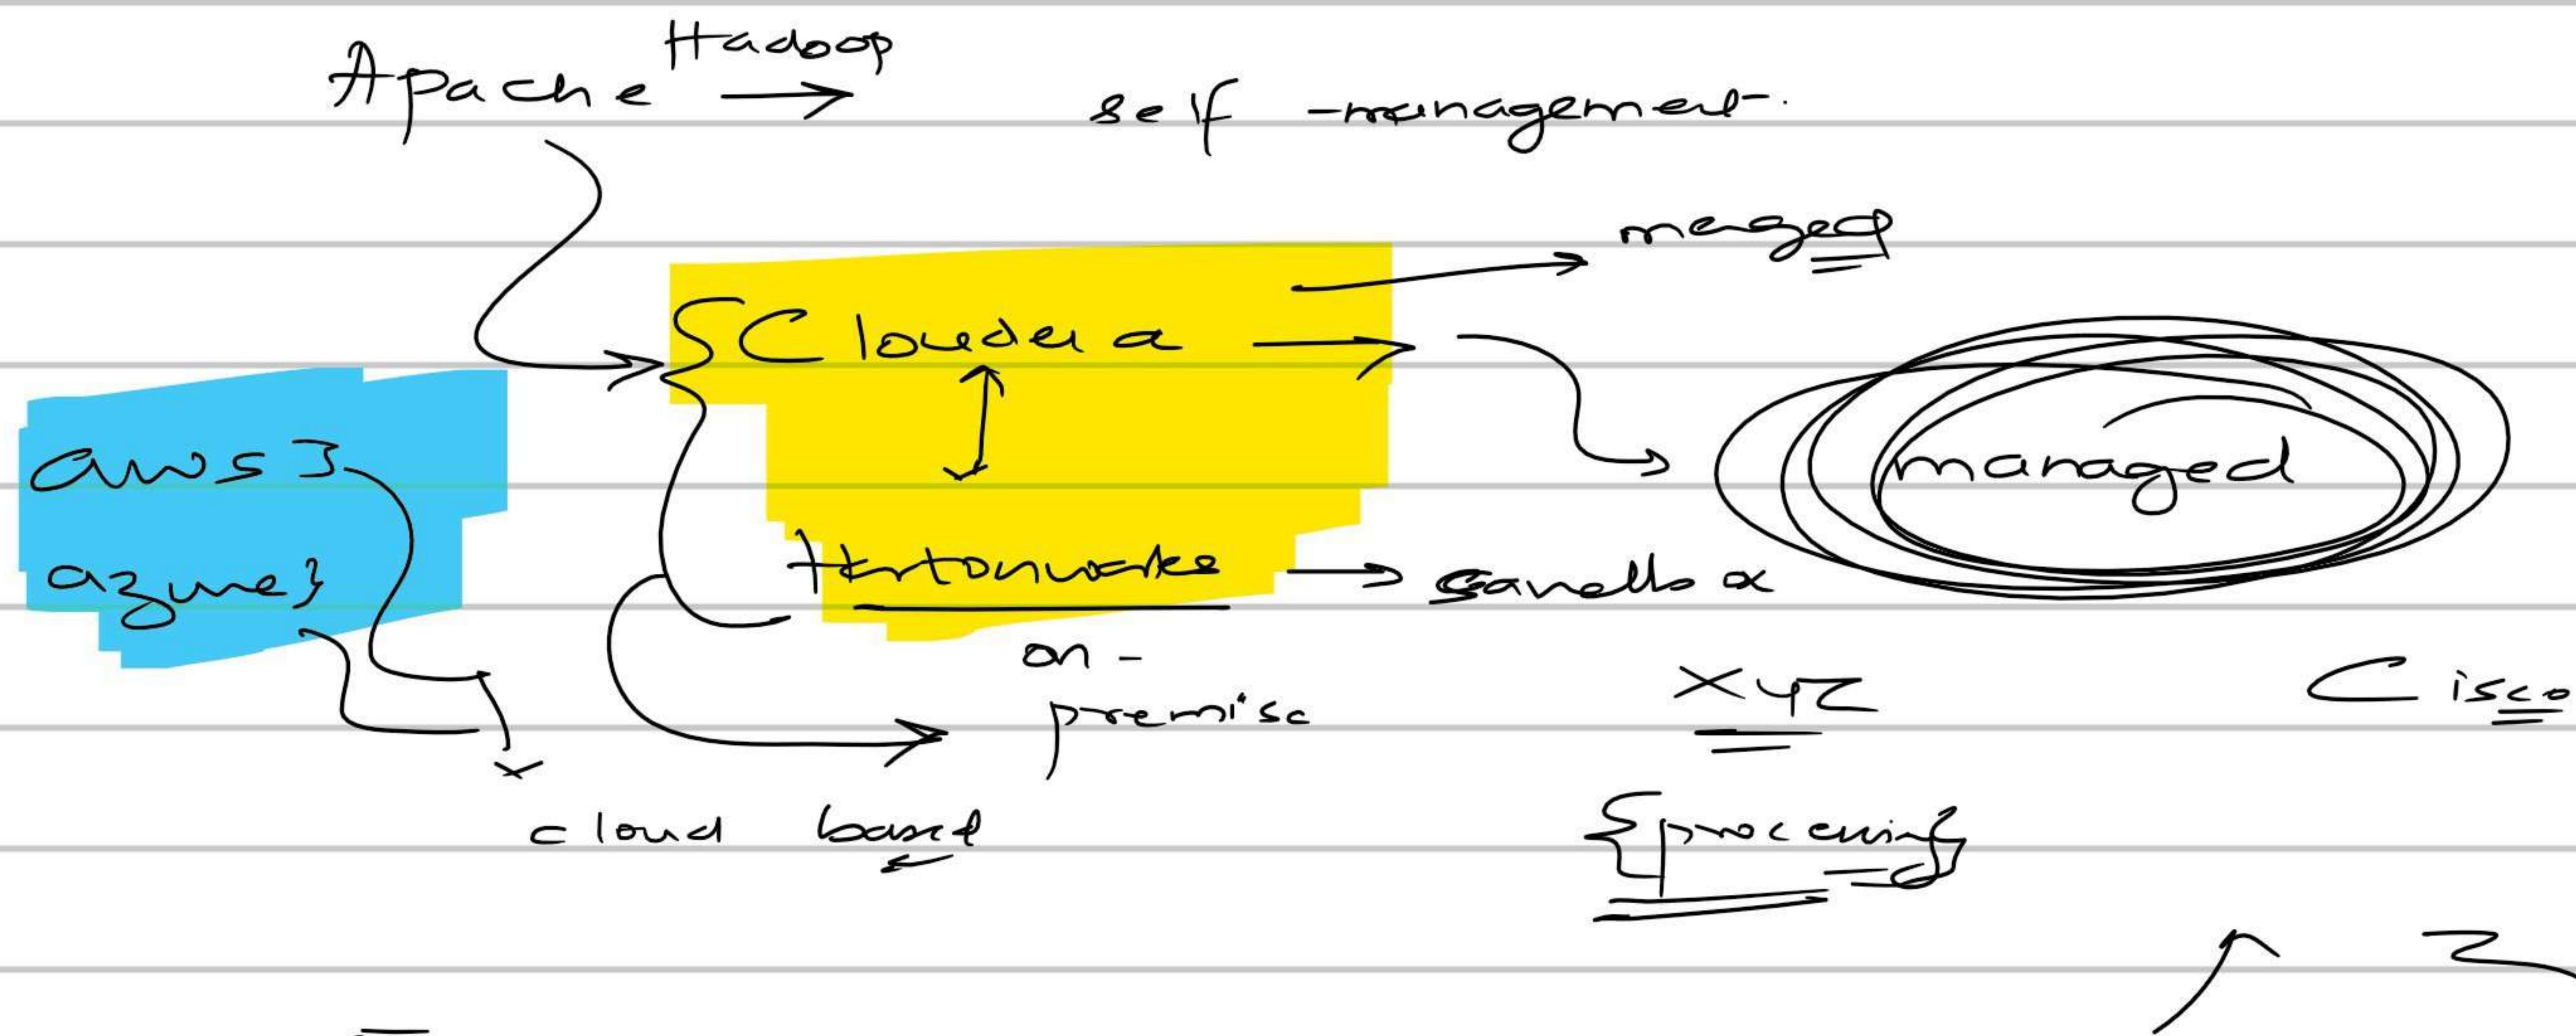


CLOUDERA

Google

yahoo!

aws



500 MB → Pen Drive

2 GB → SSD DD

Storage - 2 GB





Bay
8

Bay
16

PS
6



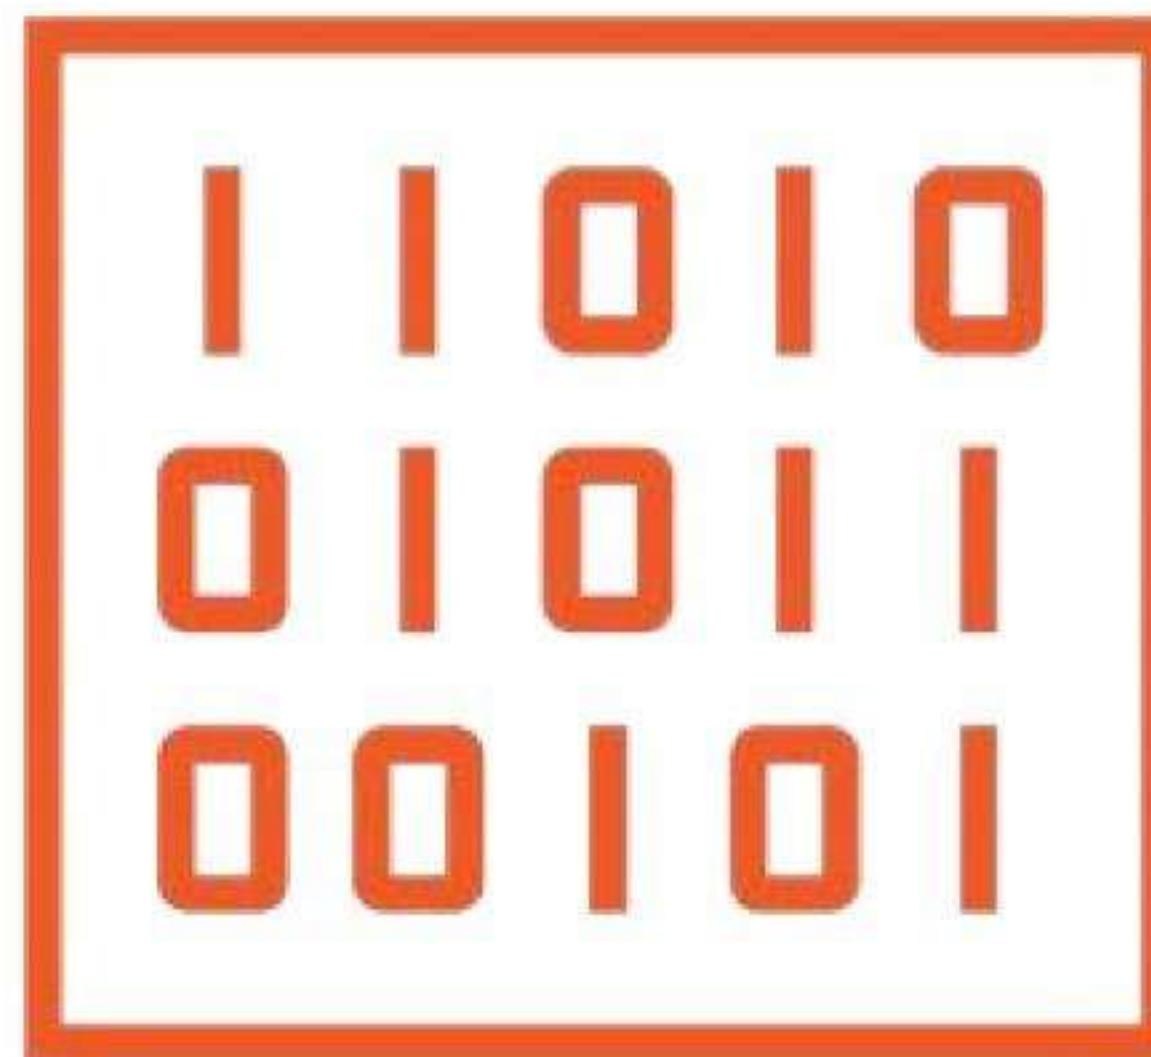


What Does Big Data Do and
HOW Does It work

How Big Data Works

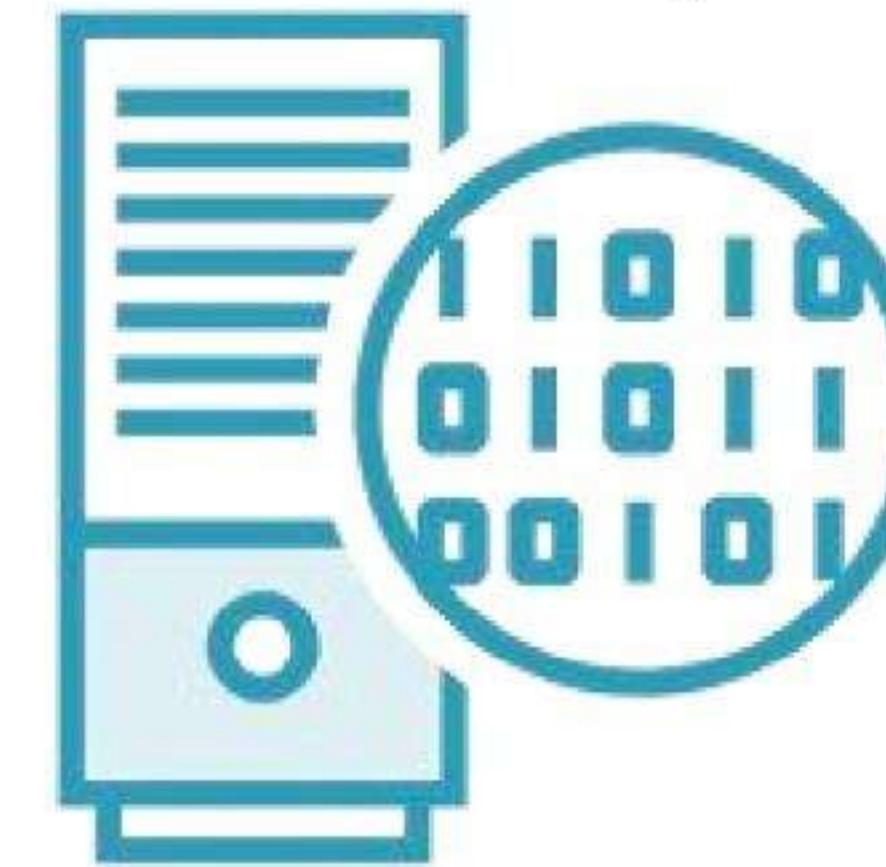
HDFS

storage

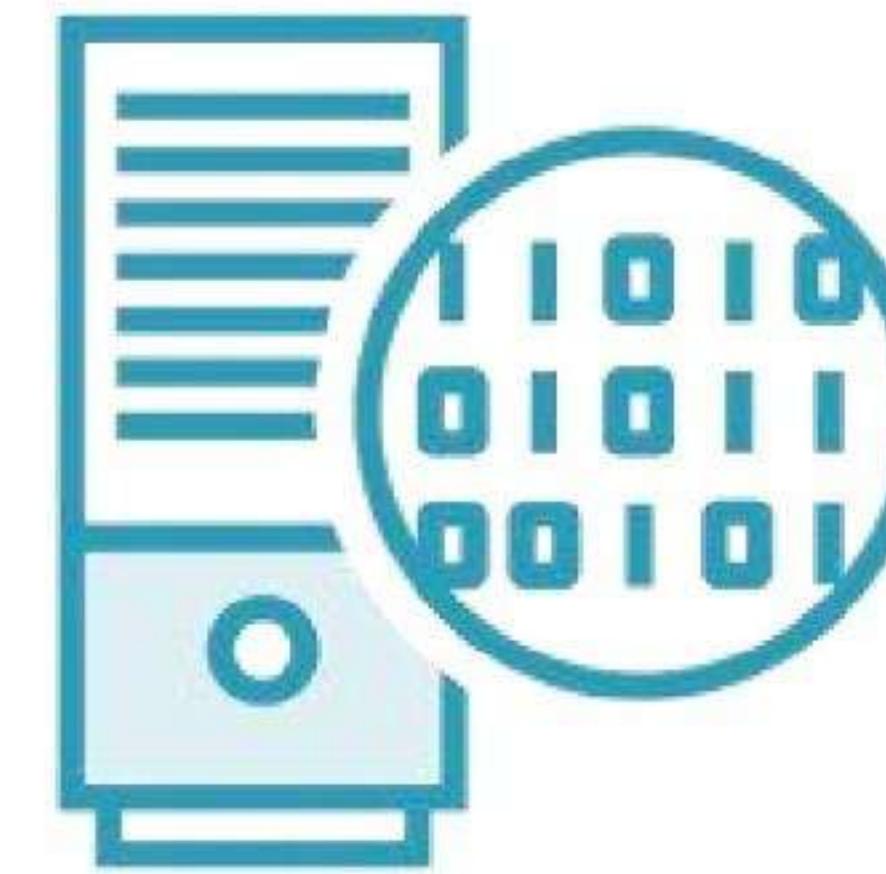
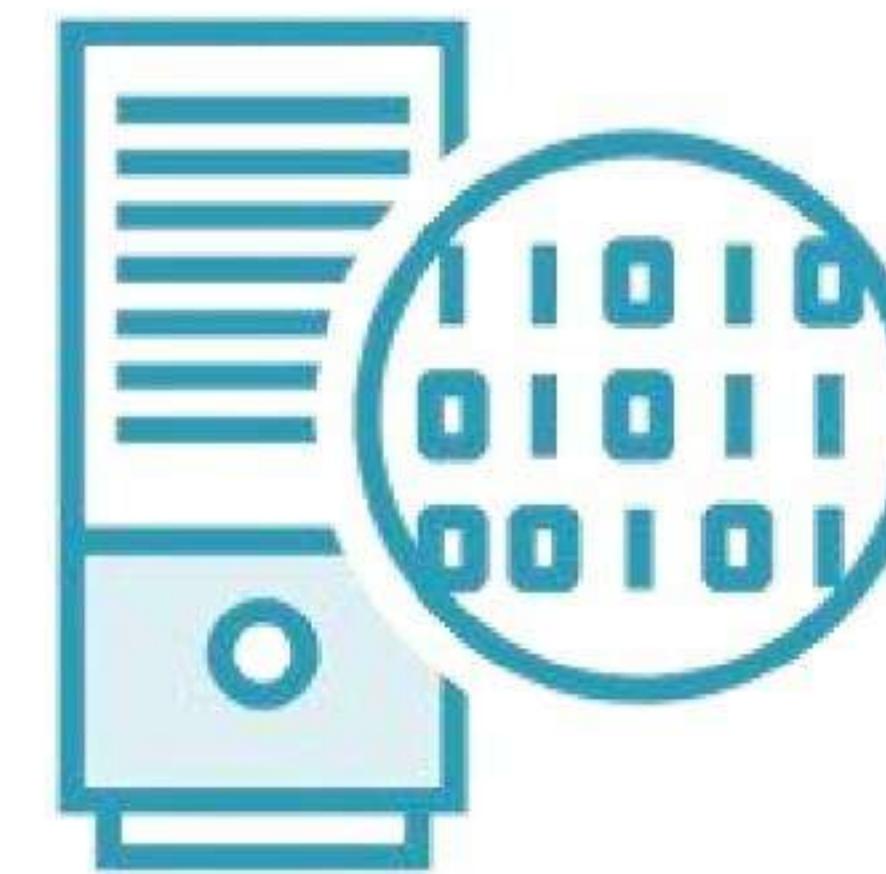


How Big Data Works

Processing
[Map Reduce]



How Big Data Works



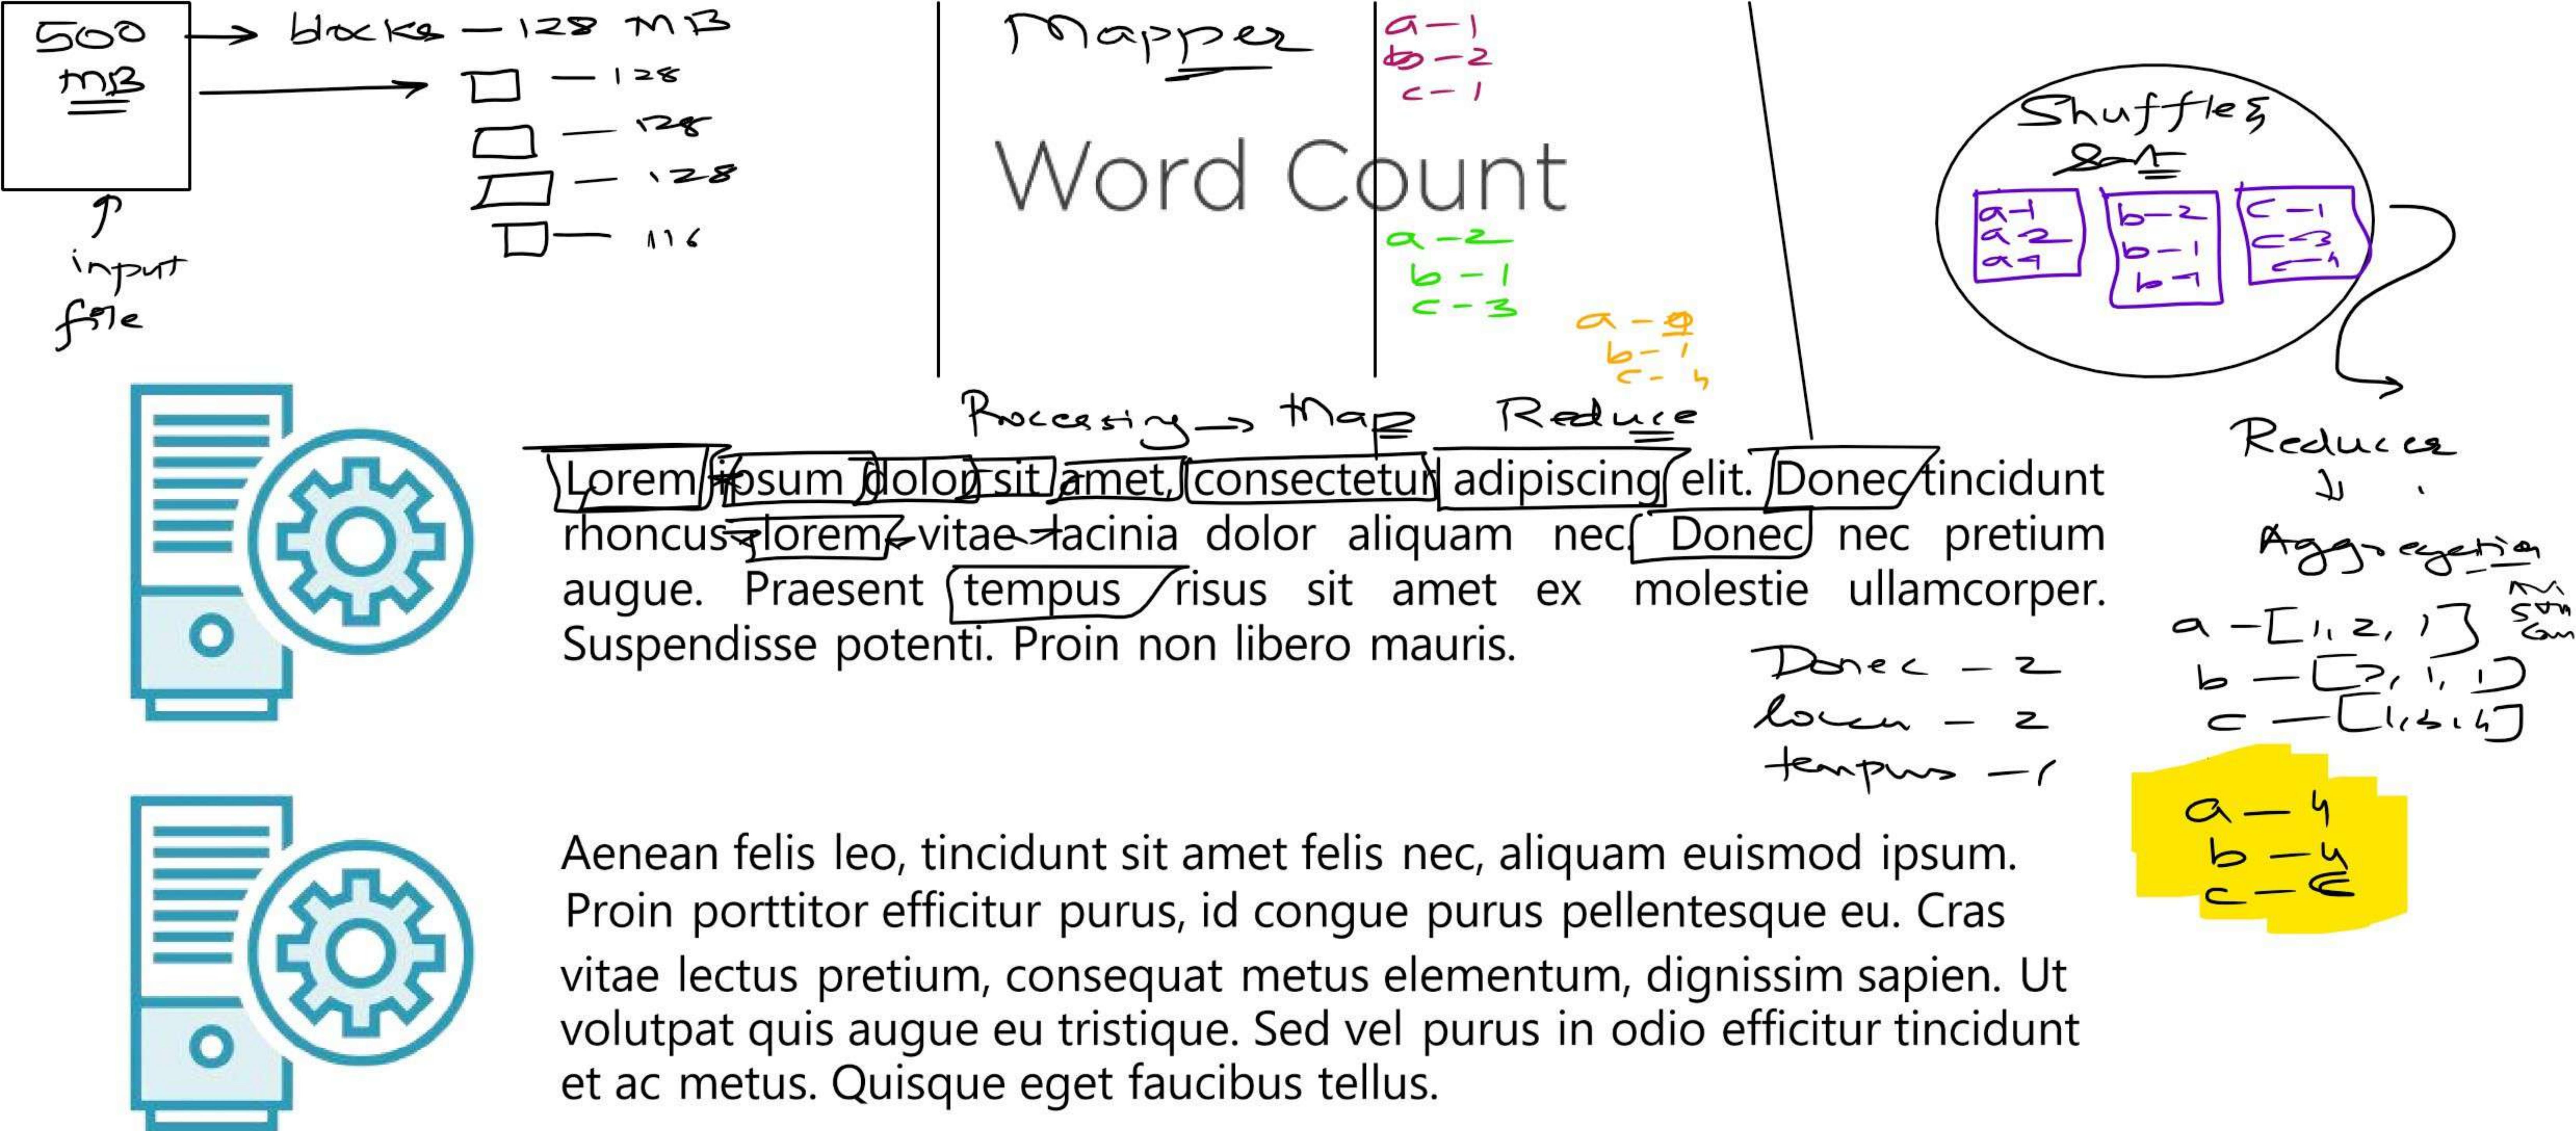
How Big Data Works

42

Word Count



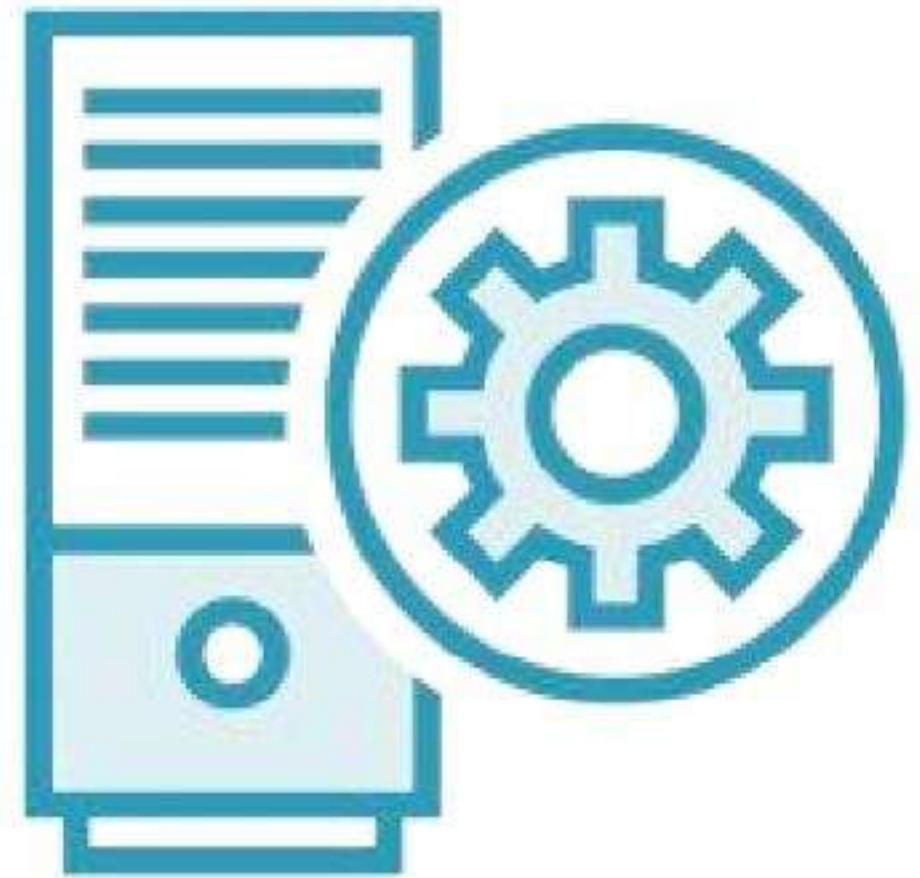
Lorem ipsum dolor sit amet, consectetur adipiscing elit. Donec tincidunt rhoncus lorem, vitae lacinia dolor aliquam nec. Donec nec pretium augue. Praesent tempus risus sit amet ex molestie ullamcorper. Suspendisse potenti. Proin non libero mauris. Aenean felis leo, tincidunt sit amet felis nec, aliquam euismod ipsum. Proin porttitor efficitur purus, id congue purus pellentesque eu. Cras vitae lectus pretium, consequat metus elementum, dignissim sapien. Ut volutpat quis augue eu tristique. Sed vel purus in odio efficitur tincidunt et ac metus. Quisque eget faucibus tellus.



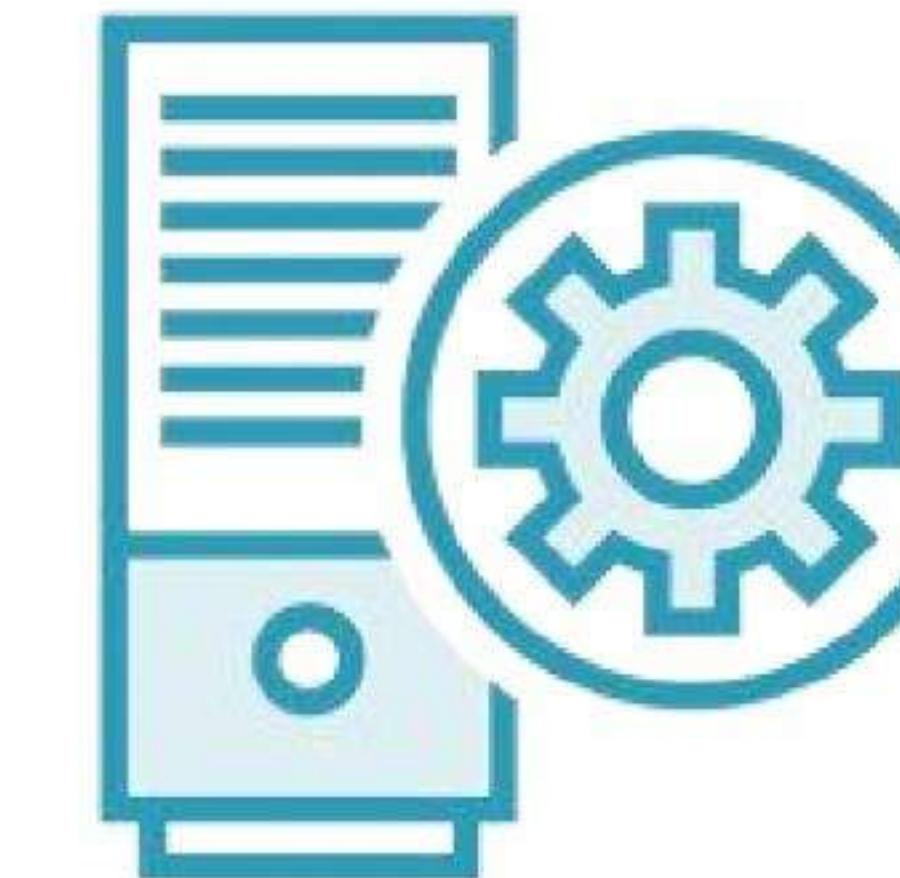
Word Count



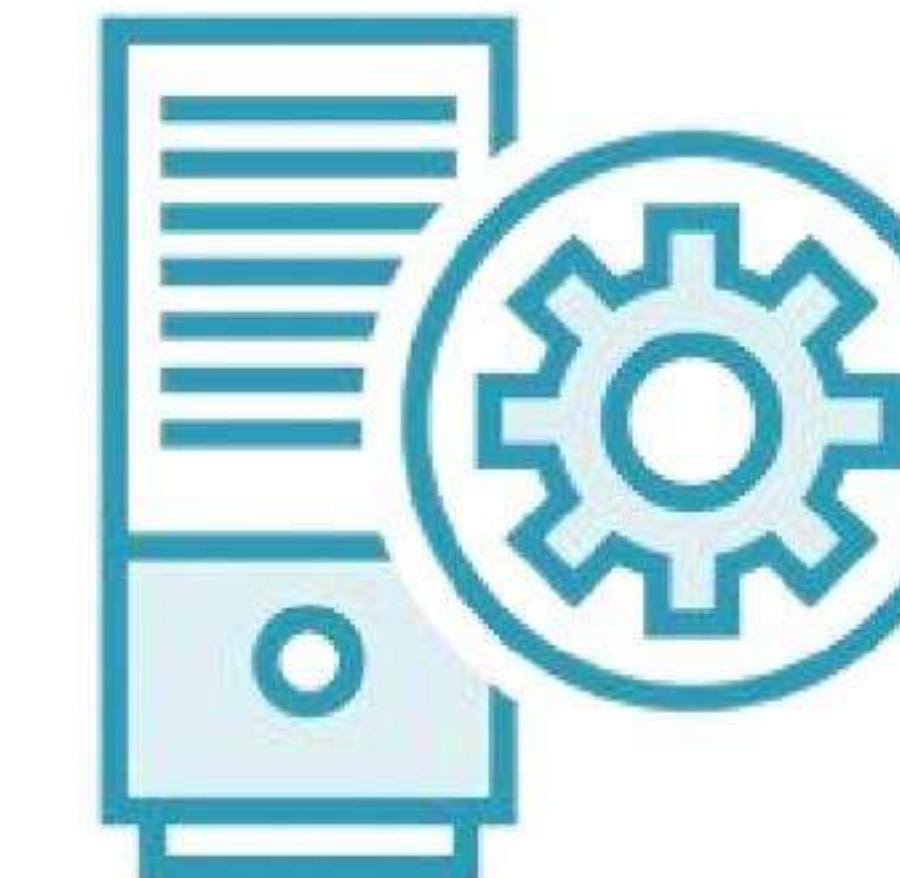
Lorem ipsum dolor sit amet, consectetur adipiscing elit. Donec tincidunt rhoncus lorem, vitae lacinia dolor aliquam nec.



Aenean felis leo, tincidunt sit amet felis nec, aliquam euismod ipsum. Proin porttitor efficitur purus, id congue purus pellentesque eu.



Donec nec pretium augue. Praesent tempus risus sit amet ex molestie ullamcorper. Suspendisse potenti. Proin non libero mauris.



Cras vitae lectus pretium, consequat metus elementum, dignissim sapien. Ut volutpat quis augue eu tristique. Sed vel purus in odio efficitur tincidunt et ac metus. Quisque eget faucibus tellus.

Big Data platforms

Modern Data platform

A future-proof architecture
for Business Analytics

Your Enterprise Data Cloud



Data Hub Clusters



Data Engineering



Data Warehouse



Operational Database



Machine Learning

Control Plane



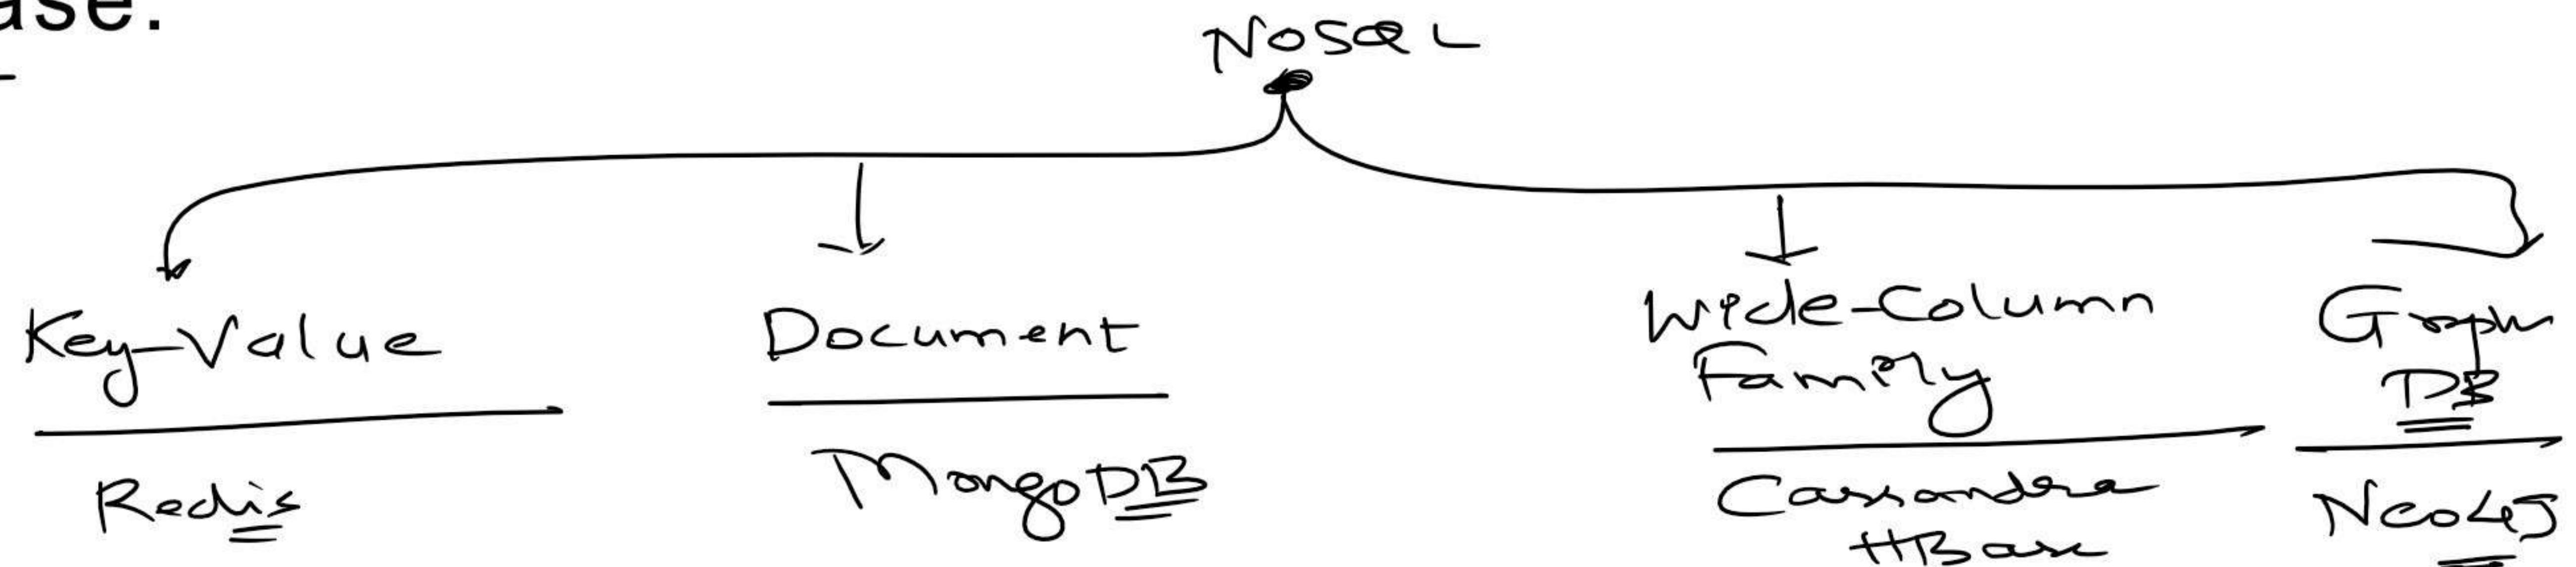
Types of Databases

NoSQL

Not Only SQL

NoSQL Database

Generic term used for any non-relational database.



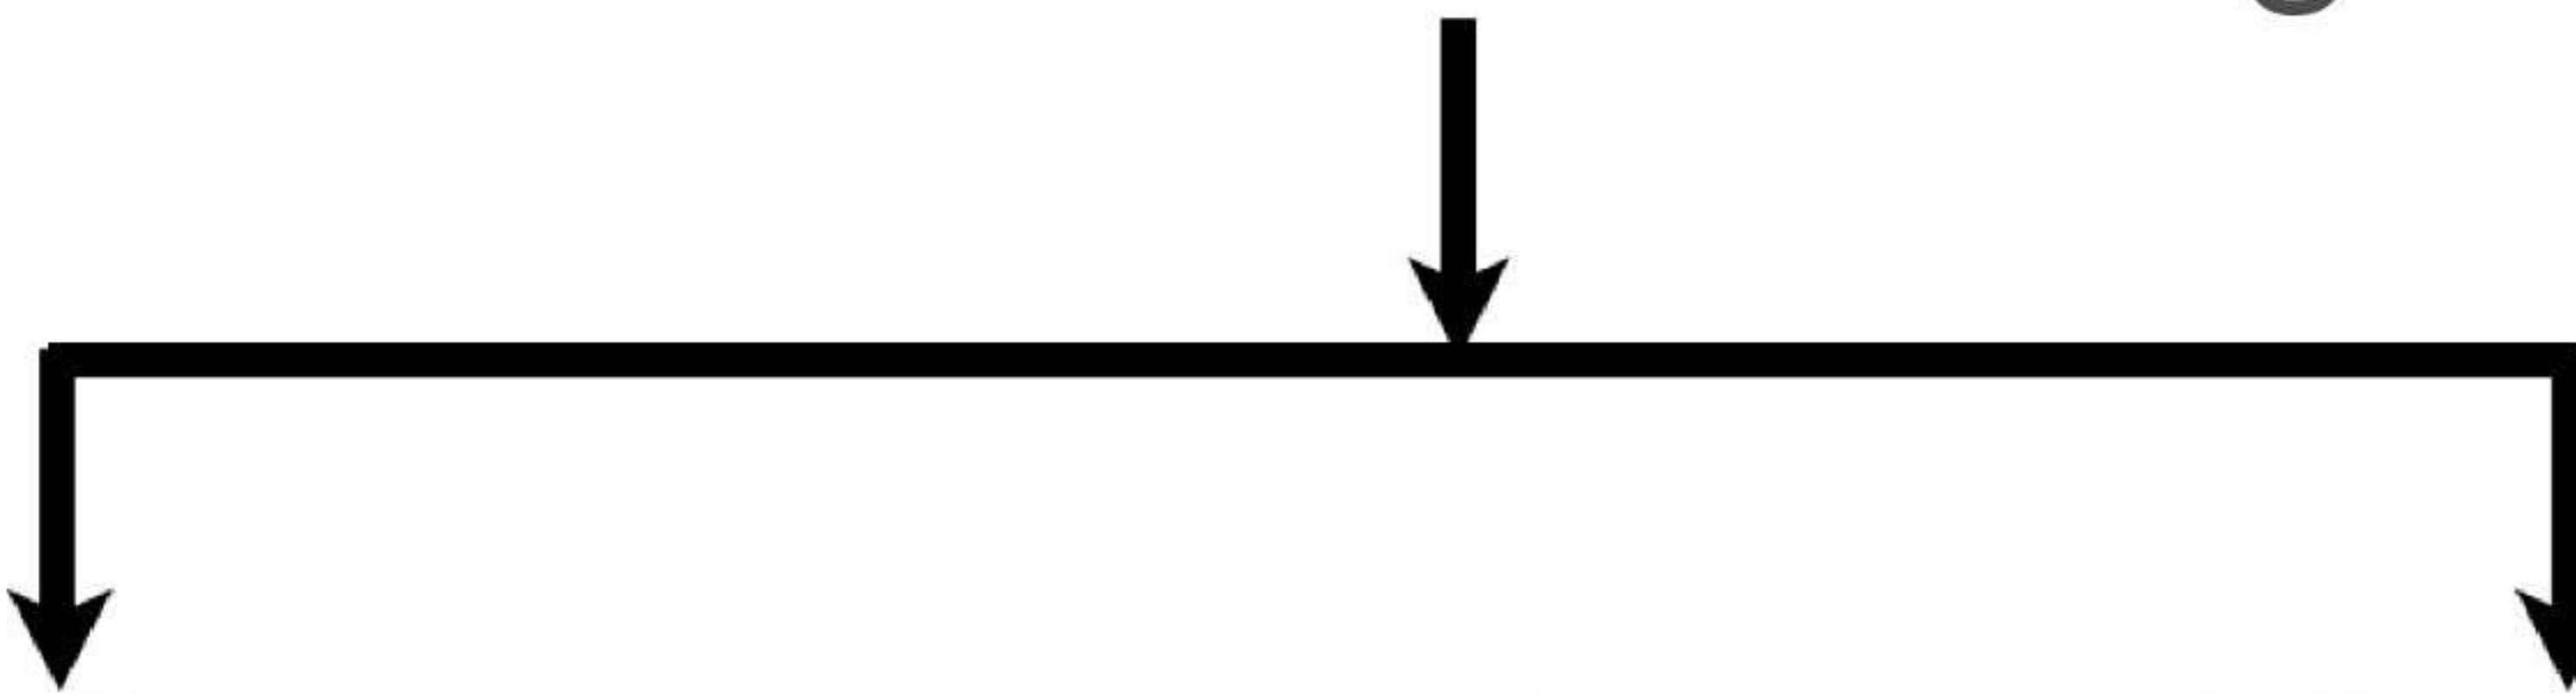
Relational Database

Generic term used for any database that stores data logically organized into relations - tables with rows and columns.

Database Technologies

NoSQL Databases

Relational Databases

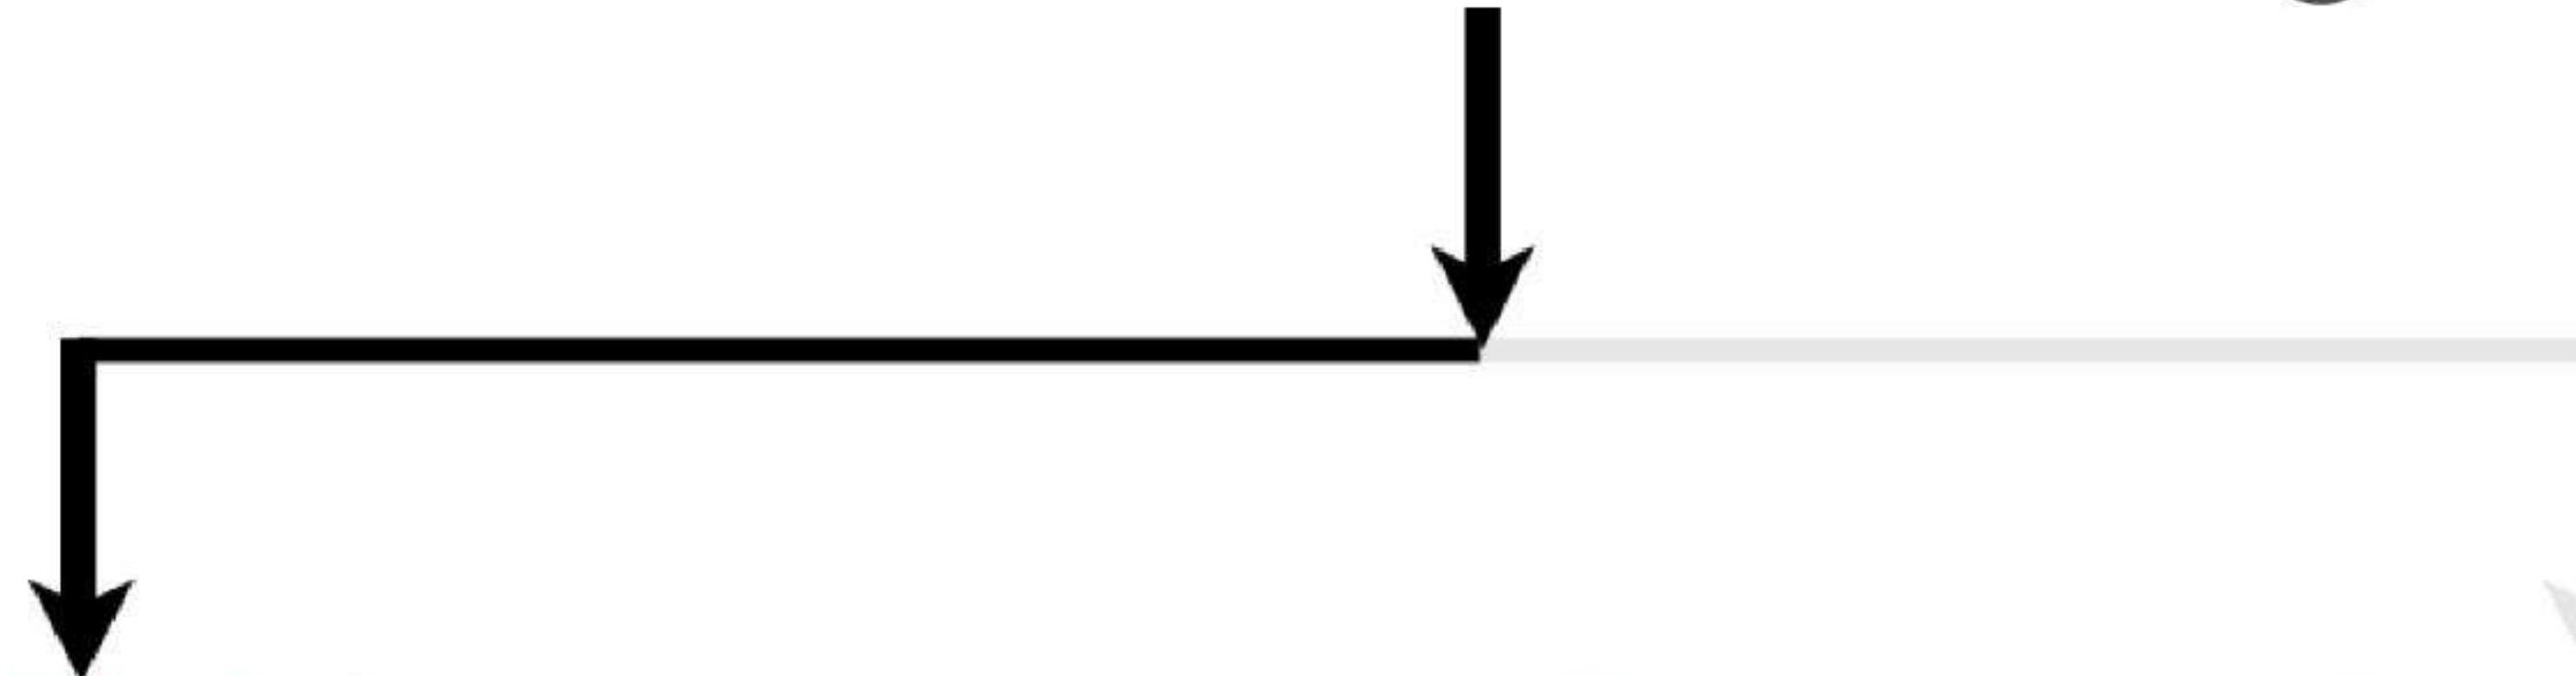


Database Technologies

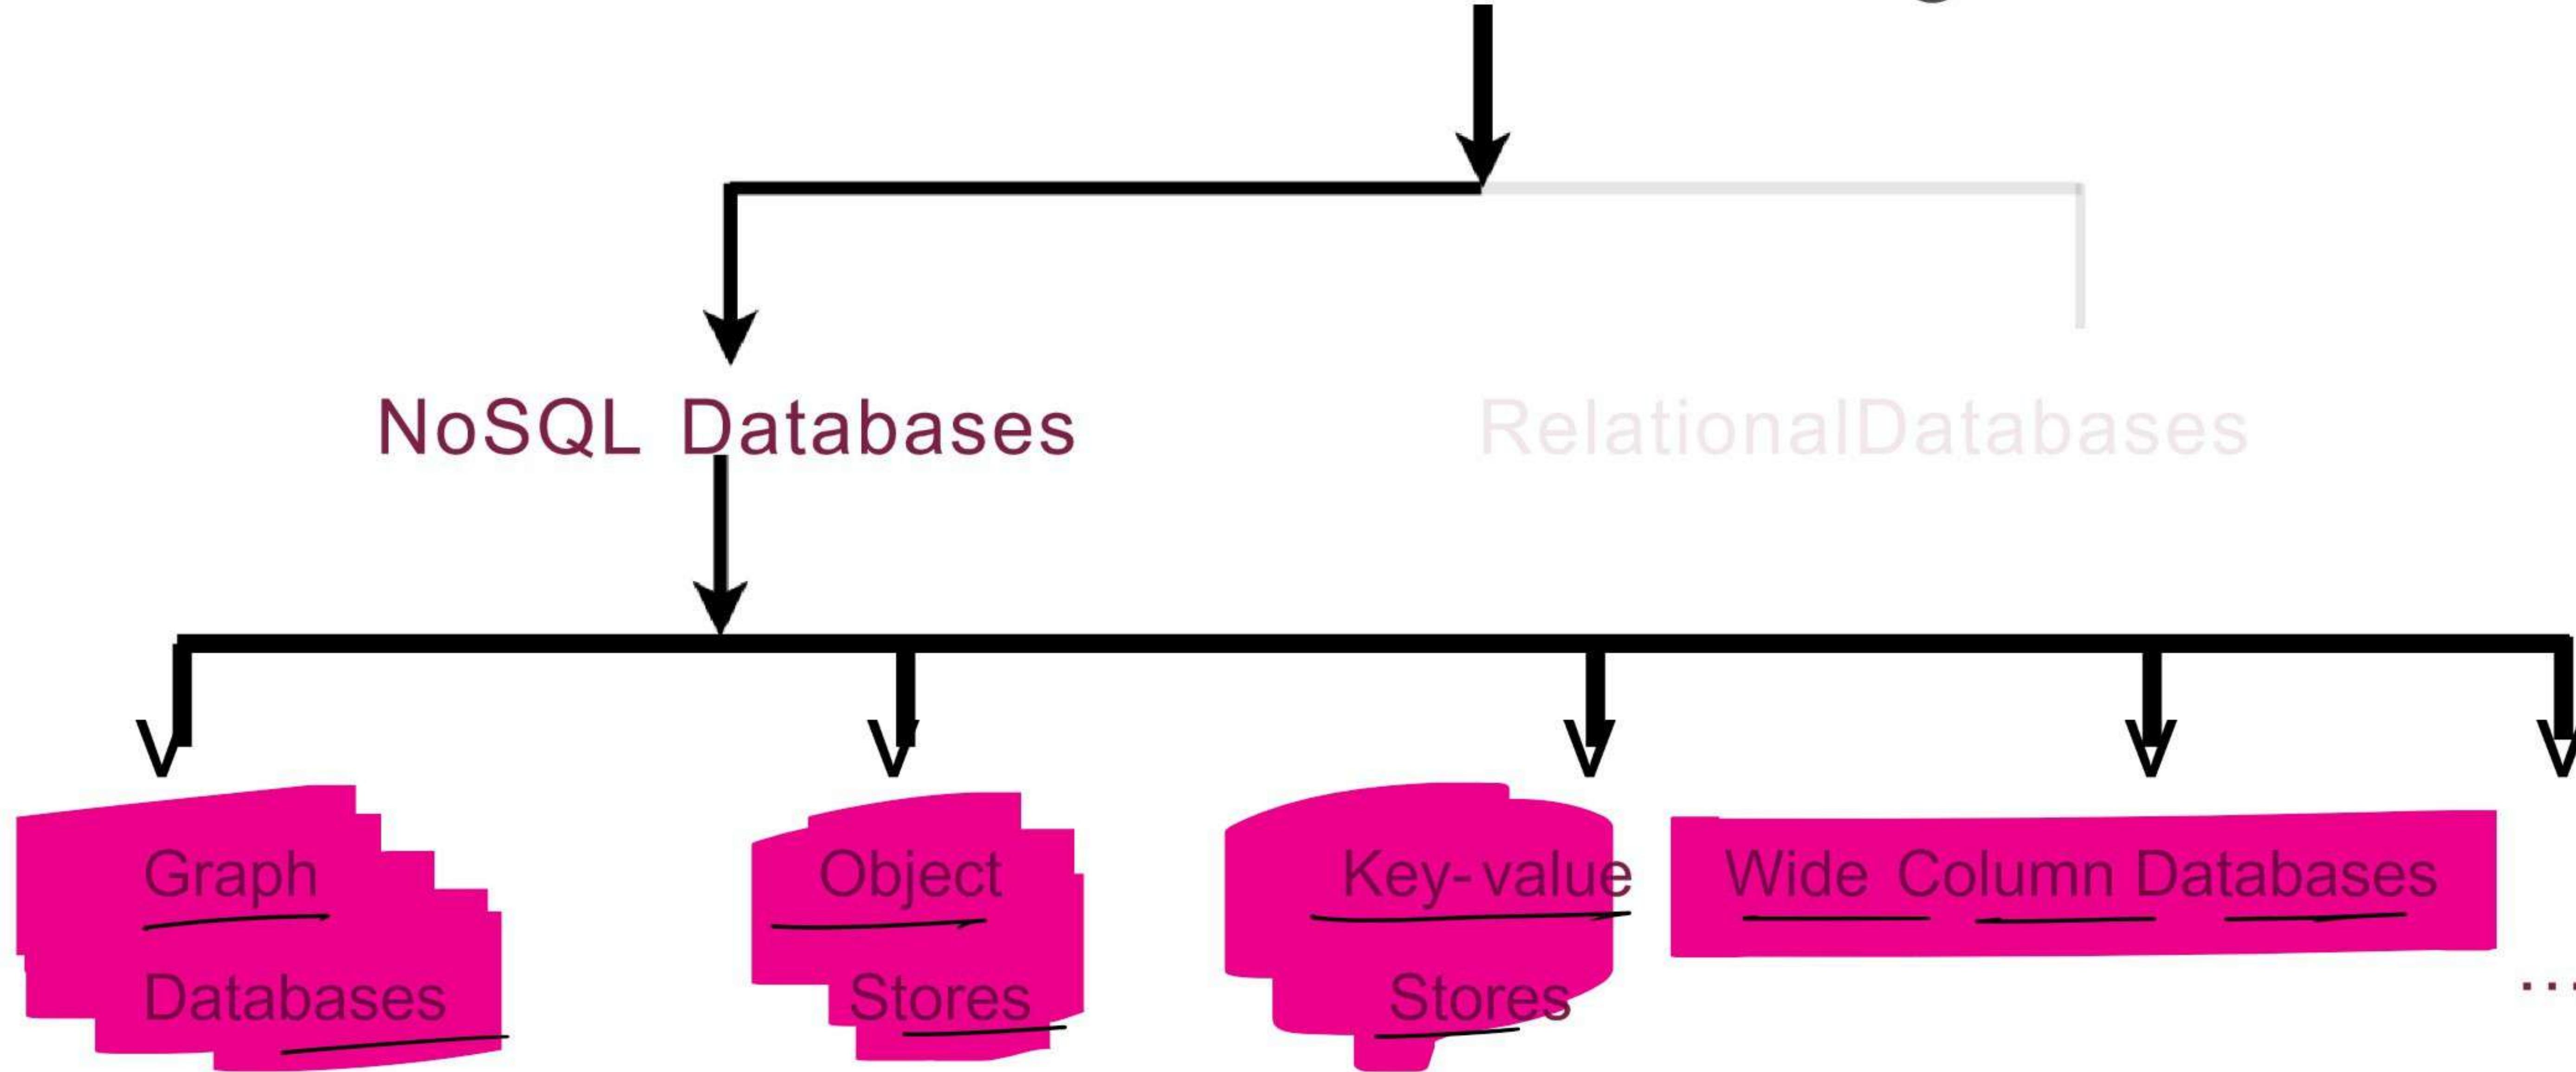
NoSQL Databases

Variety of data models

Relational Databases



Database Technologies



Structured and Unstructured Data

Problem: Creating Customer Profiles



Imagine you work at an electronics company with different sales channels

- Online store
- Physical stores
- Third-parties

Problem: Creating Customer Profiles



Customer data from online store

- User ID
- Name
- Credit-card details
- Address
- Date of birth
- Points

Customer Profile: Online Store

User ID	Name	Credit Card	Address	DOB	Points
C1	Jill	4034	<u>Sydney</u>	<u>3-Jun-1976</u>	<u>387</u>
C2	John	3323	Perth	12-Feb-1987	231
C3	Mike	2300	Adelaide	30-Dec-1990	45
C4	Lilly	9831	Hobart	7-May-1981	24

A standard structure can be expected from this source

Problem: Creating Customer Profiles



Customer data from physical stores

- Credit-card details?
- Name?
- User ID?

Different details will be available for different customers

Customer Profile: Physical Stores

User ID	Name	Credit Card	Store	Points
	Brianna	1011	Hobart	345
C23	Luca		Brisbane	128
	Ayako	8346	Alice Springs	
C43			Brisbane	137

The data is unstructured

Problem: Creating Customer Profiles



Customer data from third parties

- Fields available may vary based on third party
- Format and types will be different

Amazon → Date → java.util.Date

Xyz → Date → "String"

3 different sources

Online

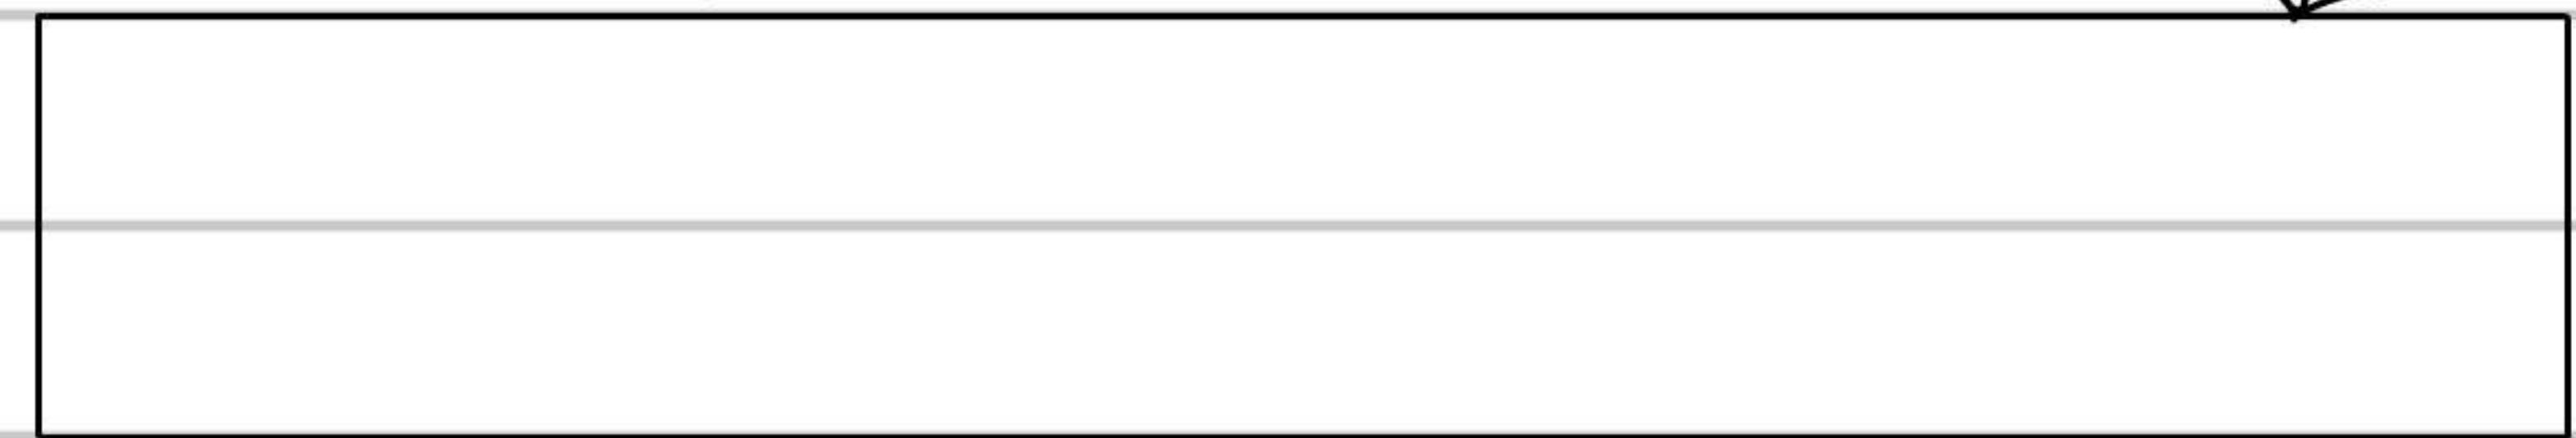
Physical

Third Party

structured

unstructured

structured but different formal -



RDBMS?

Customer Profile

	User ID	Name	Credit Card	Address	DOB	Points	Age Group
D S	C1	Jill	4034	Sydney	3-Jun-1976	387	
OS	C2	John	3323	Perth	12-Feb-1987	231	
Ps		Brianna	1011	Hobart		345	
T P		John		Perth		198	31-35

A lot of fields are
empty

Customer Profile

User ID	Name	Credit Card	Address	DOB	Points	Age Group
C1	Jill	4034	Sydney	3-Jun-1976	387	
C2	John	3323	Perth	12-Feb-1987	231	
	Brianna	1011	Hobart		345	
	John		Perth		198	31-35

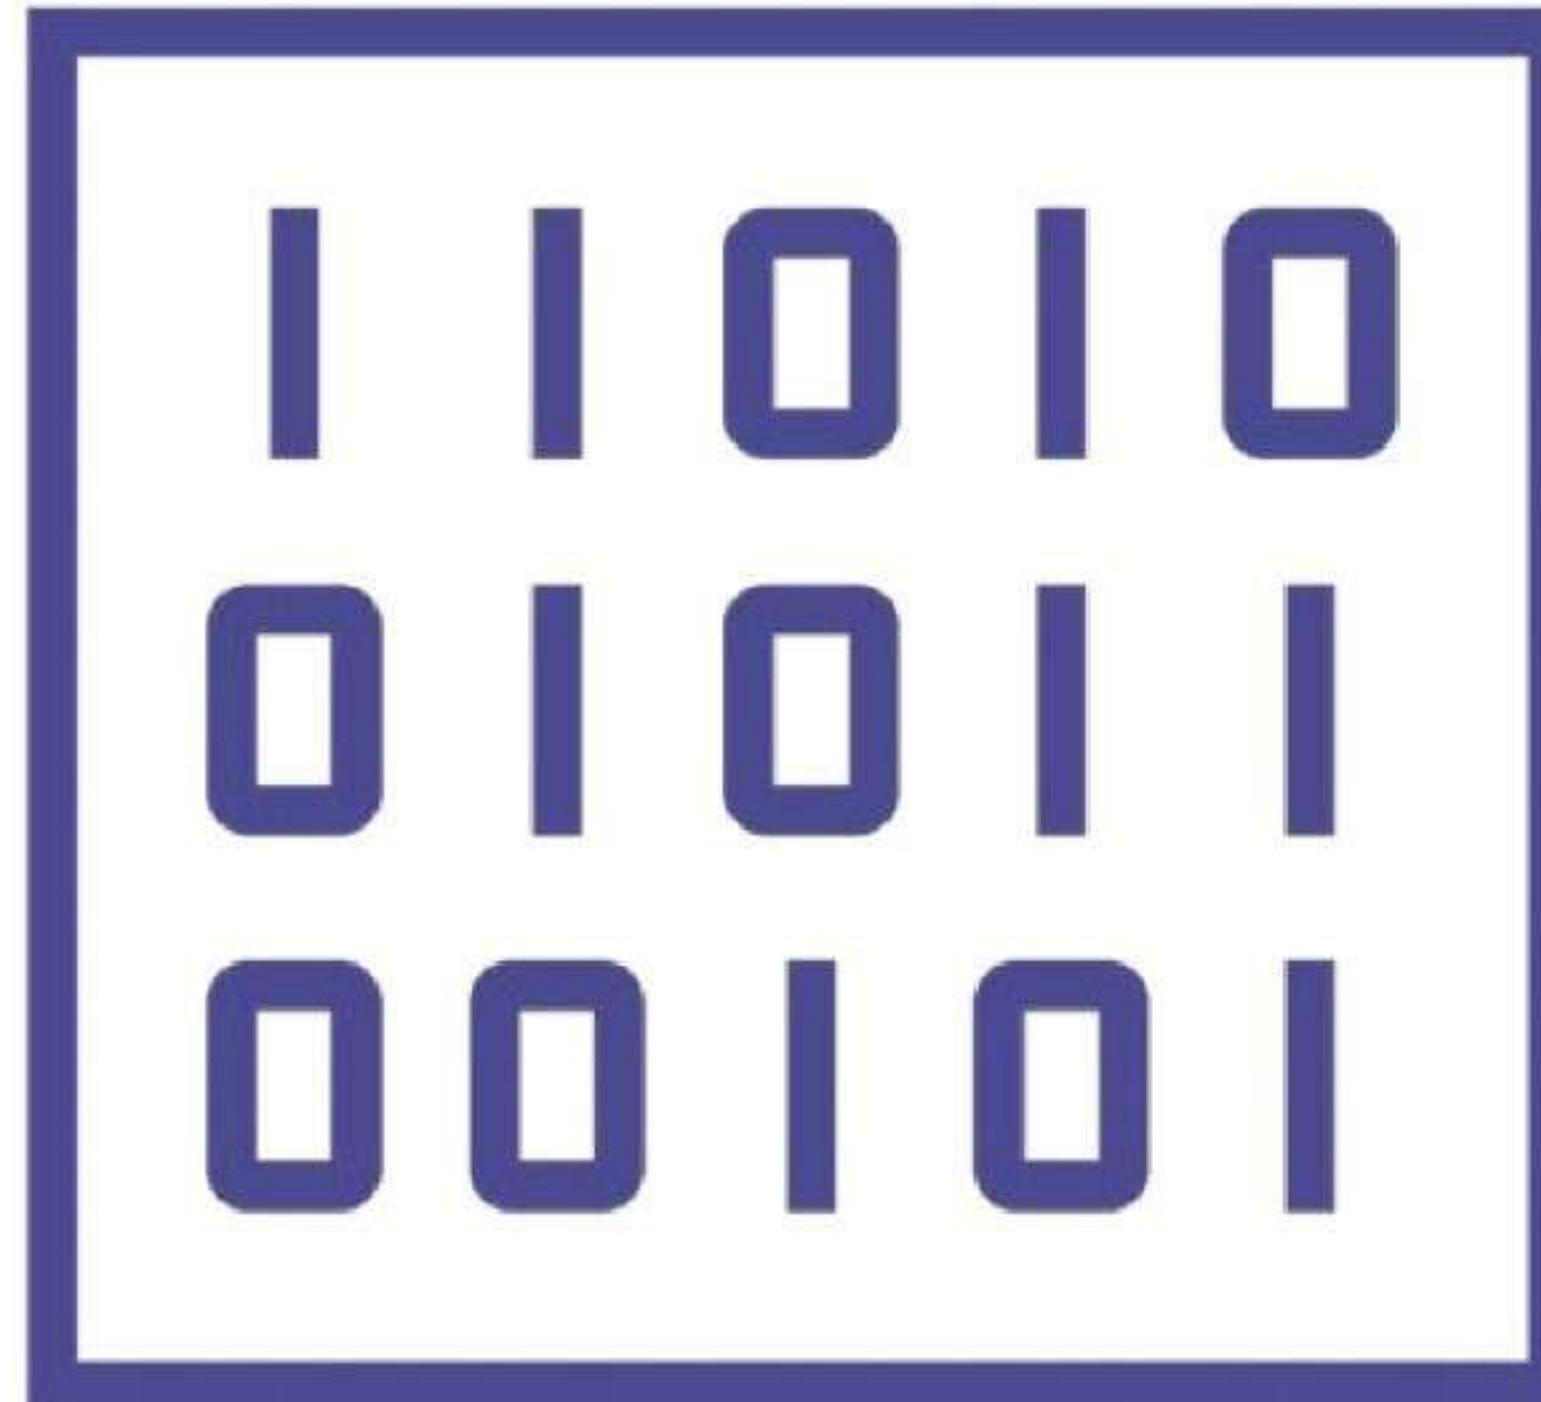
Do these refer to the same person?

Customer Profile

User ID	Name	Credit Card	Address	DOB	Points	Age Group
C1	Jill	4034	Sydney	3-Jun-1976	387	
C2	John	3323	Perth	12-Feb-1987	231	
	Brianna	1011	Hobart		345	
	John		Perth		198	31-35

This data can still be
analyzed

Unstructured Data



May originate from multiple sources

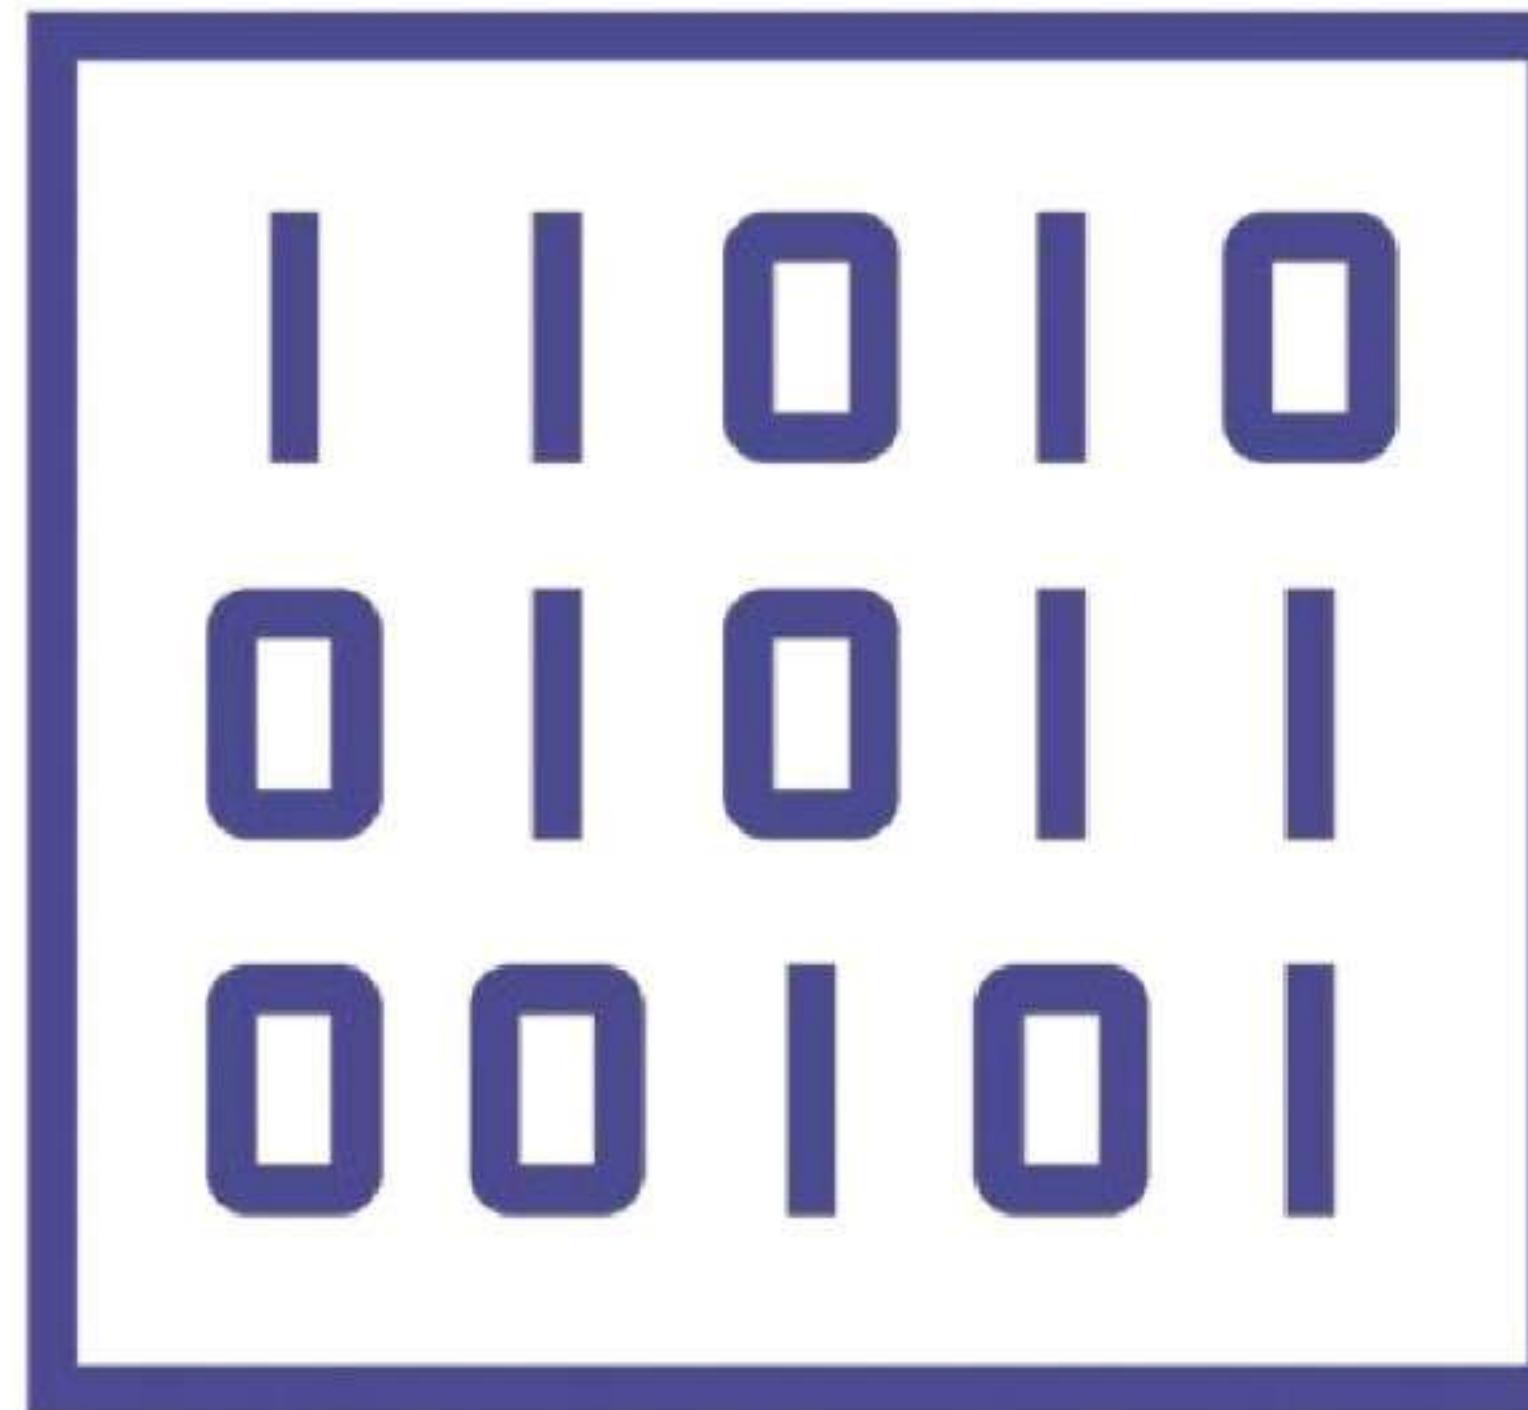
Different attributes available for same types of entities

Data types for attributes may vary

Data may be partially structured

- Some fields are always available

Unstructured Data



Relational databases are not well-suited to unstructured data

- May strictly enforce schemas

Some NoSQL databases are designed to handle this lack of structure

- E.g. document databases

An Overview of Big Data

Big Data

A field which deals with the analysis of large and complex datasets in order to extract meaningful information.

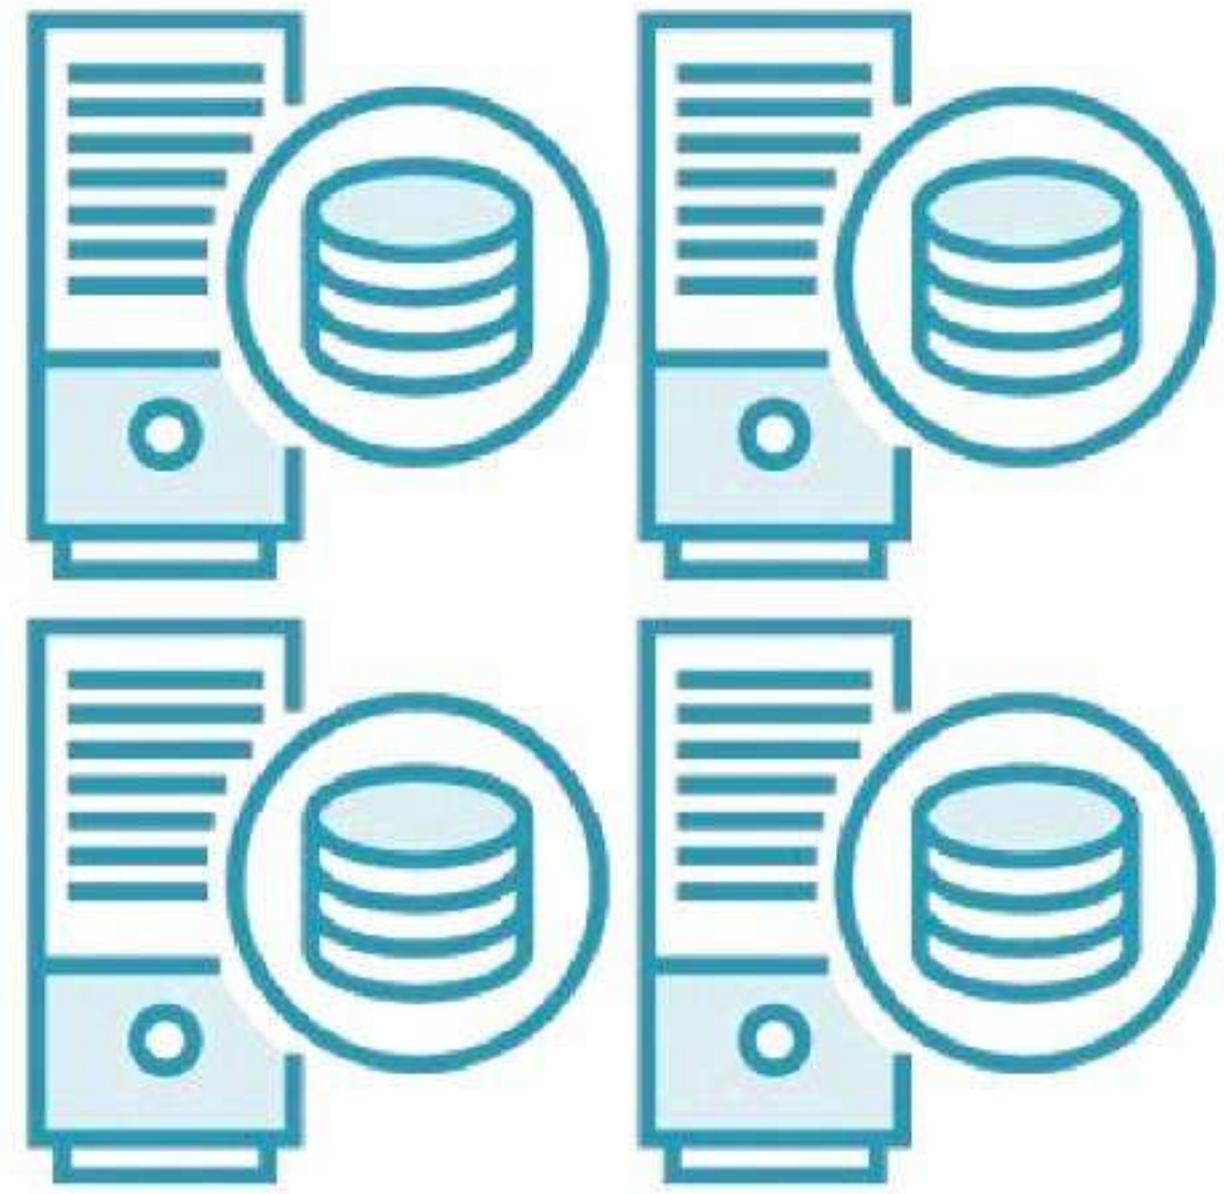
Big Data

A field which deals with the **analysis** of large and complex datasets in order to extract meaningful information.

Big Data

A field which deals with the analysis of **large and complex datasets** in order to extract meaningful information.

3 Vs of Big Data

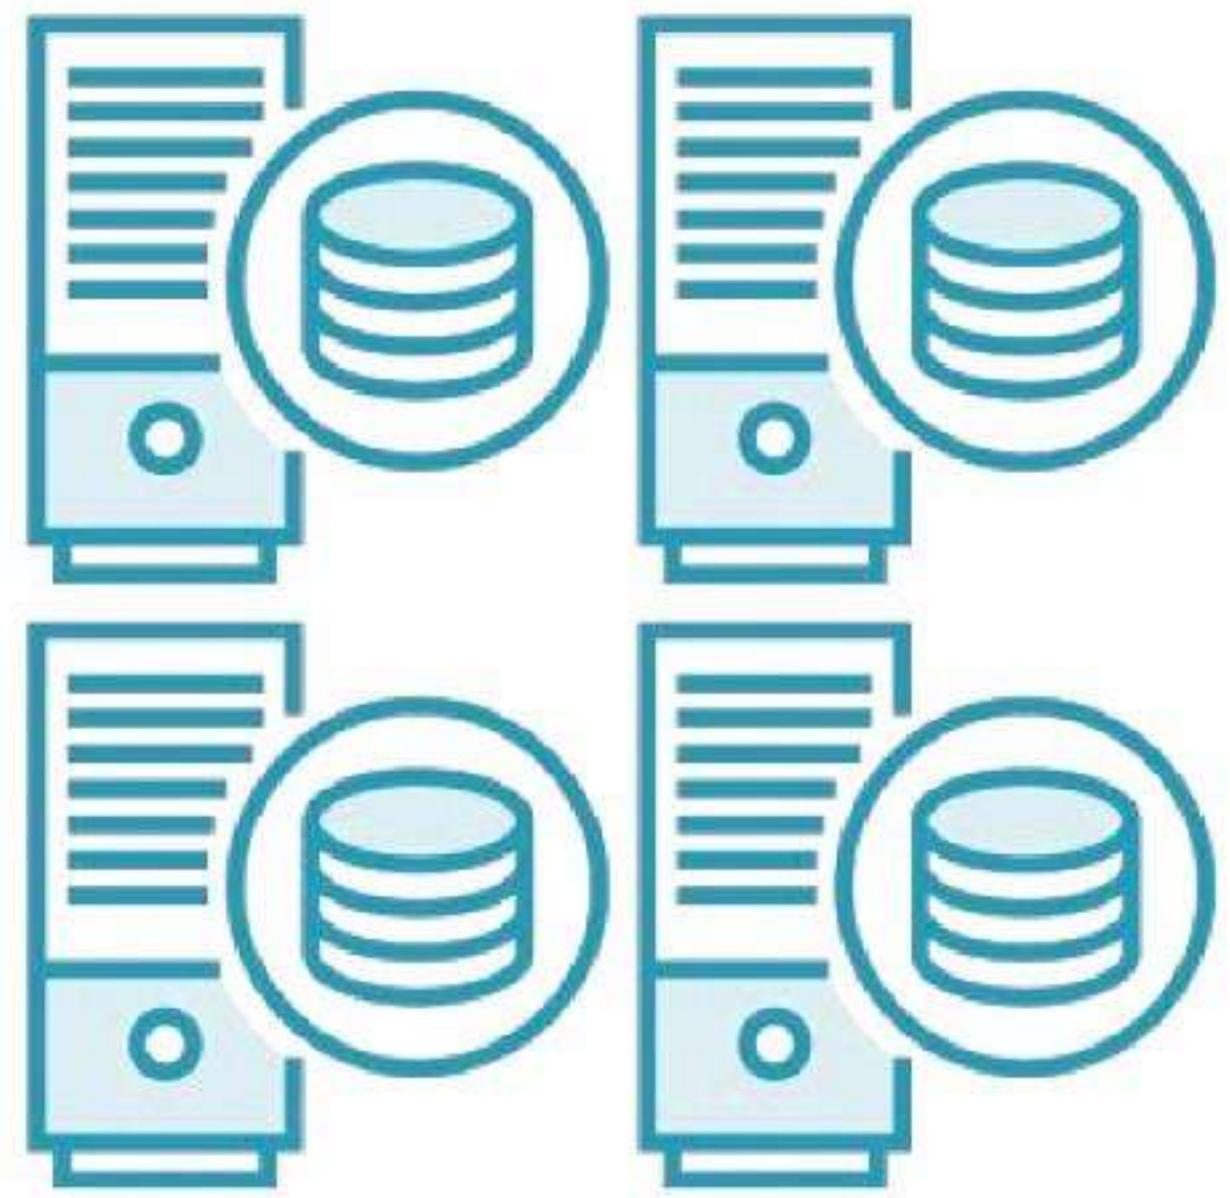


Volume: Amount of data

Variety: Number and type of sources

Velocity: Batch and streaming

Properties of Big Data Systems



Data distributed on a cluster with
multiple machines

Semi-structured or **unstructured** data

No random access to data

Data **replicated**, propagation of
updates take time

Different sources may have **different**
unknown formats

Database Use Cases



Data Storage

Data Security

Transactions

Data Analysis

Transactional and Analytical Processing

Transactional Processing

Ensure correctness of individual entries

Access to recent data, from the last
few hours or days

Updates data

Fast real-time access

Usually a single data source

Analytical Processing

Analyzes large batches of data

Access to older data going back months, or
even years

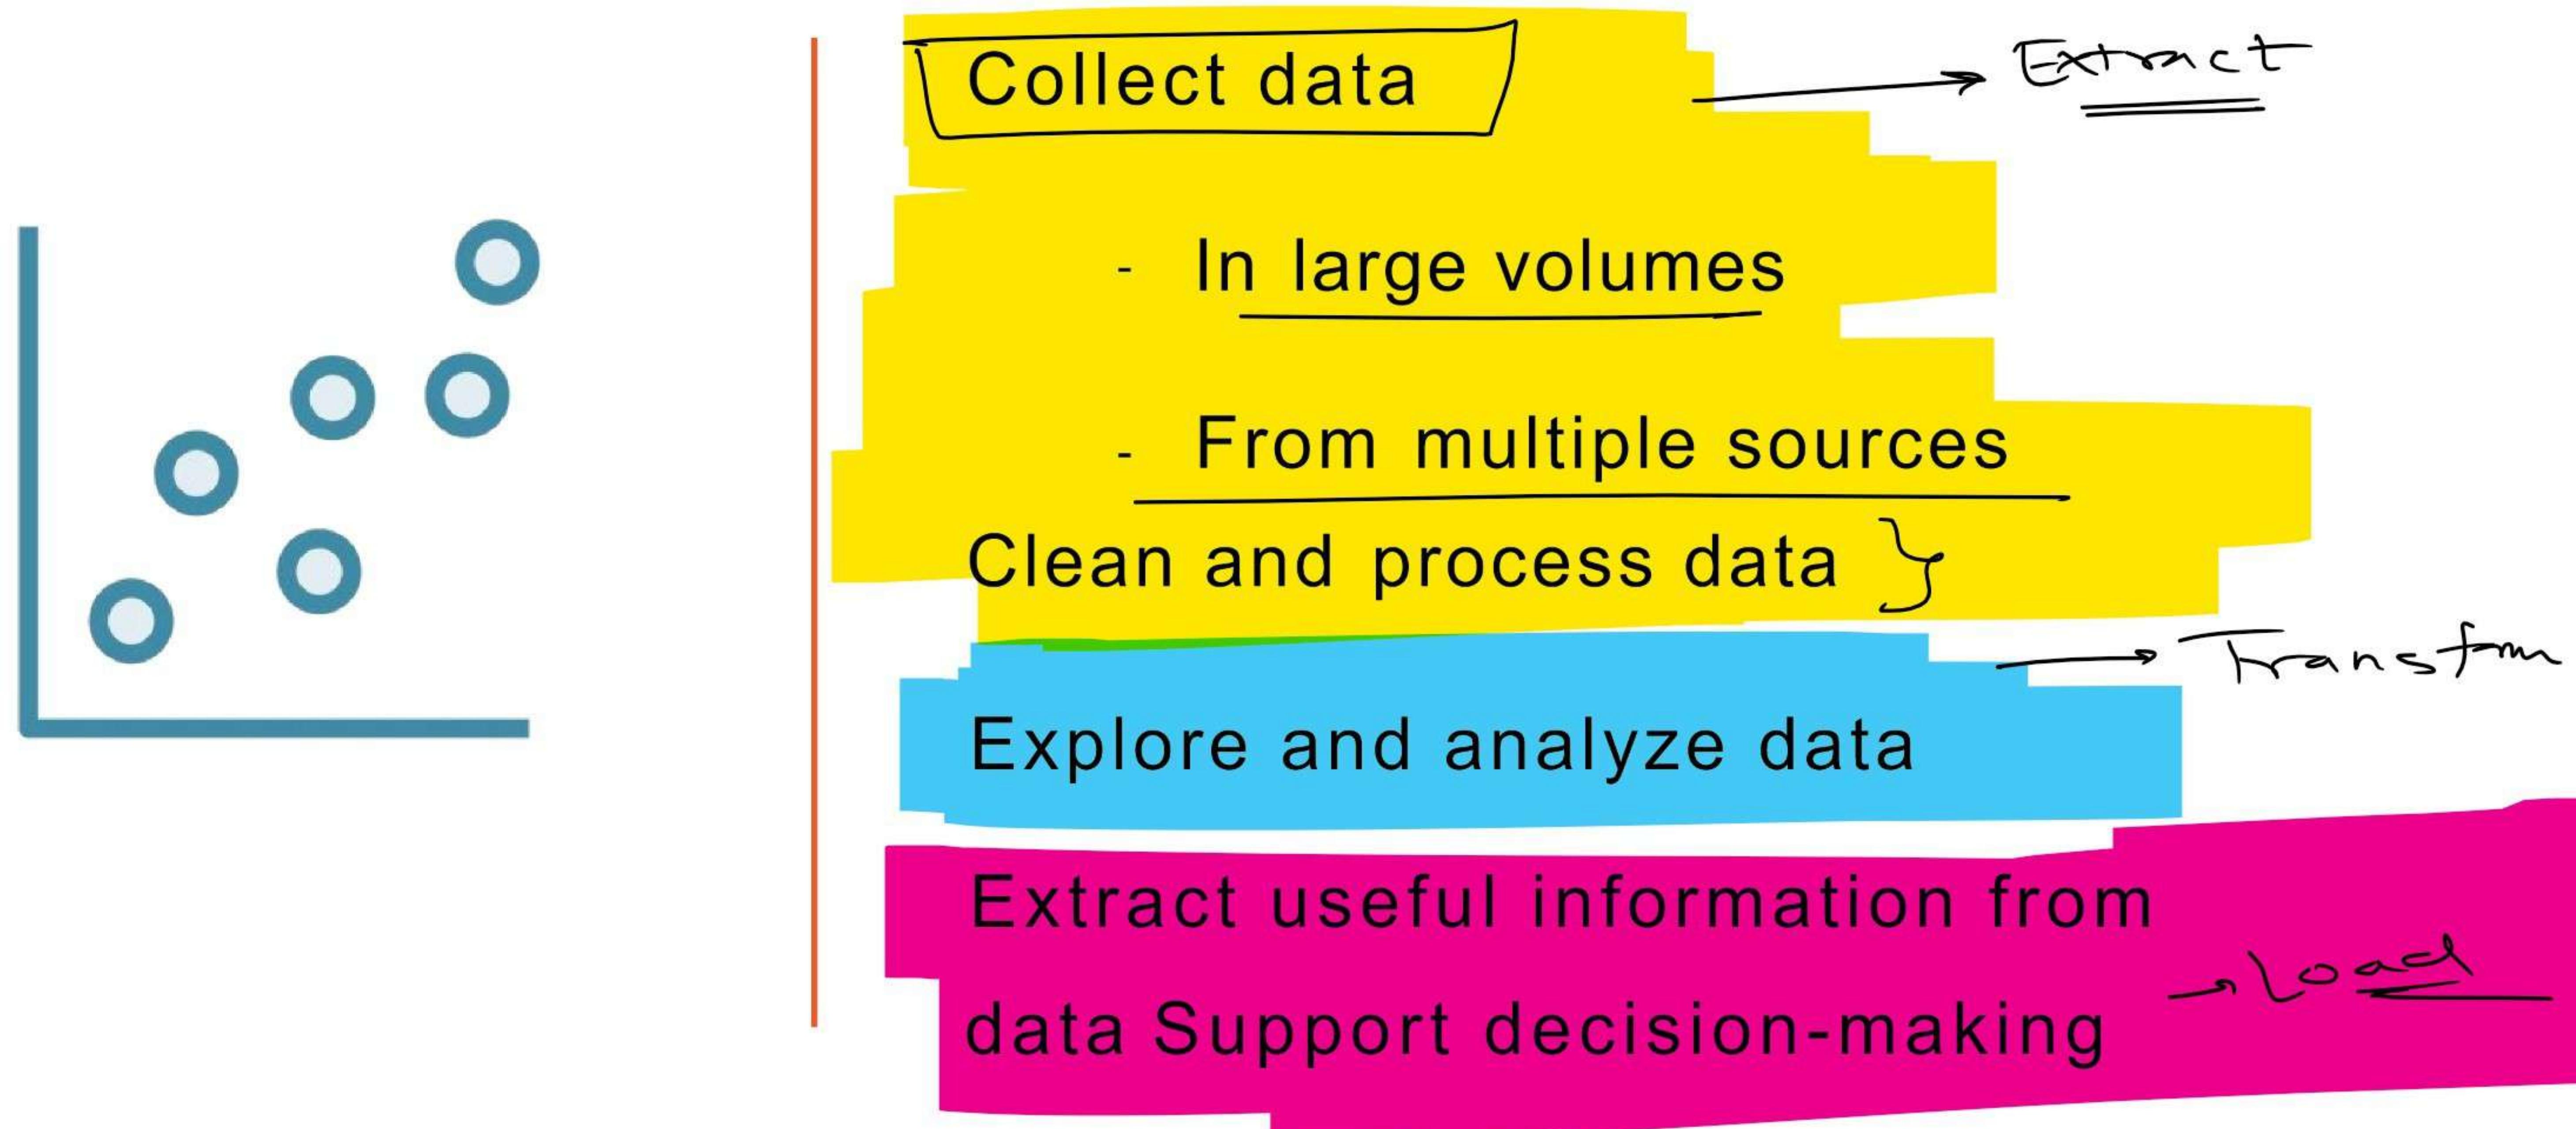
Mostly reads data

Long running jobs

Multiple data sources

Extract ~~Transform~~
ETL → Load

Data Analysis and Big Data



NoSQL databases are
more
suitable for Big Data
processing than RDBMS

NoSQL Databases and Big Data

5 V's → Velocity, Variety, Volume, Value, ~~Variety~~ → Cassandra

Features of NoSQL Databases

Variety
Semi-structured Data

Volume
Large Datasets

High Availability

Analytical Queries

Variety
Real-time and Streaming

Caching and Prototyping

Features of NoSQL Databases

Semi-structured
Data

Large Datasets

High Availability

Analytical Queries

Real-time and
Streaming

Caching and
Prototyping

These use cases map to 3 properties of big data

Features of NoSQL Databases

Variety

Volume

High Availability

Analytical Queries

Velocity

Caching and
Prototyping

The 3 Vs of big data

Features of NoSQL Databases

Variety

Volume

High Availability

Analytical Queries

Velocity

Caching and
Prototyping

This can be ensured with a distributed system

Features of NoSQL Databases

Variety

Volume

High Availability

Analytical Queries

Velocity

Caching and
Prototyping

Queries meant to understand data in the aggregate

Features of NoSQL Databases

Variety

Volume

High Availability

Analytical Queries

Velocity

Caching and
Prototyping

Contrasts with traditional use case for RDBMS

Relational vs. NoSQL Databases

Relational Databases

Vertical scaling - often not distributed systems

Normalized data - small, inter-related tables

Optimized for efficient storage - avoid redundancy, hierarchy, nesting

ACID

Atomicity, Consistency, Isolation, Durability

NoSQL Databases

Horizontal scaling - are almost always distributed systems

Denormalized data - large, self-contained collections

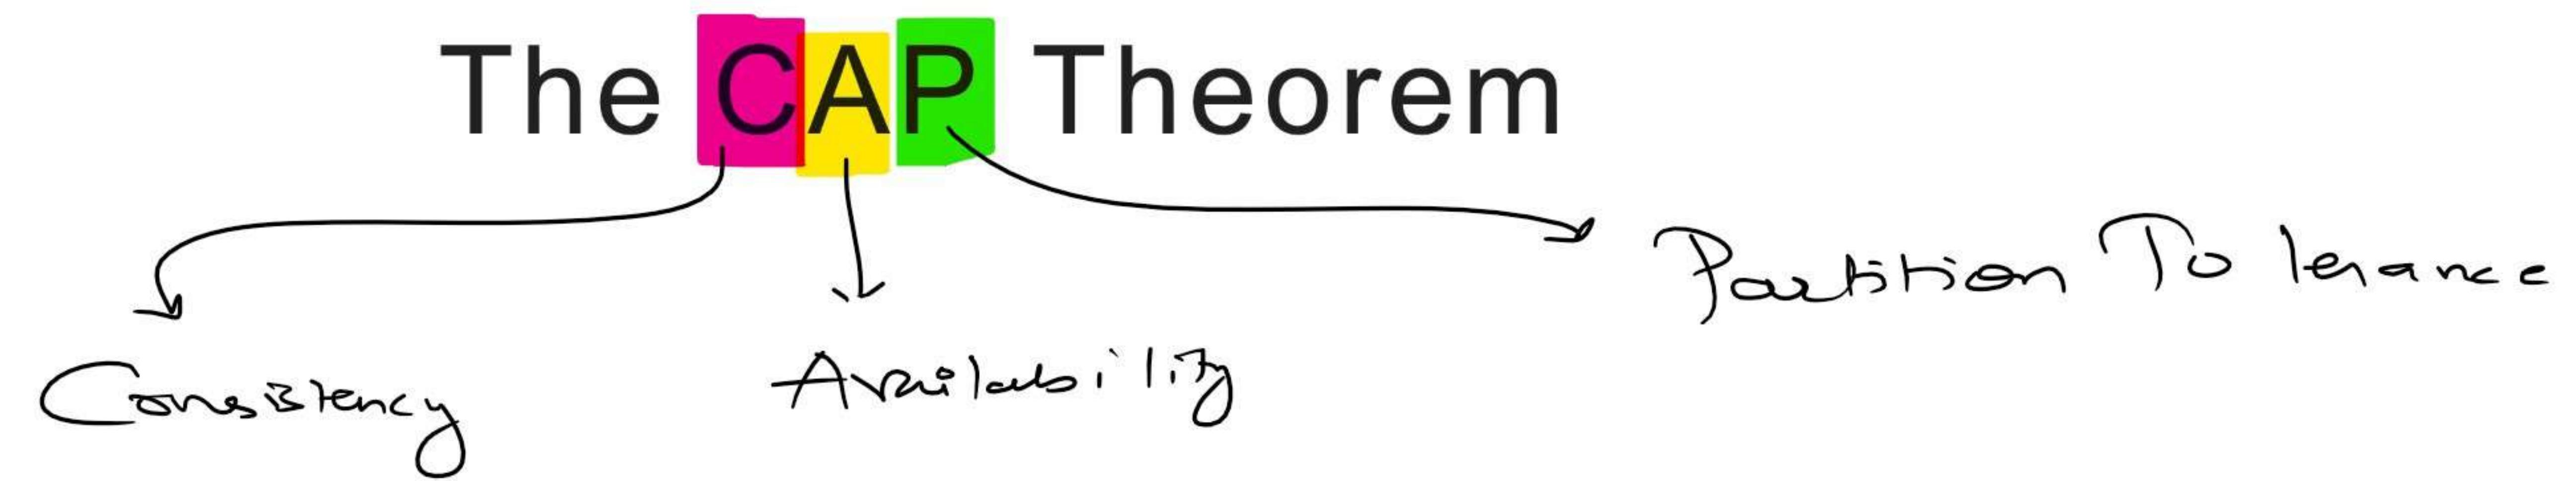
Optimized for efficient network access - use hierarchy, redundancy, nesting

BASE

Basically
A vailable
Soft State
Eventual Consistency.

Good, Cheap & Fast ?

→ Pick any 2 ?



CAP Theorem



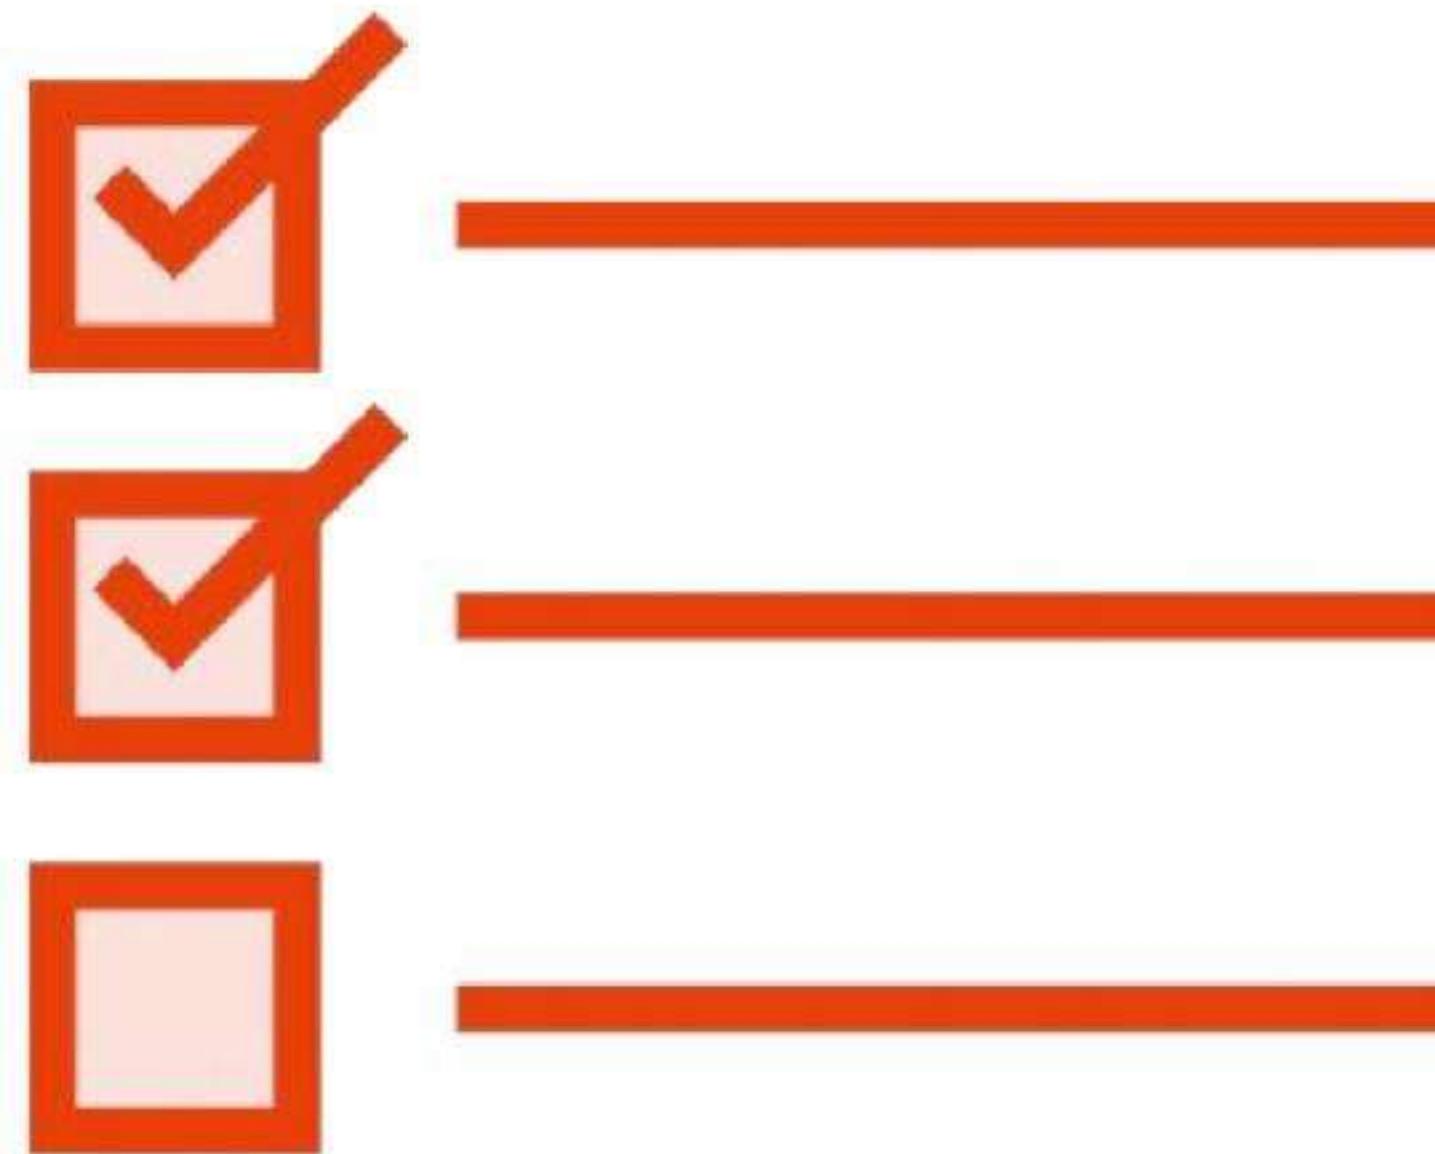
A distributed database must
choose
2 out of the 3 CAP guarantees:

- Consistency
- Availability
- Partition Tolerance

The CAP Guarantees

C
A
P

by id allotment



Consistency: Every read will receive most recent write, else error

Availability: Every request will receive a non-error response

Partition Tolerance: Network failures between nodes will be "managed gracefully" "gracefully"

C → Hard Working
A → Loyal
P → Alpine

CP (Better delay than wrong data) Consistency → data accuracy may deny the request if data is not consistent
AP financial inventory

CA → RDBMS

Availability (Always ready, even if config) AP
Prioritize uptime, even if mean to show static data
social media logs, metrics

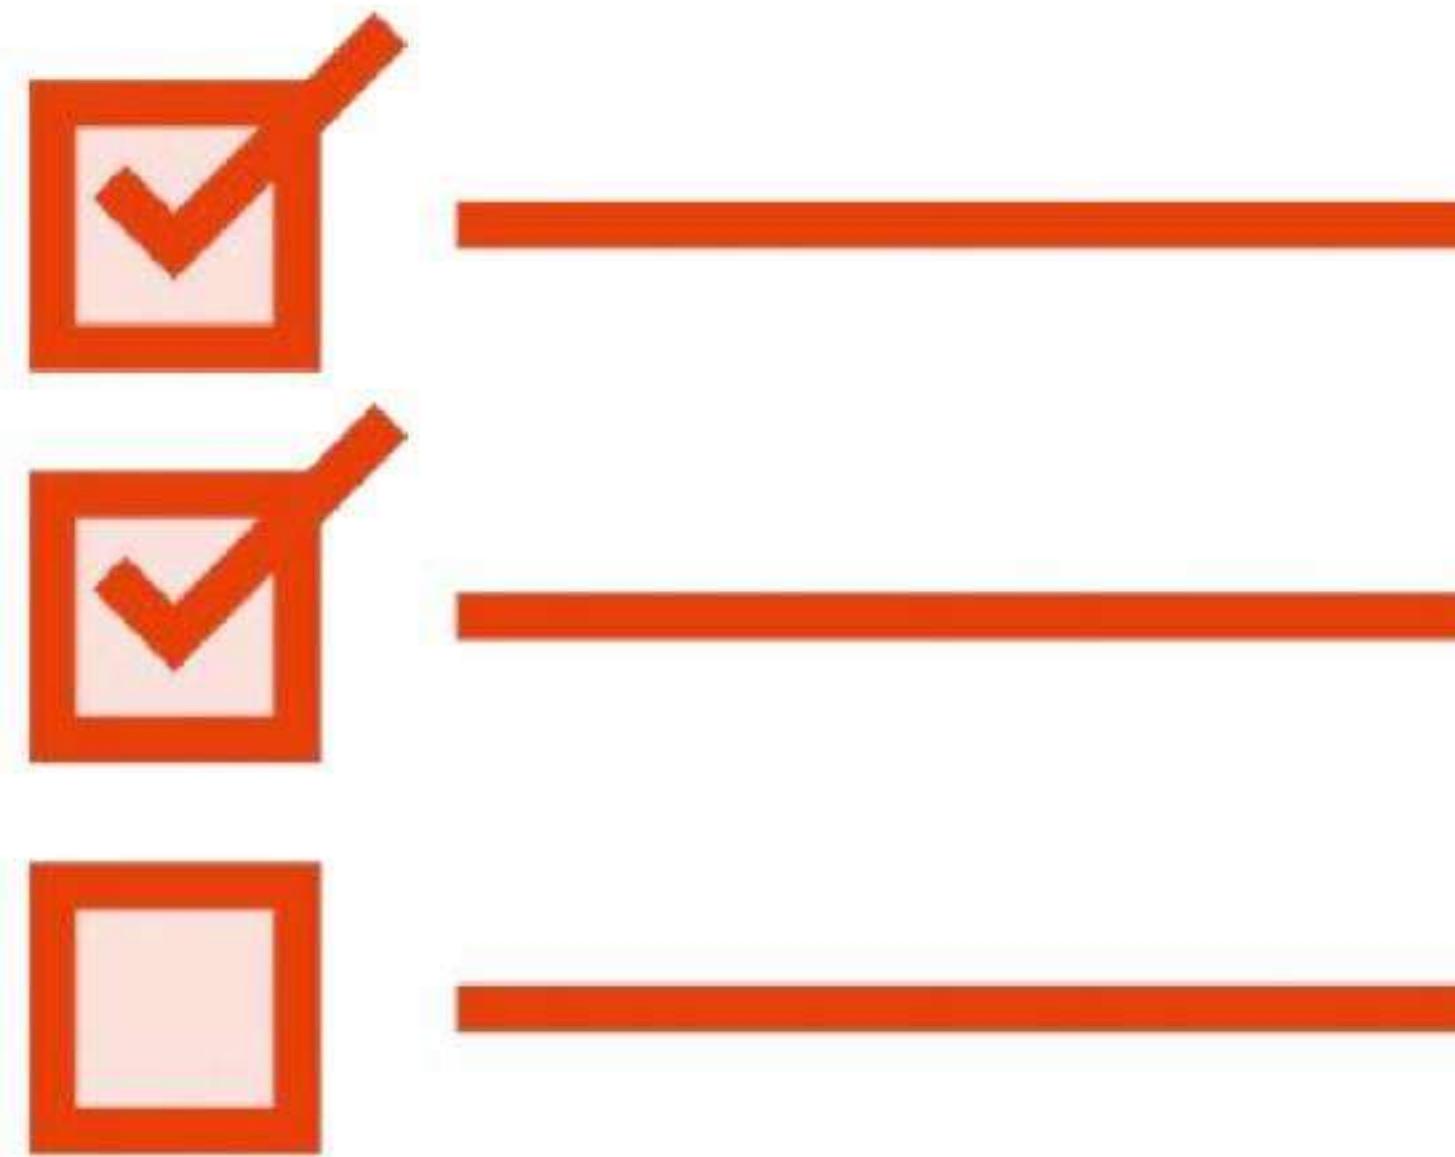
Partition Tolerance

Flight Ticket — 

P ↗

C	P	+Base Mongo DB <u>Cassandra</u> (strong consistency)
A	P	Cassandra (default eventual consistency) DynamoDB CouchDB

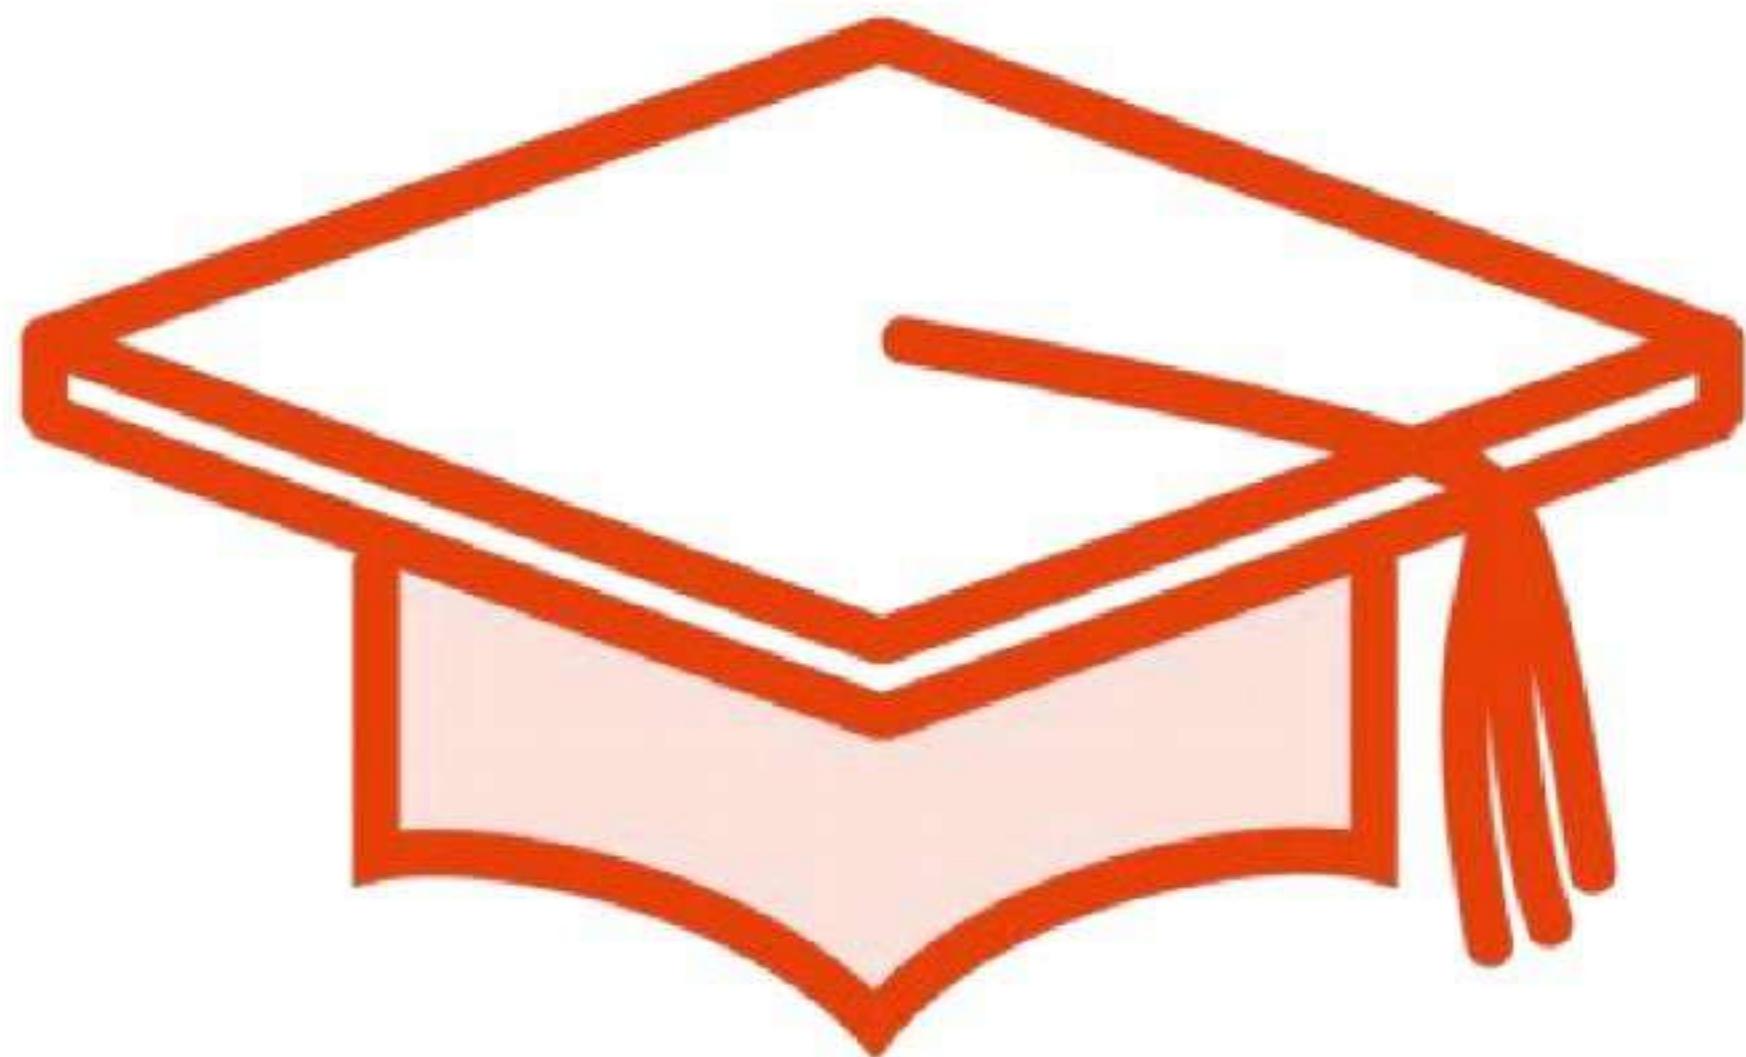
Partition Tolerance



When network failures partition the distributed system, it must either:

- Cancel the operation
 - Compromises availability but maintains consistency
- Allow the operation
 - Compromises consistency but maintain availability

The CAP Theorem



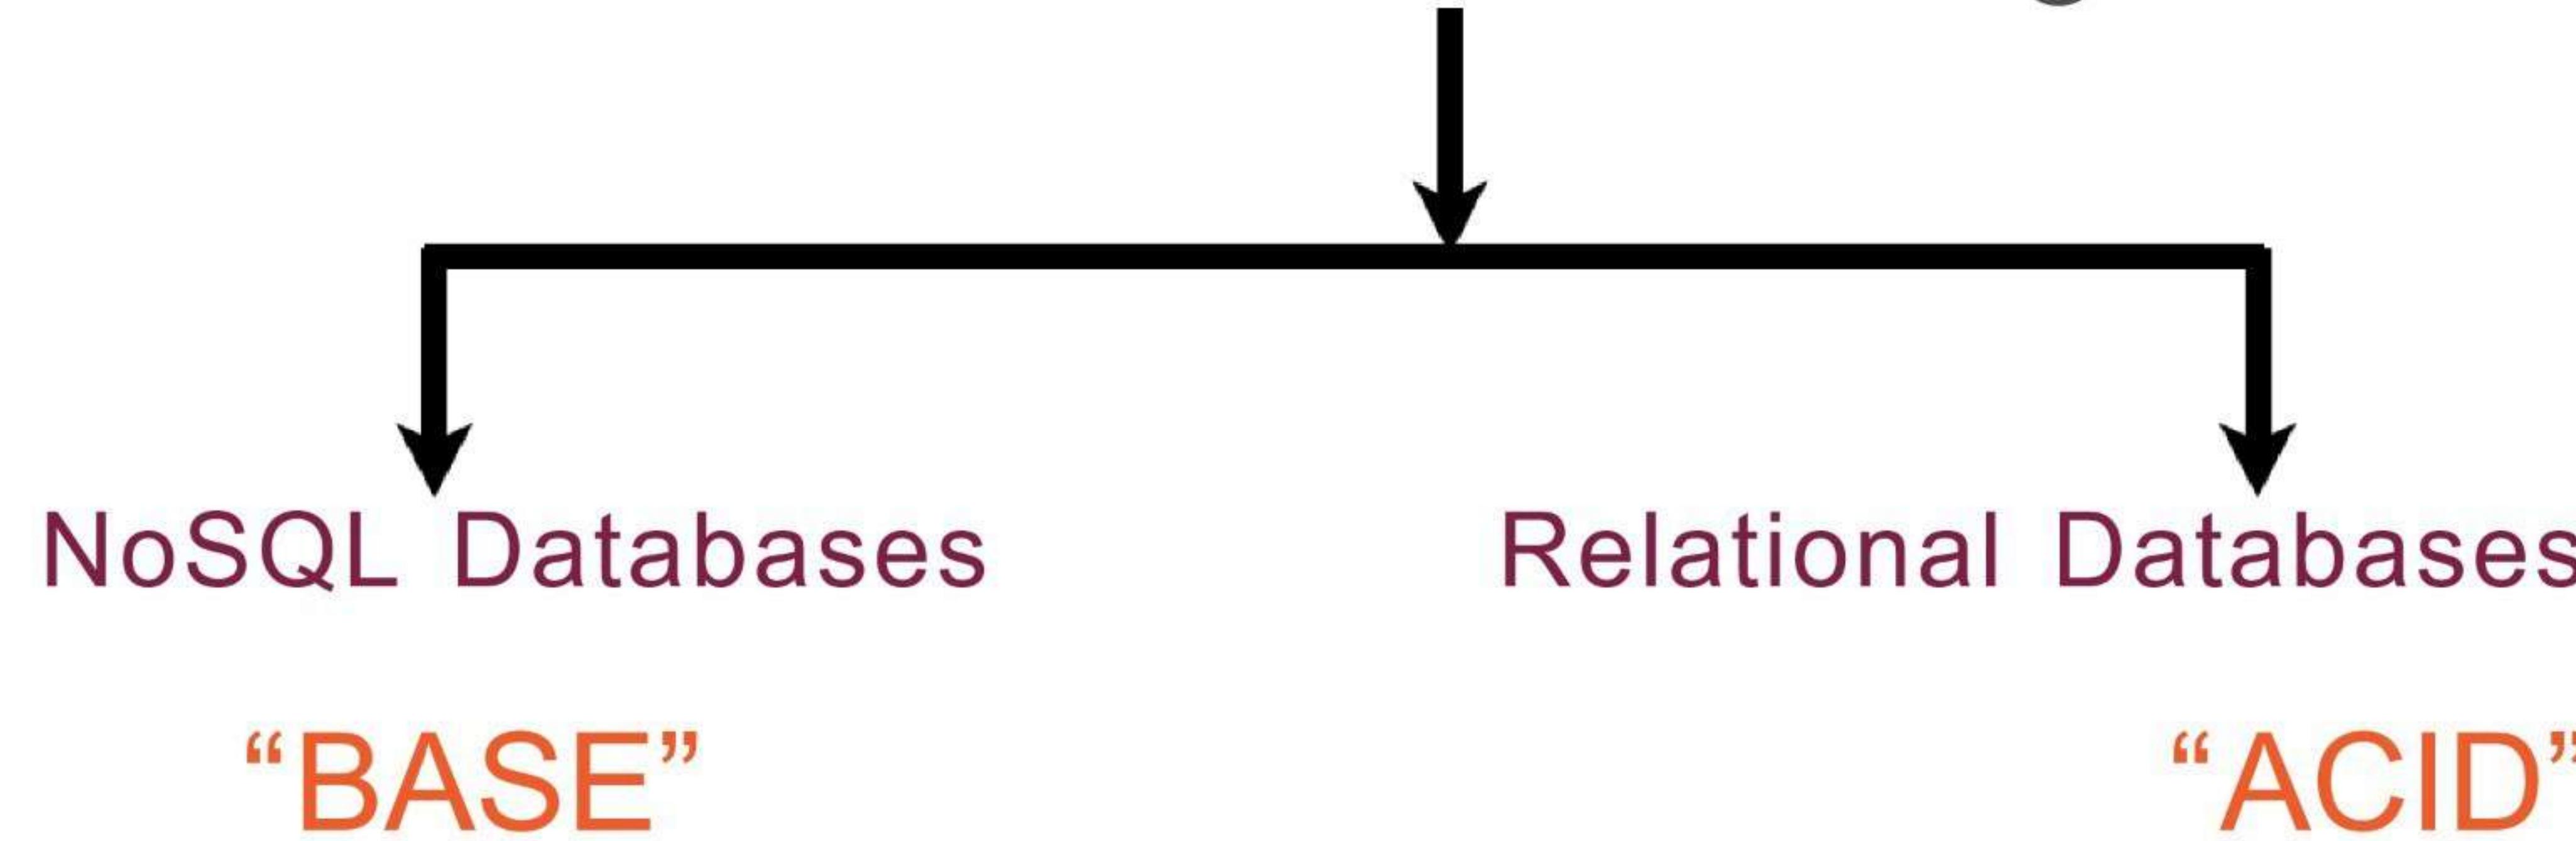
Not possible for a distributed database to achieve all 3 CAP guarantees

Big data technologies are invariably distributed database systems

- Horizontal scaling

Hence must choose what CAP guarantees to provide

Database Technologies



BASE vs. ACID

NoSQL Databases

Choose availability over consistency

BASE: Basically Available, Soft state, Eventual consistency

Writes are faster - don't wait for system to become fully consistent before returning

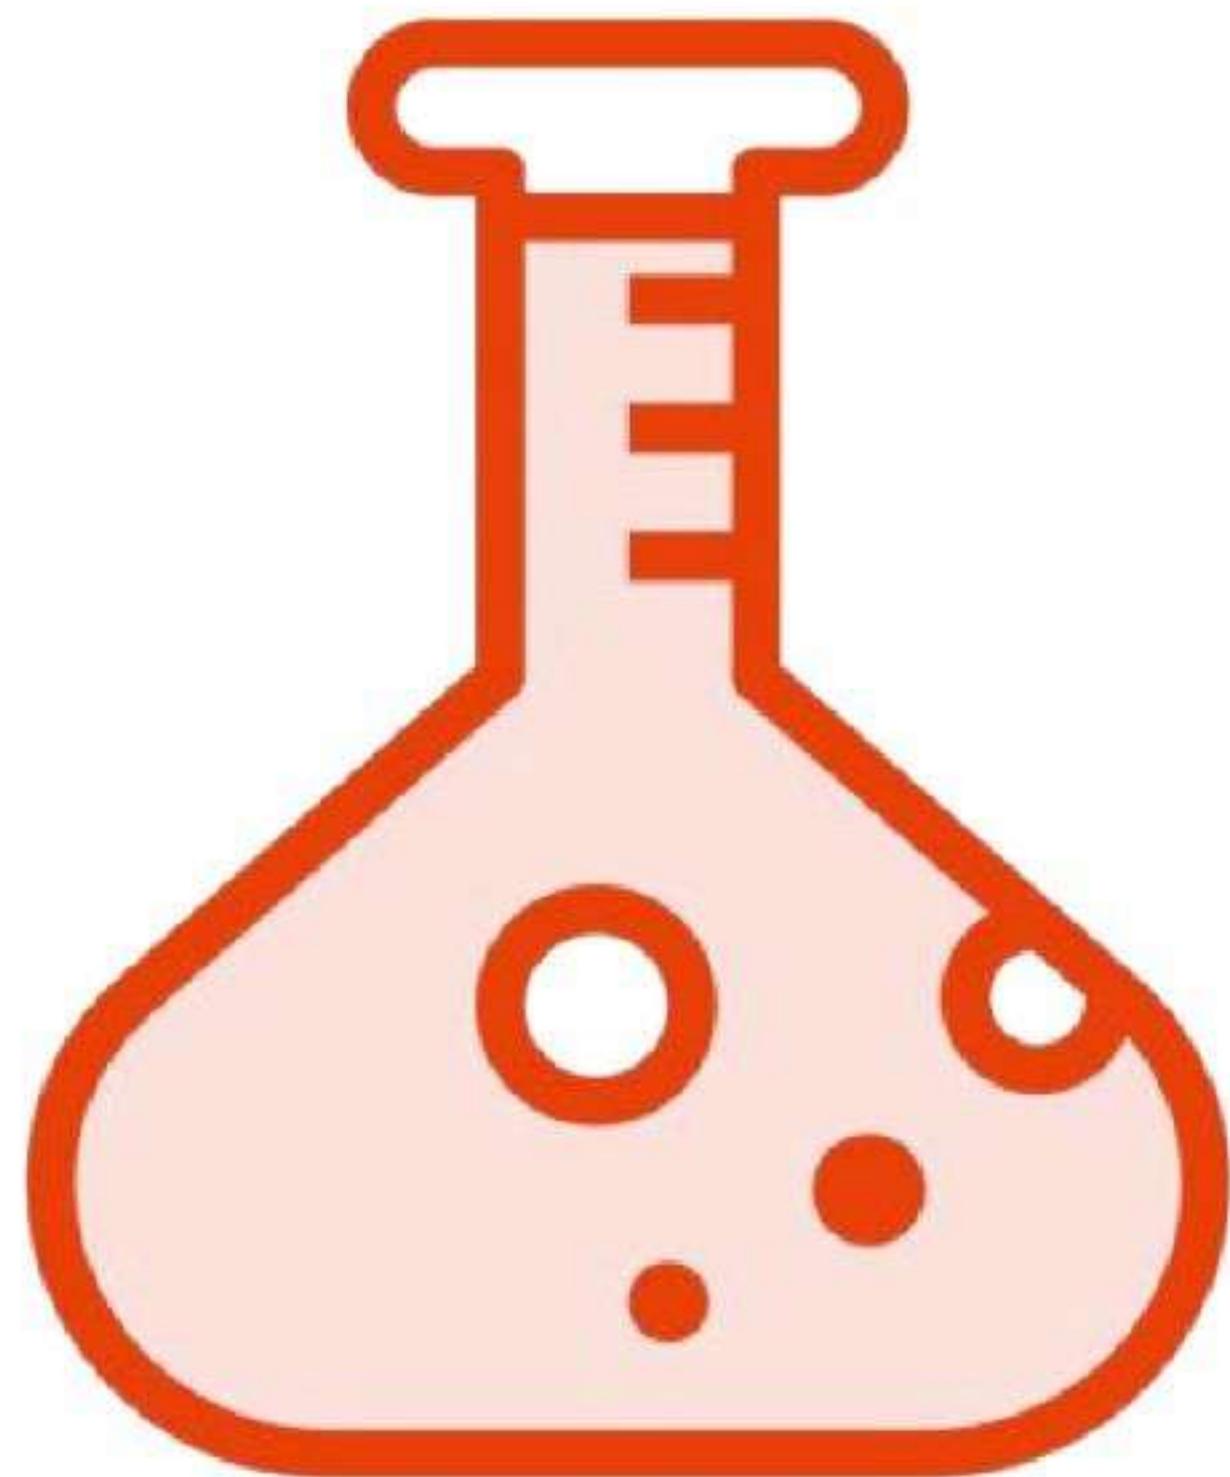
Relational Databases

Choose consistency over availability

ACID: Atomicity, Consistency, Isolation, Durability

Writes are slower - wait for full consistency of system before returning

The BASE Philosophy



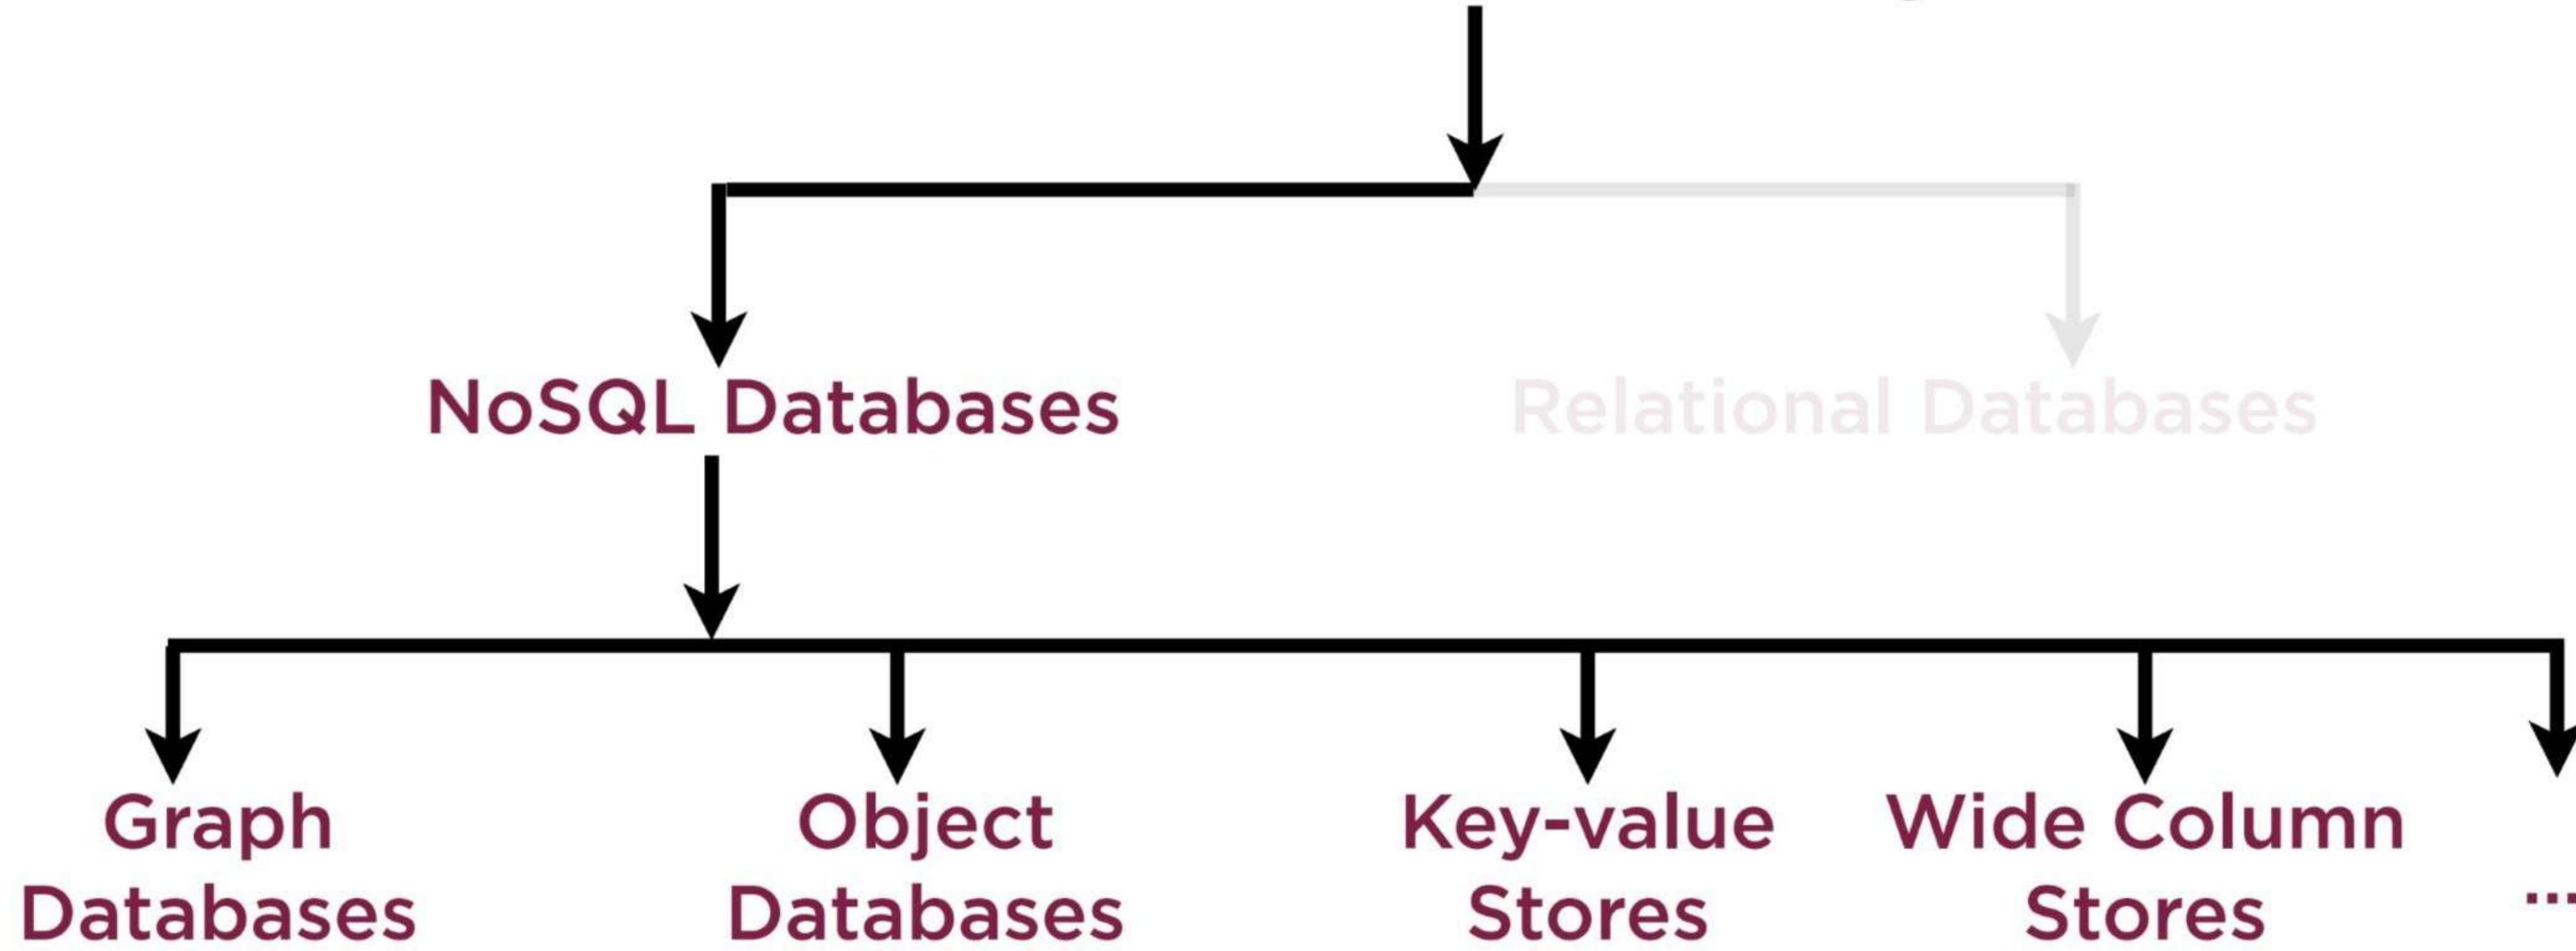
Basically Available: System will stay “up” using replication/sharding

Soft state: No guarantee that data is consistent - reads can yield stale data

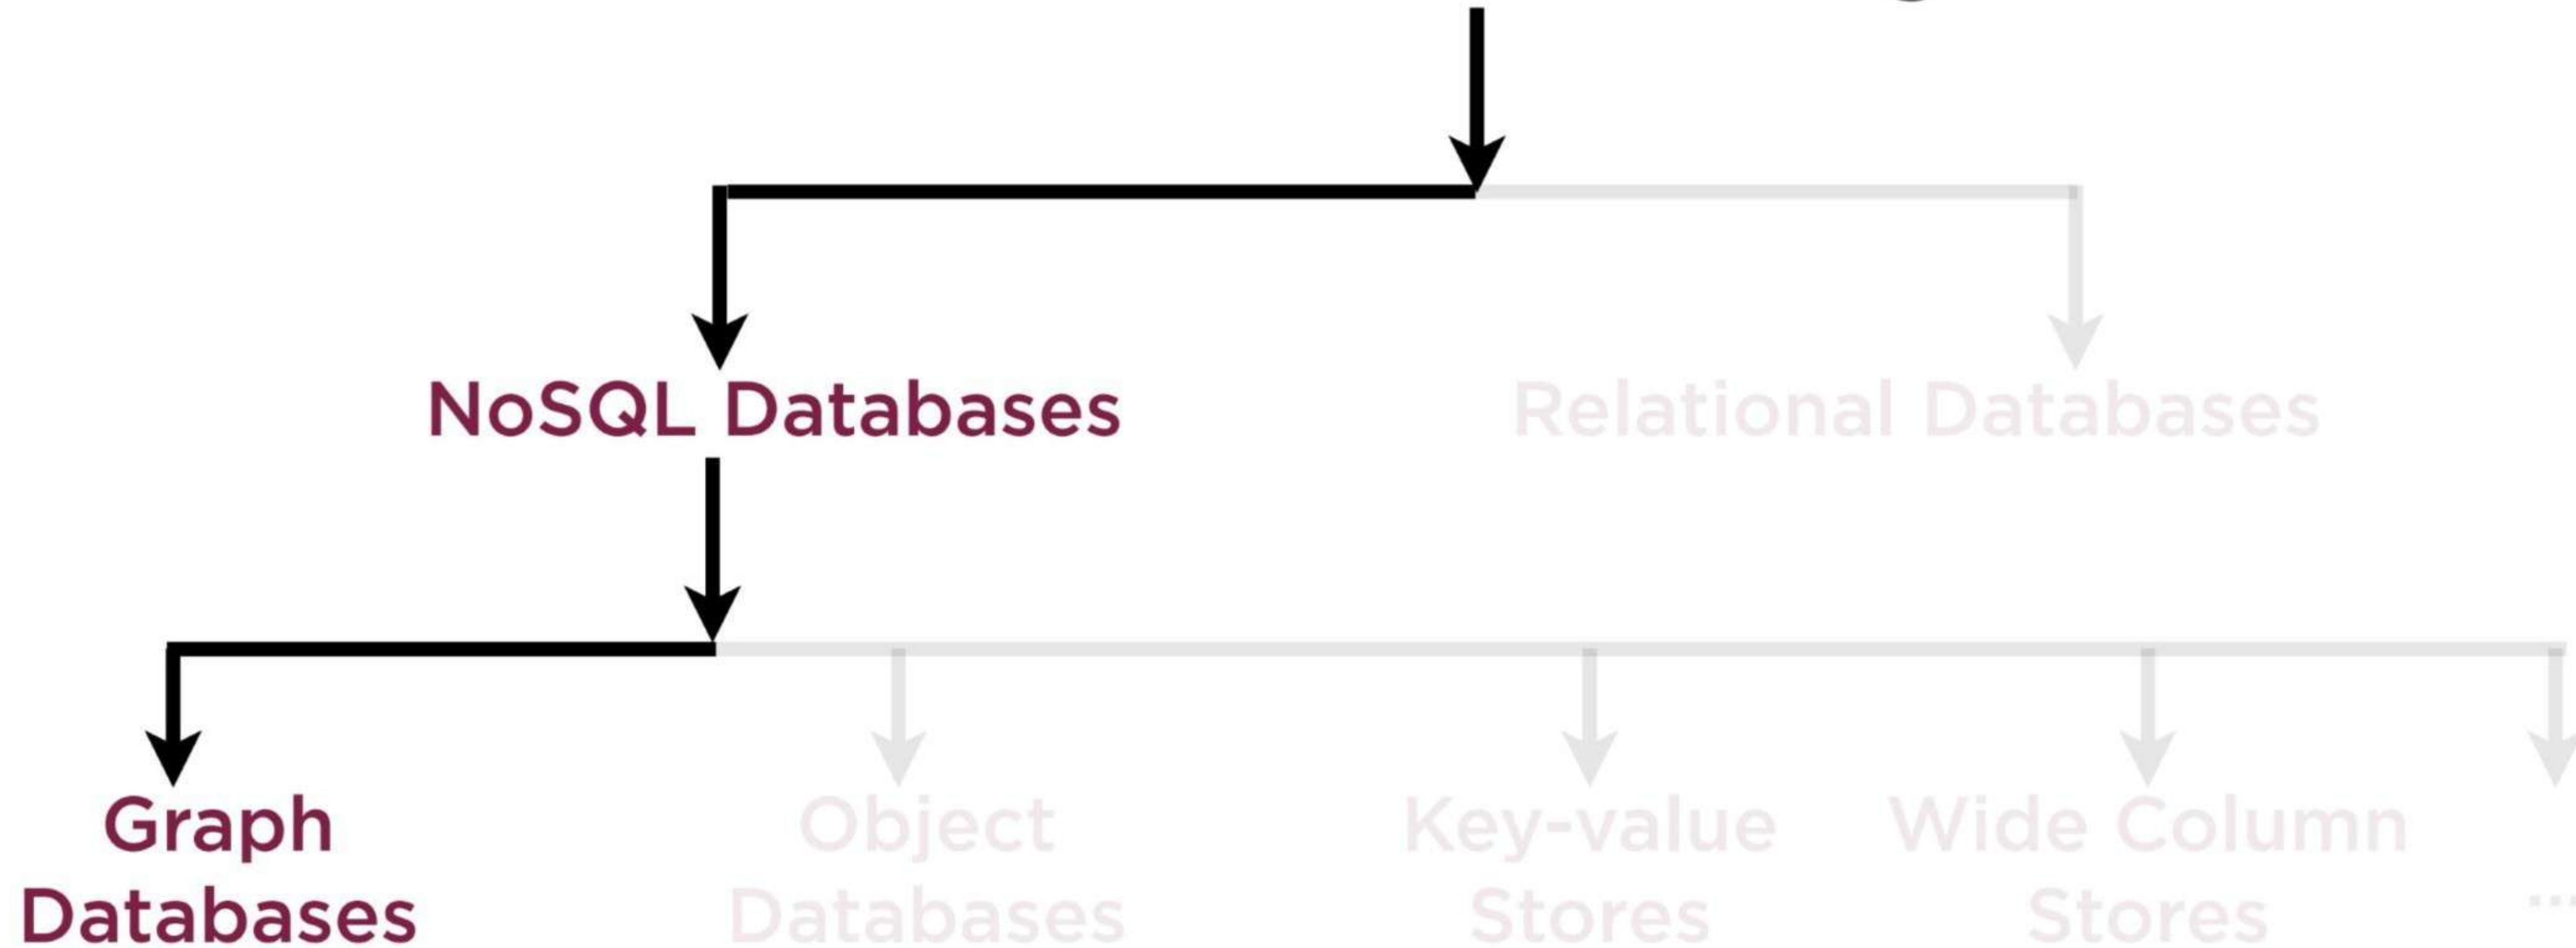
Eventual consistency: If we wait long enough, reads will yield consistent data

- No guarantee on how long that is

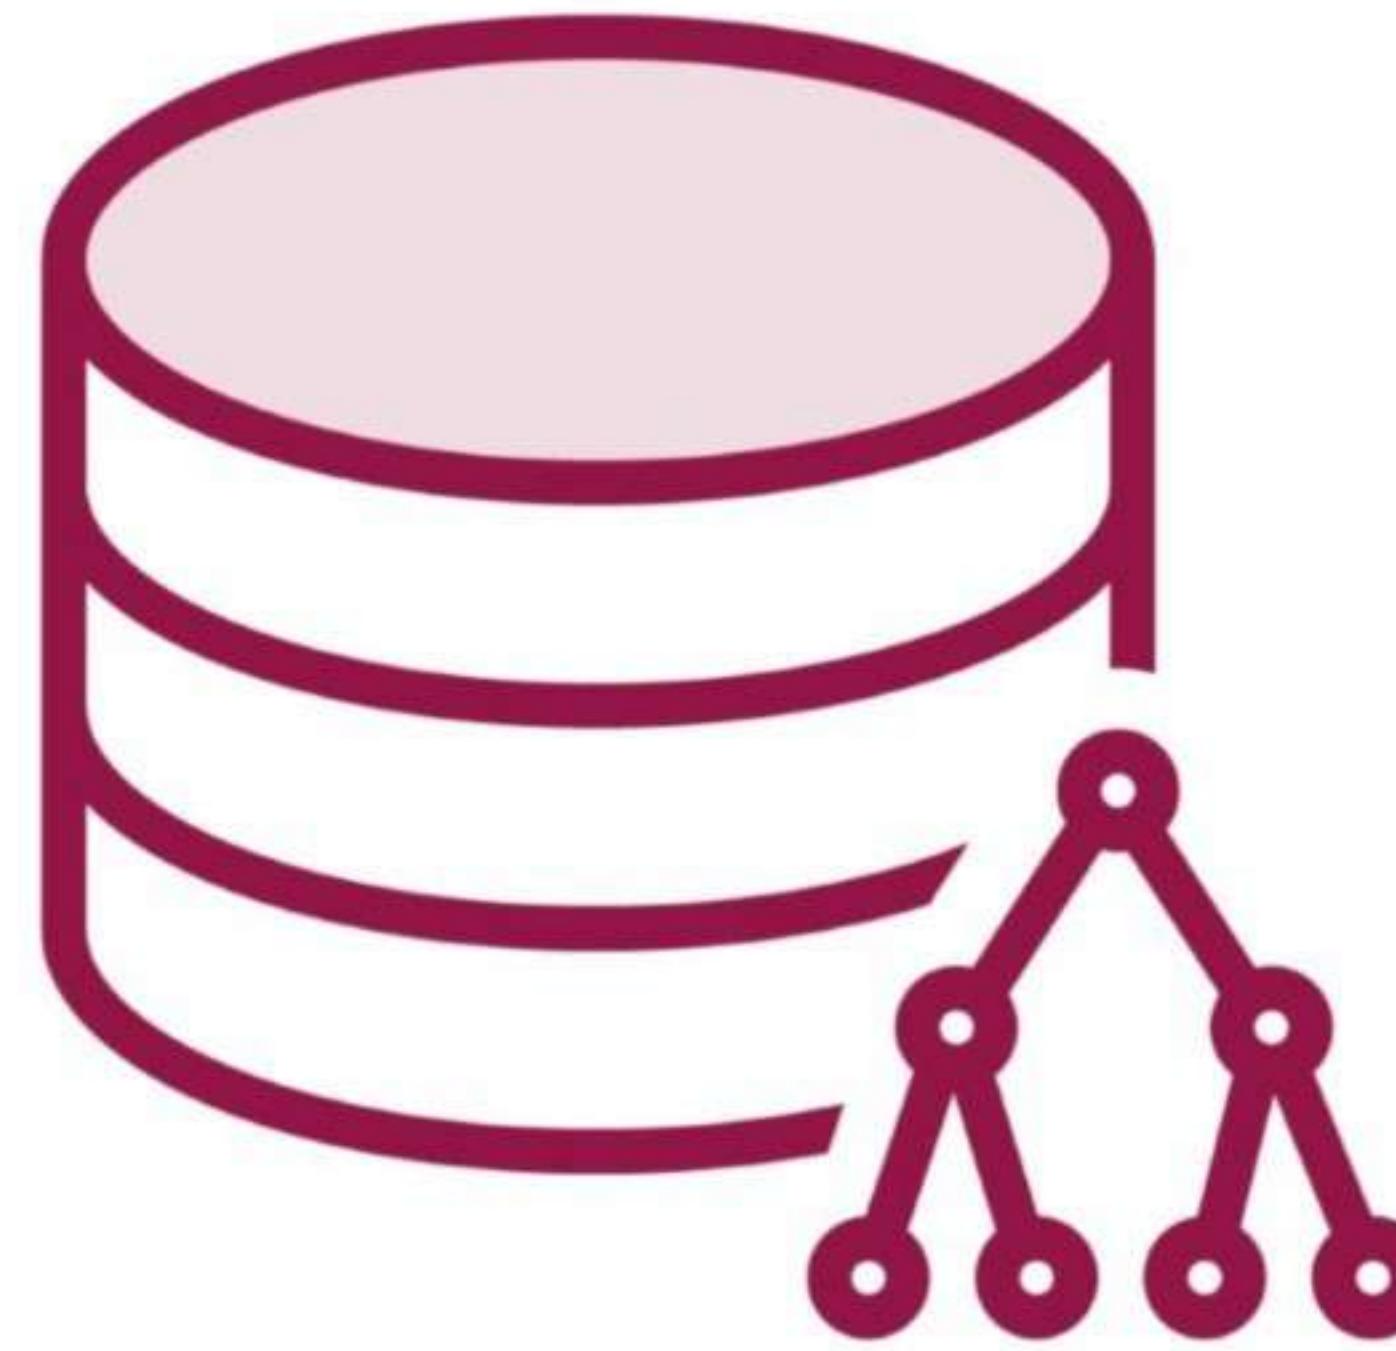
Database Technologies



Database Technologies



Graph Databases



Data organized into graphs

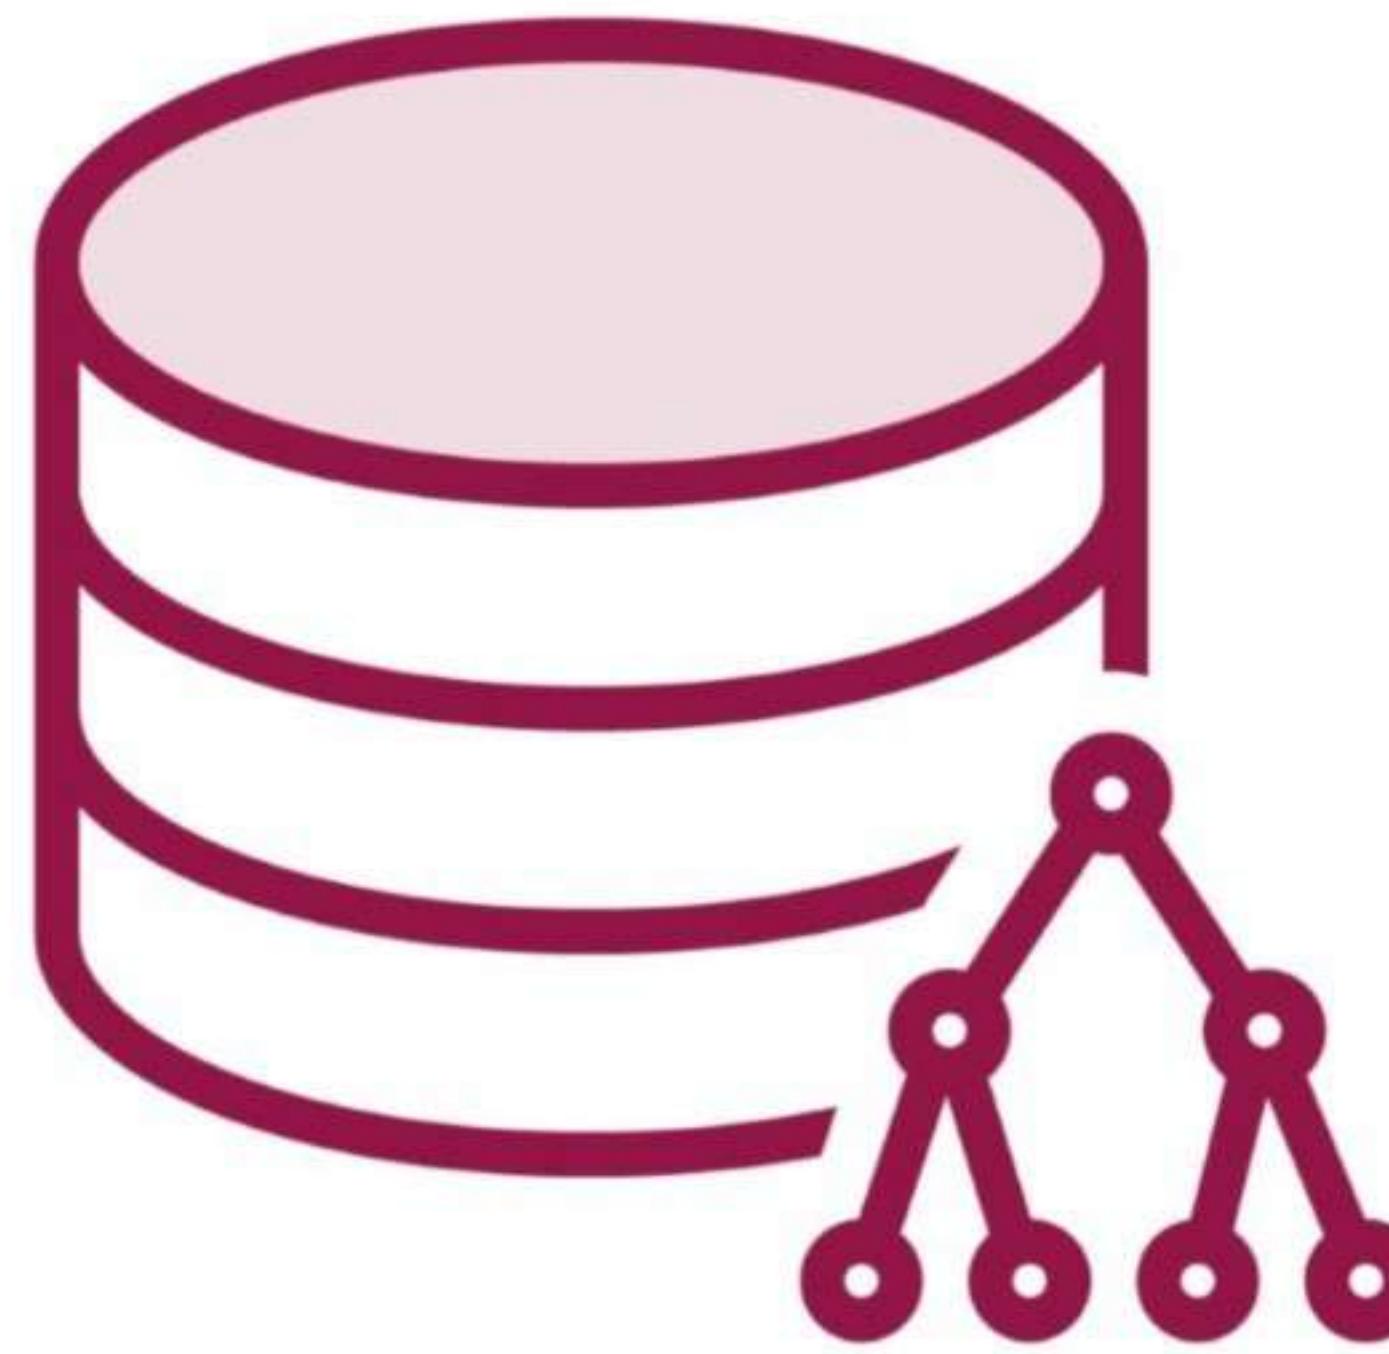
Emphasize relationships over entities

Nodes and edges

- Nodes for entities
- Edges for relationships

DGraph
IBM Graph
Amazon
DB

Graph Databases

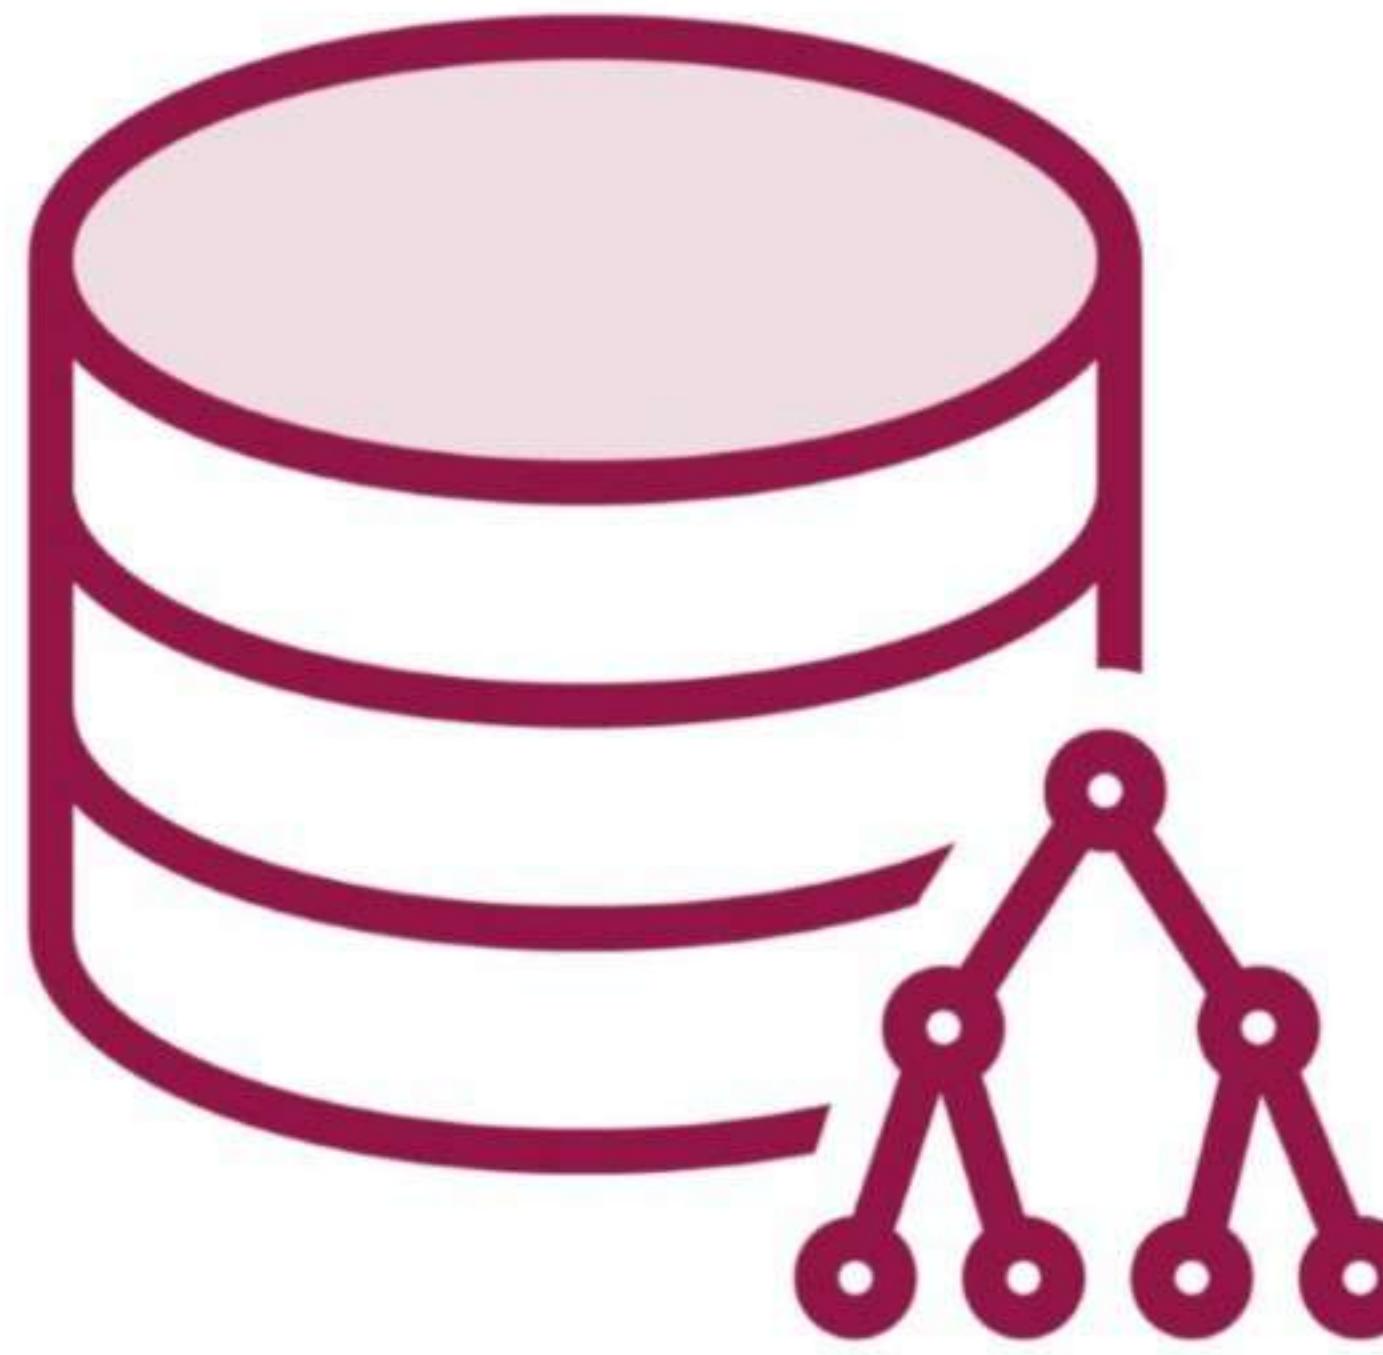


Support for semantic queries

- Query based on associations, context

Quickly retrieve complex hierarchical structures

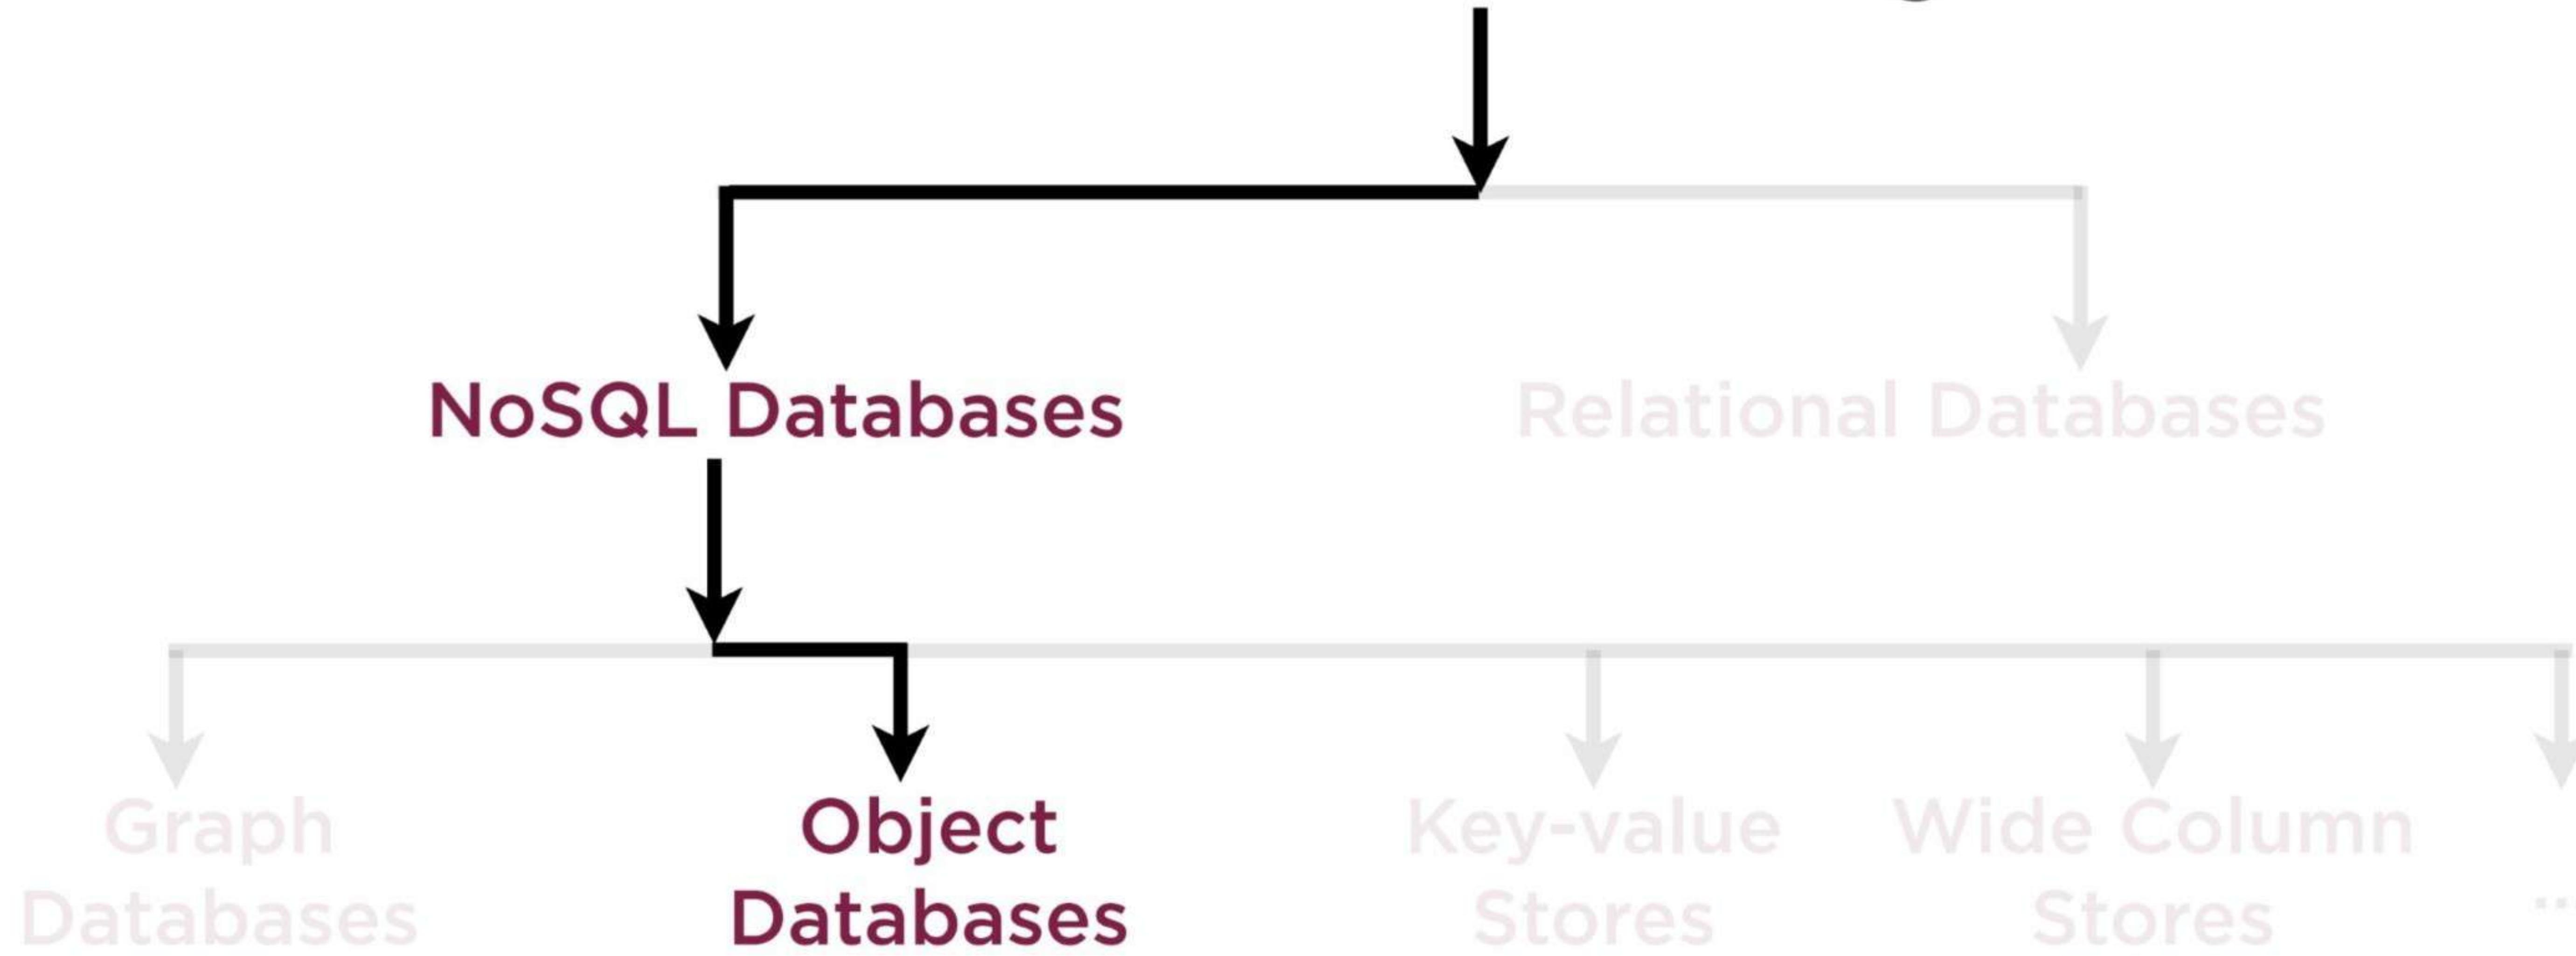
Graph Databases



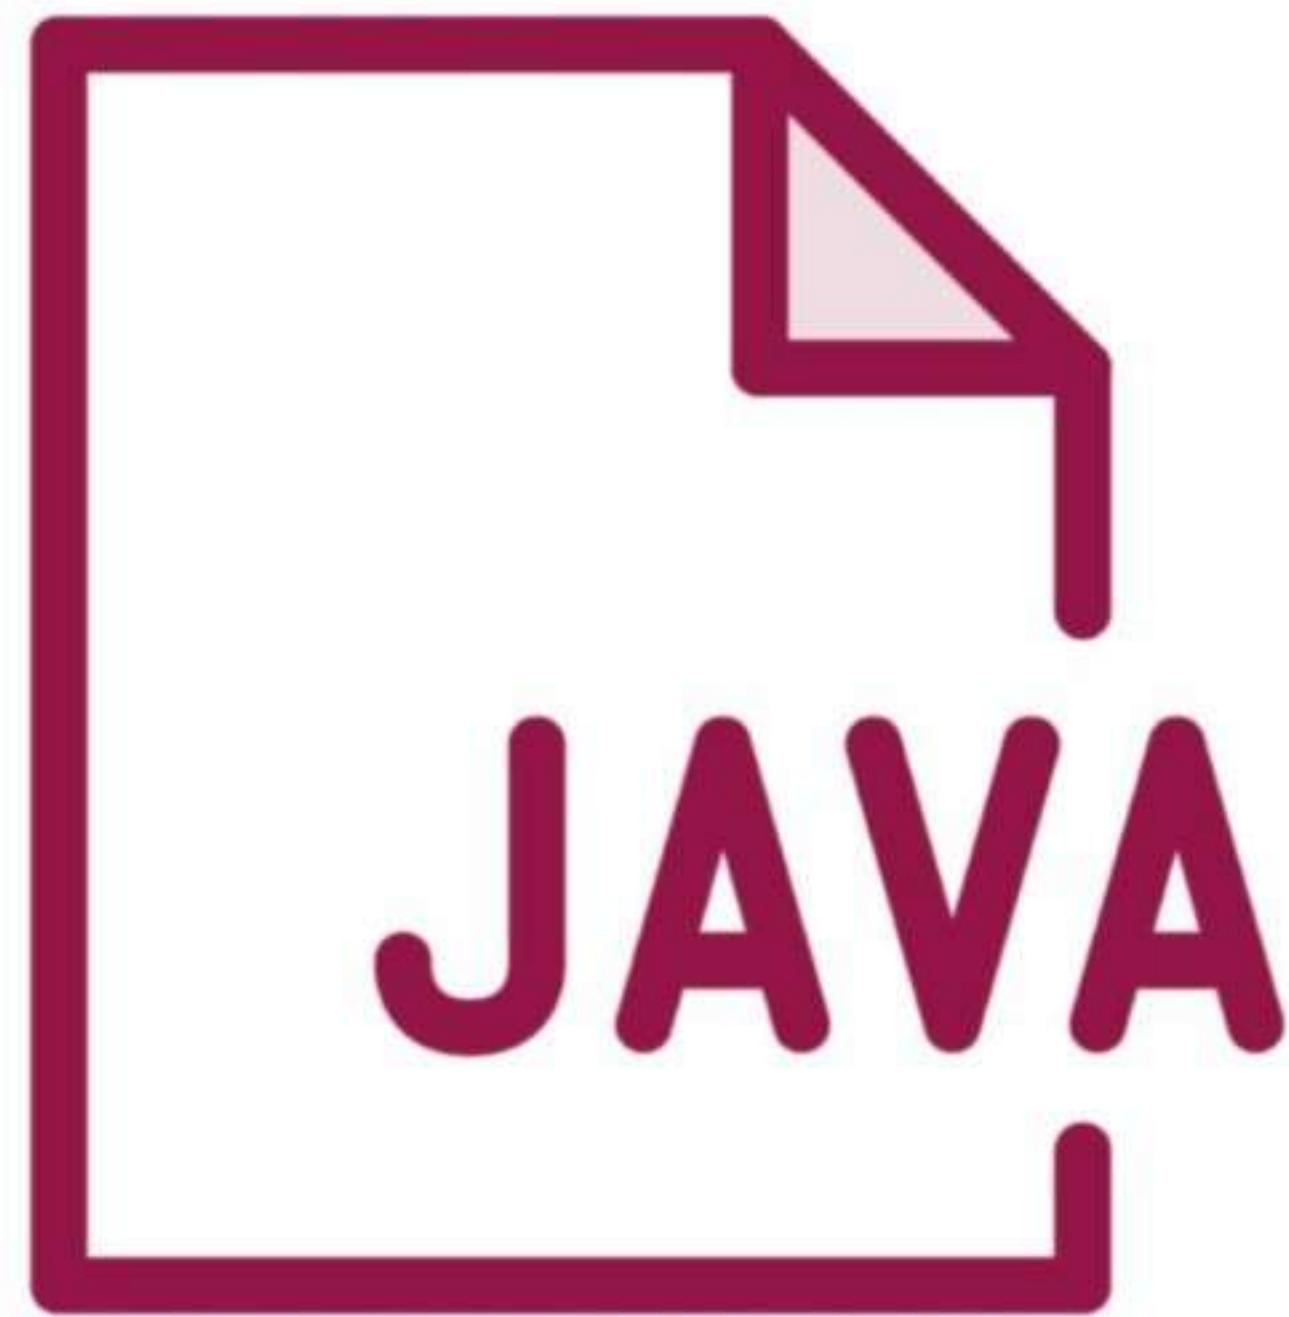
No standard query language

- SQL is ill-suited to graph databases
- Major barrier to adoption

Database Technologies



Object Databases



Programming languages usually model data using classes, objects

Relational databases model data using tables, rows

“Object-Relational Impedance Mismatch”

Object Databases attempt to solve this

Object Databases



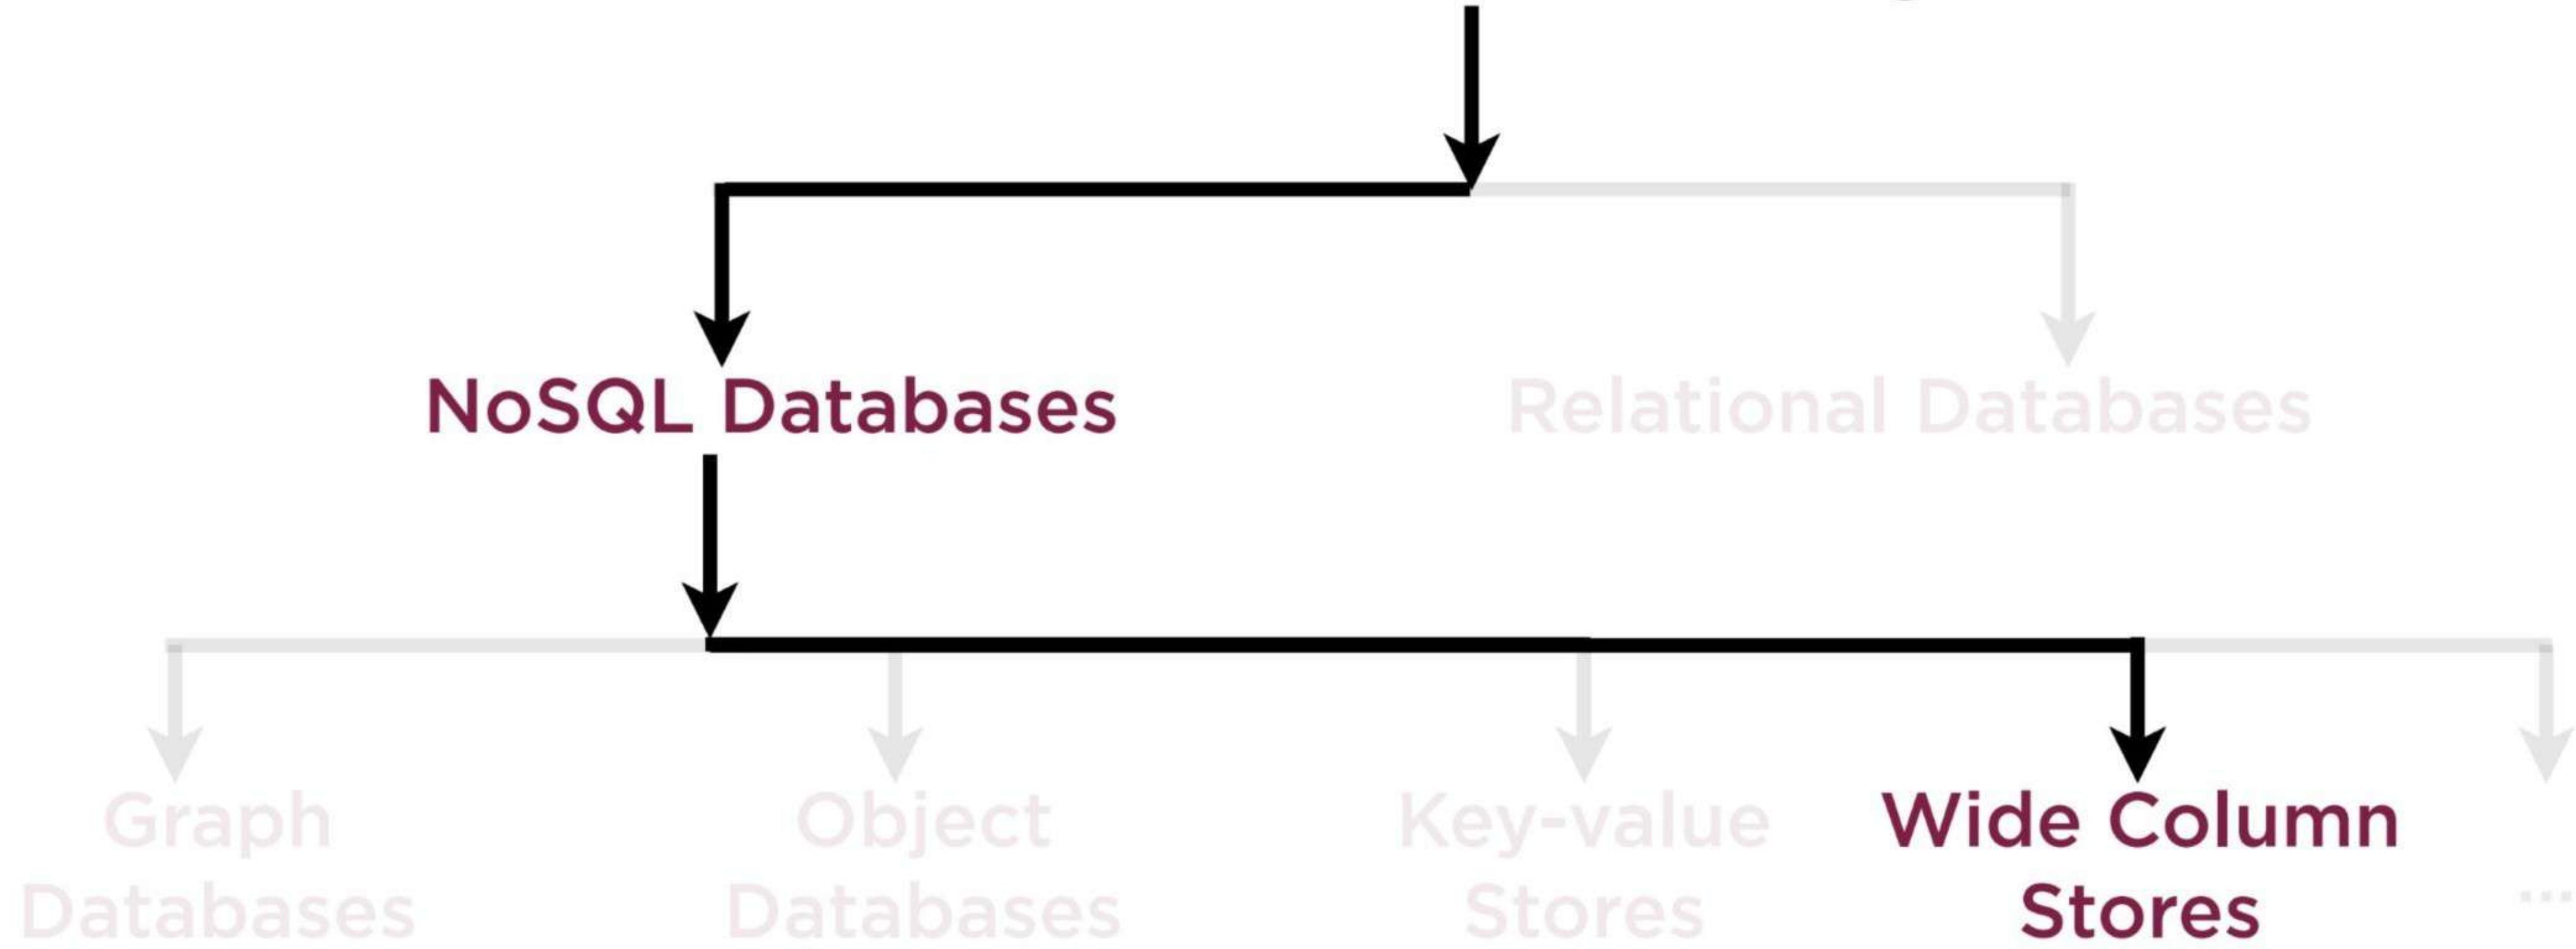
No standard language

- SQL ill-suited to object databases
- Major barrier to adoption

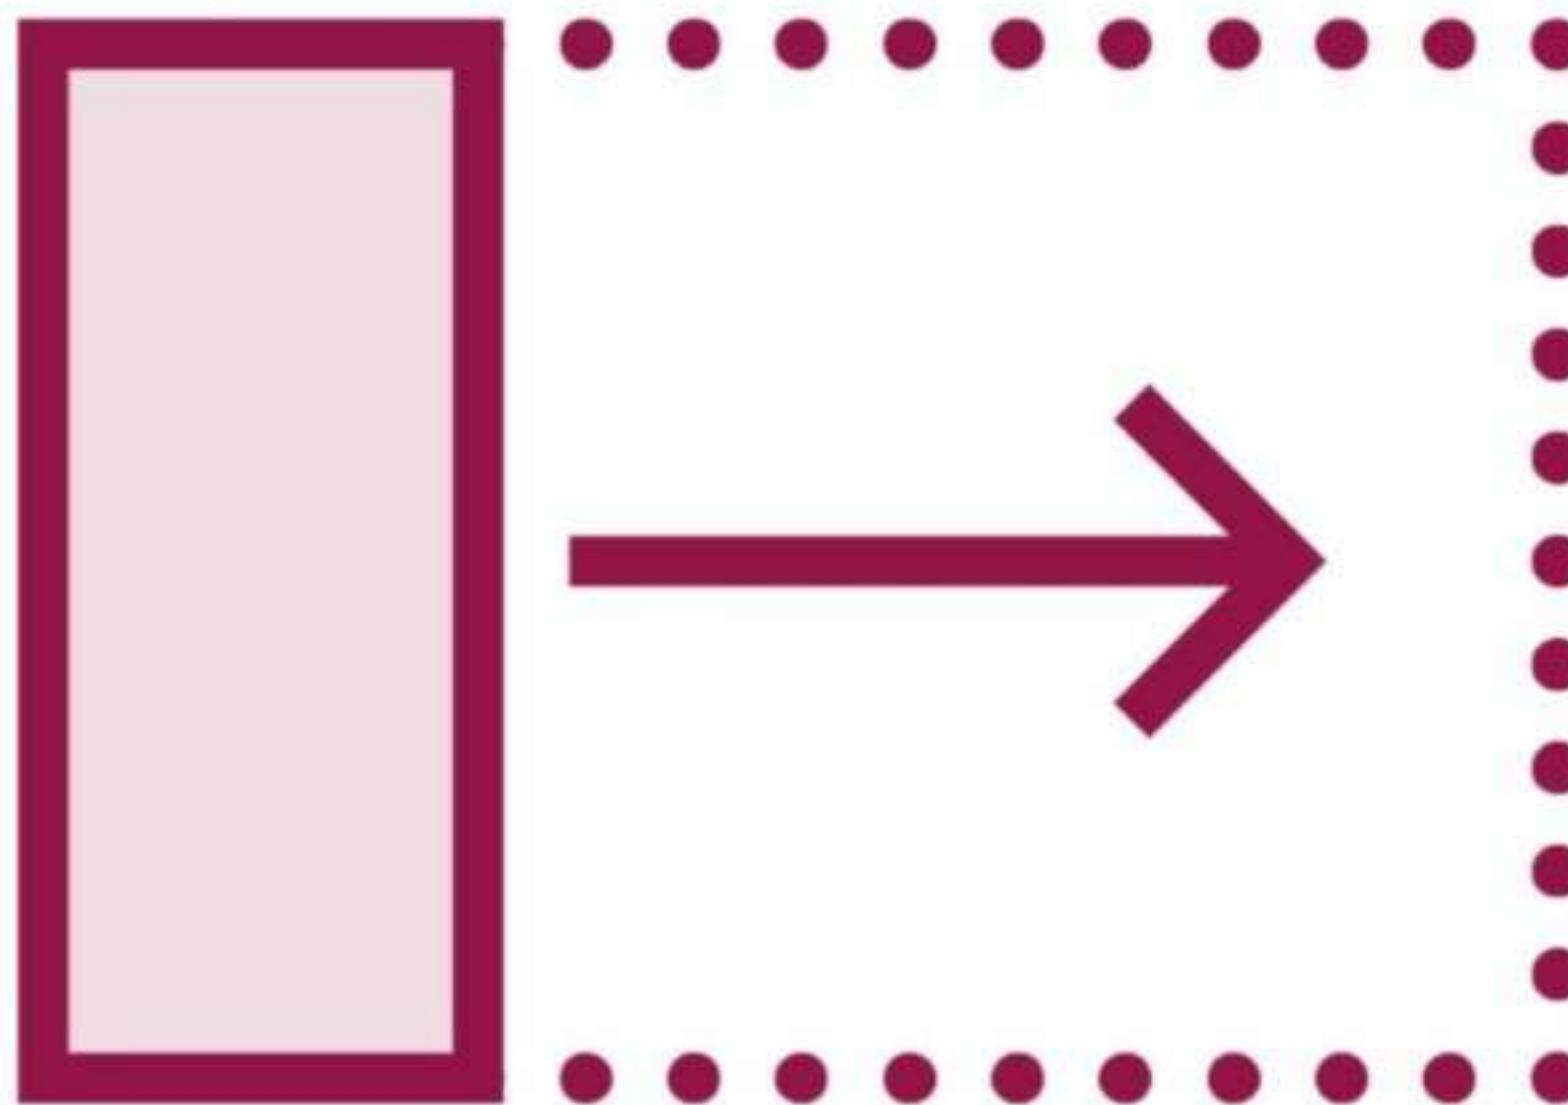
Related to ORM frameworks

- Hibernate
- JPA

Database Technologies



Wide Column Databases



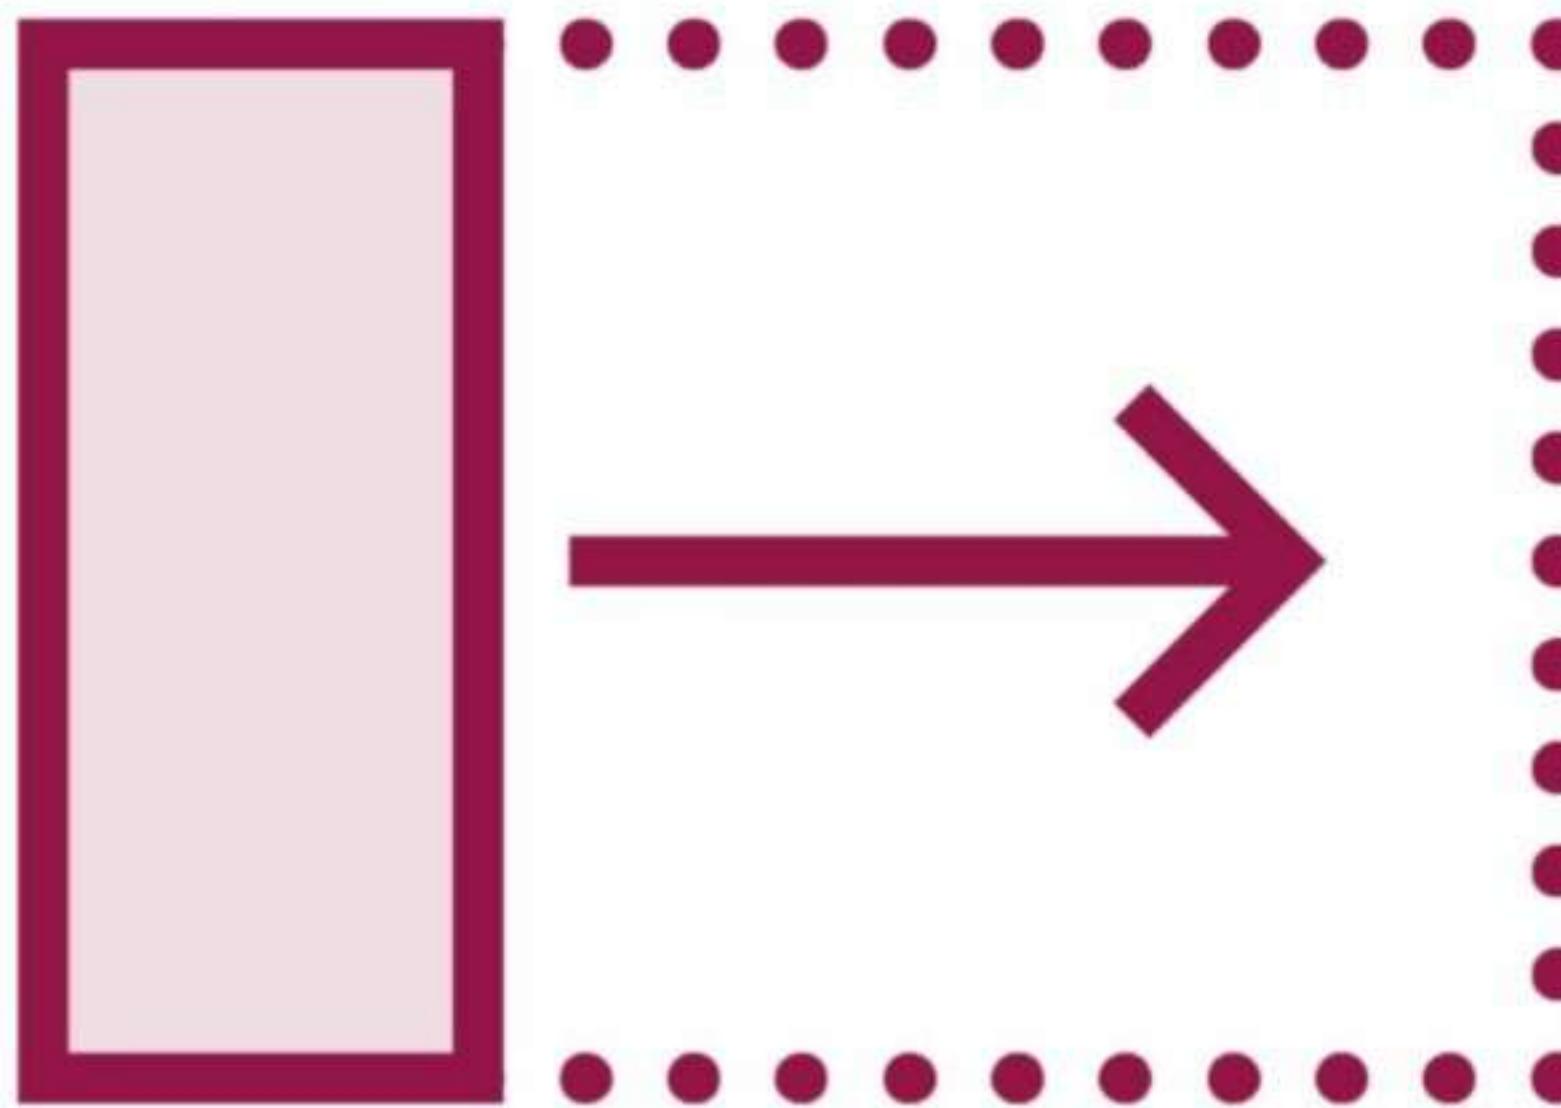
Relational databases feature fixed schemas

Altering schemas to add/remove columns is onerous

NULL values occupy significant space

Wide Column databases address these weaknesses

Wide Column Databases



Several wide column databases have achieved widespread popularity

- HBase
- Cassandra

Syntax closer to SQL

- Has helped drive adoption

Relational Data Model: Wide Tables

Id	To	Type	Content
1	mike	offer	Mobile offer
2	john	sale	Redmi sale
3	jill	order	Order delivered
4	megan	sale	Clothes sale

As columns are added, table gets wider

Relational Data Model: Wide Tables

Id	To	Type	Content	Value
1	mike	offer	Mobile offer	NULL
2	john	sale	Redmi sale	NULL
3	jill	order	Order delivered	NULL
4	megan	sale	Clothes sale	\$53

Addition of columns can lead to several null values

From Wide To Long

Id	To	Type	Content	Value
1	mike	offer	Mobile offer	NULL
2	john	sale	Redmi sale	NULL
3	jill	order	Order delivered	NULL
4	megan	sale	Clothes sale	\$53



Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Mobile offer
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	megan
4	Type	sale
4	Content	Clothes sale
4	Value	\$53

From Wide To Long

Id	To	Type	Content	Value
1	mike	offer	Mobile offer	NULL
2	john	sale	Redmi sale	NULL
3	jill	order	Order delivered	NULL
4	megan	sale	Clothes sale	\$53



Id	Column	Value
1	To	mike
1	Type	offer
1	Content	Mobile offer
2	To	john
2	Type	sale
2	Content	Redmi sale
3	To	jill
3	Type	order
3	Content	Order delivered
4	To	megan
4	Type	sale
4	Content	Clothes sale
4	Value	\$53

Introduction to Cassandra

Cassandra - A Decentralized

Avinash Lakshman
Facebook

Roshant
Malik

Inbox Search

By Carolyn Abram on Thursday, June 19, 2008 at 8:16pm

This month, we're rolling out Inbox search, a tool which allows you to search through all of your messages either by name of the person who sent it, or by a keyword that shows up in the text. Ever had a friend send you a Facebook message suggesting a good restaurant? Or maybe you were sent an address or phone number of an old friend. Then time passes and your Inbox fills up, and by the time you're ready to go out to eat with your old friend you have to page through hundreds of Facebook messages to find the one you need. Inbox Search aims to fix this problem.

Compose Message

Search Inbox

Apache Cassandra

open source, distributed database,

designed to handle large amount of data across commodity servers

NoSQL

→ Highly Available

No SPOF
Single Point of Failure

Peer to Peer Network

Major Focus → Performance

Cluster based

commodity cluster



The image shows a Mac OS X desktop environment with two browser windows open. The top window is a Google search results page for the query "Bigtable". The search results include links to "Bigtable paper" and "Bigtable paper pdf". The bottom window is a PDF document titled "Dynamo: Amazon's Highly Available Key-value Store" by a team of authors from Amazon. The PDF is hosted on www.allthingsdistributed.com/files/amazon-dynamo-sosp2007.pdf. The desktop background features a dark, abstract pattern.

static.googleusercontent.com

Paul

static.googleusercontent.com/media/research.google.com/en/us/archive/bigtabl...

Bigtable

Fay Chang, J
Mik

www.allthingsdistributed.com

Paul

www.allthingsdistributed.com/files/amazon-dynamo-sosp2007.pdf

Dynamo: Amazon's Highly Available Key-value Store

Giuseppe DeCandia, Deniz Hastorun, Madan Jampani, Gunavardhan Kakulapati,
Avinash Lakshman, Alex Pilchin, Swaminathan Sivasubramanian, Peter Vosshall
and Werner Vogels

Amazon.com

EXTRACT

One of the lessons our organization has learned from operating



Manage massive amounts of data, fast, without losing sleep

[Download Cassandra](#)

[Cassandra 3.11.4 Changelog](#)

What is Cassandra?

The Apache Cassandra database is the right choice when you need scalability and high availability without compromising performance. [Linear scalability](#) and proven fault-tolerance on commodity hardware or cloud infrastructure make it the perfect platform for mission-critical data. Cassandra's support for replicating across multiple datacenters is best-in-class, providing lower latency for your users and the peace of mind of knowing that you can survive regional outages.

<https://cassandra.apache.org/>

Migrating Netflix from Data X

www.slideshare.net/adrianco/migrating-netflix-from-oracle-to-global-cassandra

slideshare | Search

Home Leadership Technology Education M

Replacing Datacenter Oracle with Global Apache Cassandra

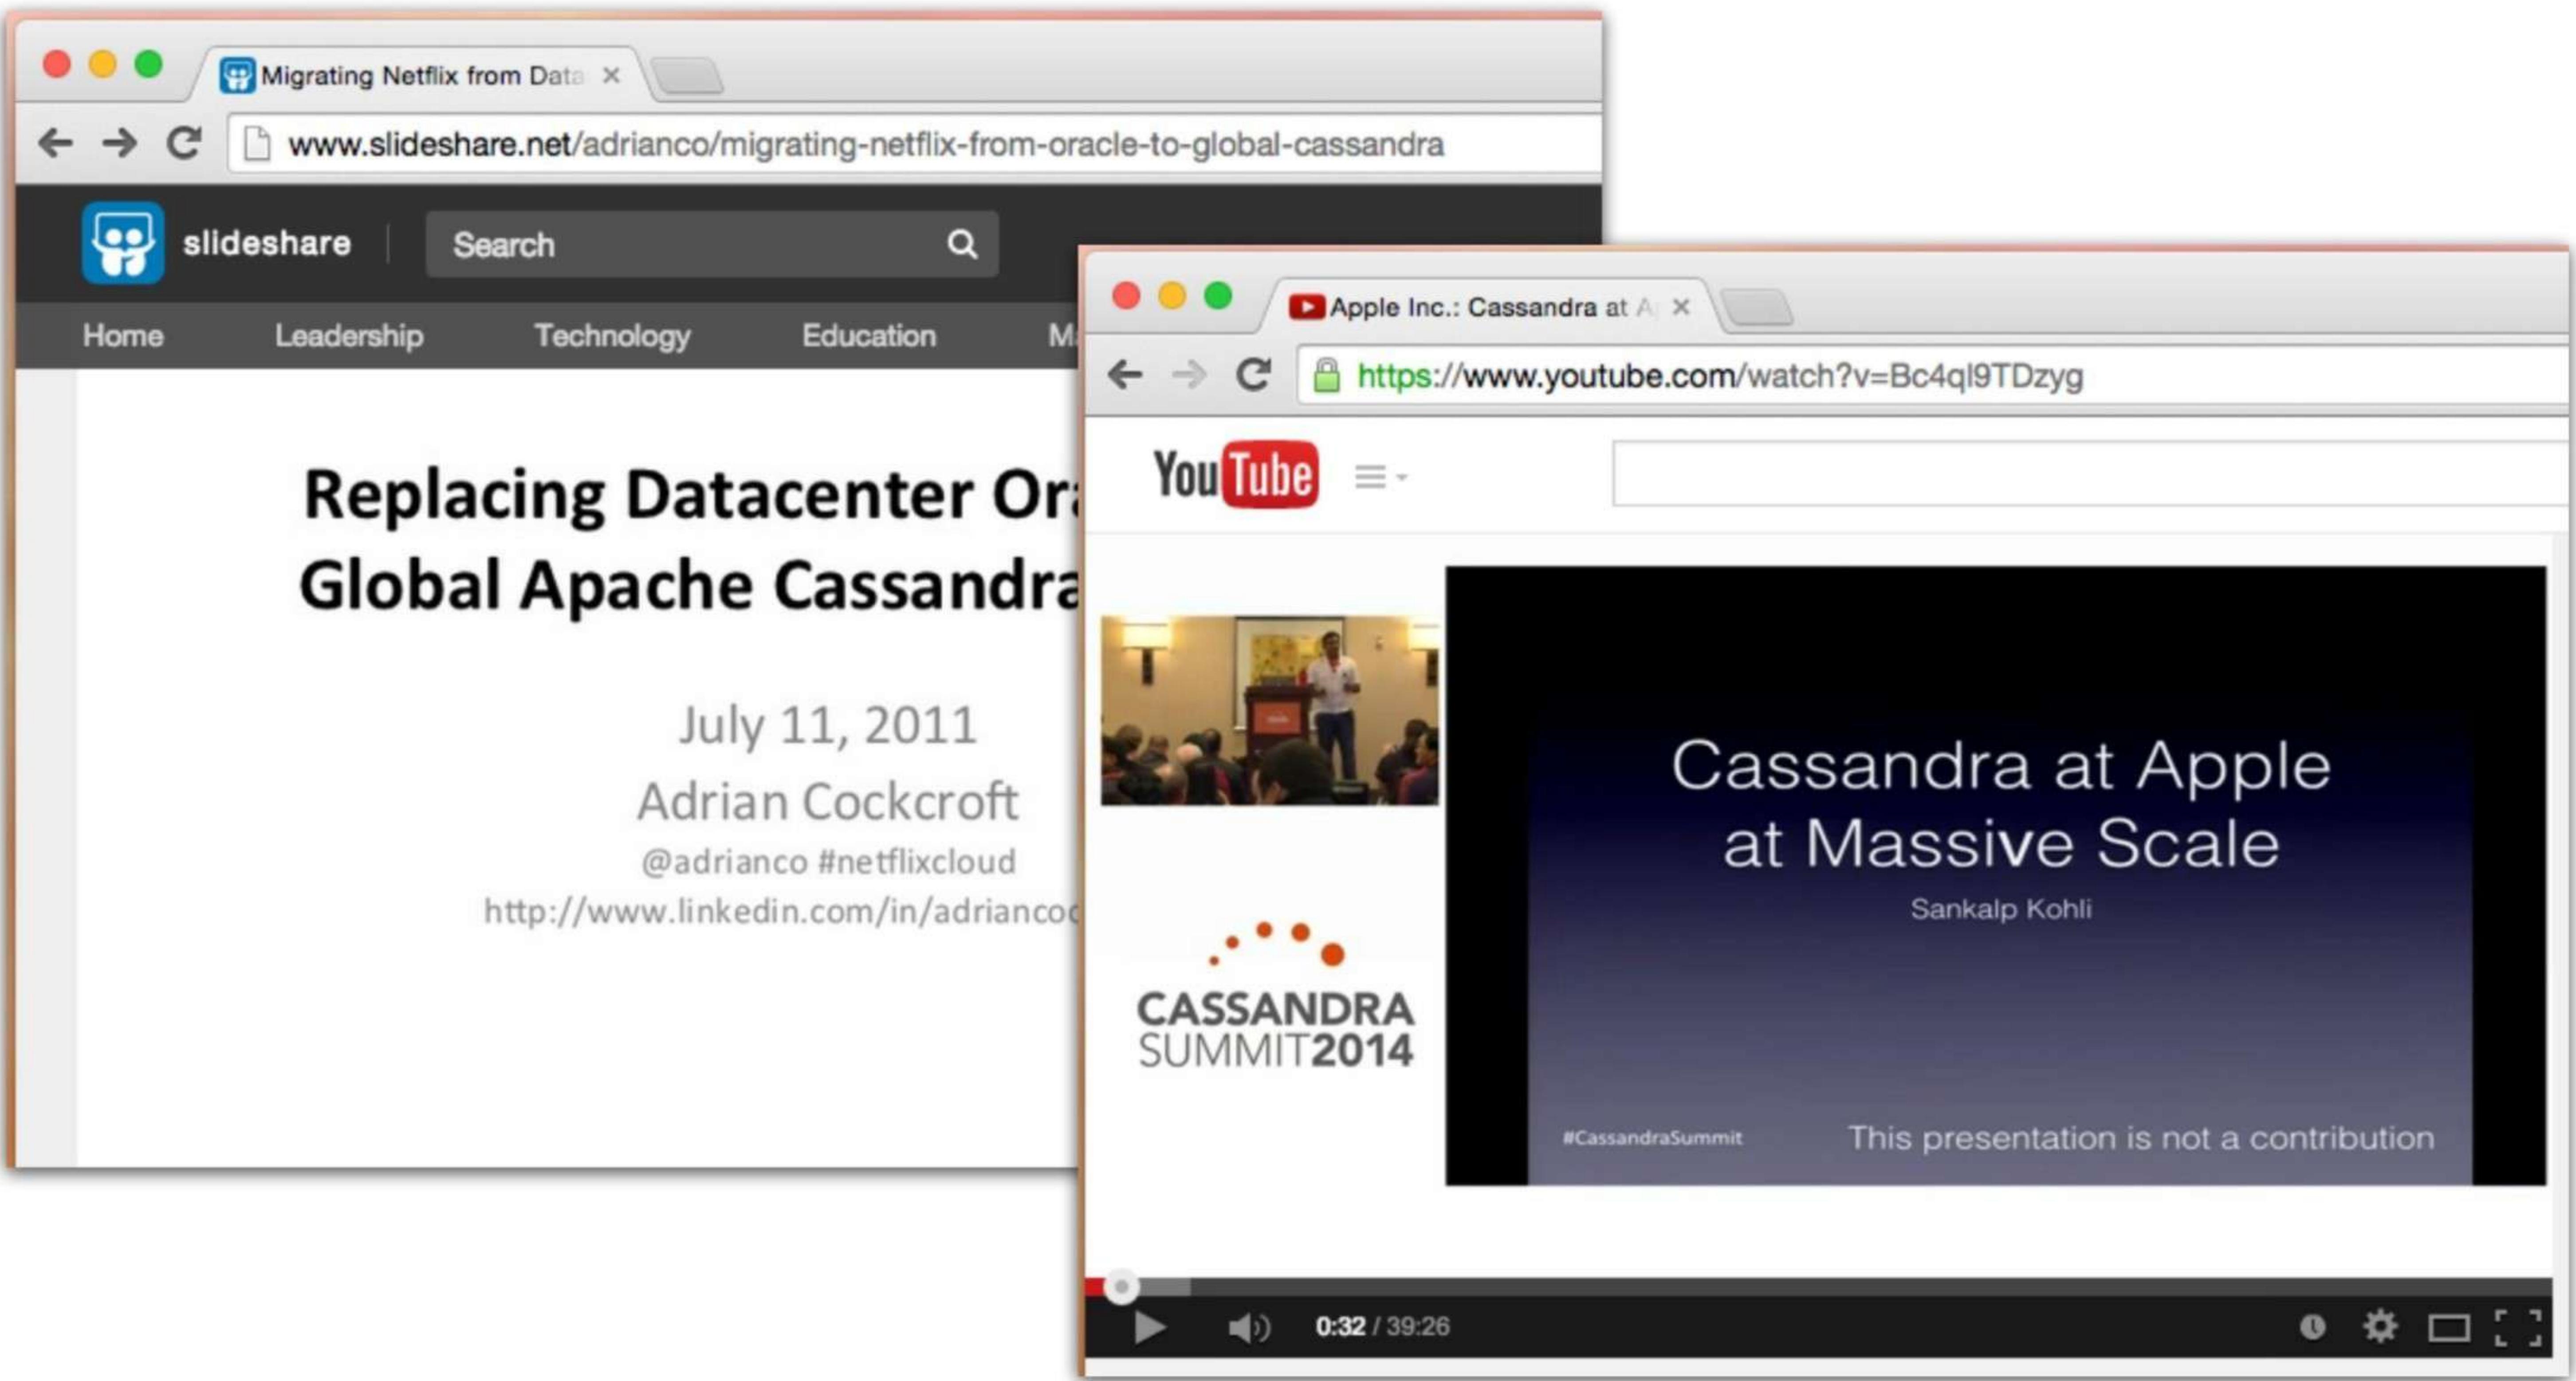
July 11, 2011

Adrian Cockcroft
@adrianco #netflixcloud
<http://www.linkedin.com/in/adrianco>

Apple Inc.: Cassandra at A X

https://www.youtube.com/watch?v=Bc4ql9TDzyg

YouTube



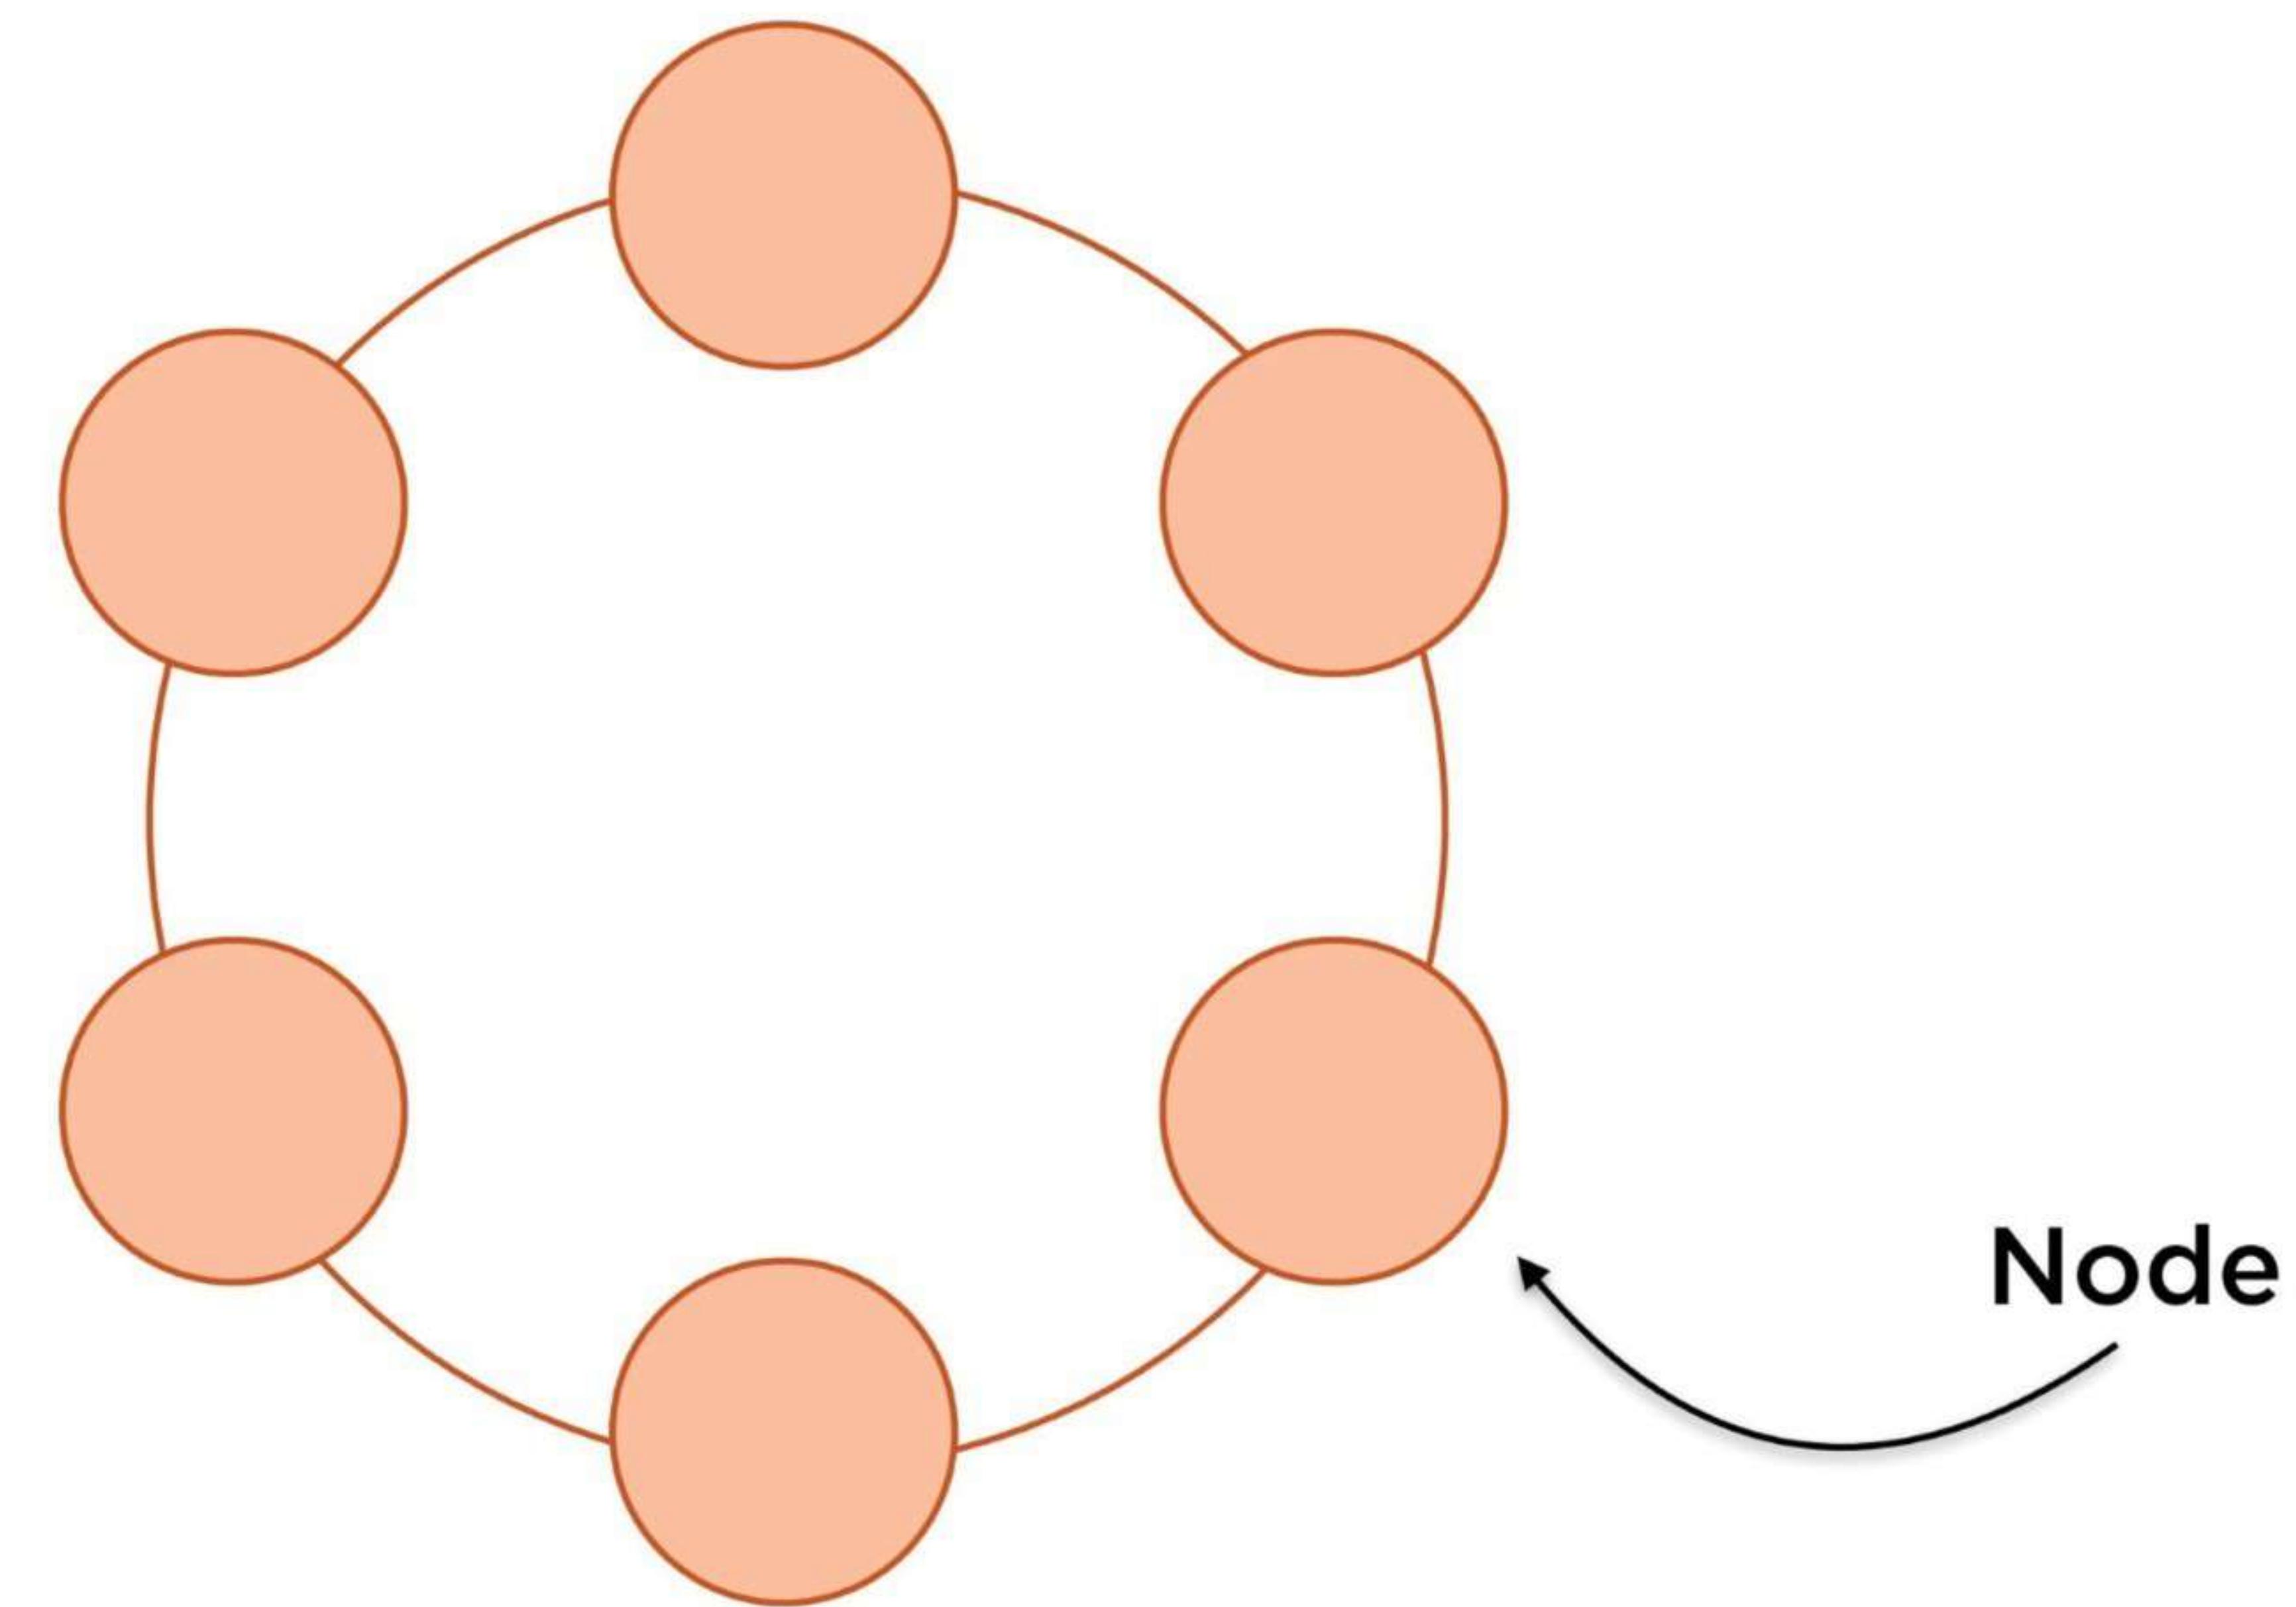
Cassandra at Apple at Massive Scale

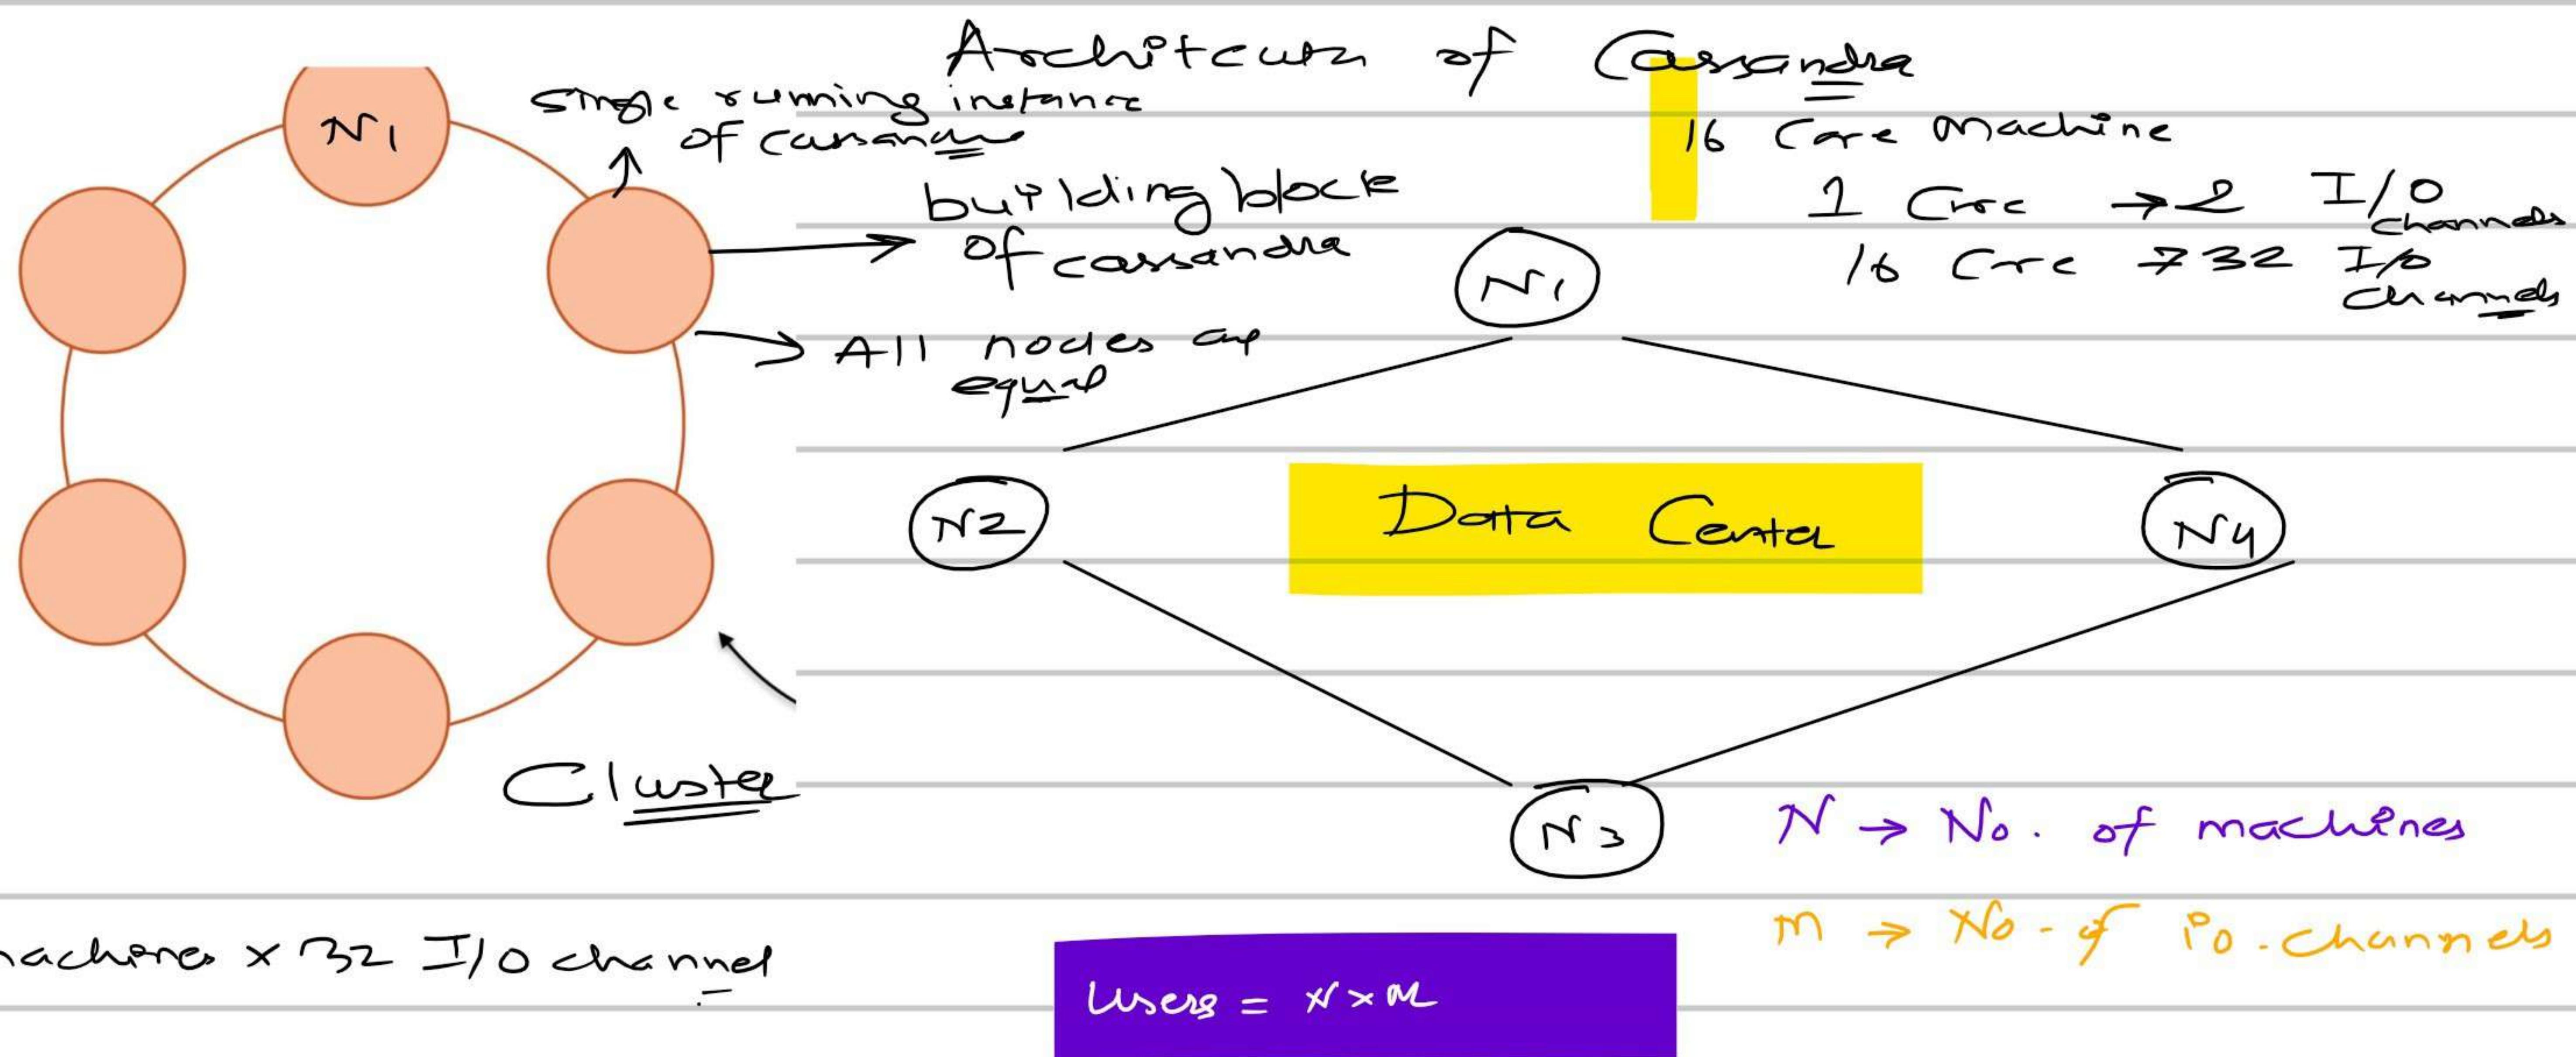
Sankalp Kohli

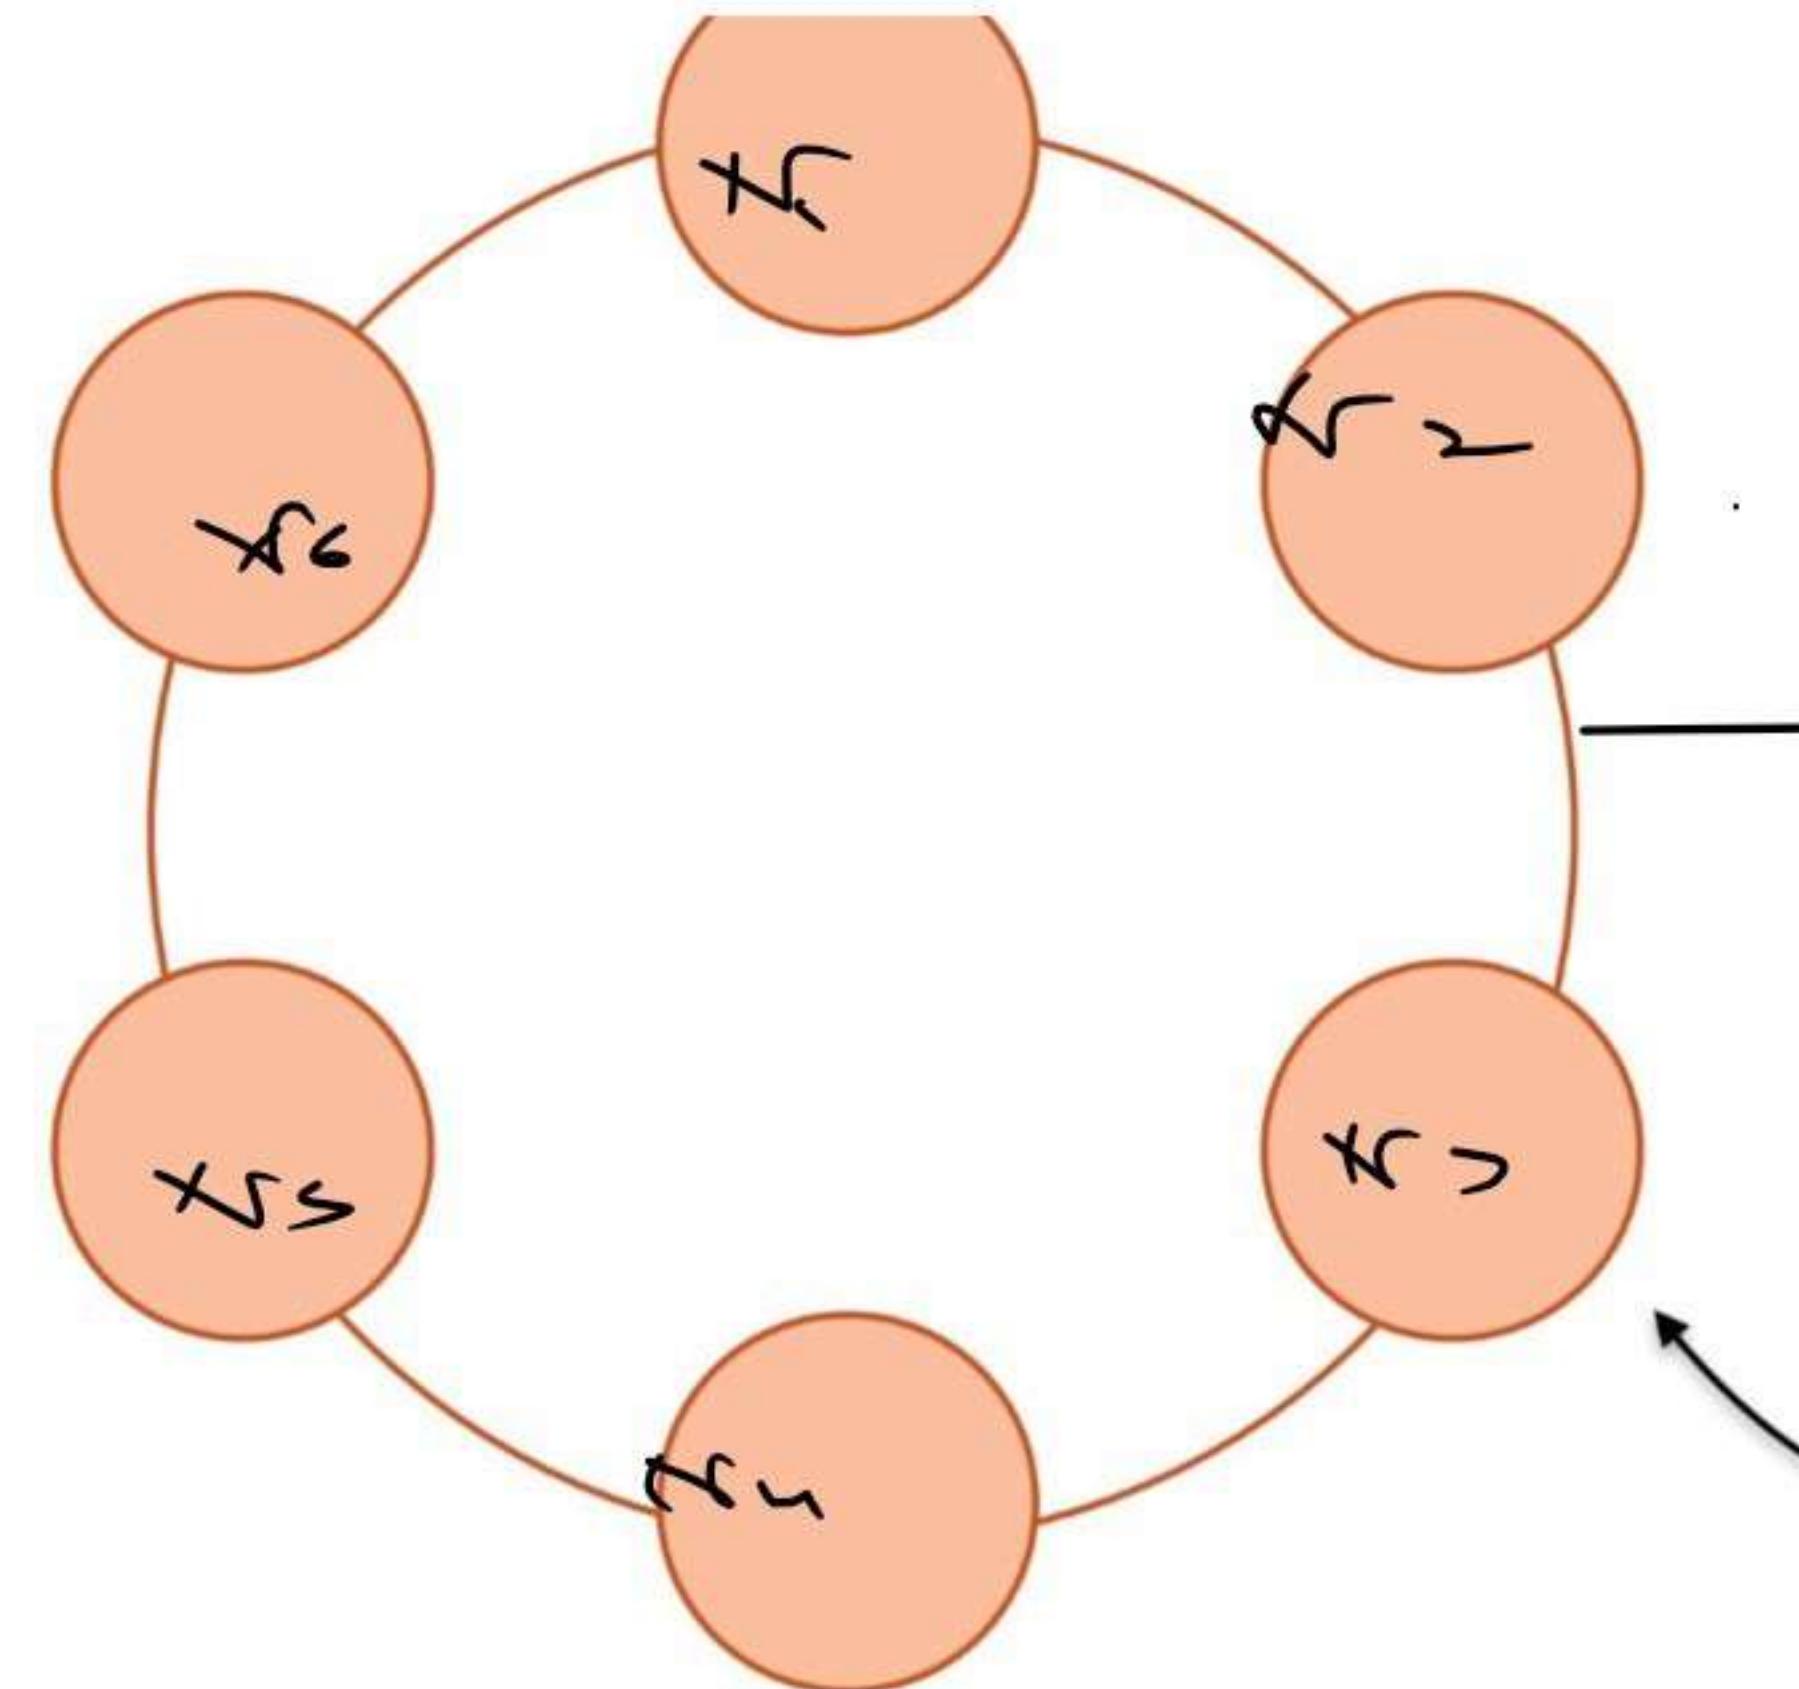
#CassandraSummit

This presentation is not a contribution

0:32 / 39:26







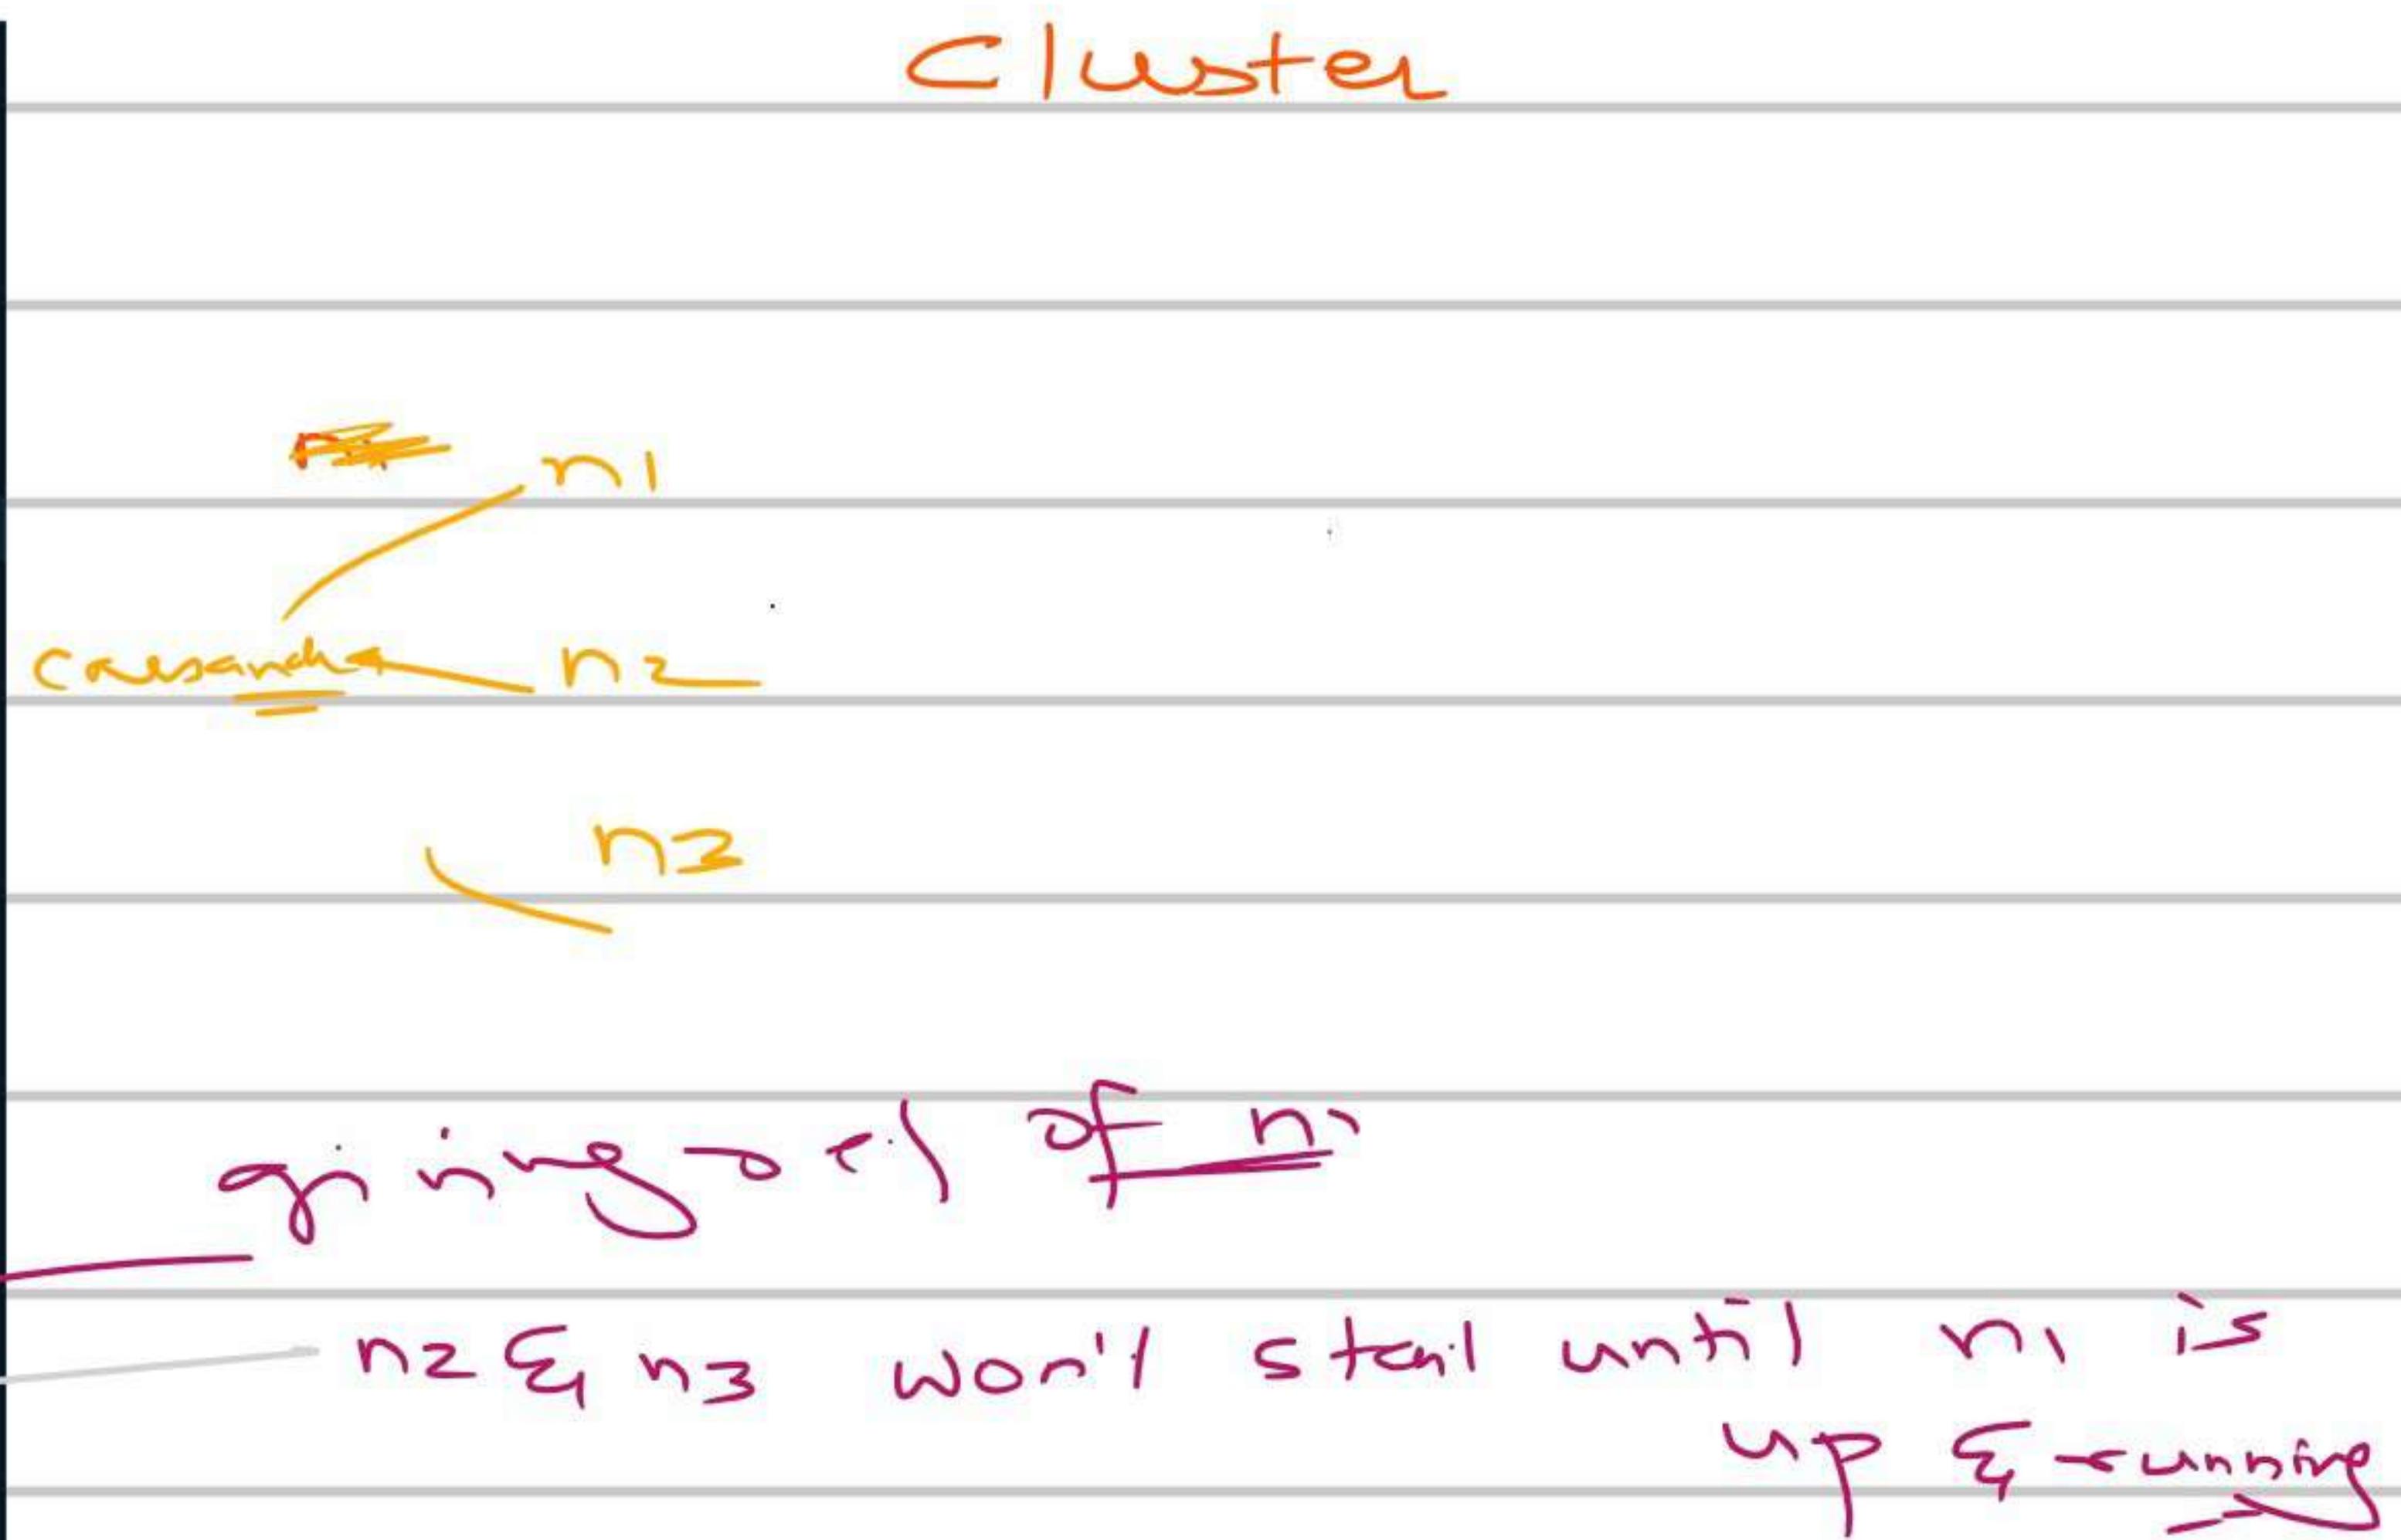
Cluster

span

→ collection of interconnected nodes

→ appear as one single logical db to client

```
>Run All Services  
services:  
  > Run Service  
    n1:  
      image: cassandra:3.11  
      networks:  
        - cluster  
    > Run Service  
    n2:  
      image: cassandra:3.11  
      networks:  
        - cluster  
      environment:  
        - CASSANDRA_SEEDS=n1  
      depends_on:  
        - n1  
    > Run Service  
    n3:  
      image: cassandra:3.11  
      networks:  
        - cluster  
      environment:  
        - CASSANDRA_SEEDS=n1  
      depends_on:  
        - n1  
    networks:  
      cluster:
```



How is data structured in Cassandra



Partition → set of rows which have the same partition key

QUERY FIRST APPROACH

before we design
the database;
we should think
of queries &
should design
the database
accordingly.

- Data should be evenly spread across the cluster
 - Minimize the partition size

No limit on storage

Query
and
Structure

Accommodate the data in such a way that the query can be easily accommodated

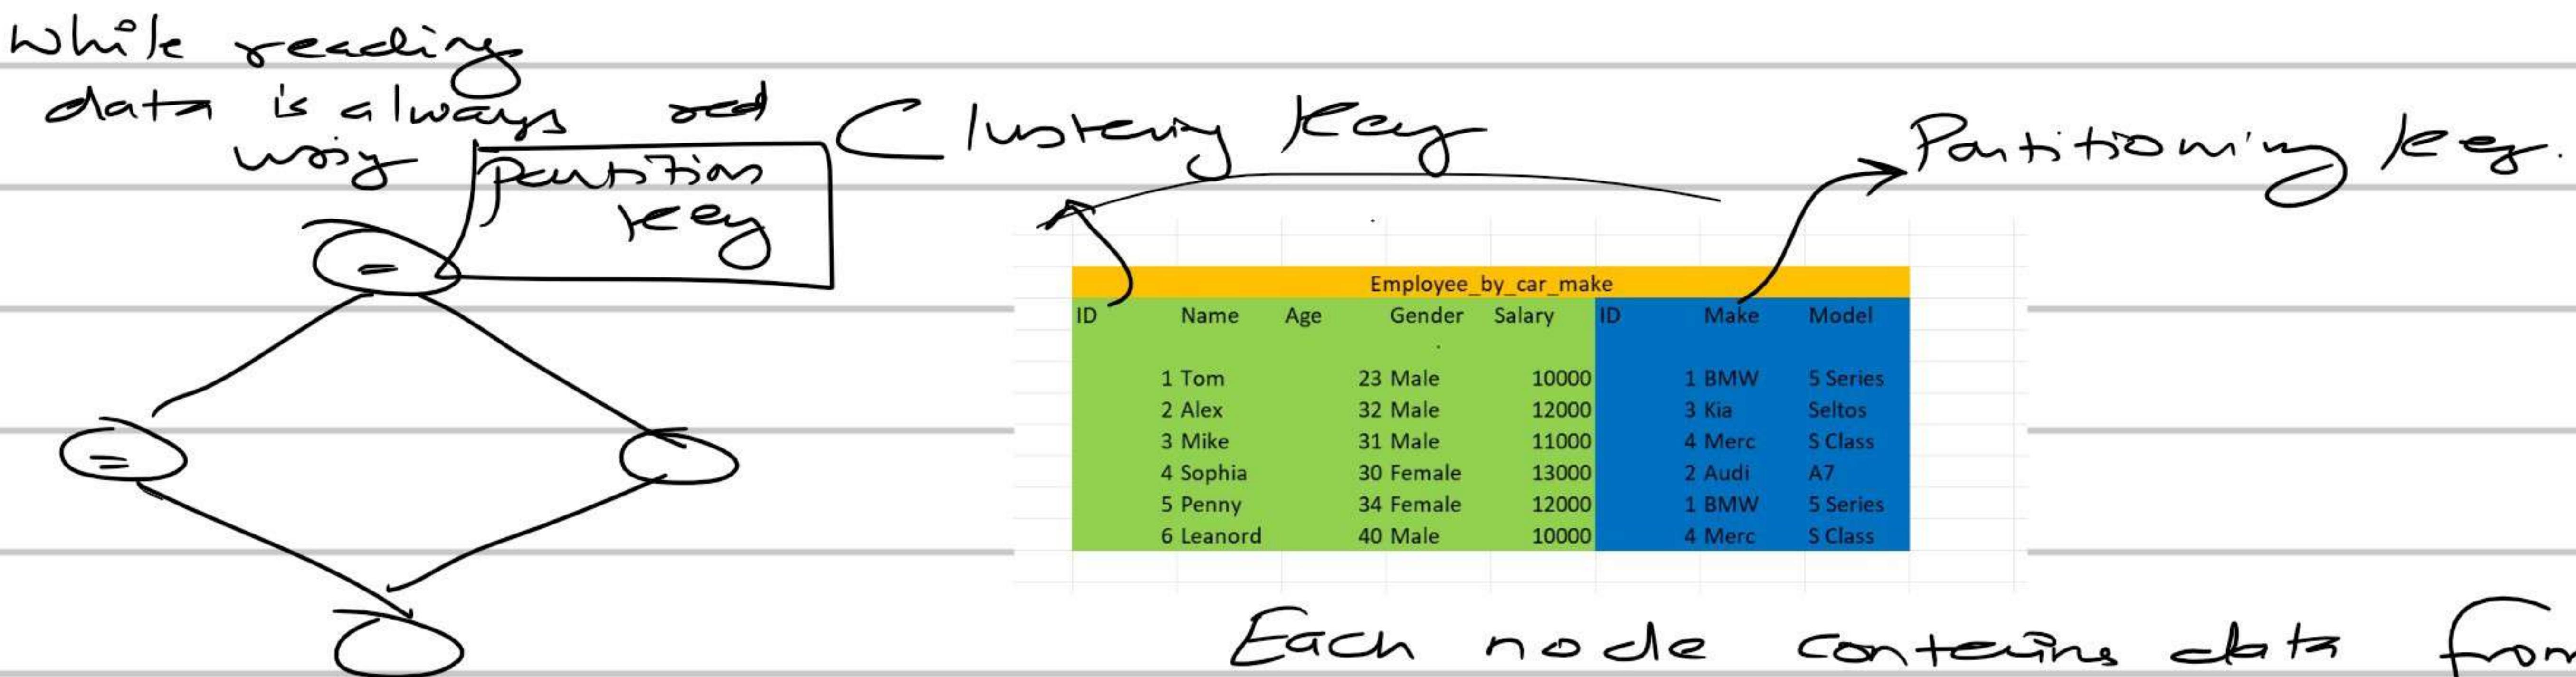
Primary Key in Cassandra

Primary Key → Partitioning + Clustering
Key

① Single Primary Key [Partition Key]

② Composite Primary Key [Partition Key, Clustering Key]
[Partition Key, CK, CK]

③ Composite Partition Primary Key [(PK, PK), CK]



Data Types Cassandra

Basic

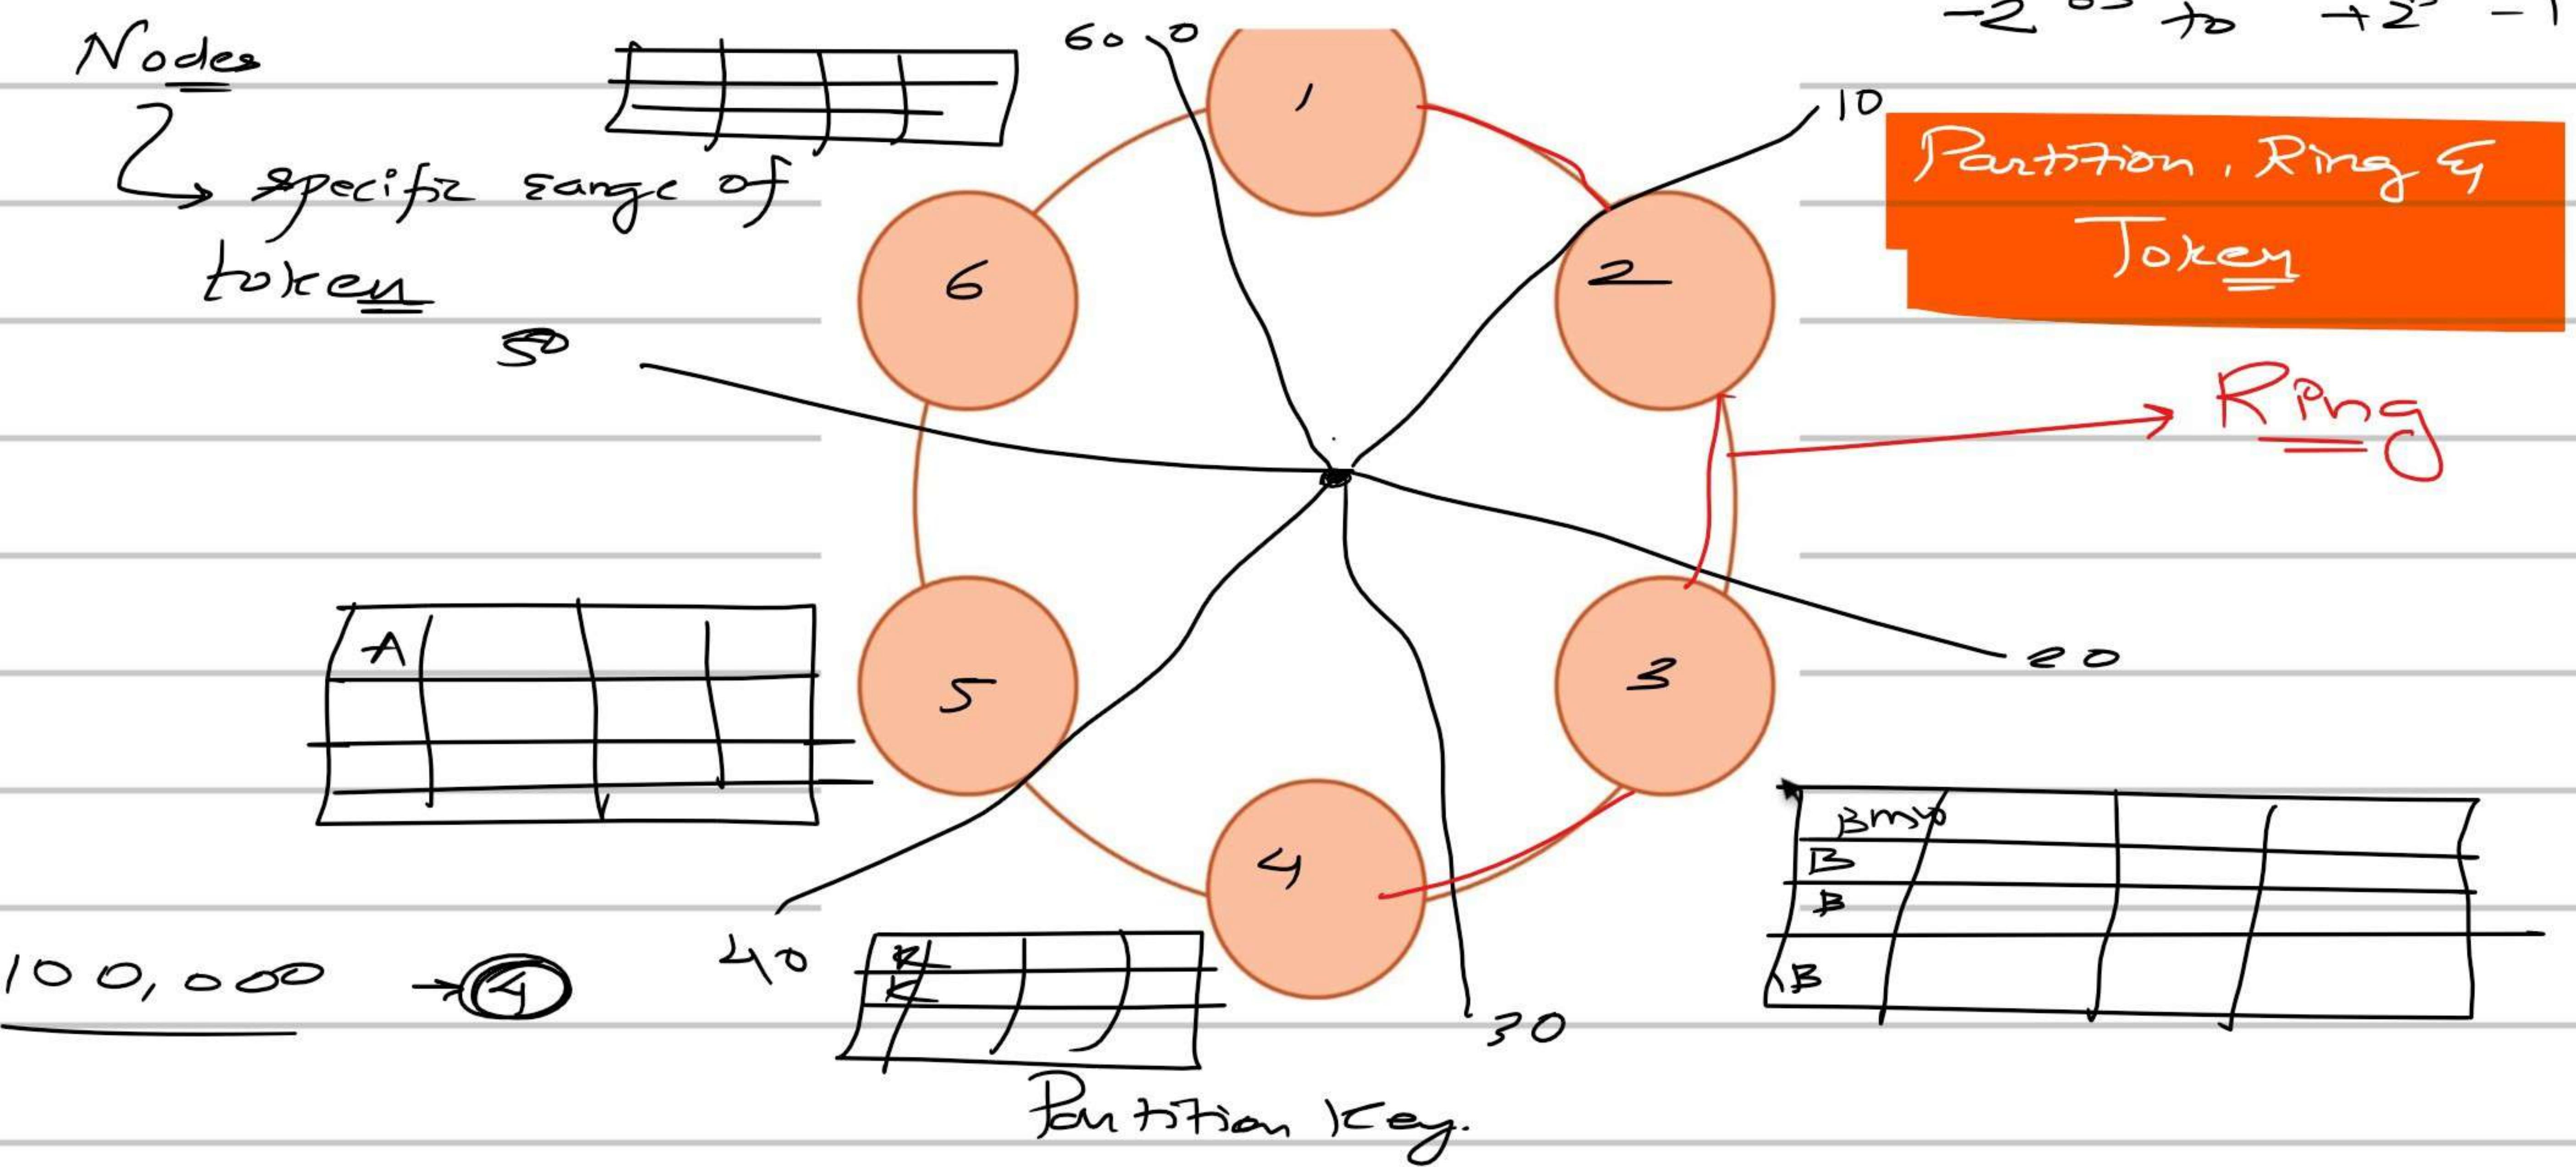
→ int, float, double, varchar, text, blob, bigint,
boolean, decimal, timestamp

Collection

→ list, set & map

UD

→ User Defined



Employee_by_car_make							
ID	Name	Age	Gender	Salary	ID	Make	Model
1 Tom		23 Male		10000	1 BMW	5 Series	
2 Alex		32 Male		12000	3 Kia	Seltos	
3 Mike		31 Male		11000	4 Merc	S Class	
4 Sophia		30 Female		13000	2 Audi	A7	
5 Penny		34 Female		12000	1 BMW	5 Series	
6 Leonard		40 Male		10000	4 Merc	S Class	

Pk → Hash Function → Token value

BMW → HF → 23

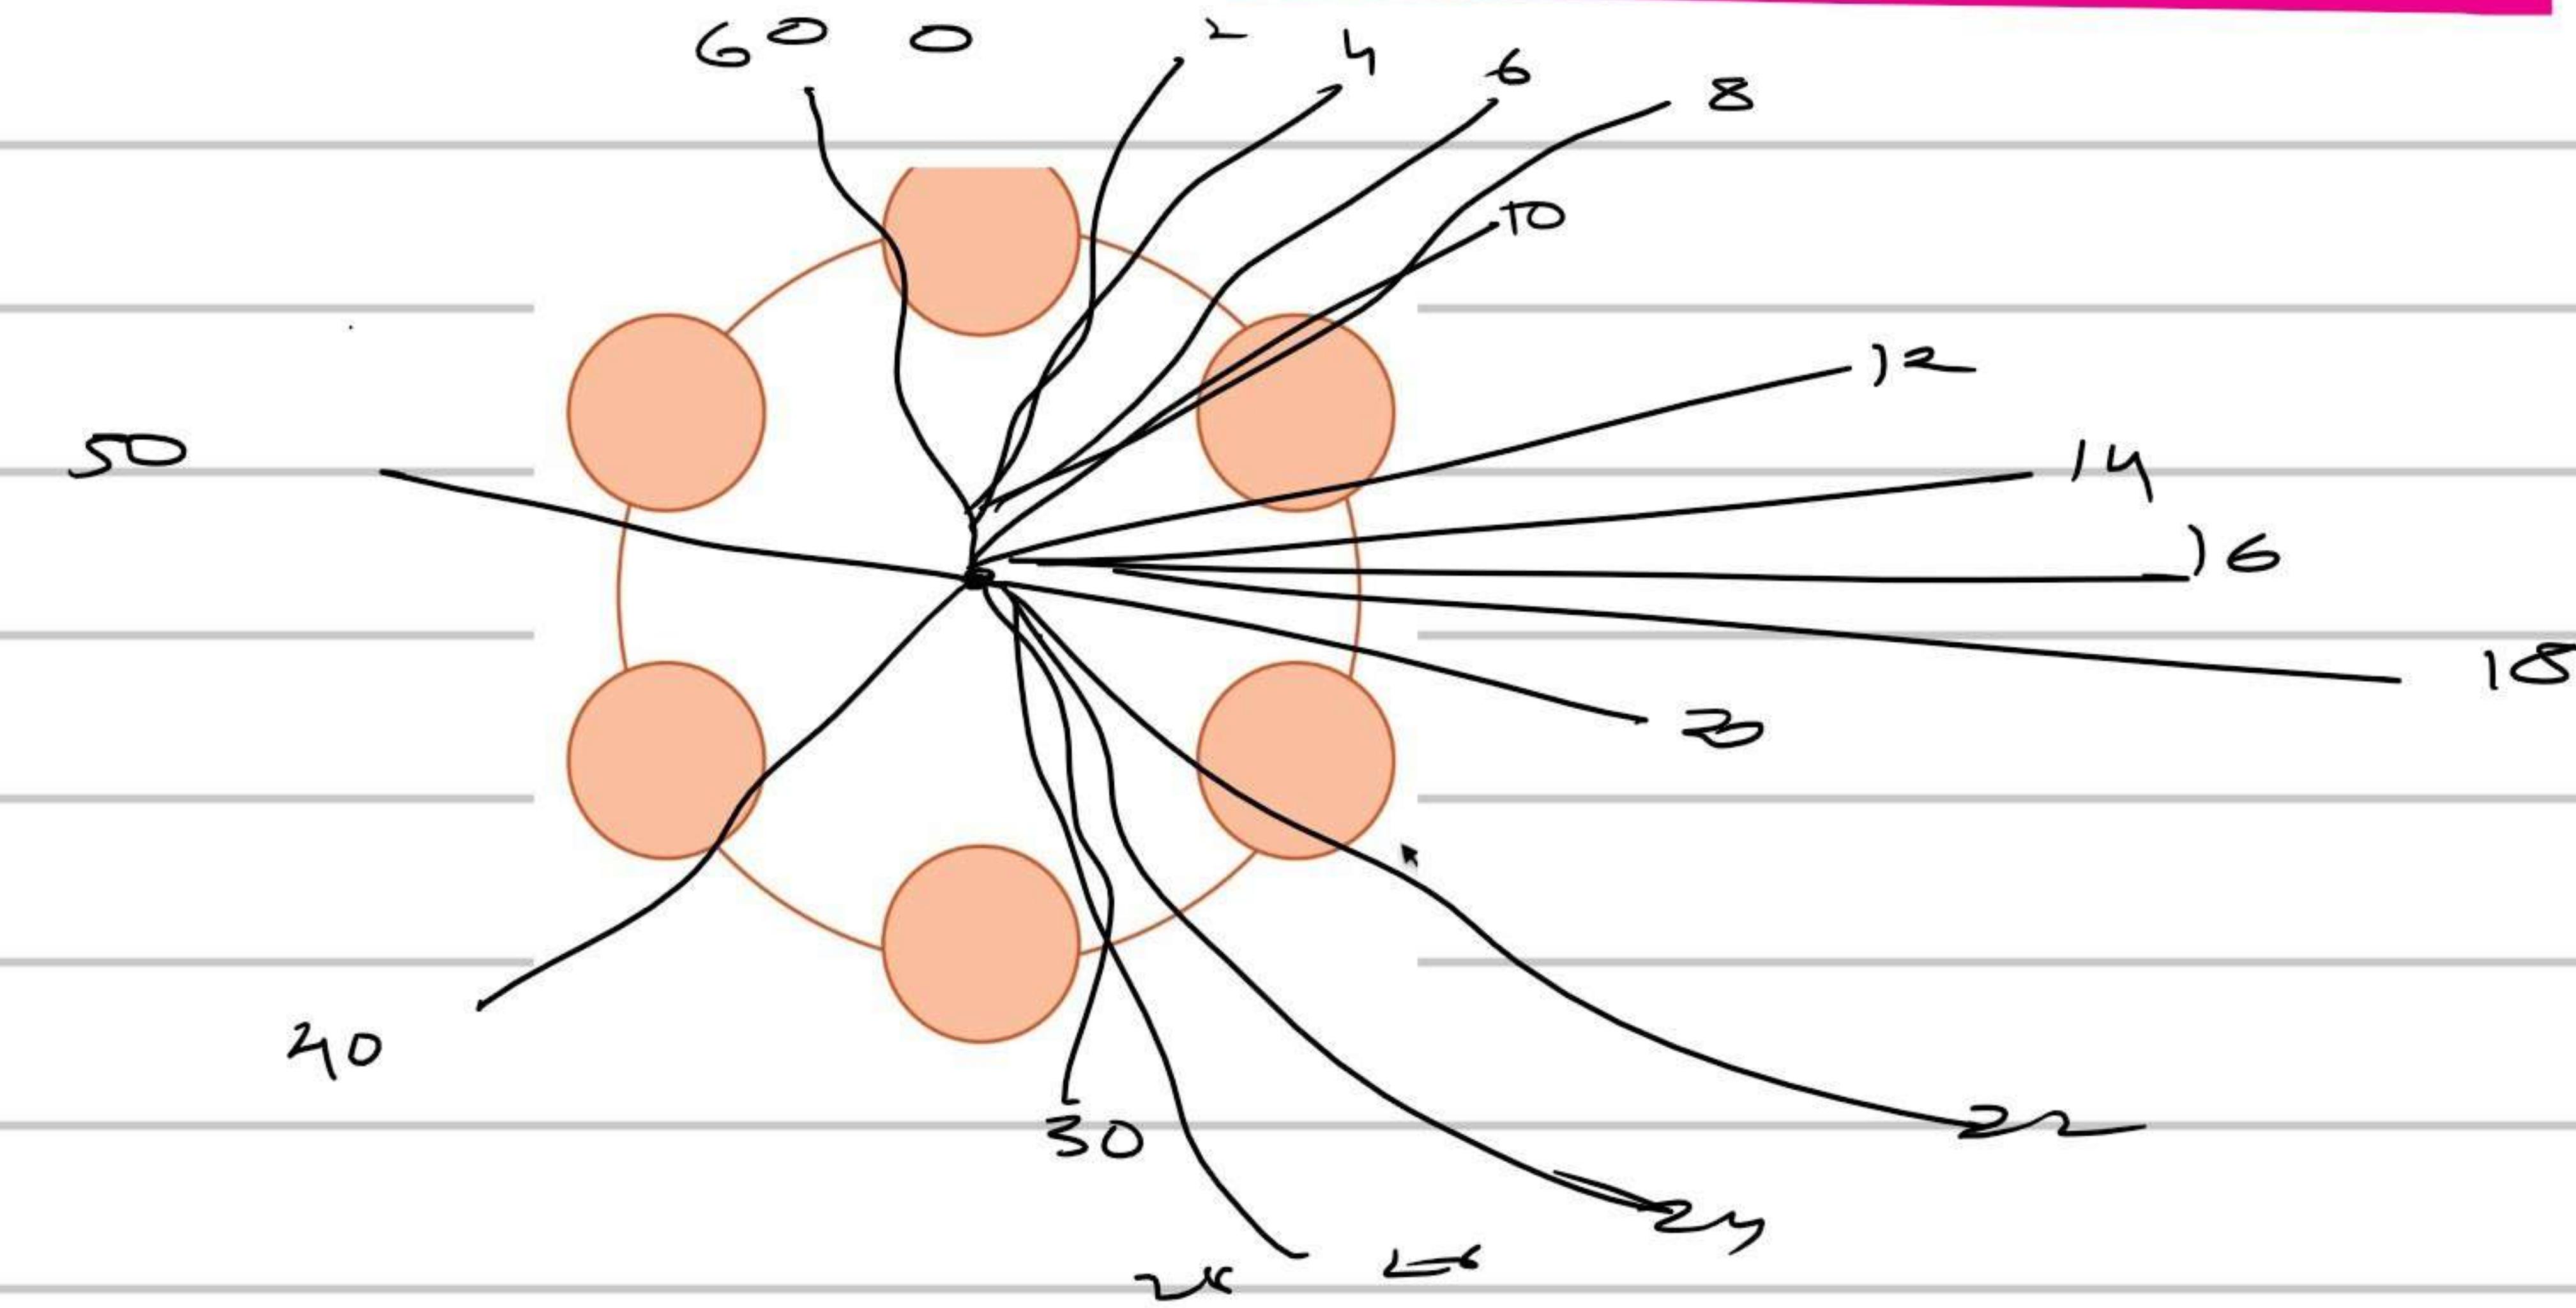
Audi → HF → 45

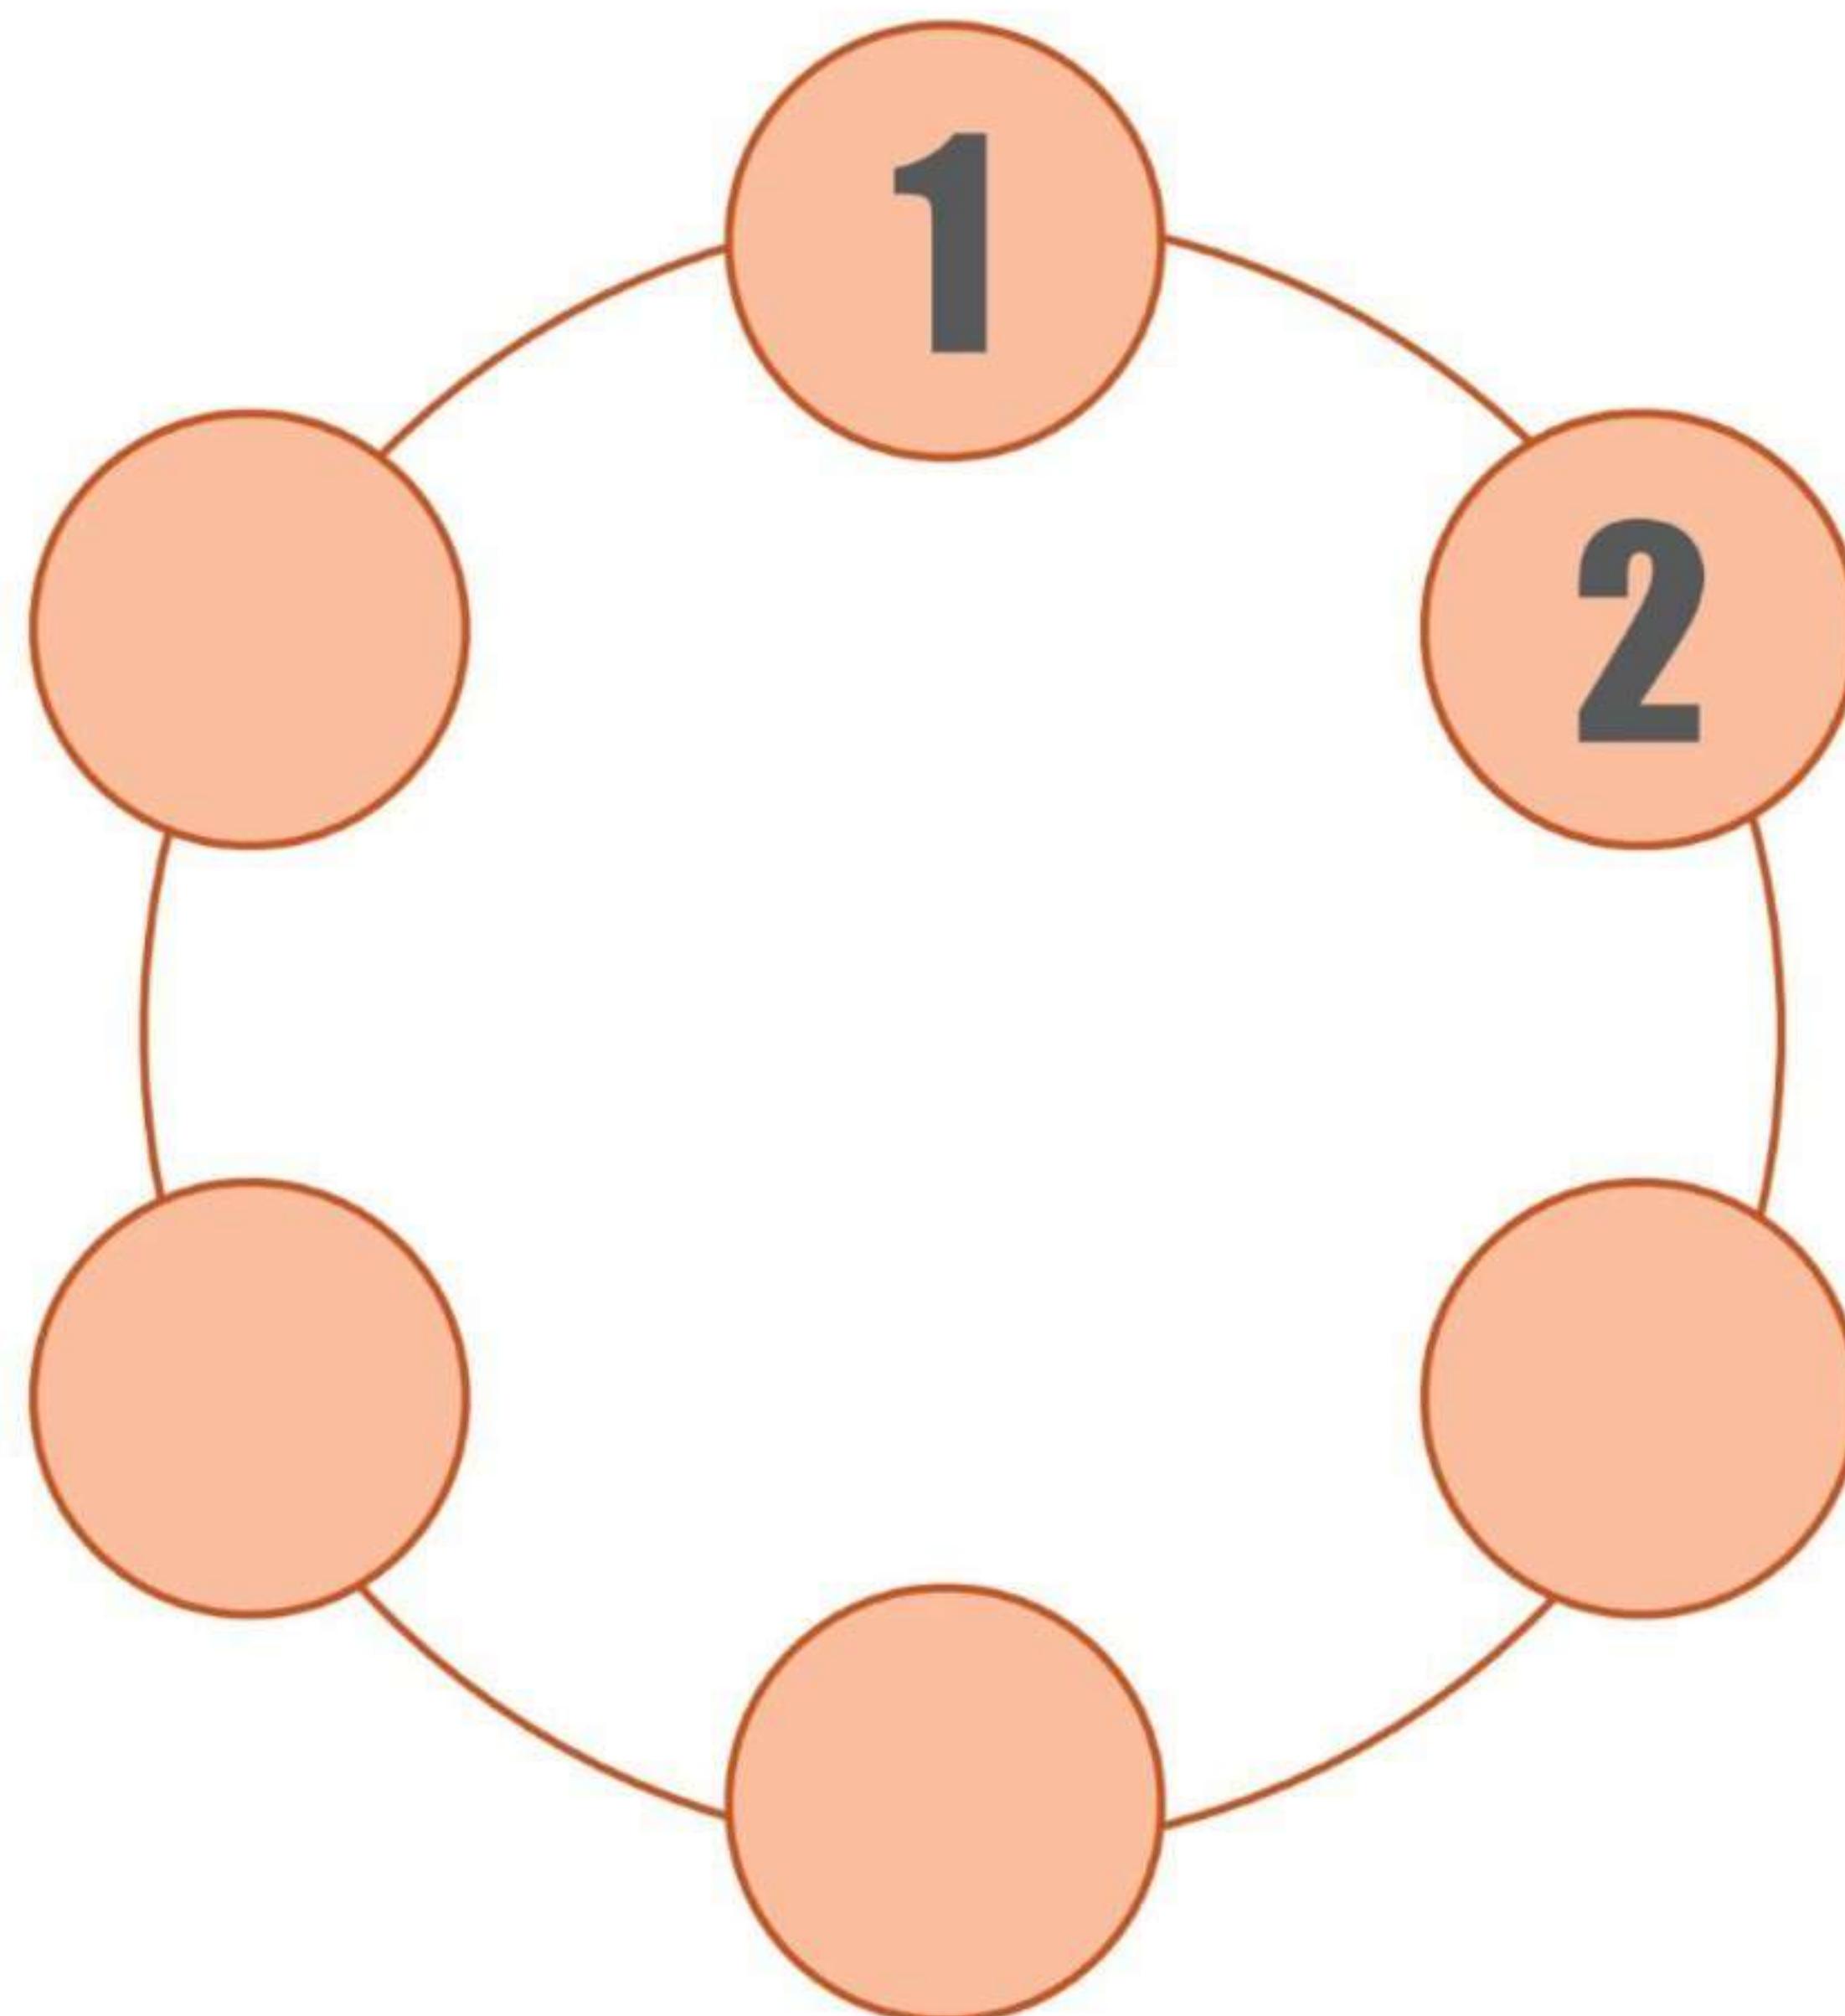
Kia → HF → 39

Merc → HF → 53

1 Node =
 256 Nodes

Virtual Nodes



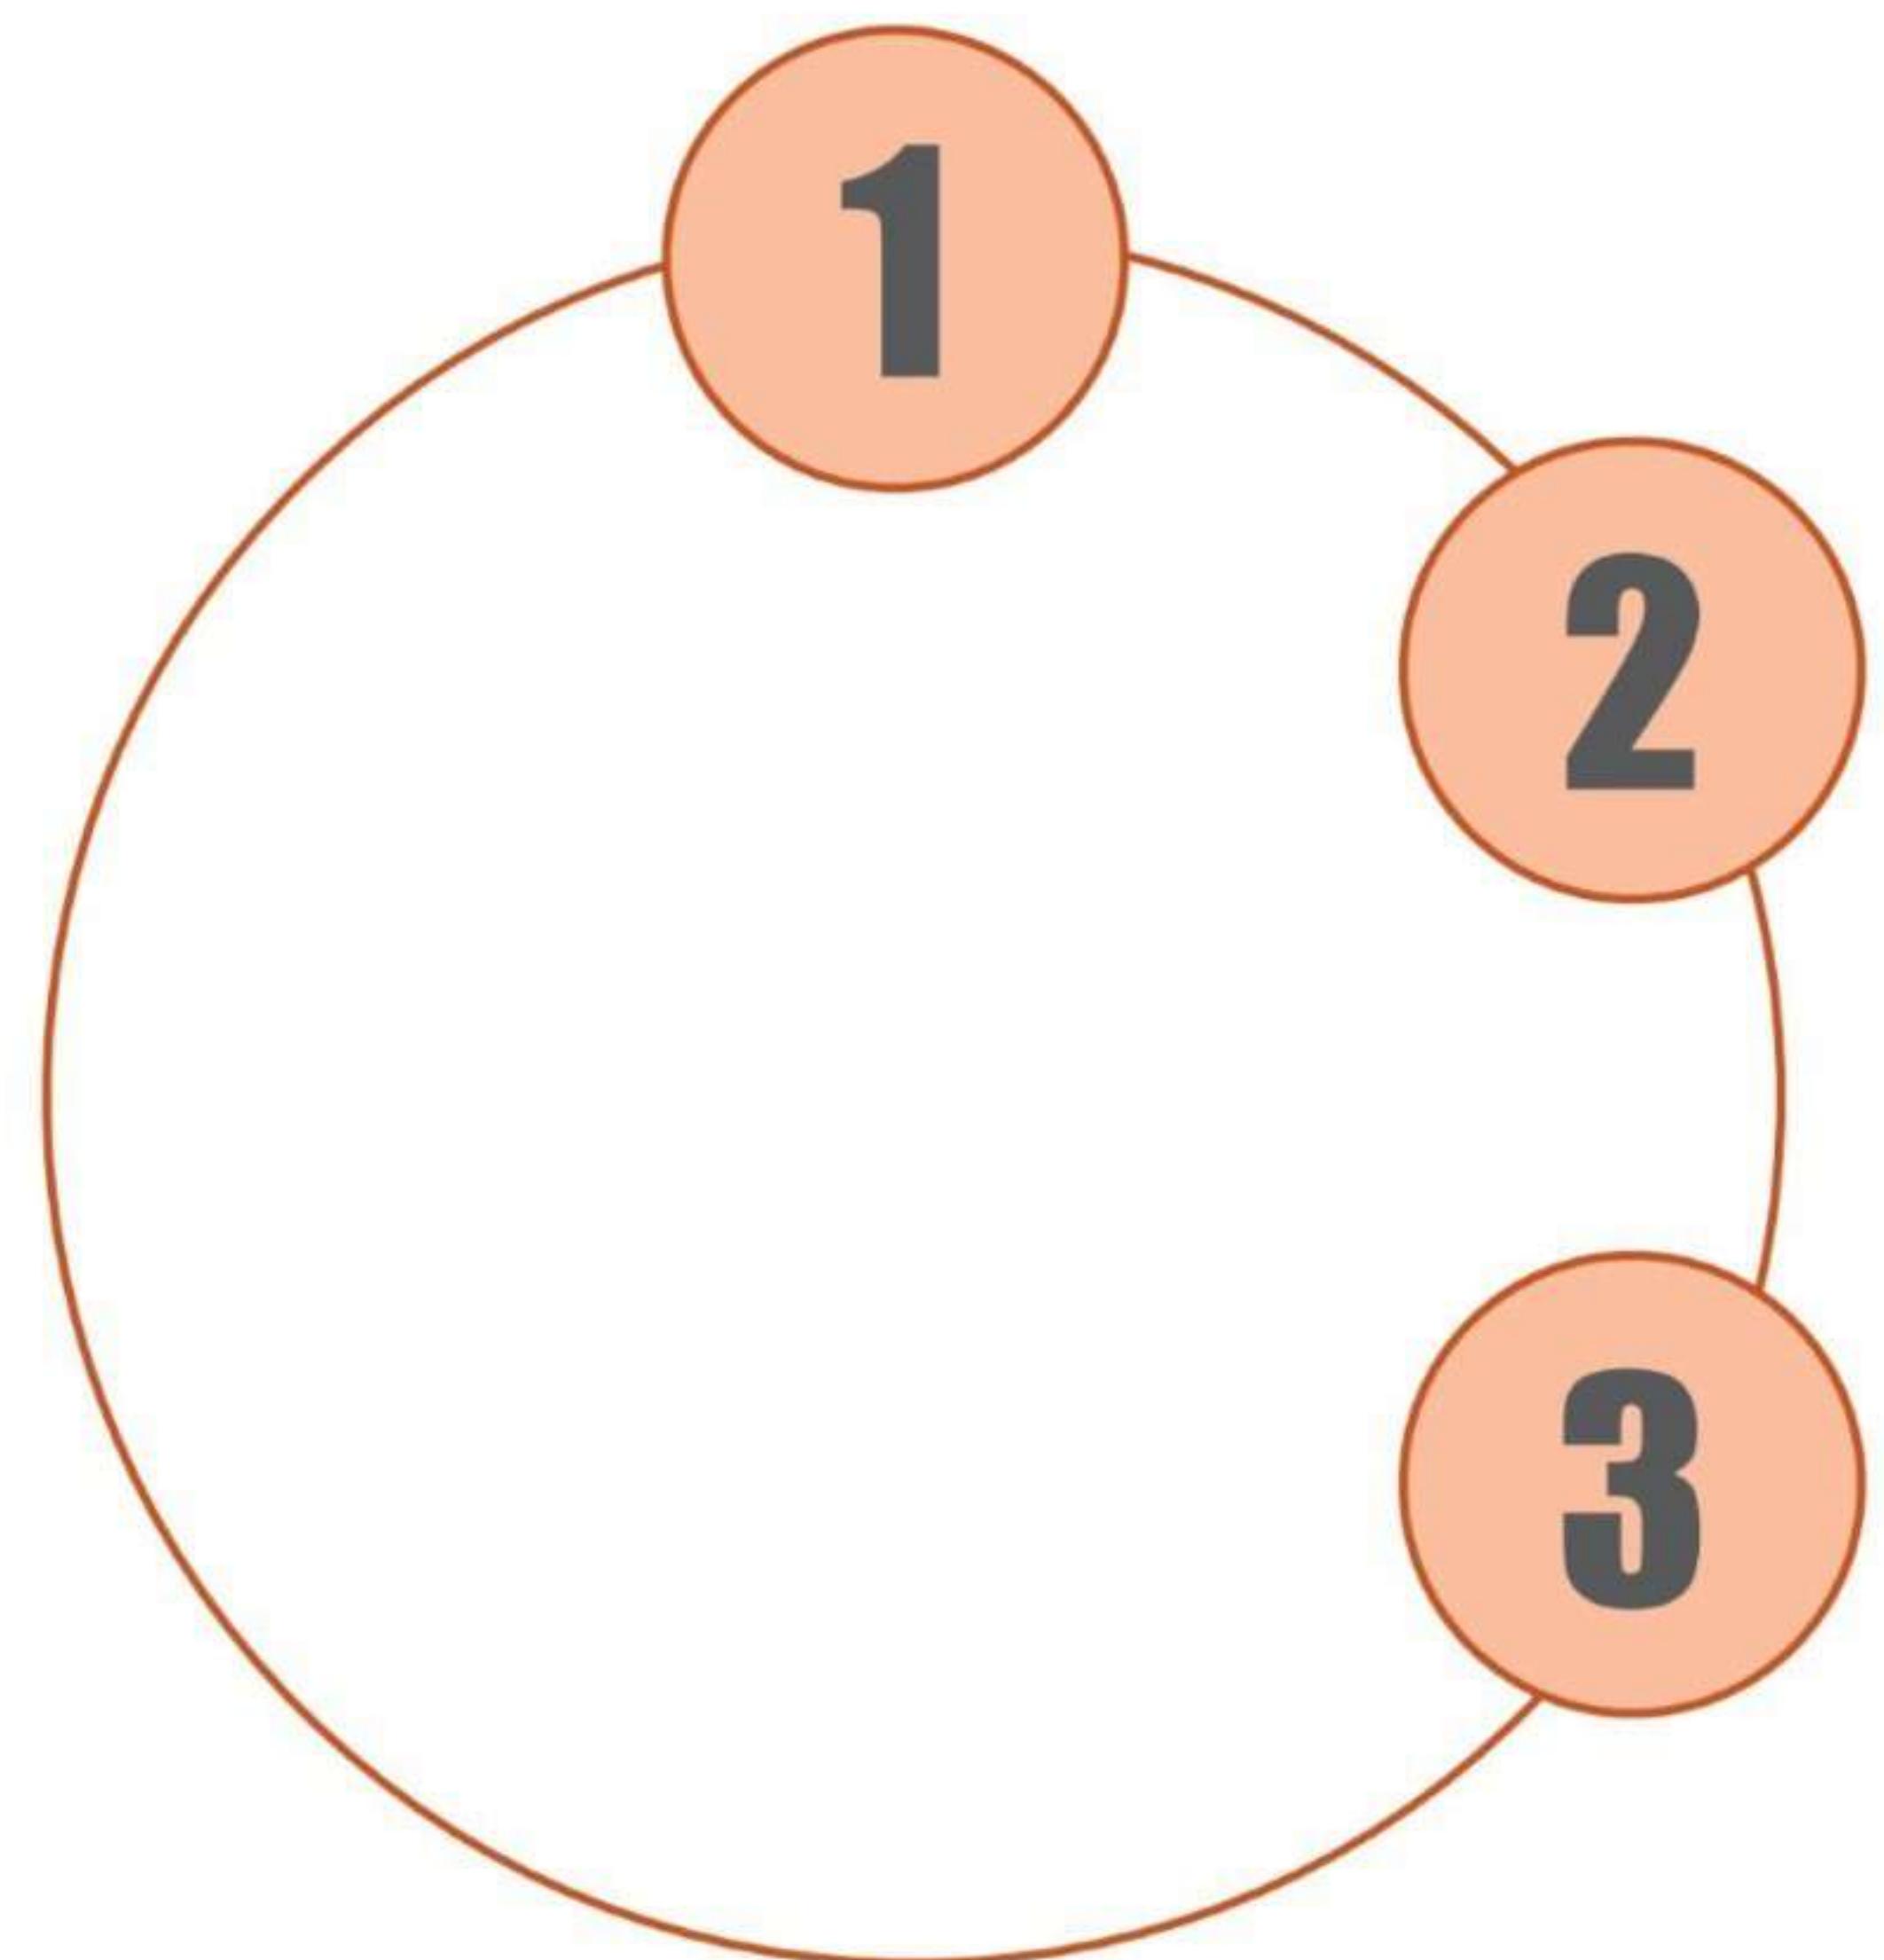


Storing Data in Cassandra

Range of all possible token values

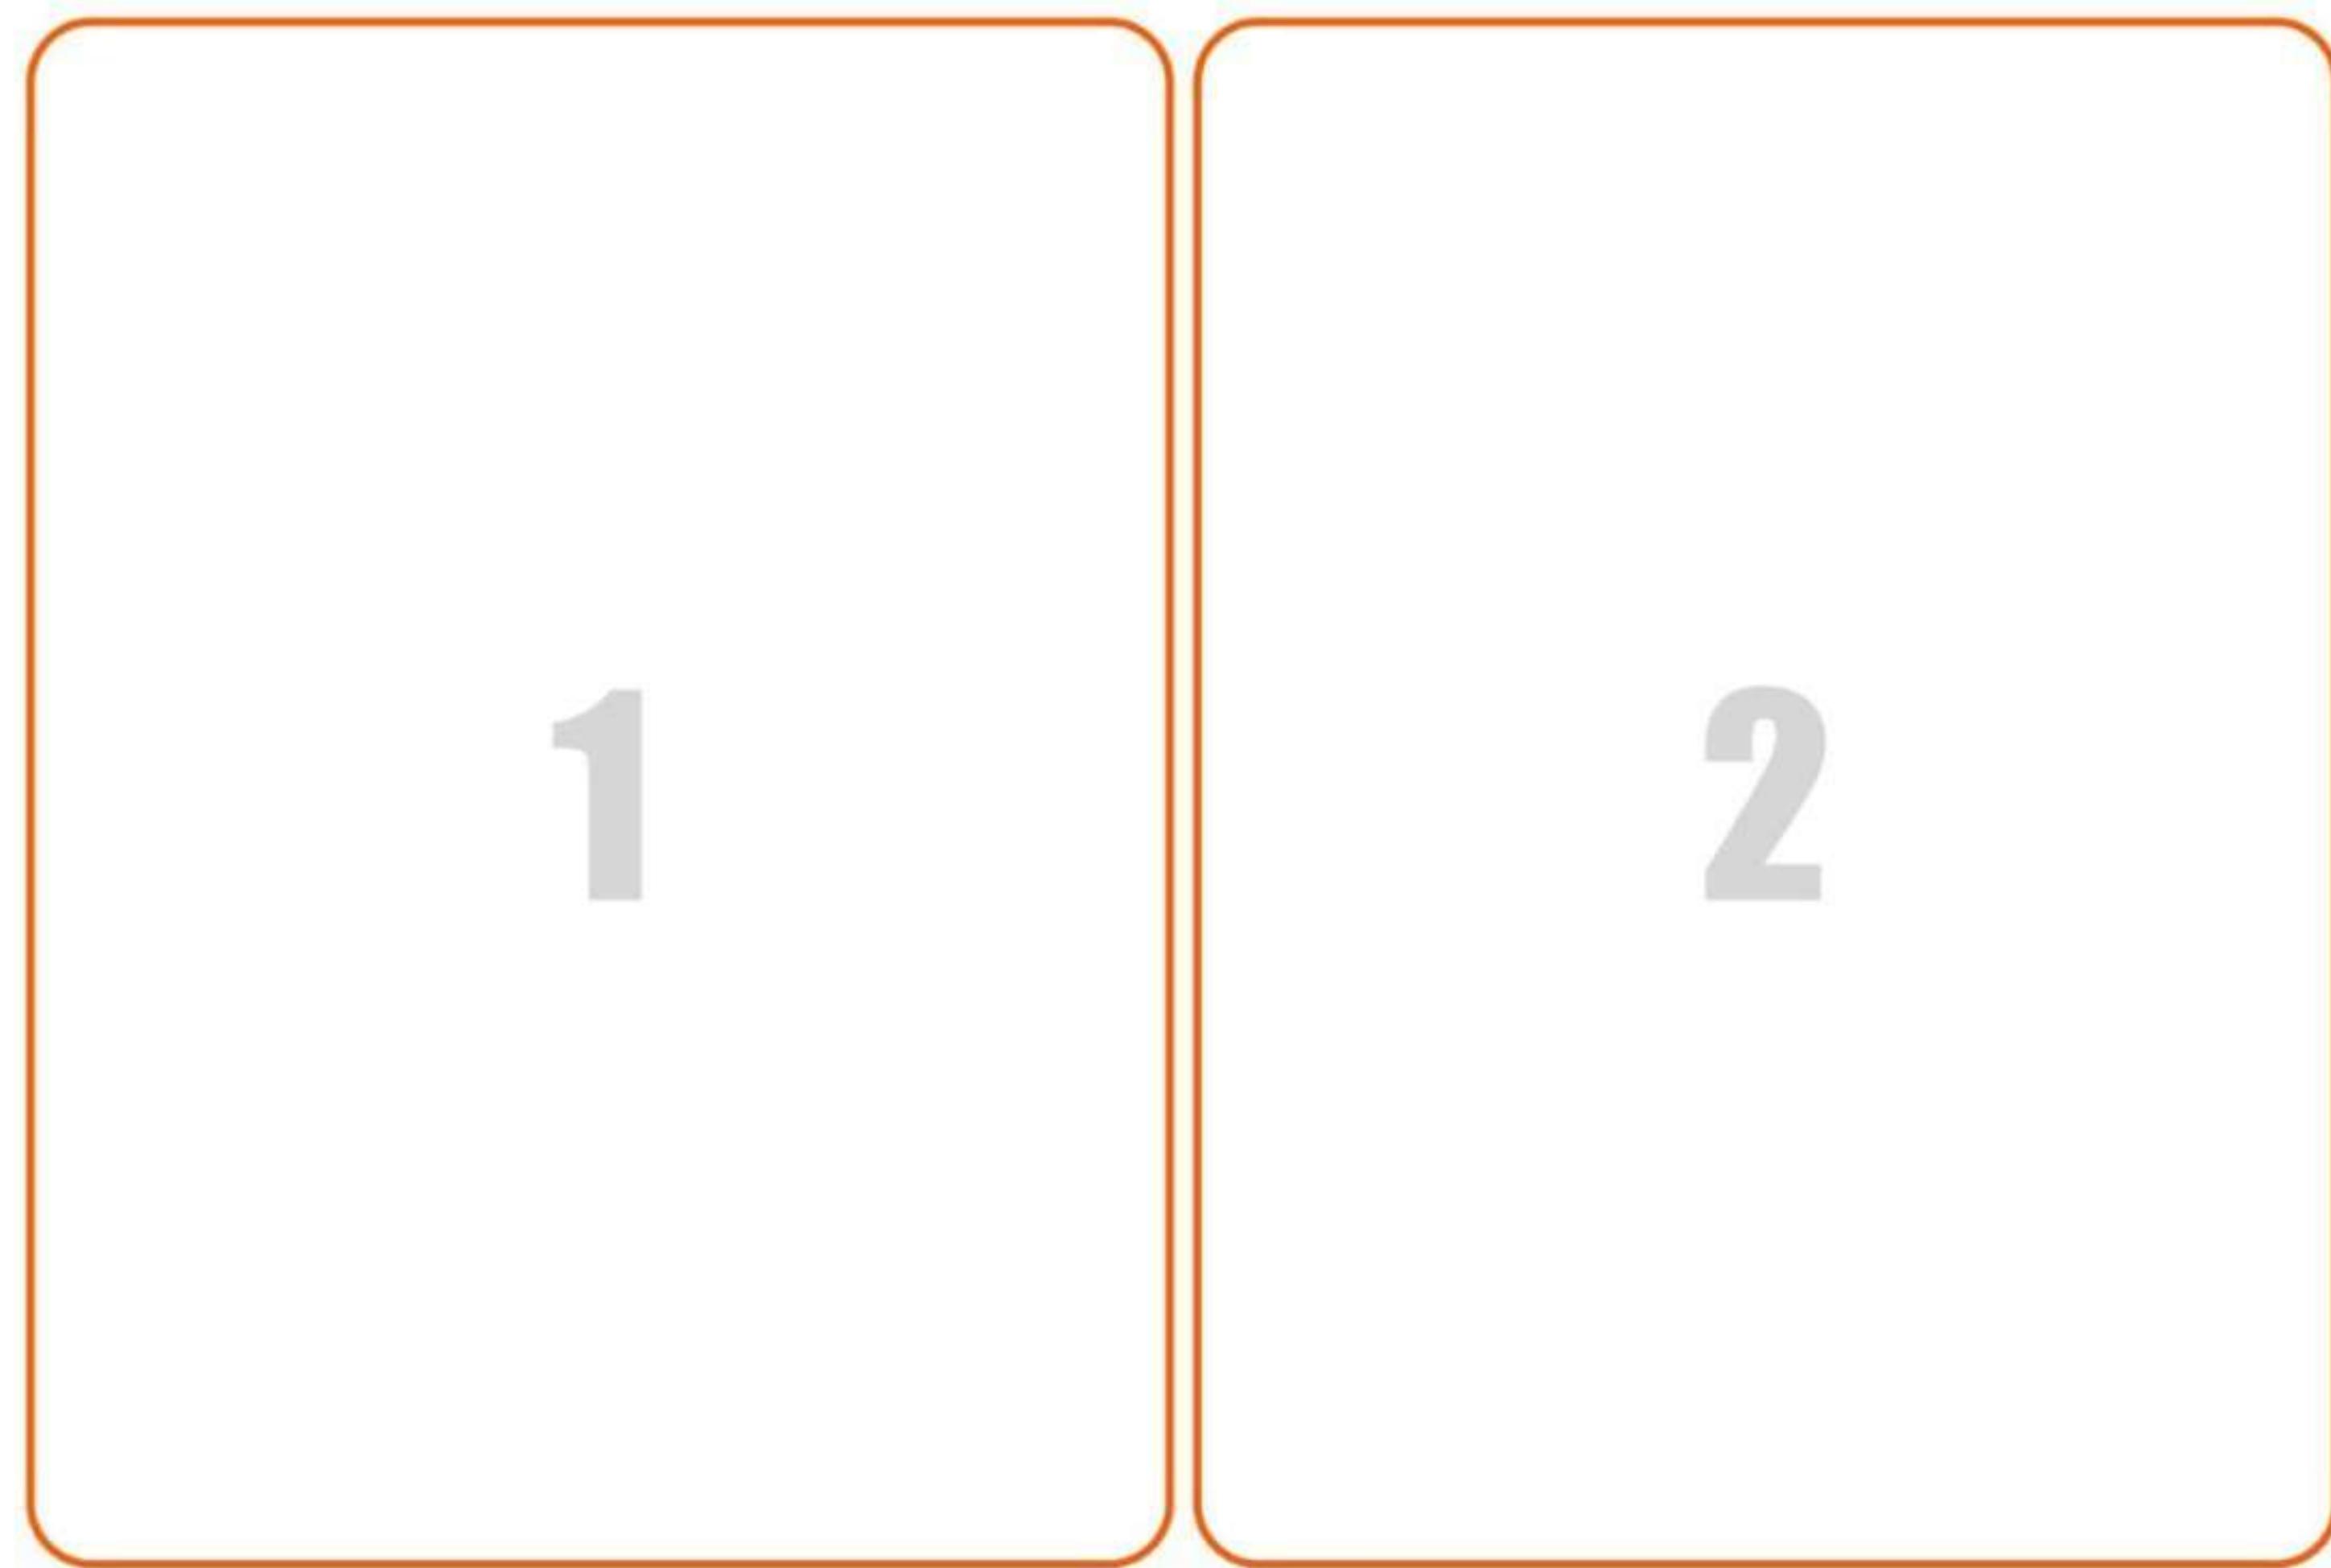
1

-2^{63} to $+2^{63}$

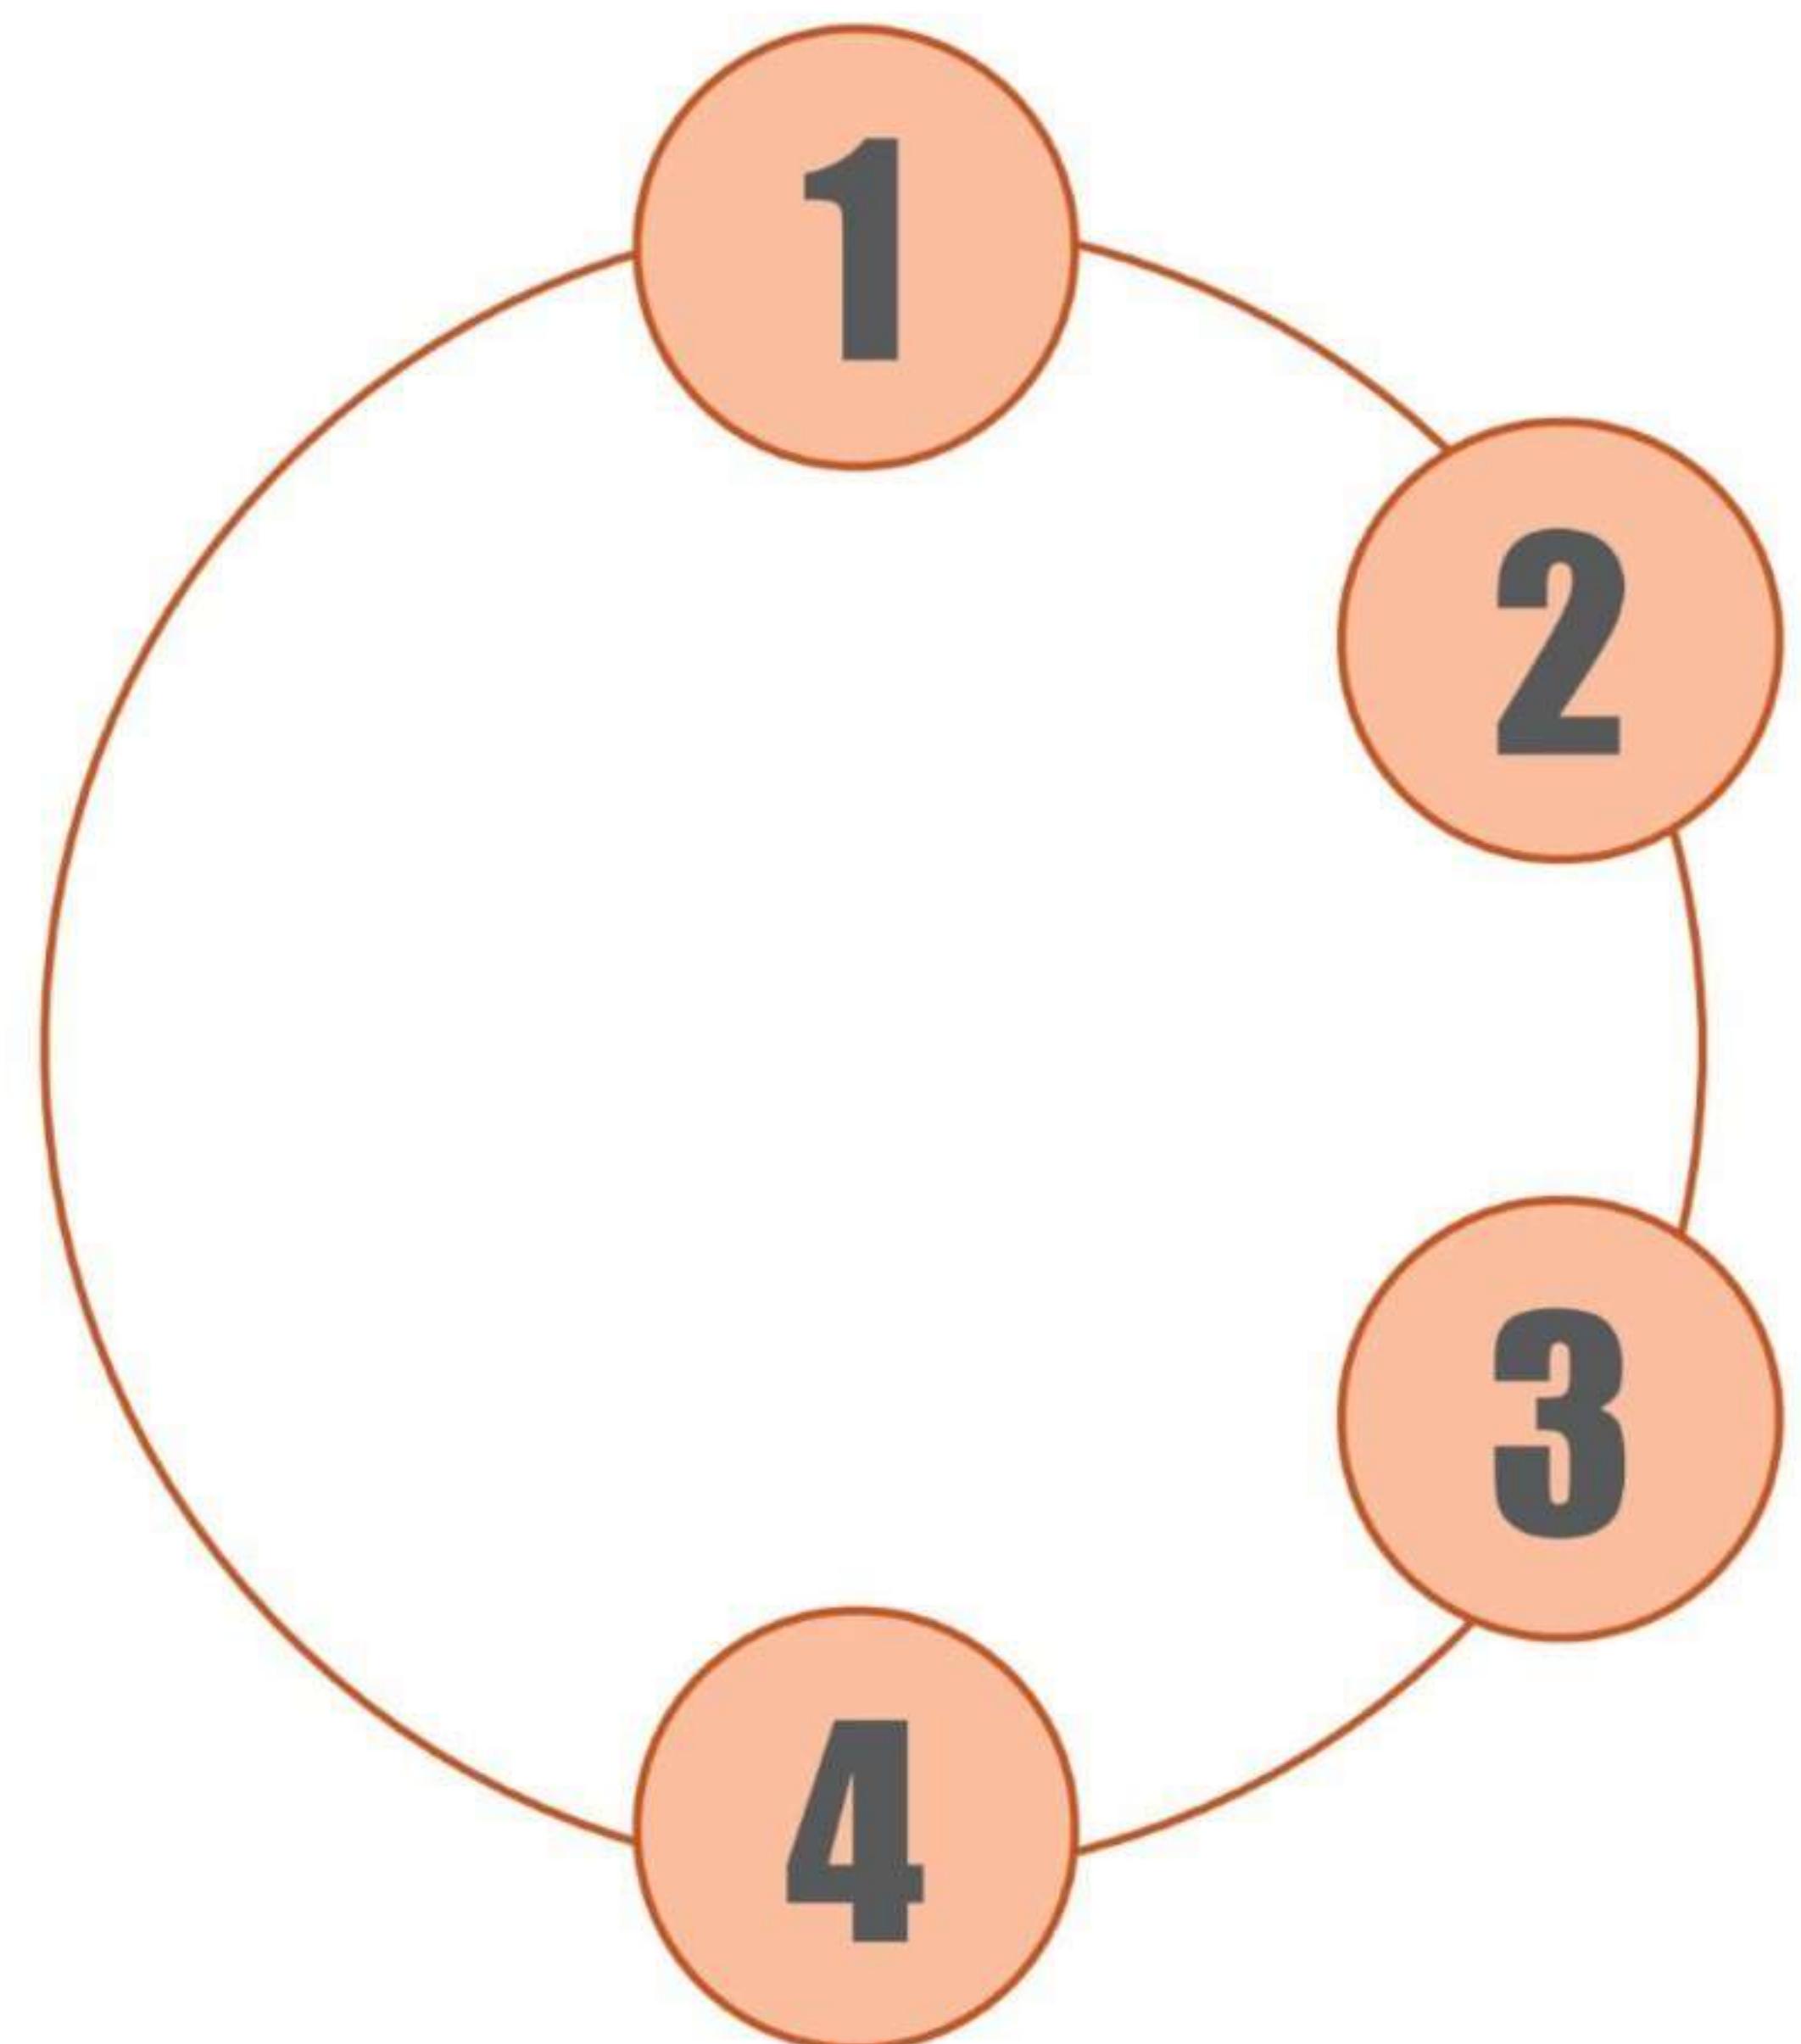


Storing Data in Cassandra

Range of all possible token values

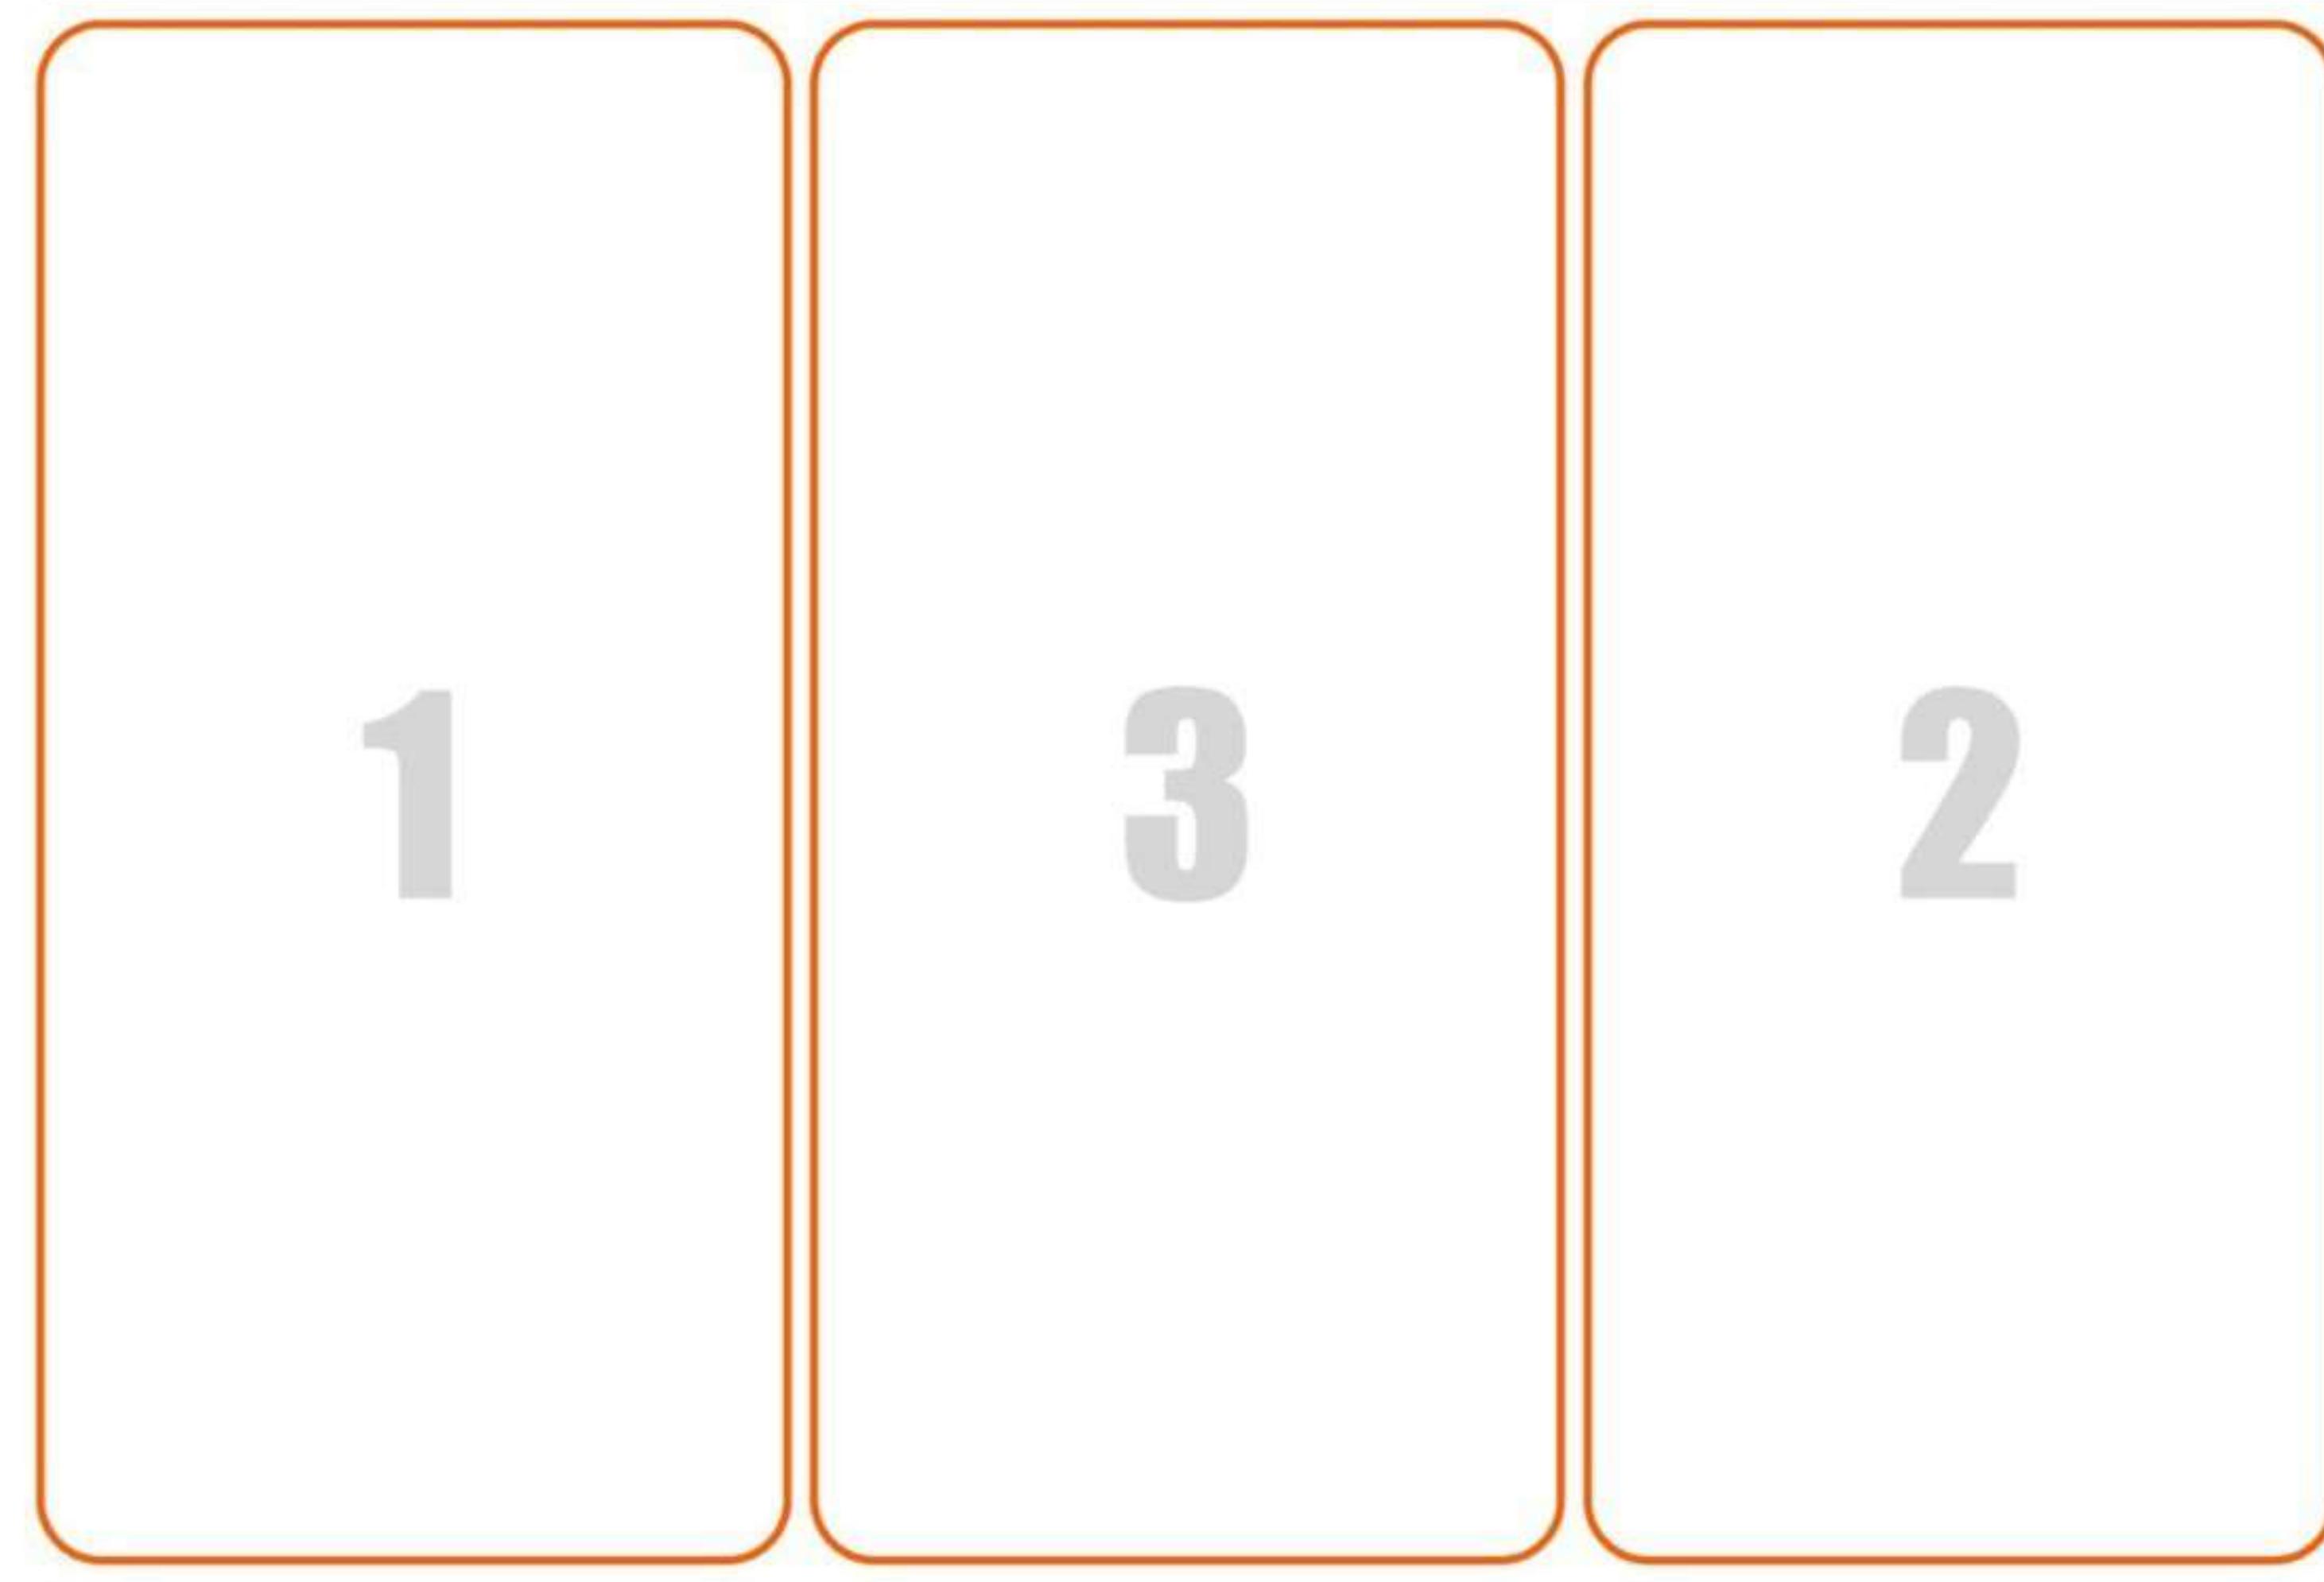


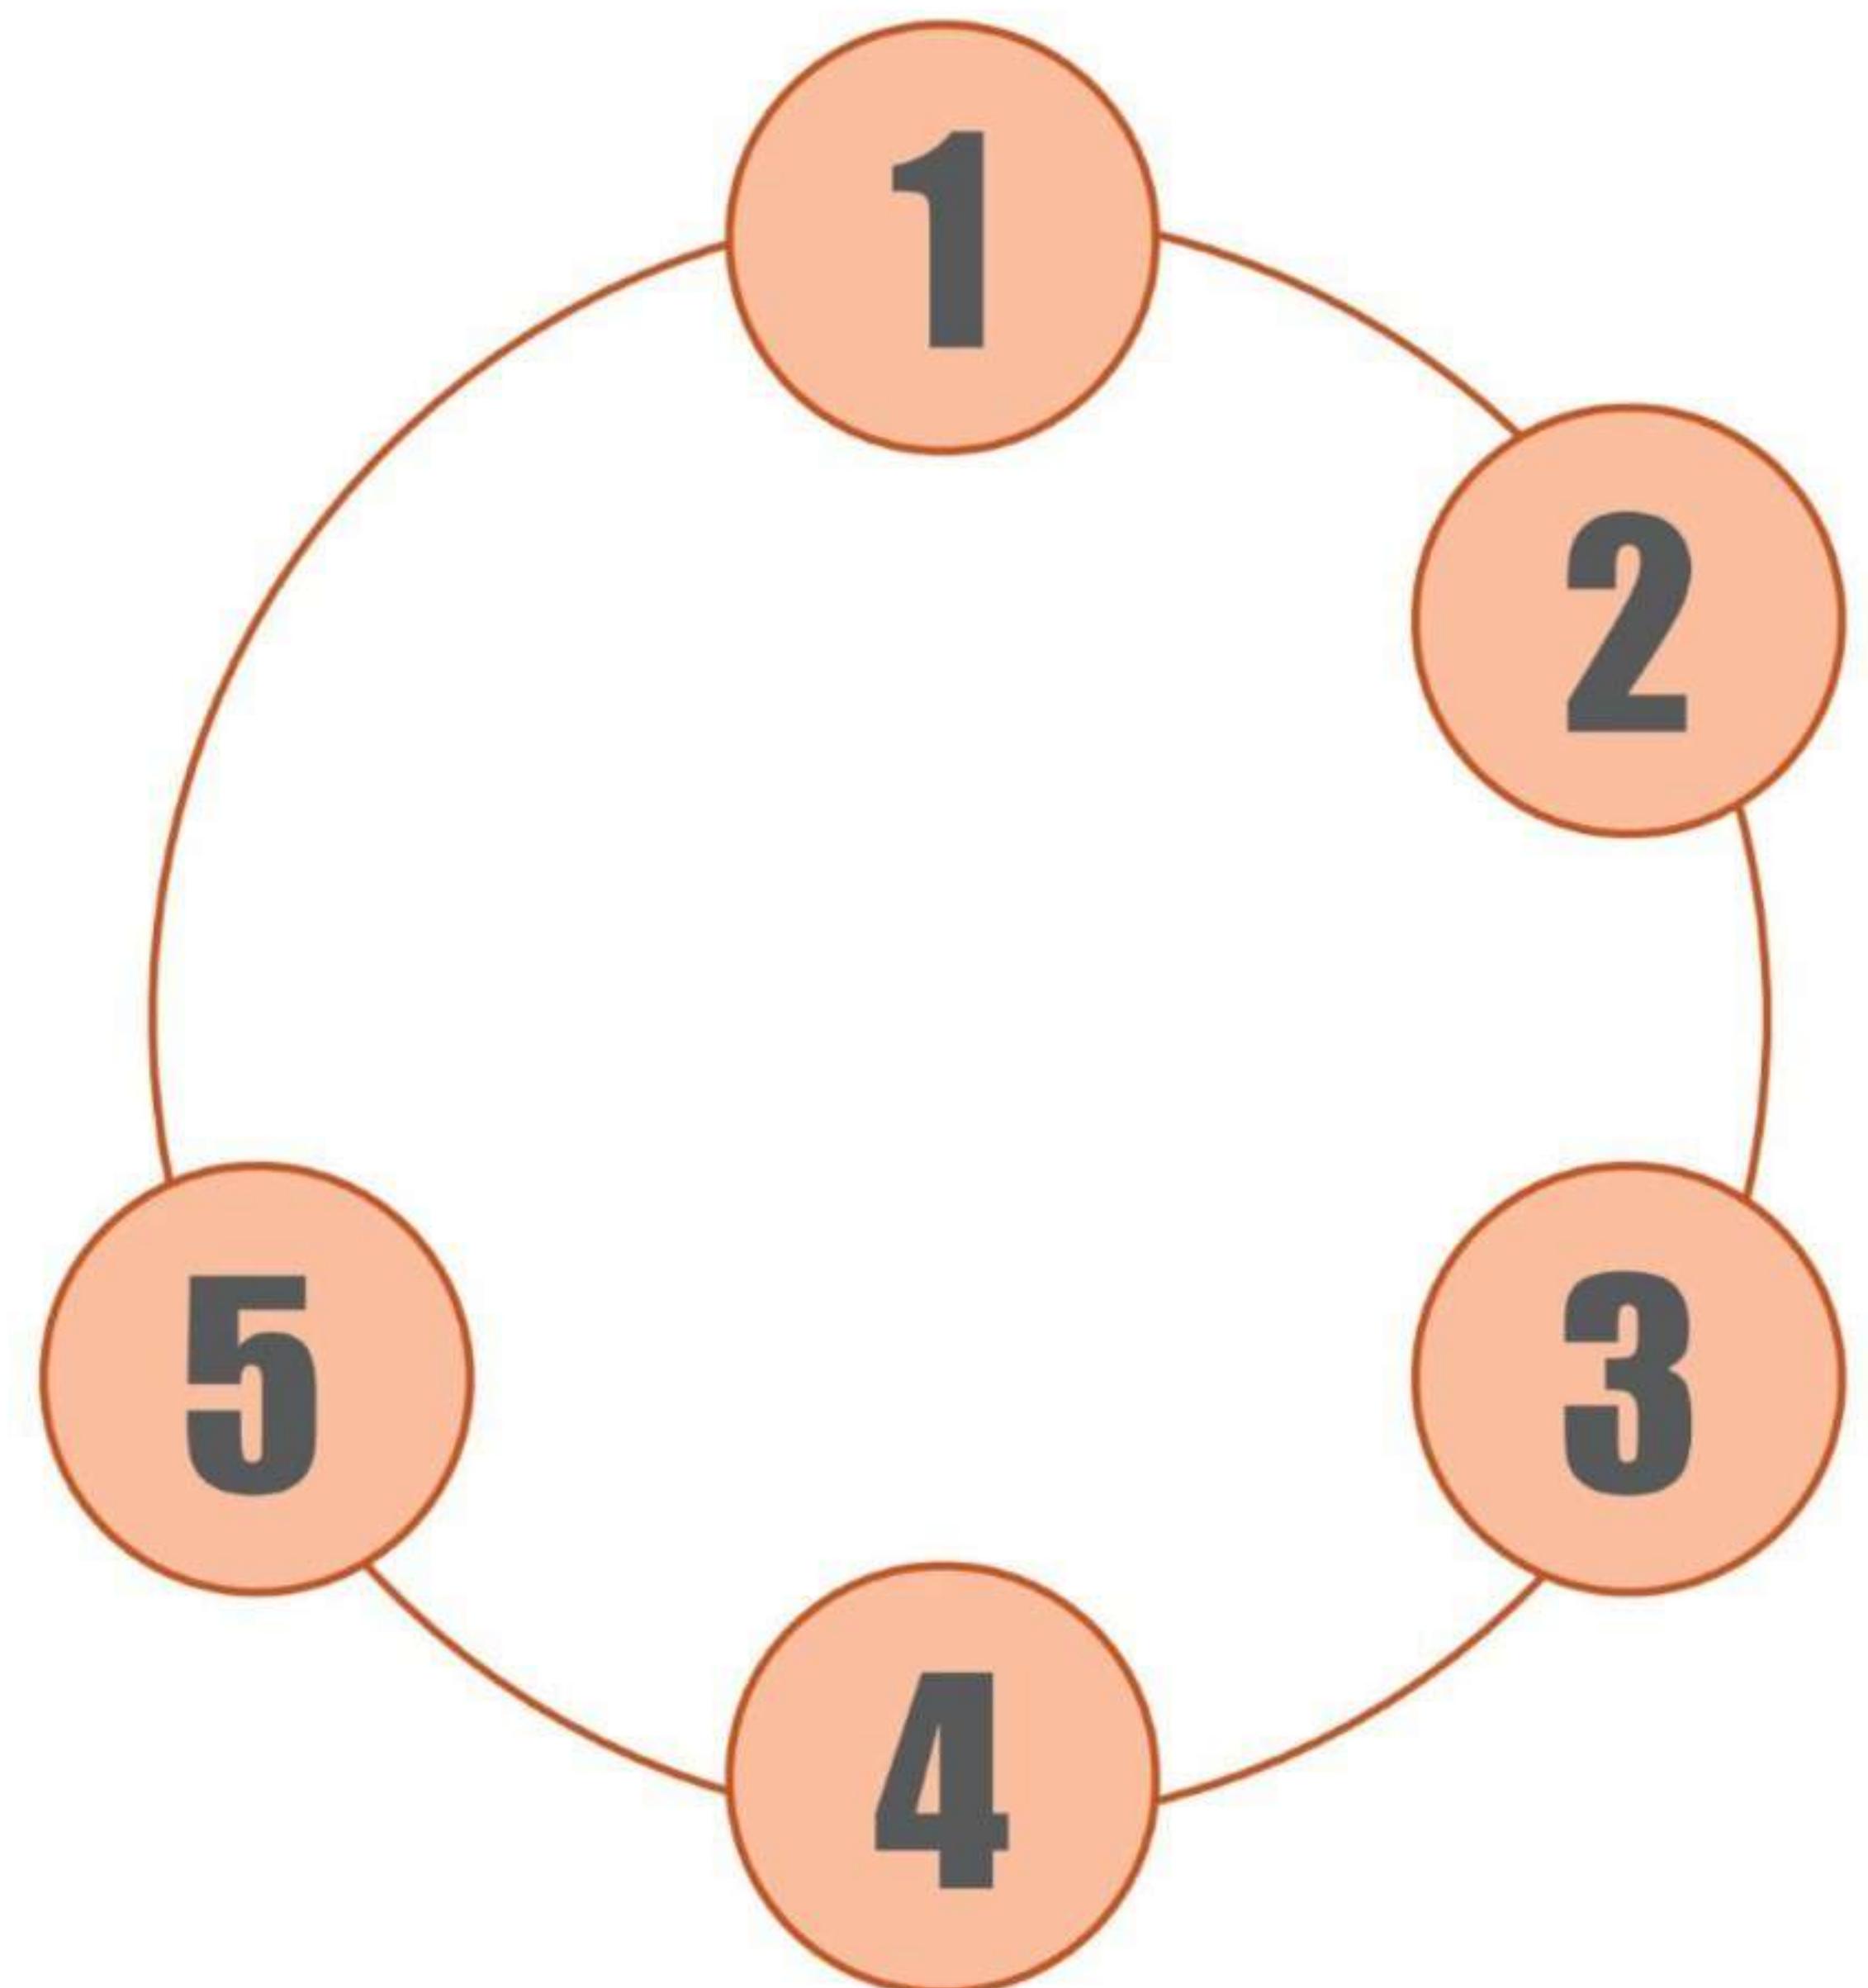
-2^{63} to $+2^{63}$



Storing Data in Cassandra

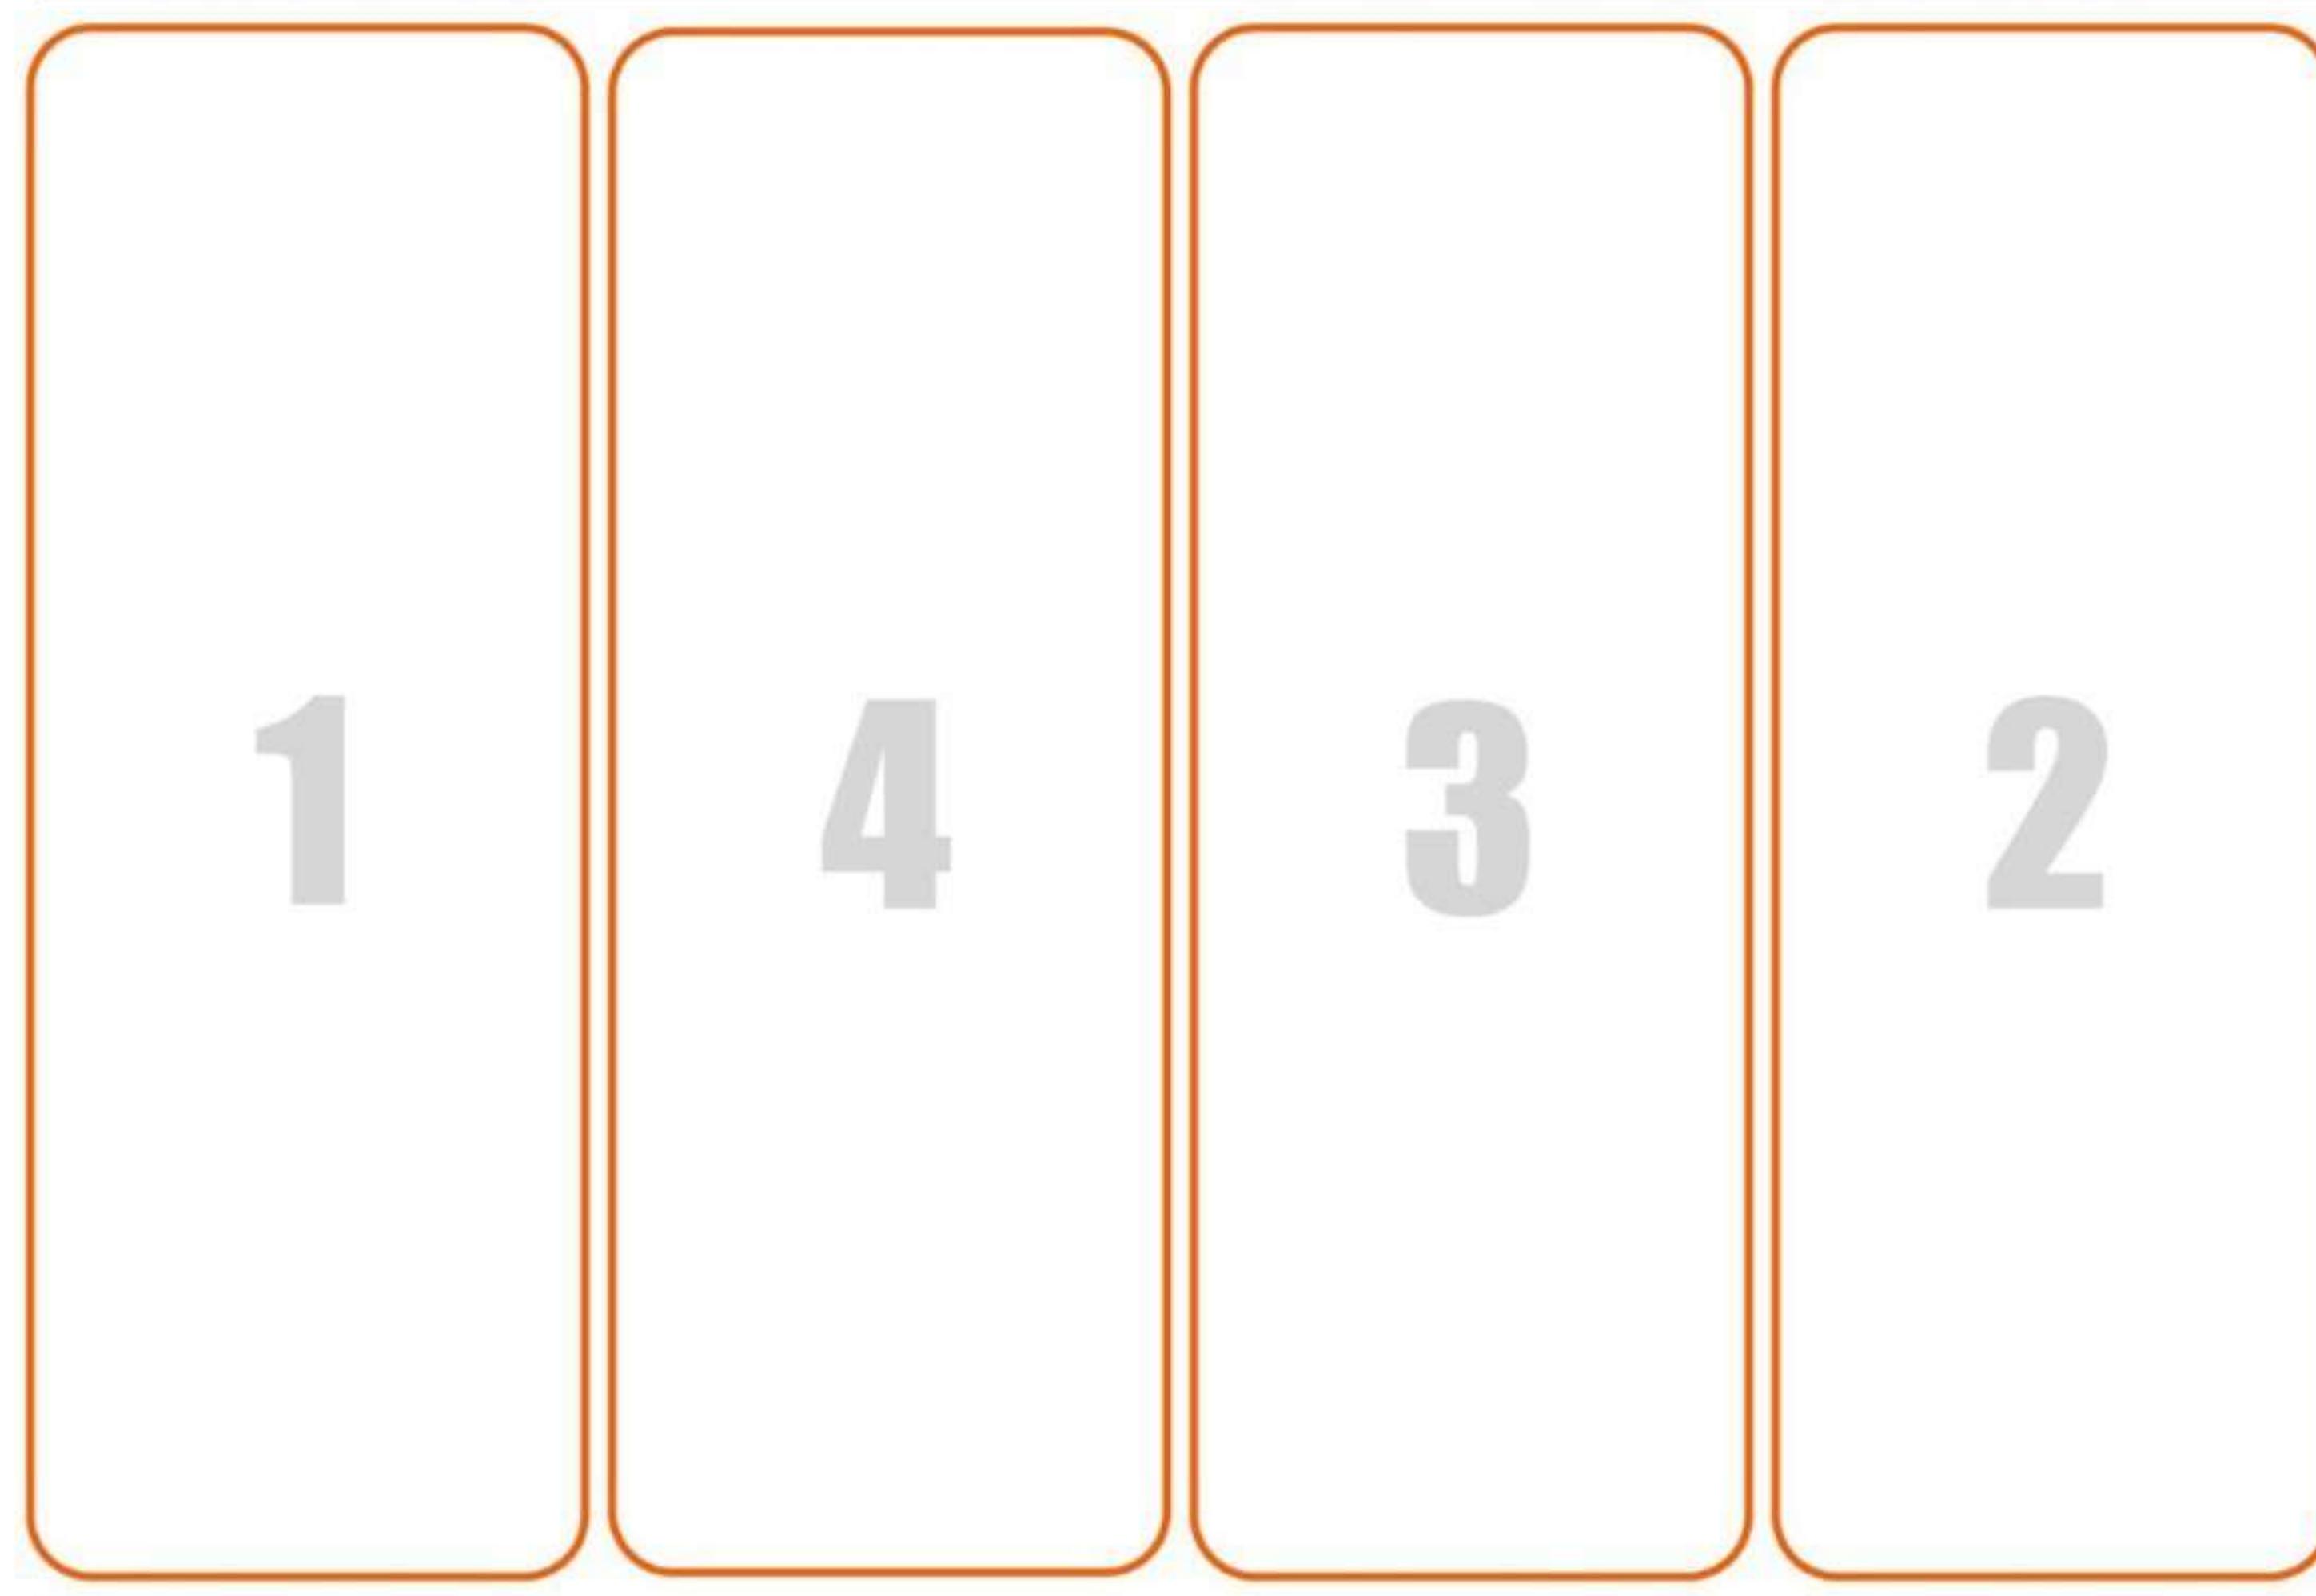
Range of all possible token values



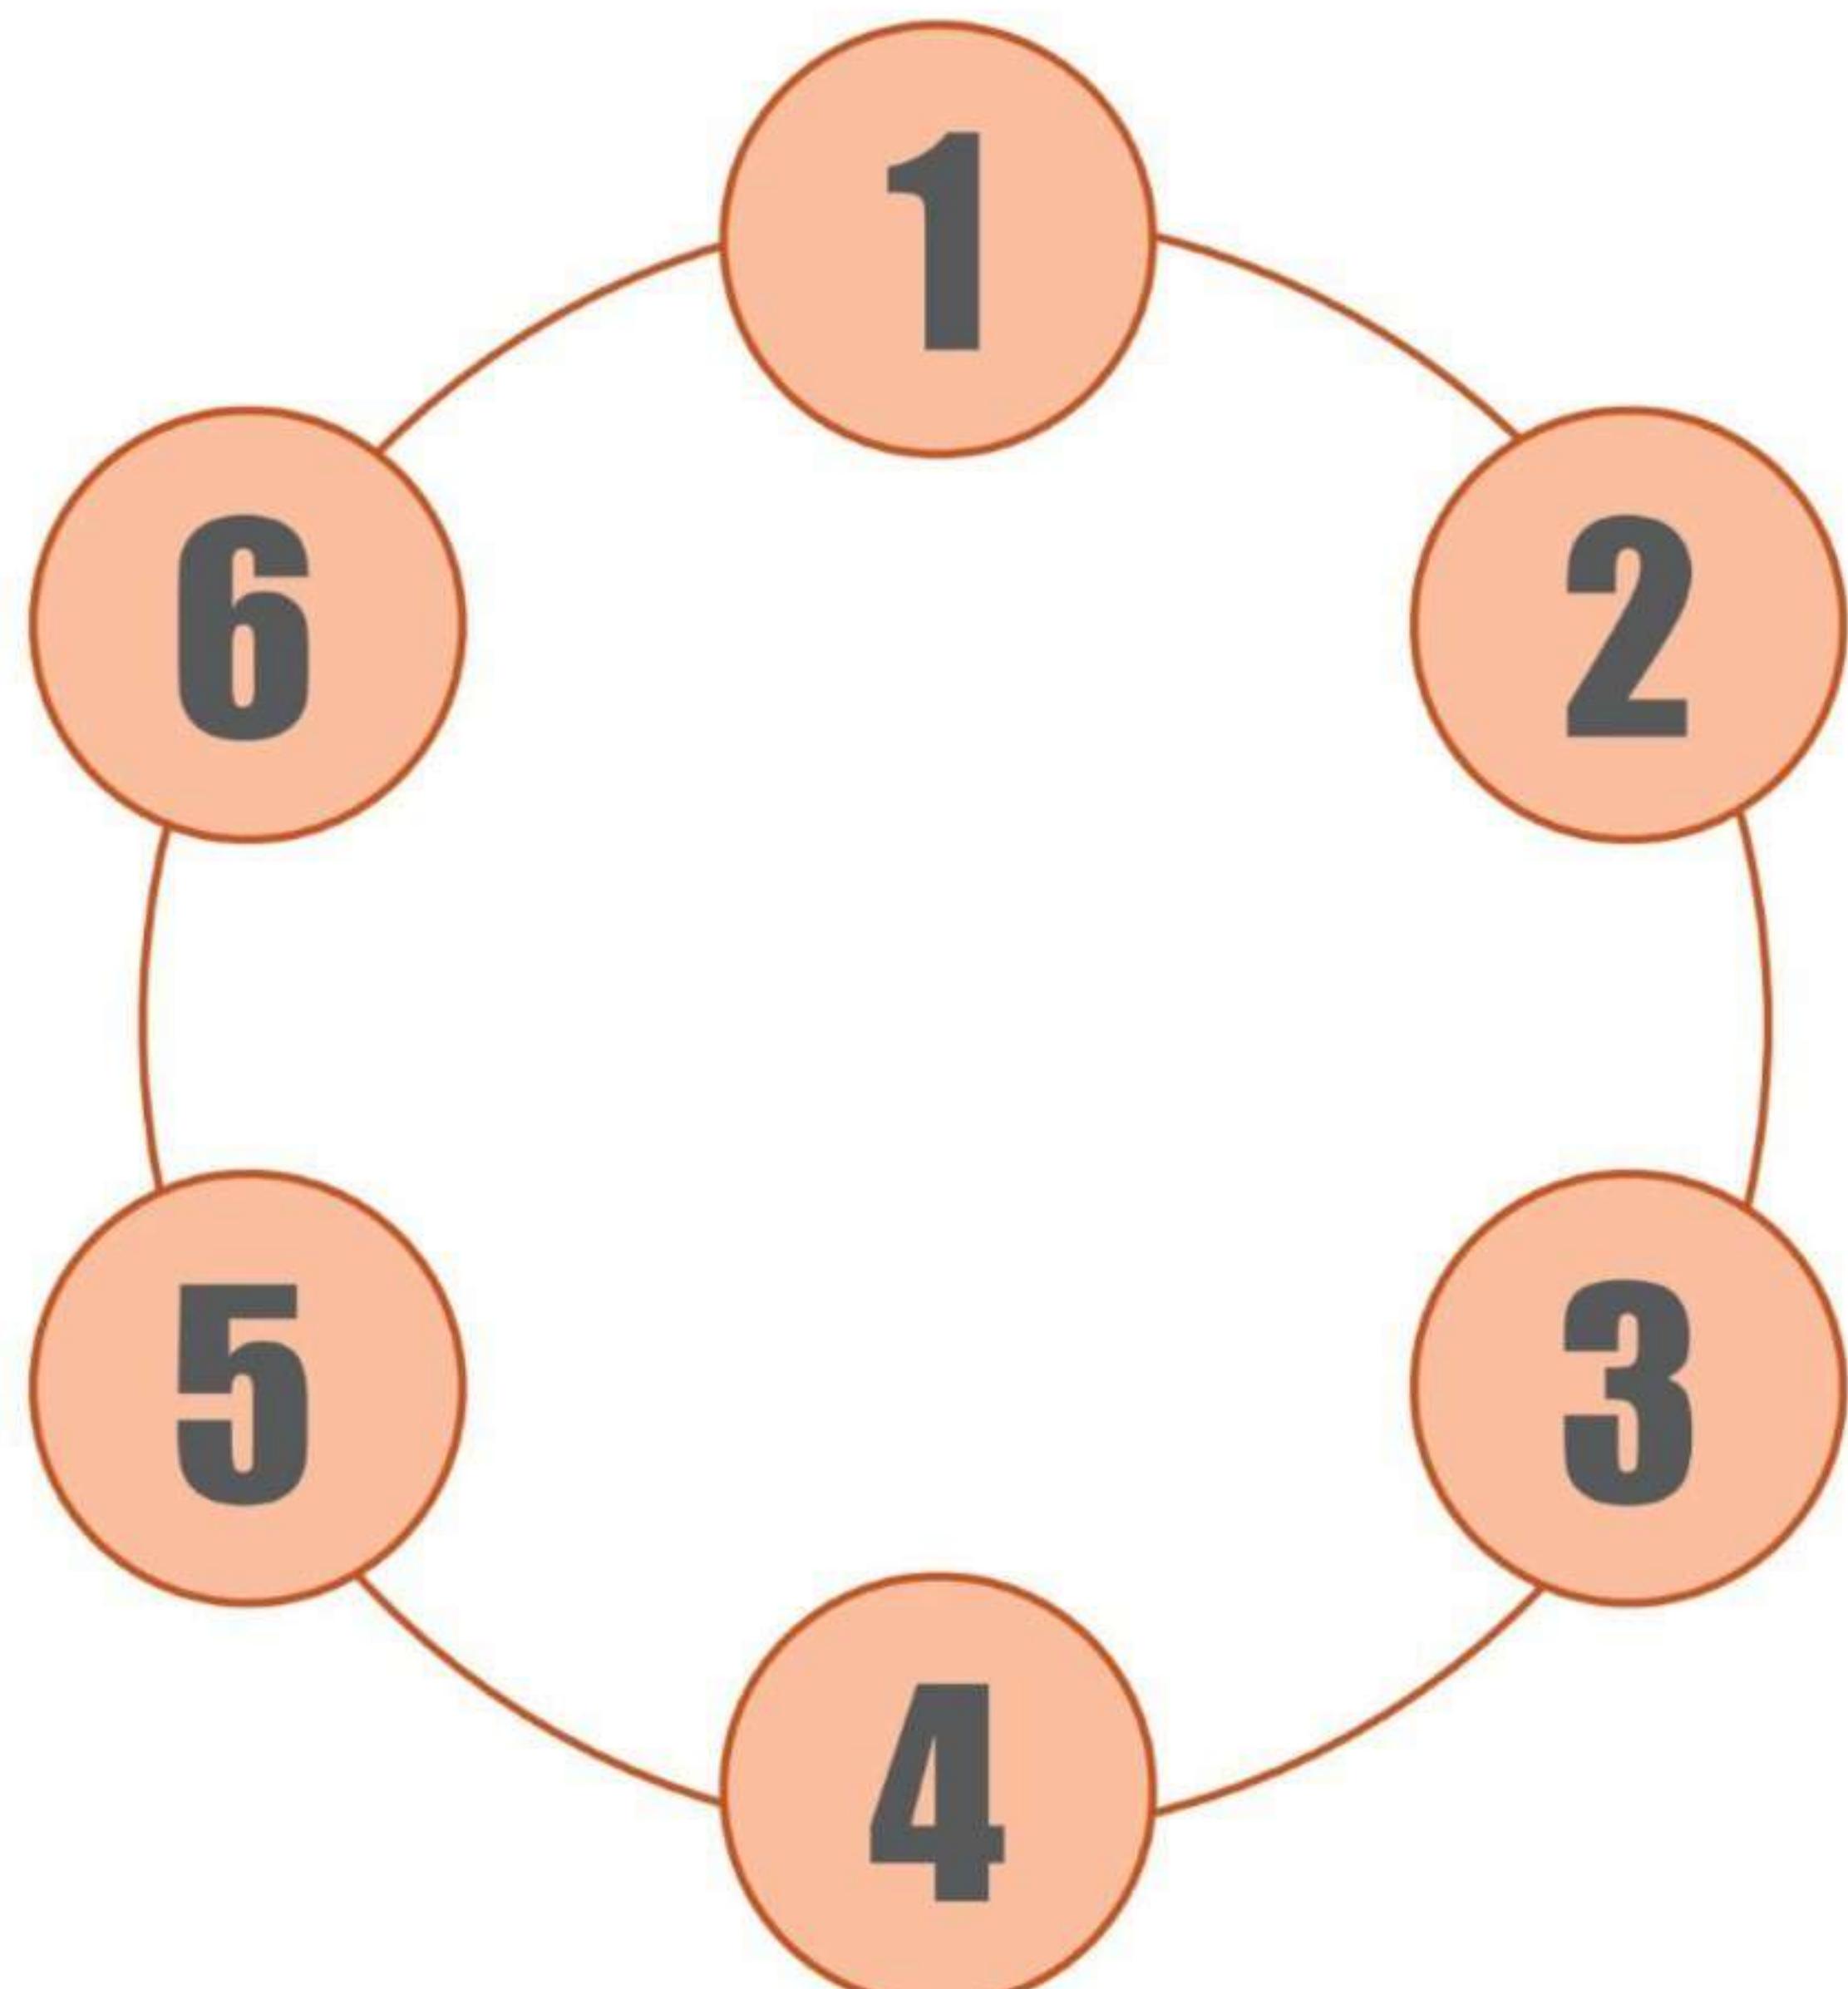


Storing Data in Cassandra

Range of all possible token values

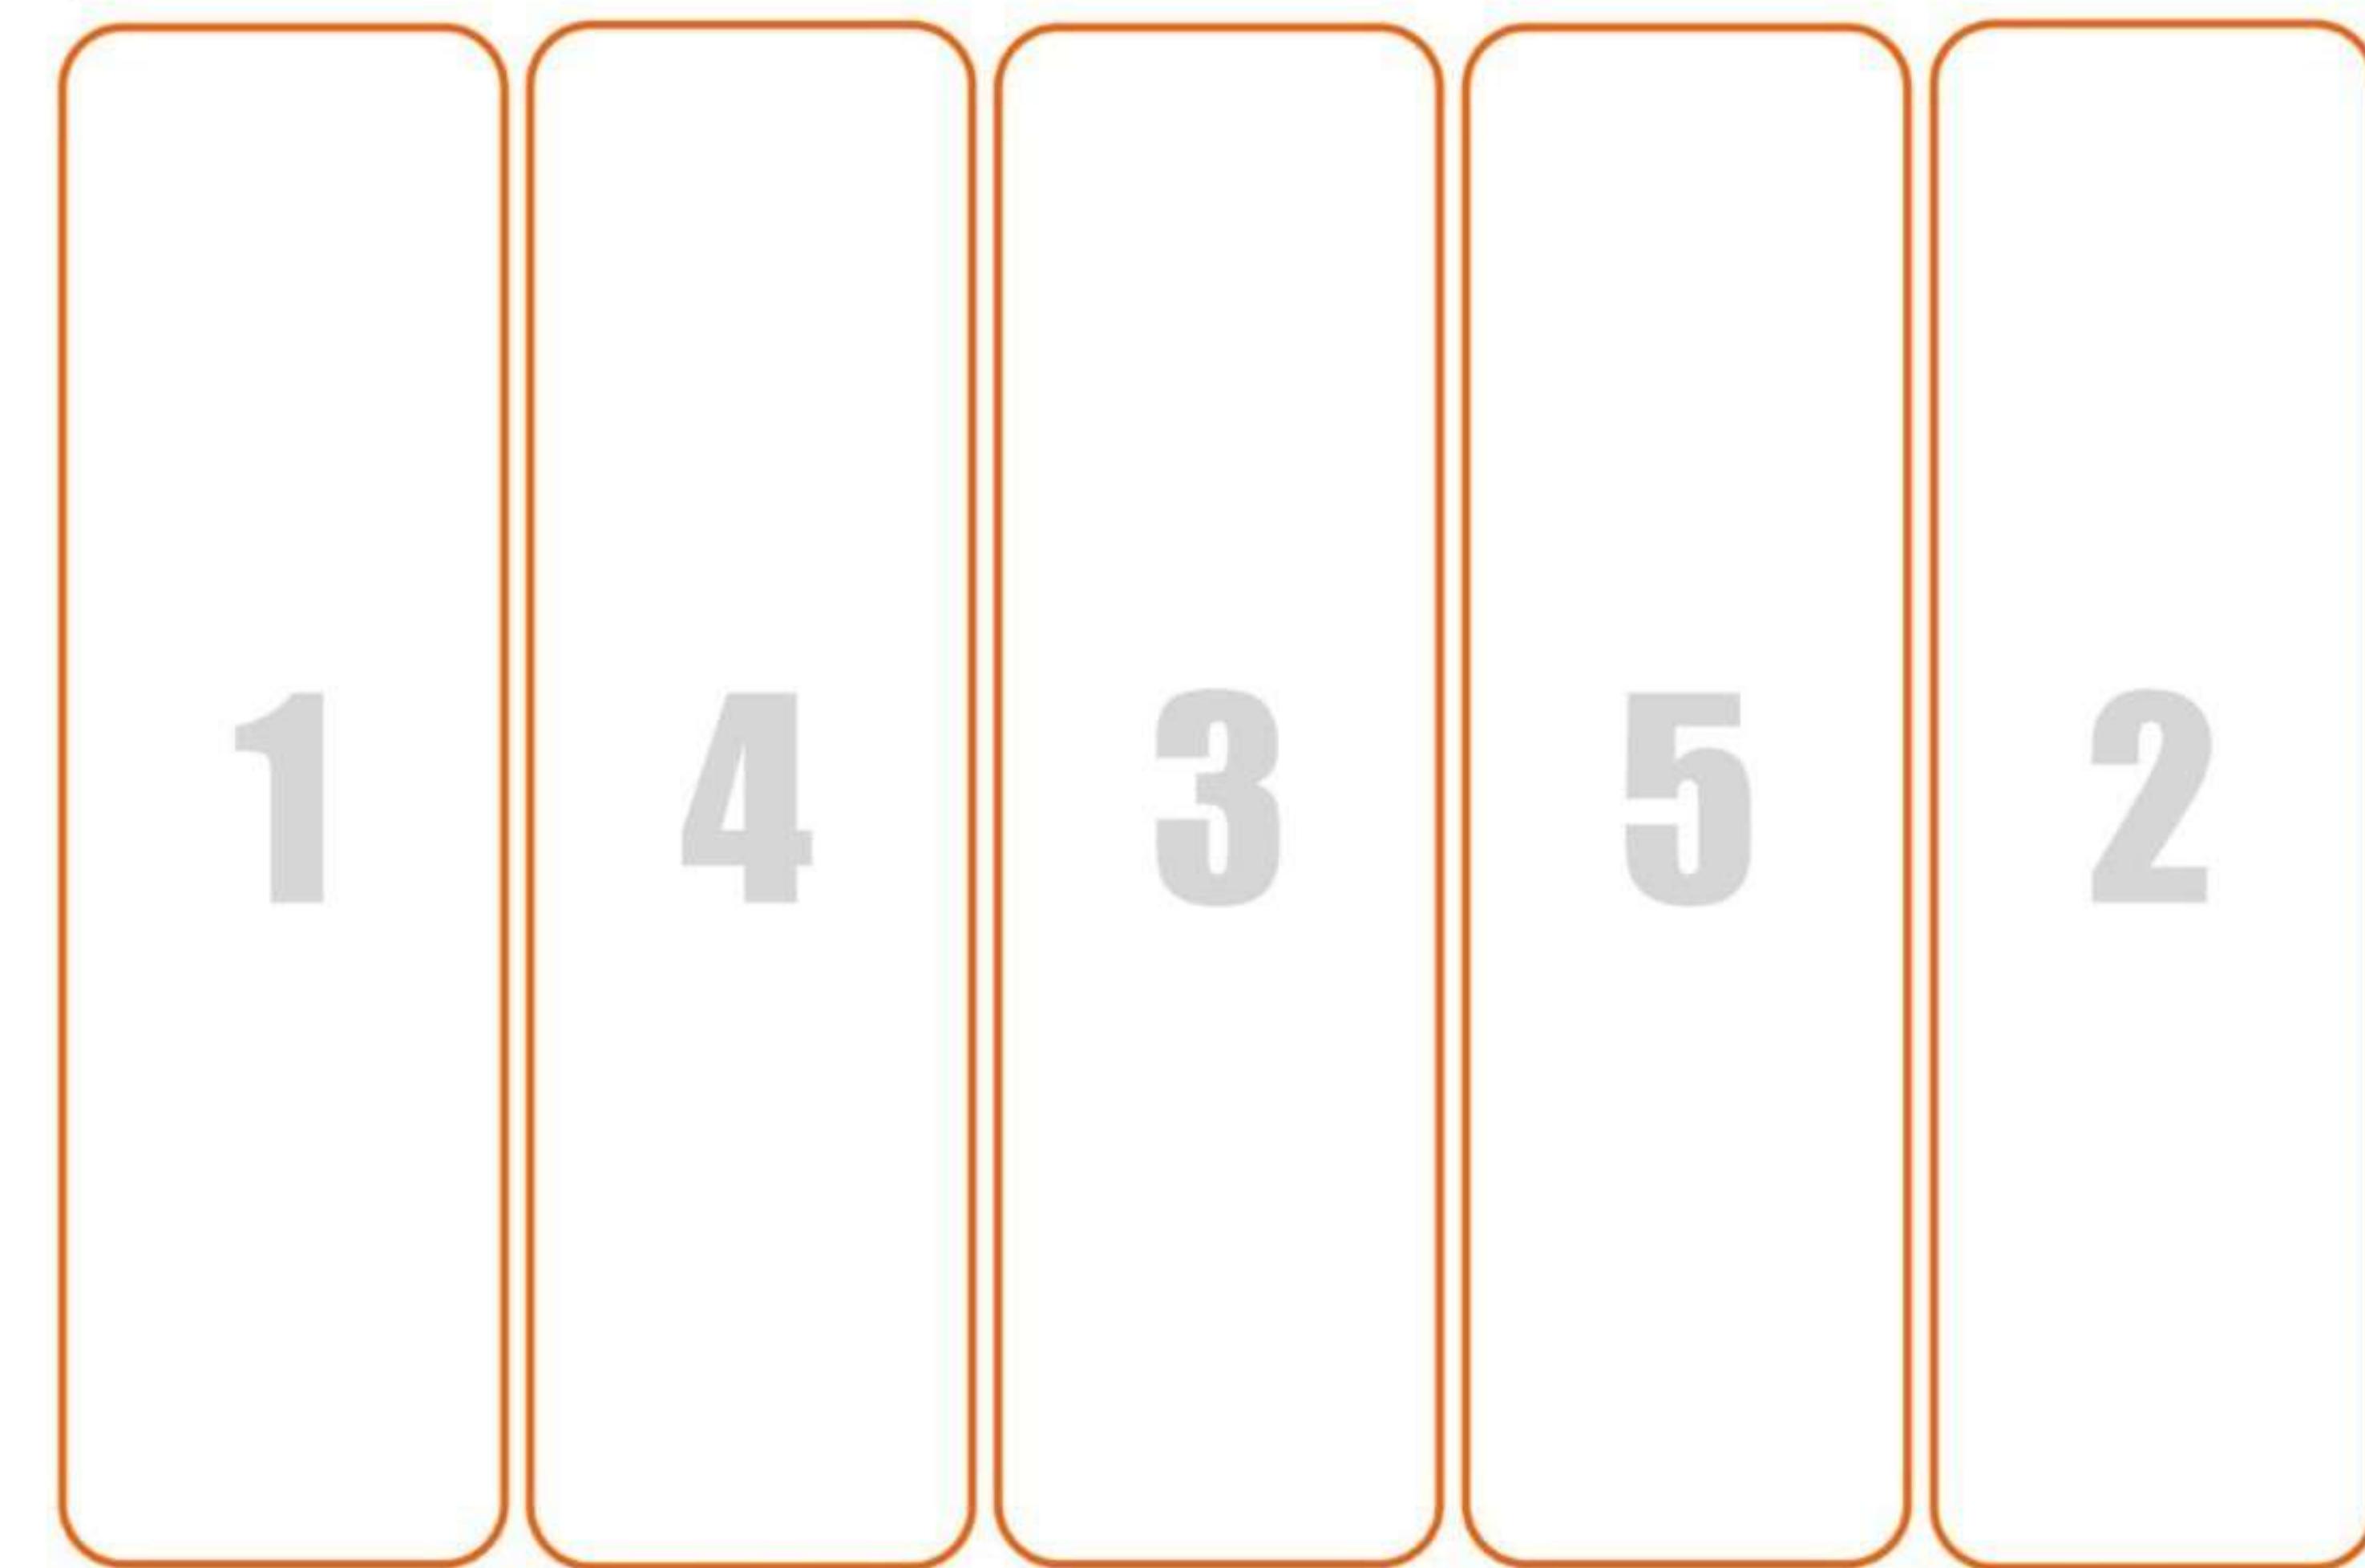


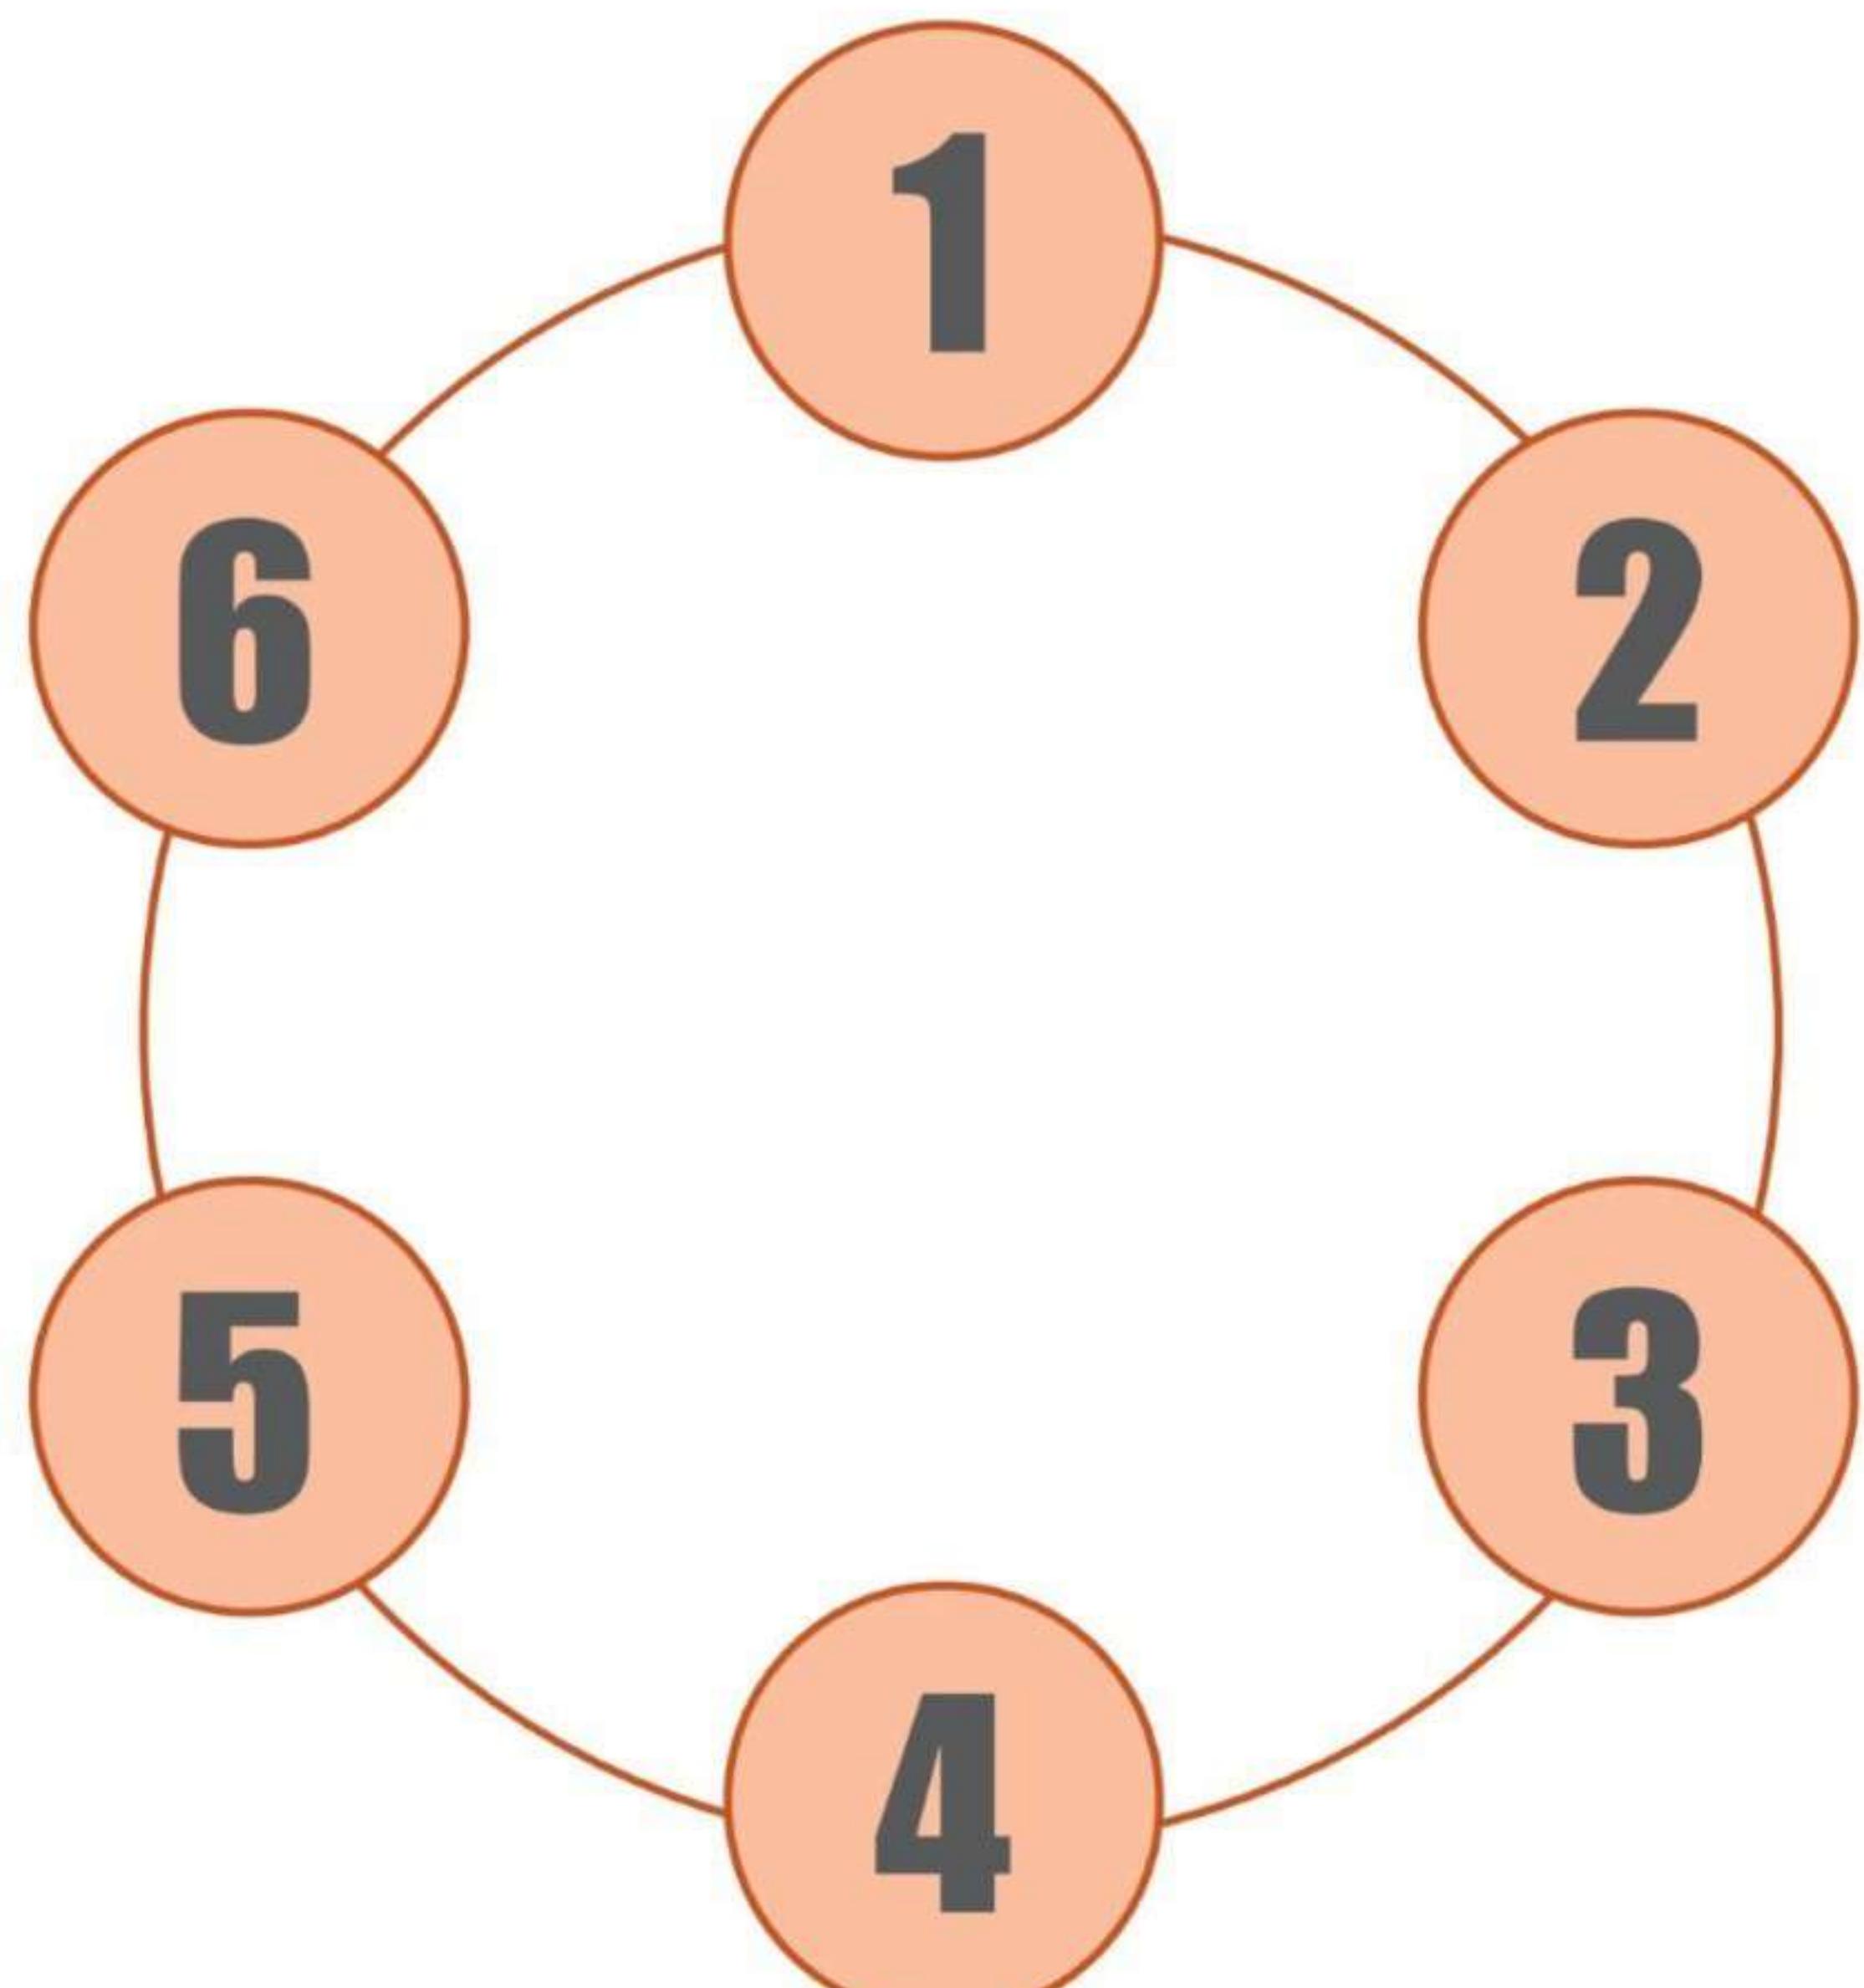
-2^{63} to $+2^{63}$



Storing Data in Cassandra

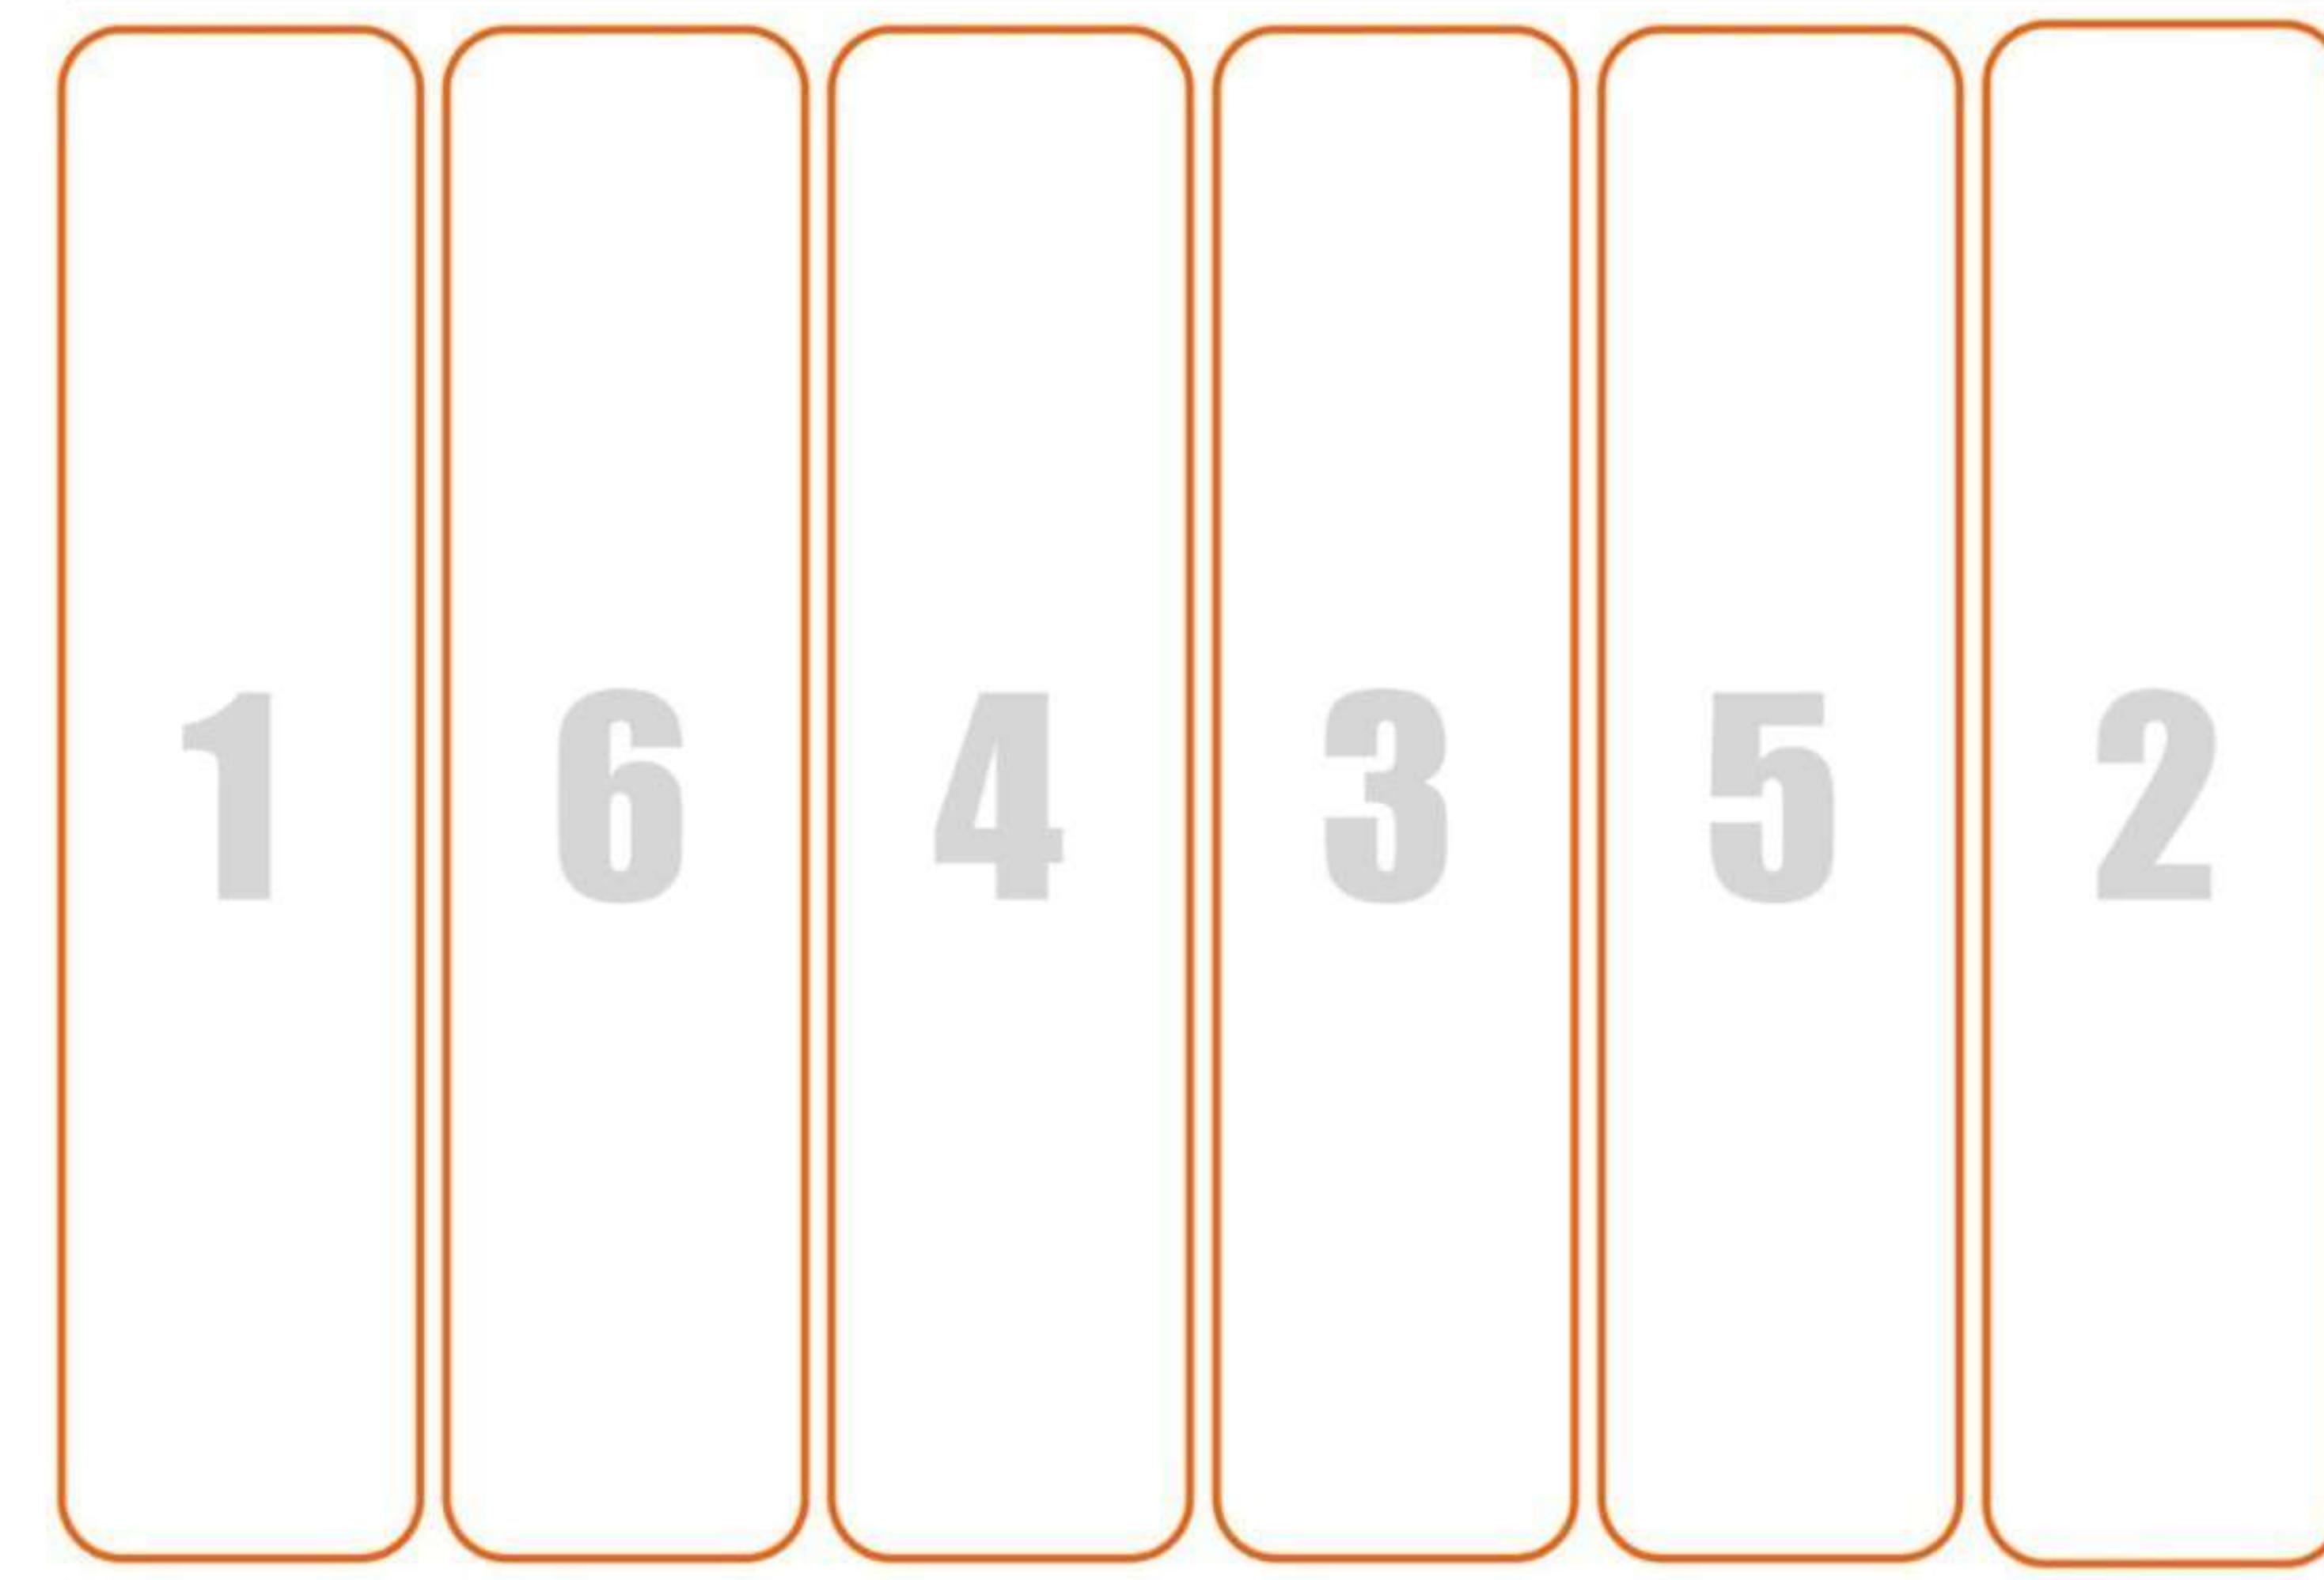
Range of all possible token values



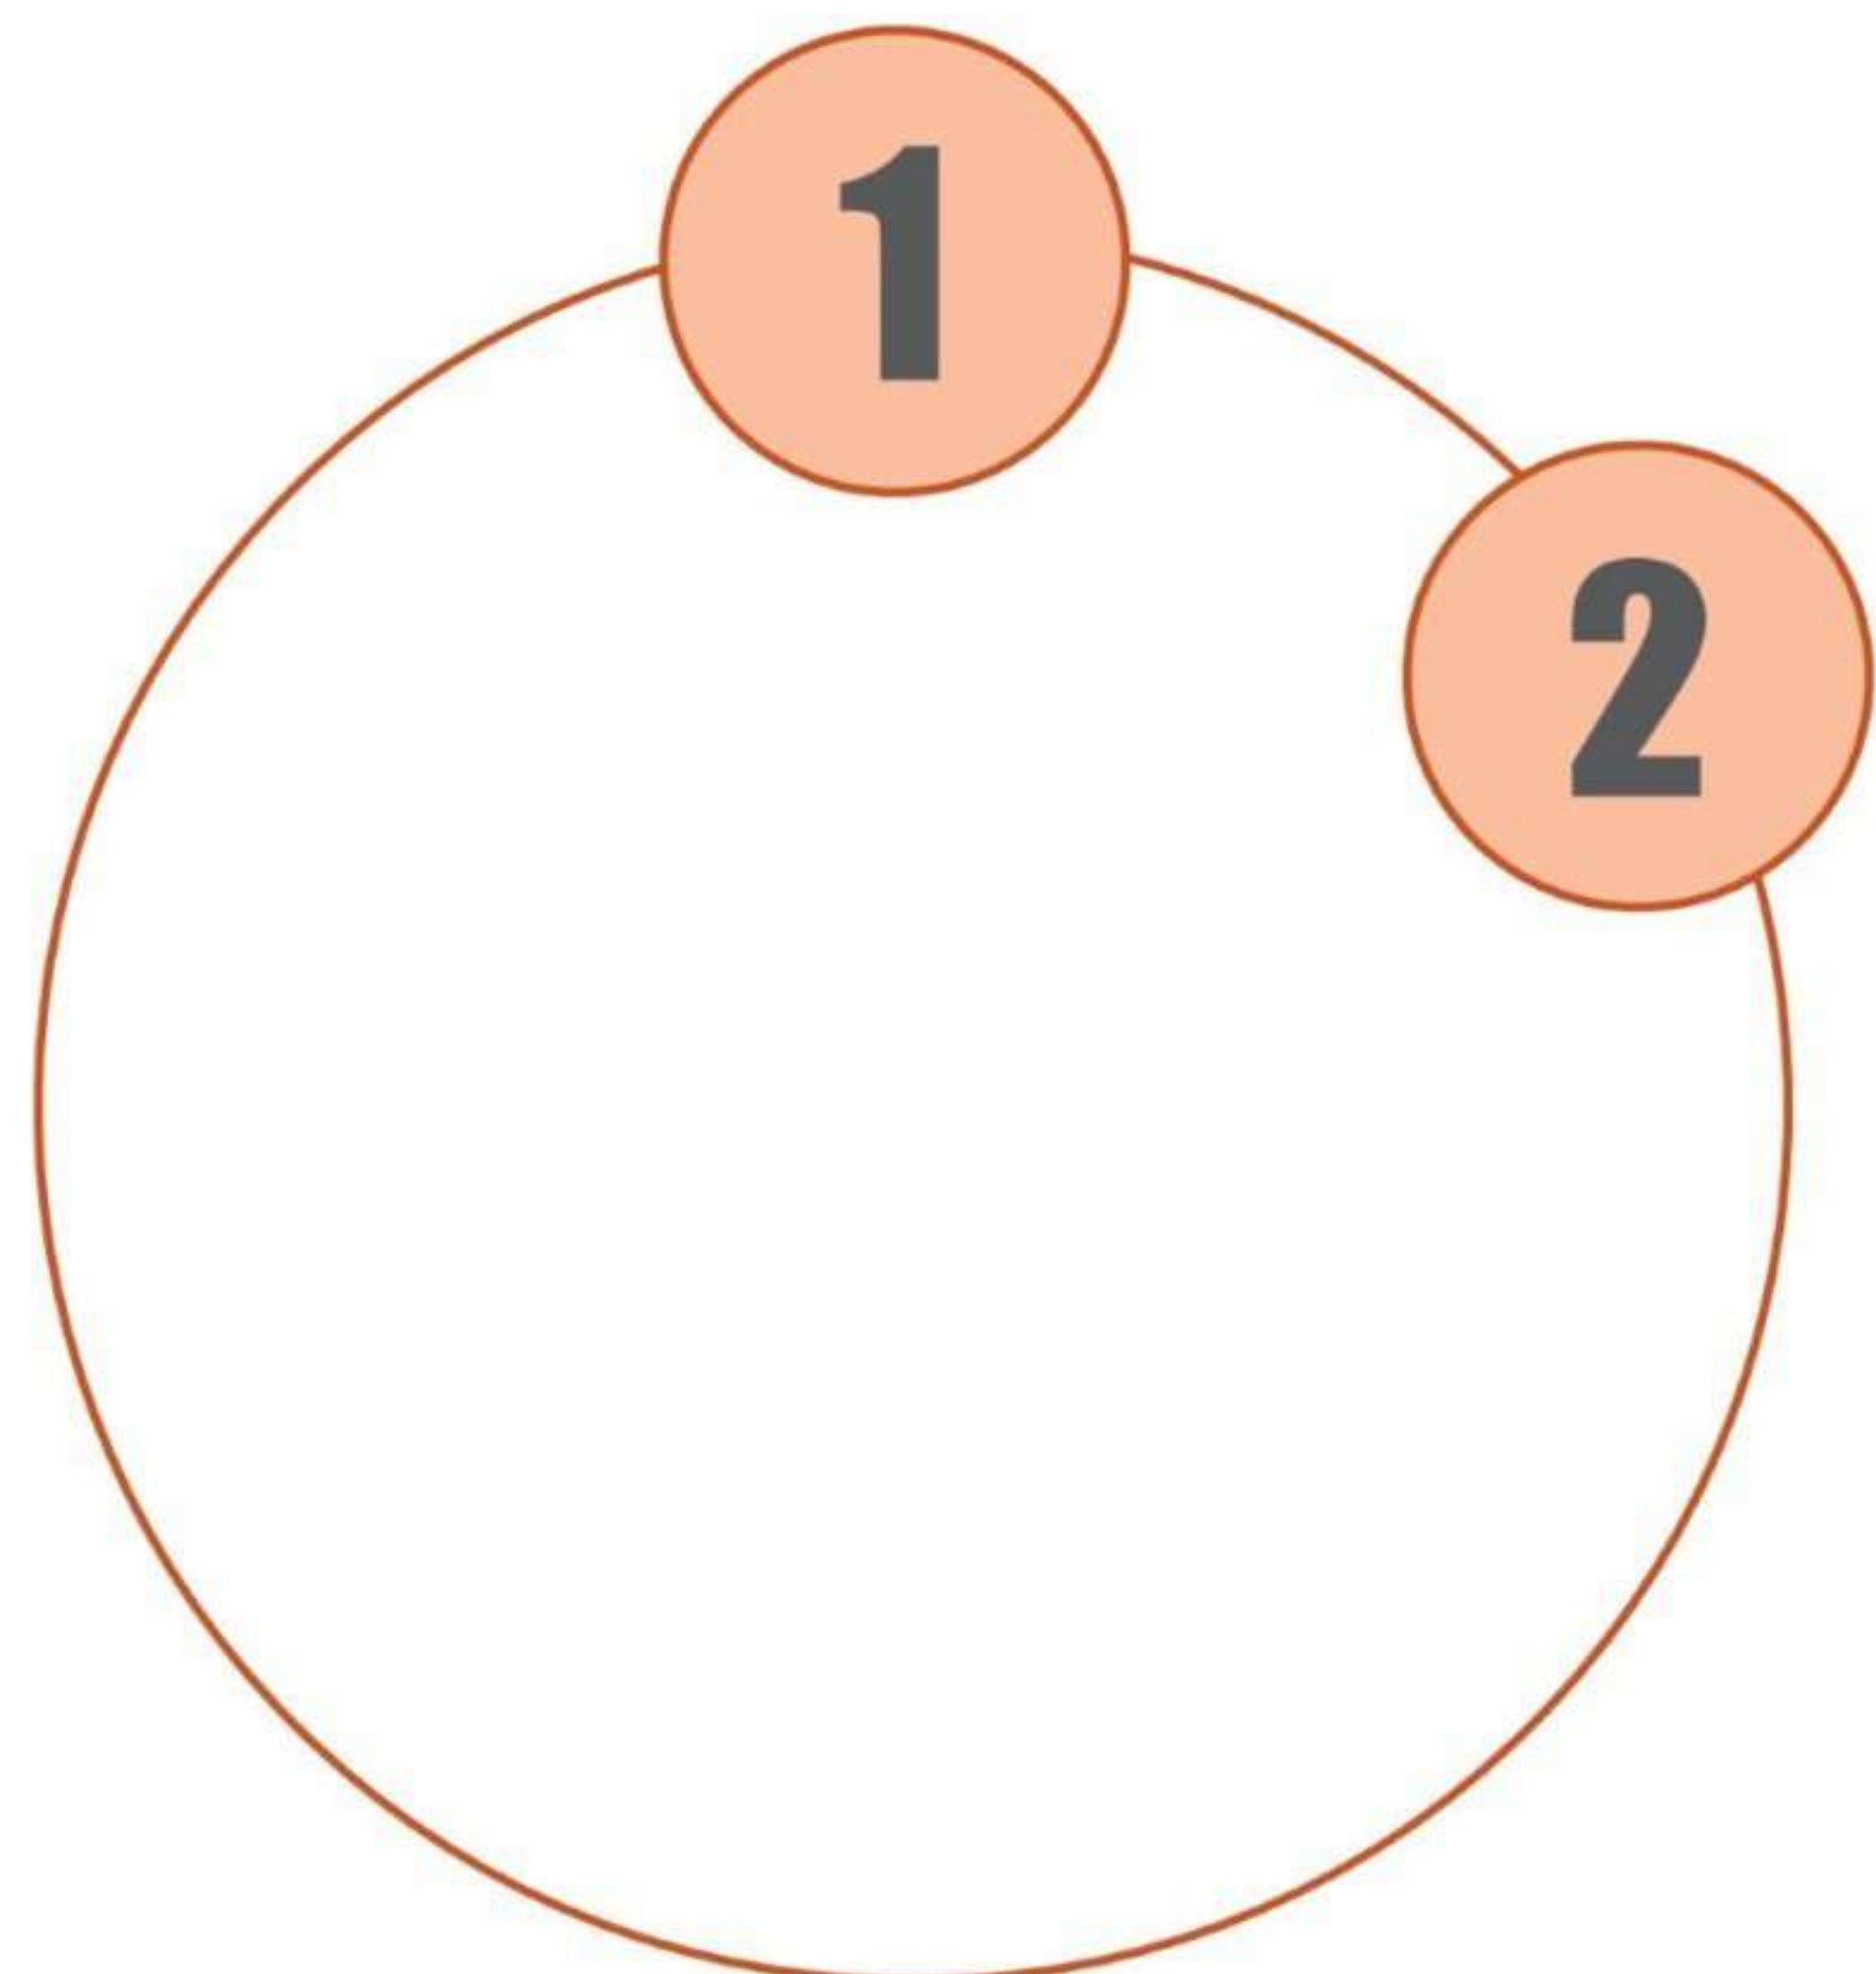


Storing Data in Cassandra

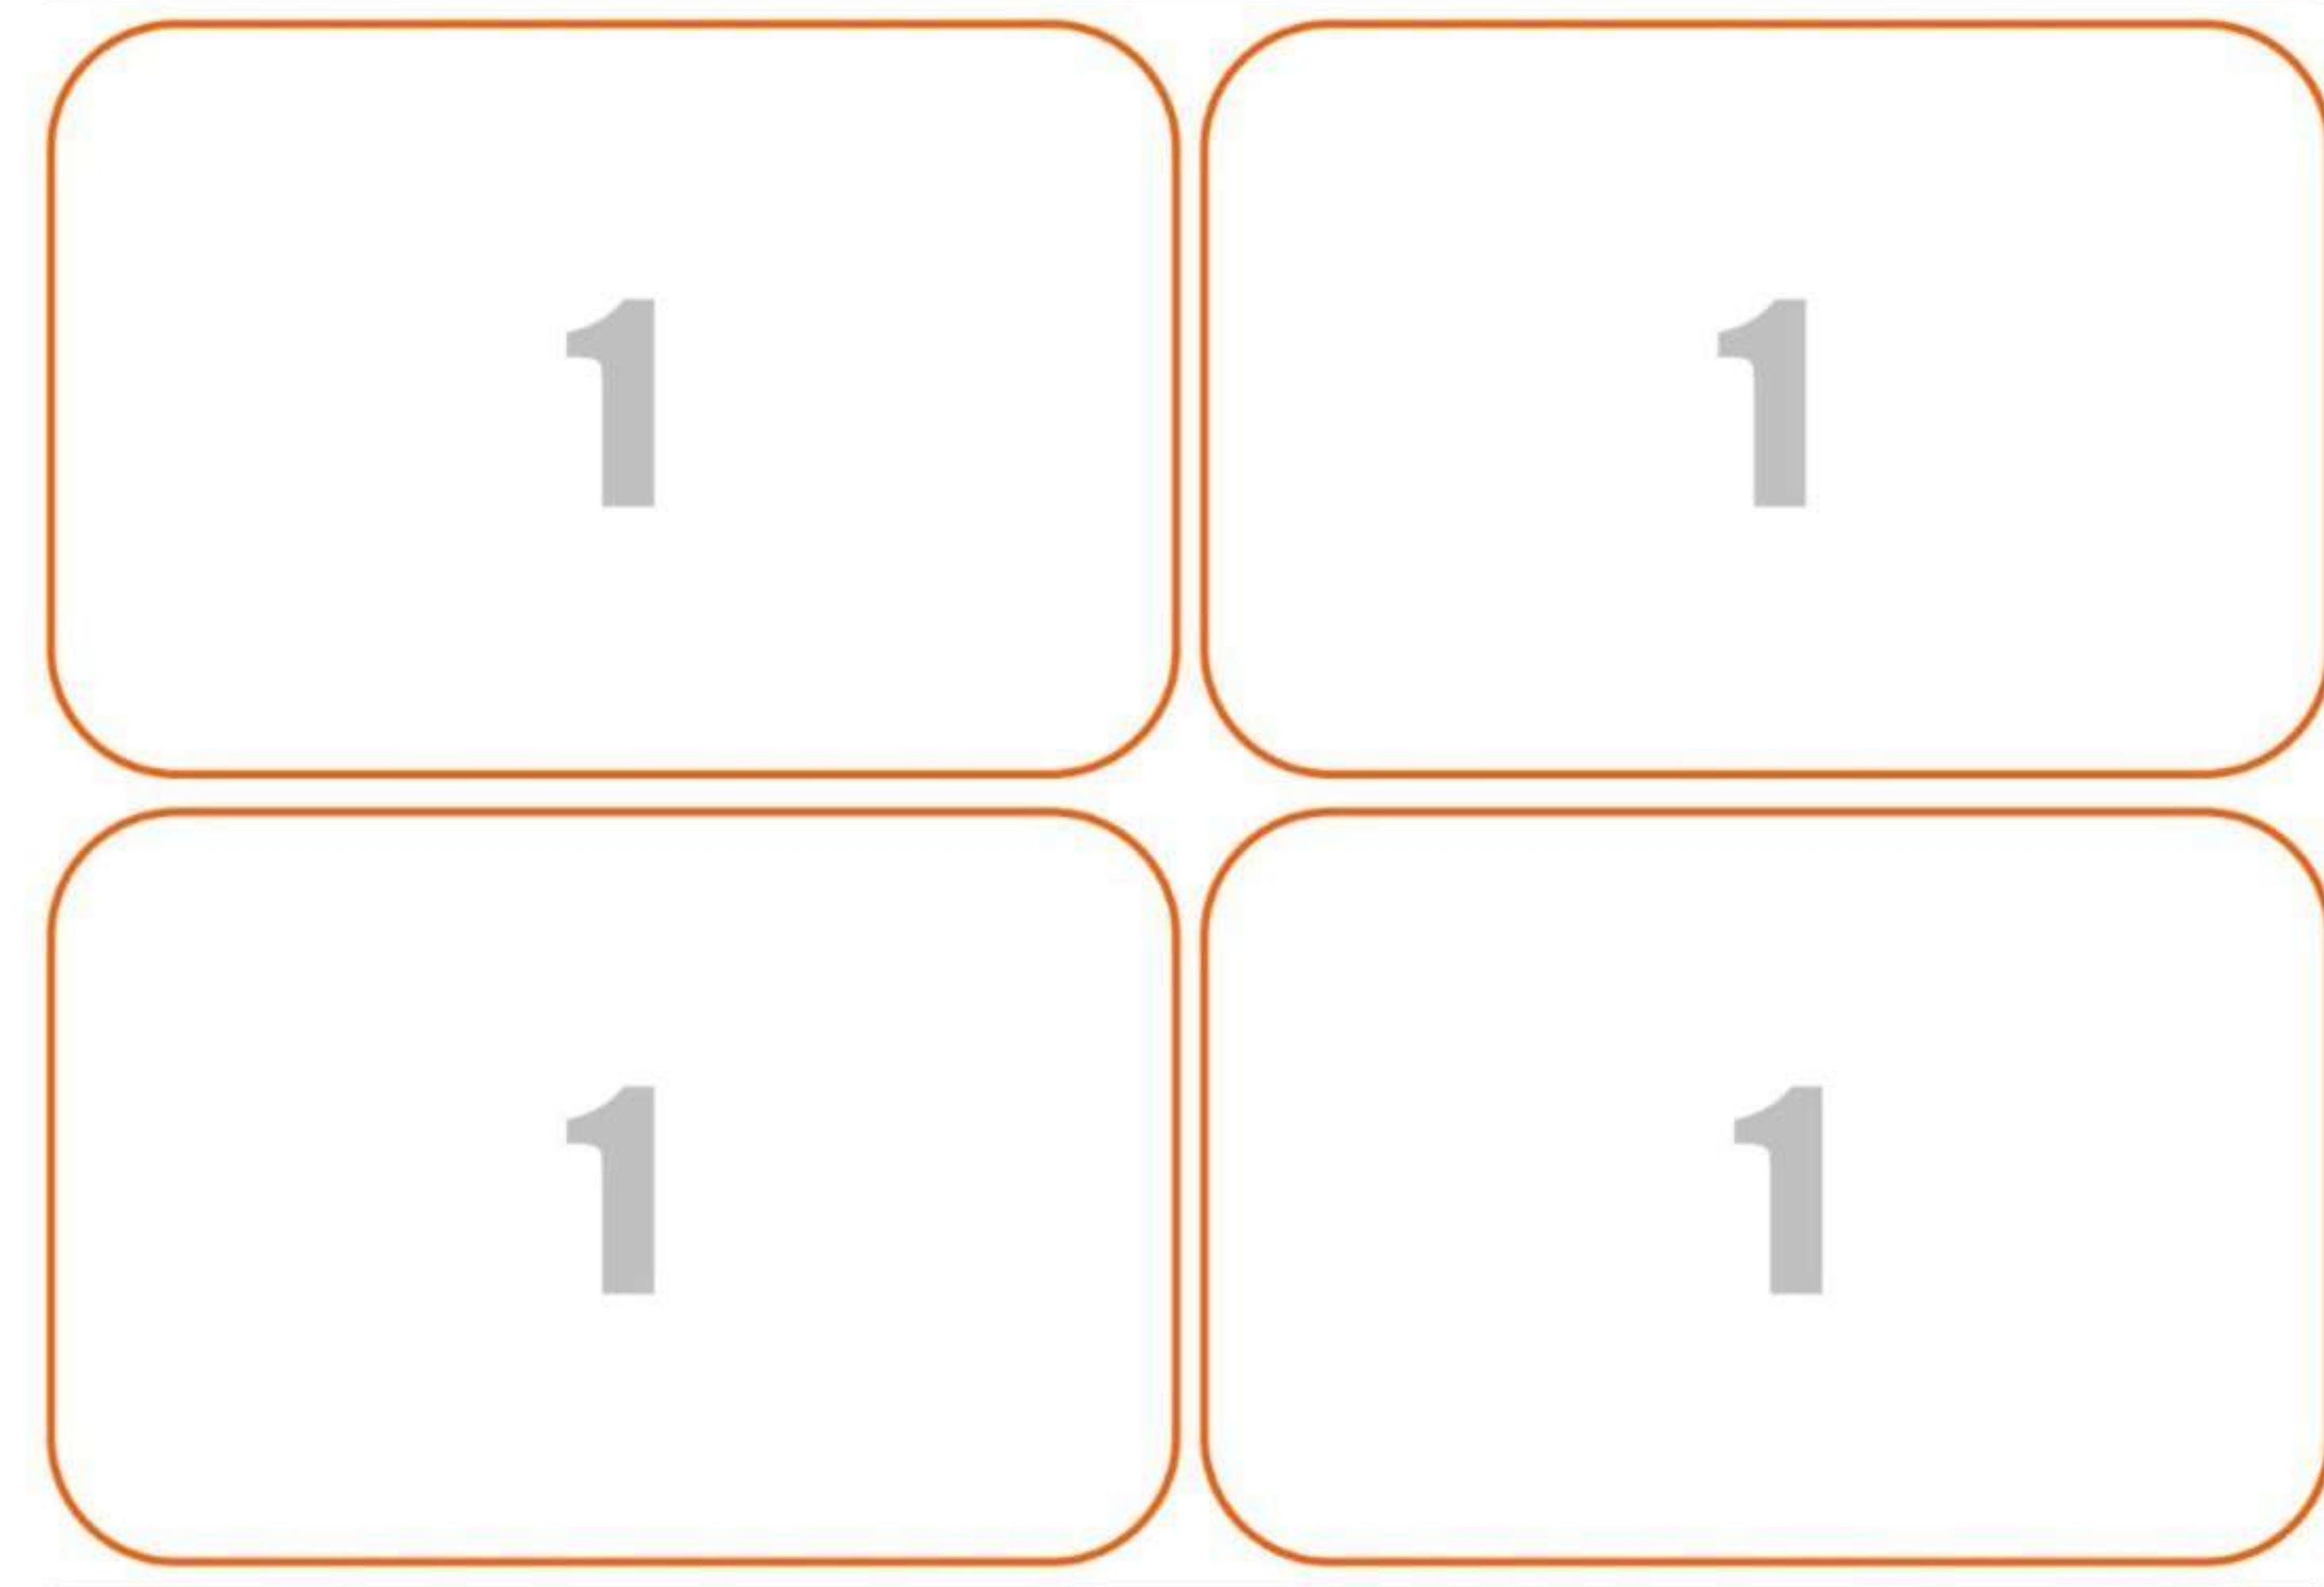
Range of all possible token values



Storing Data with Virtual Nodes

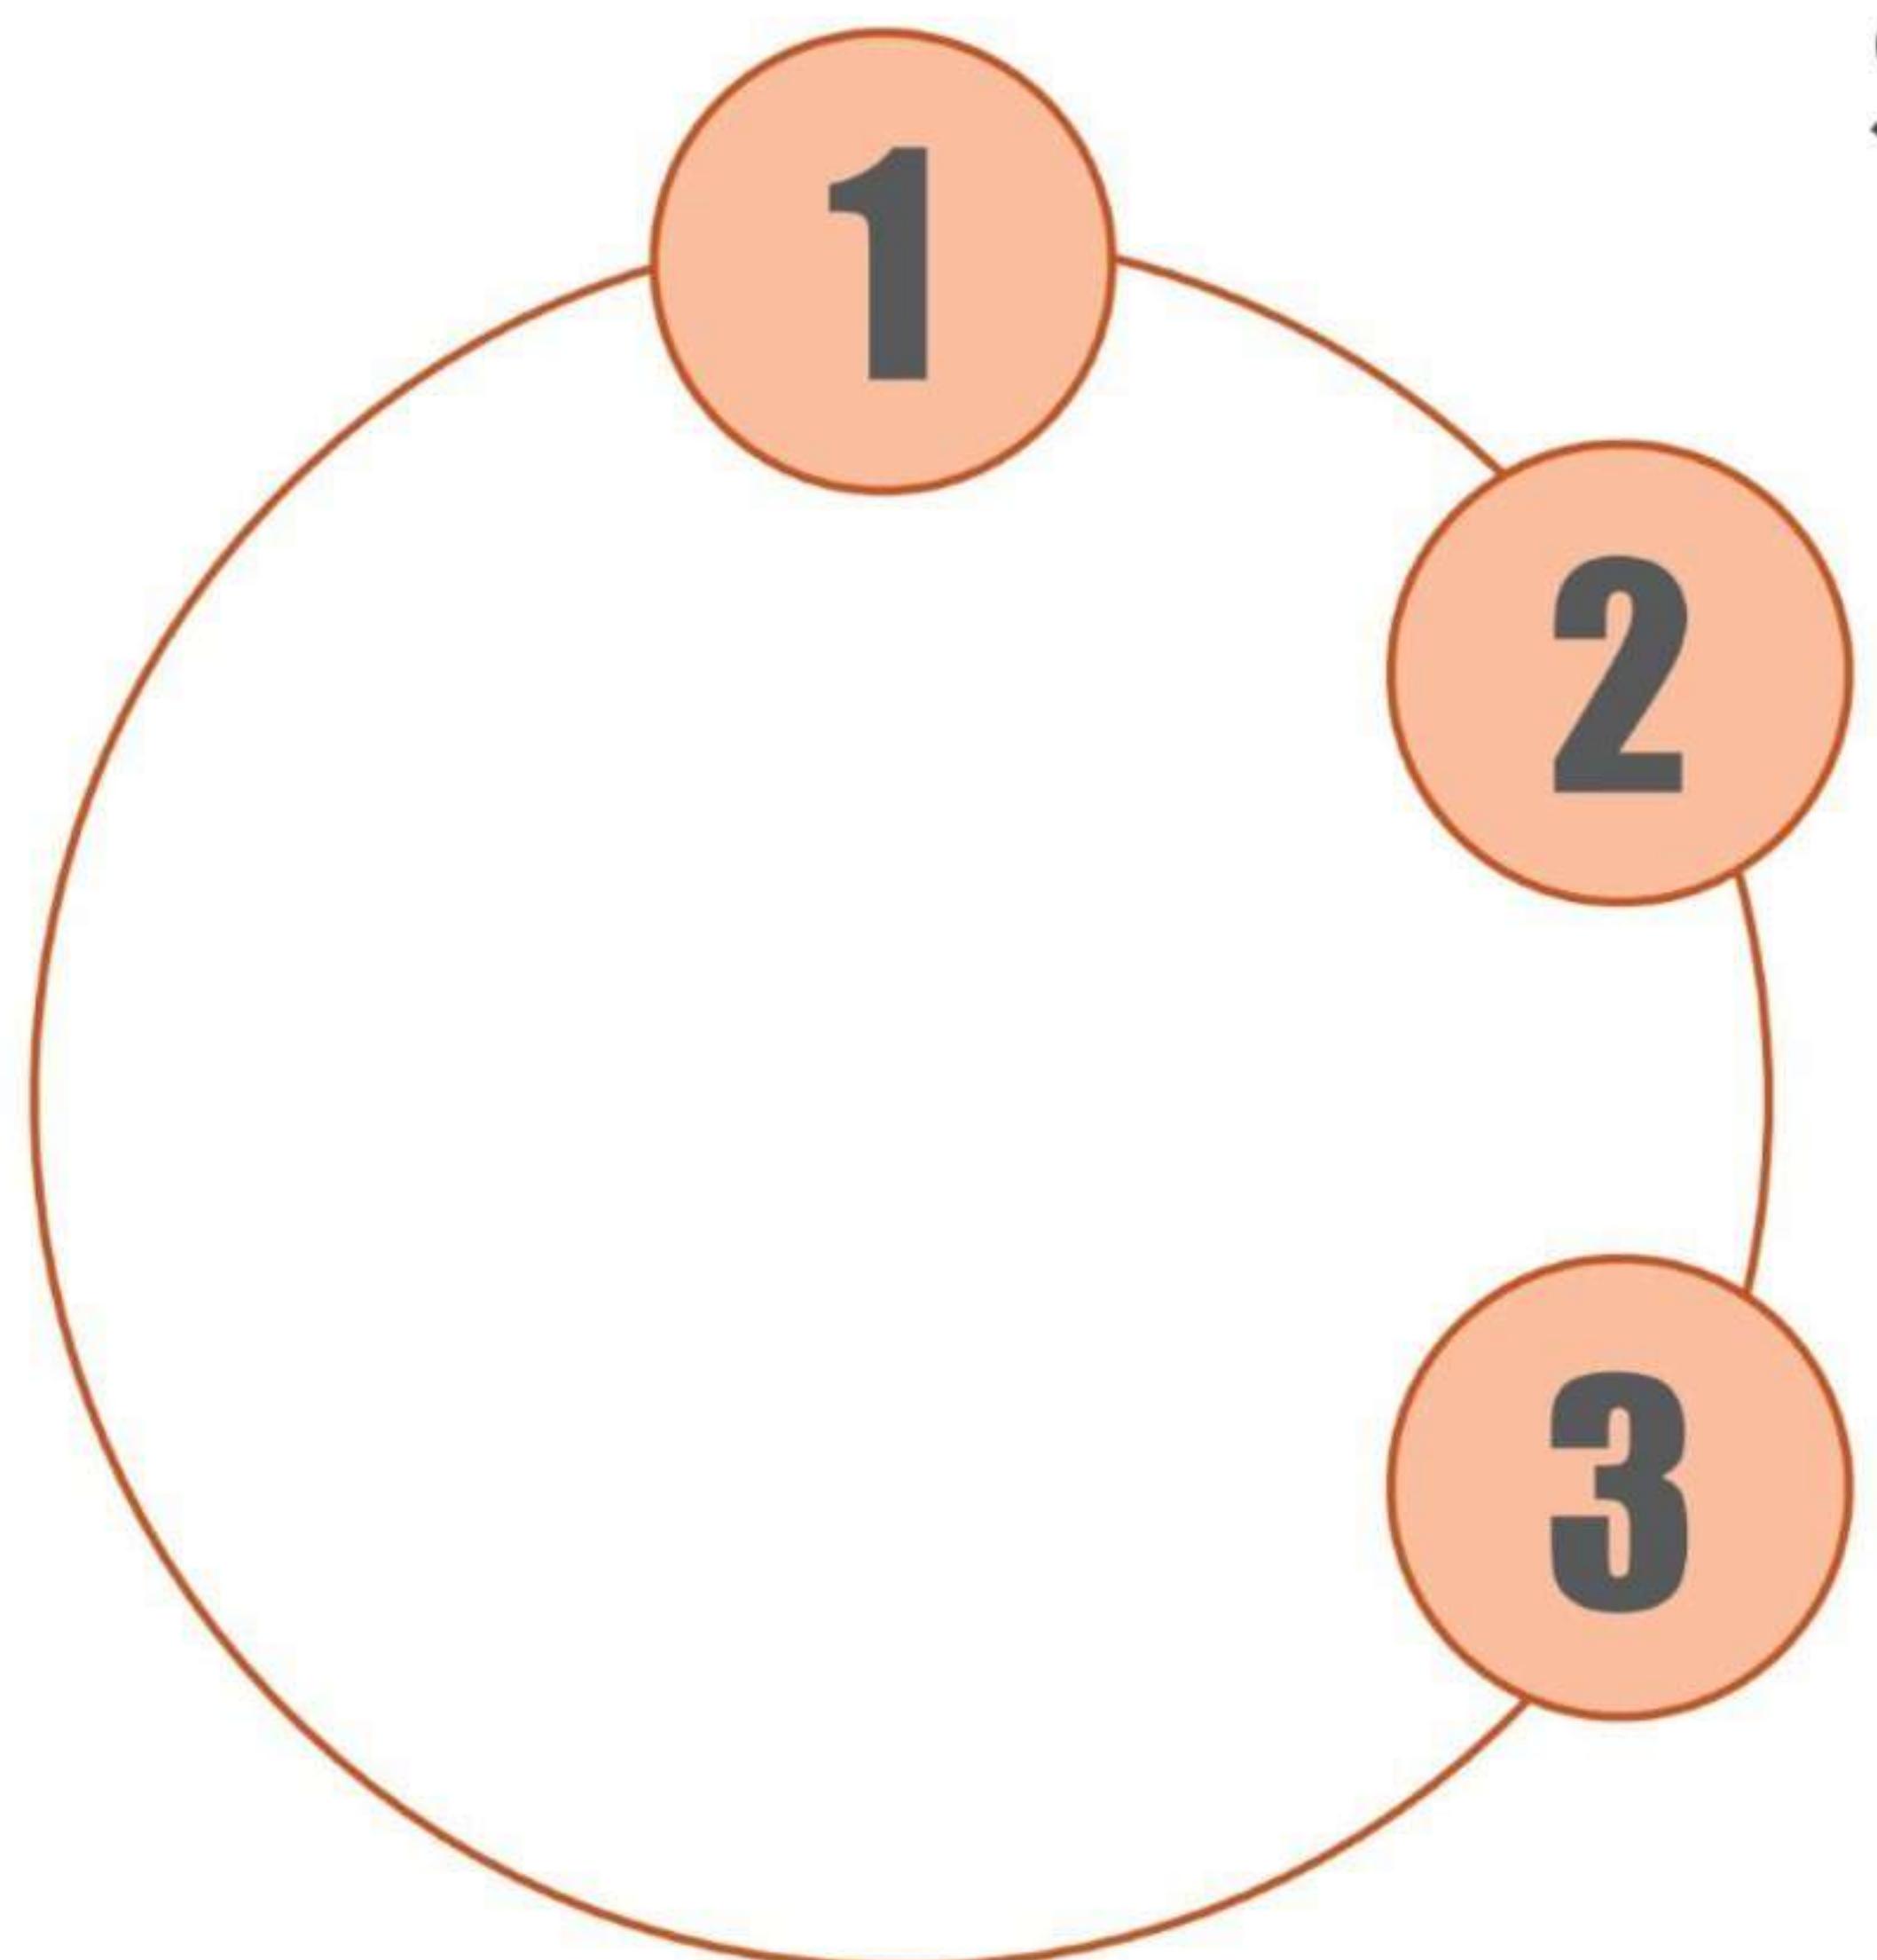


Range of all possible token values

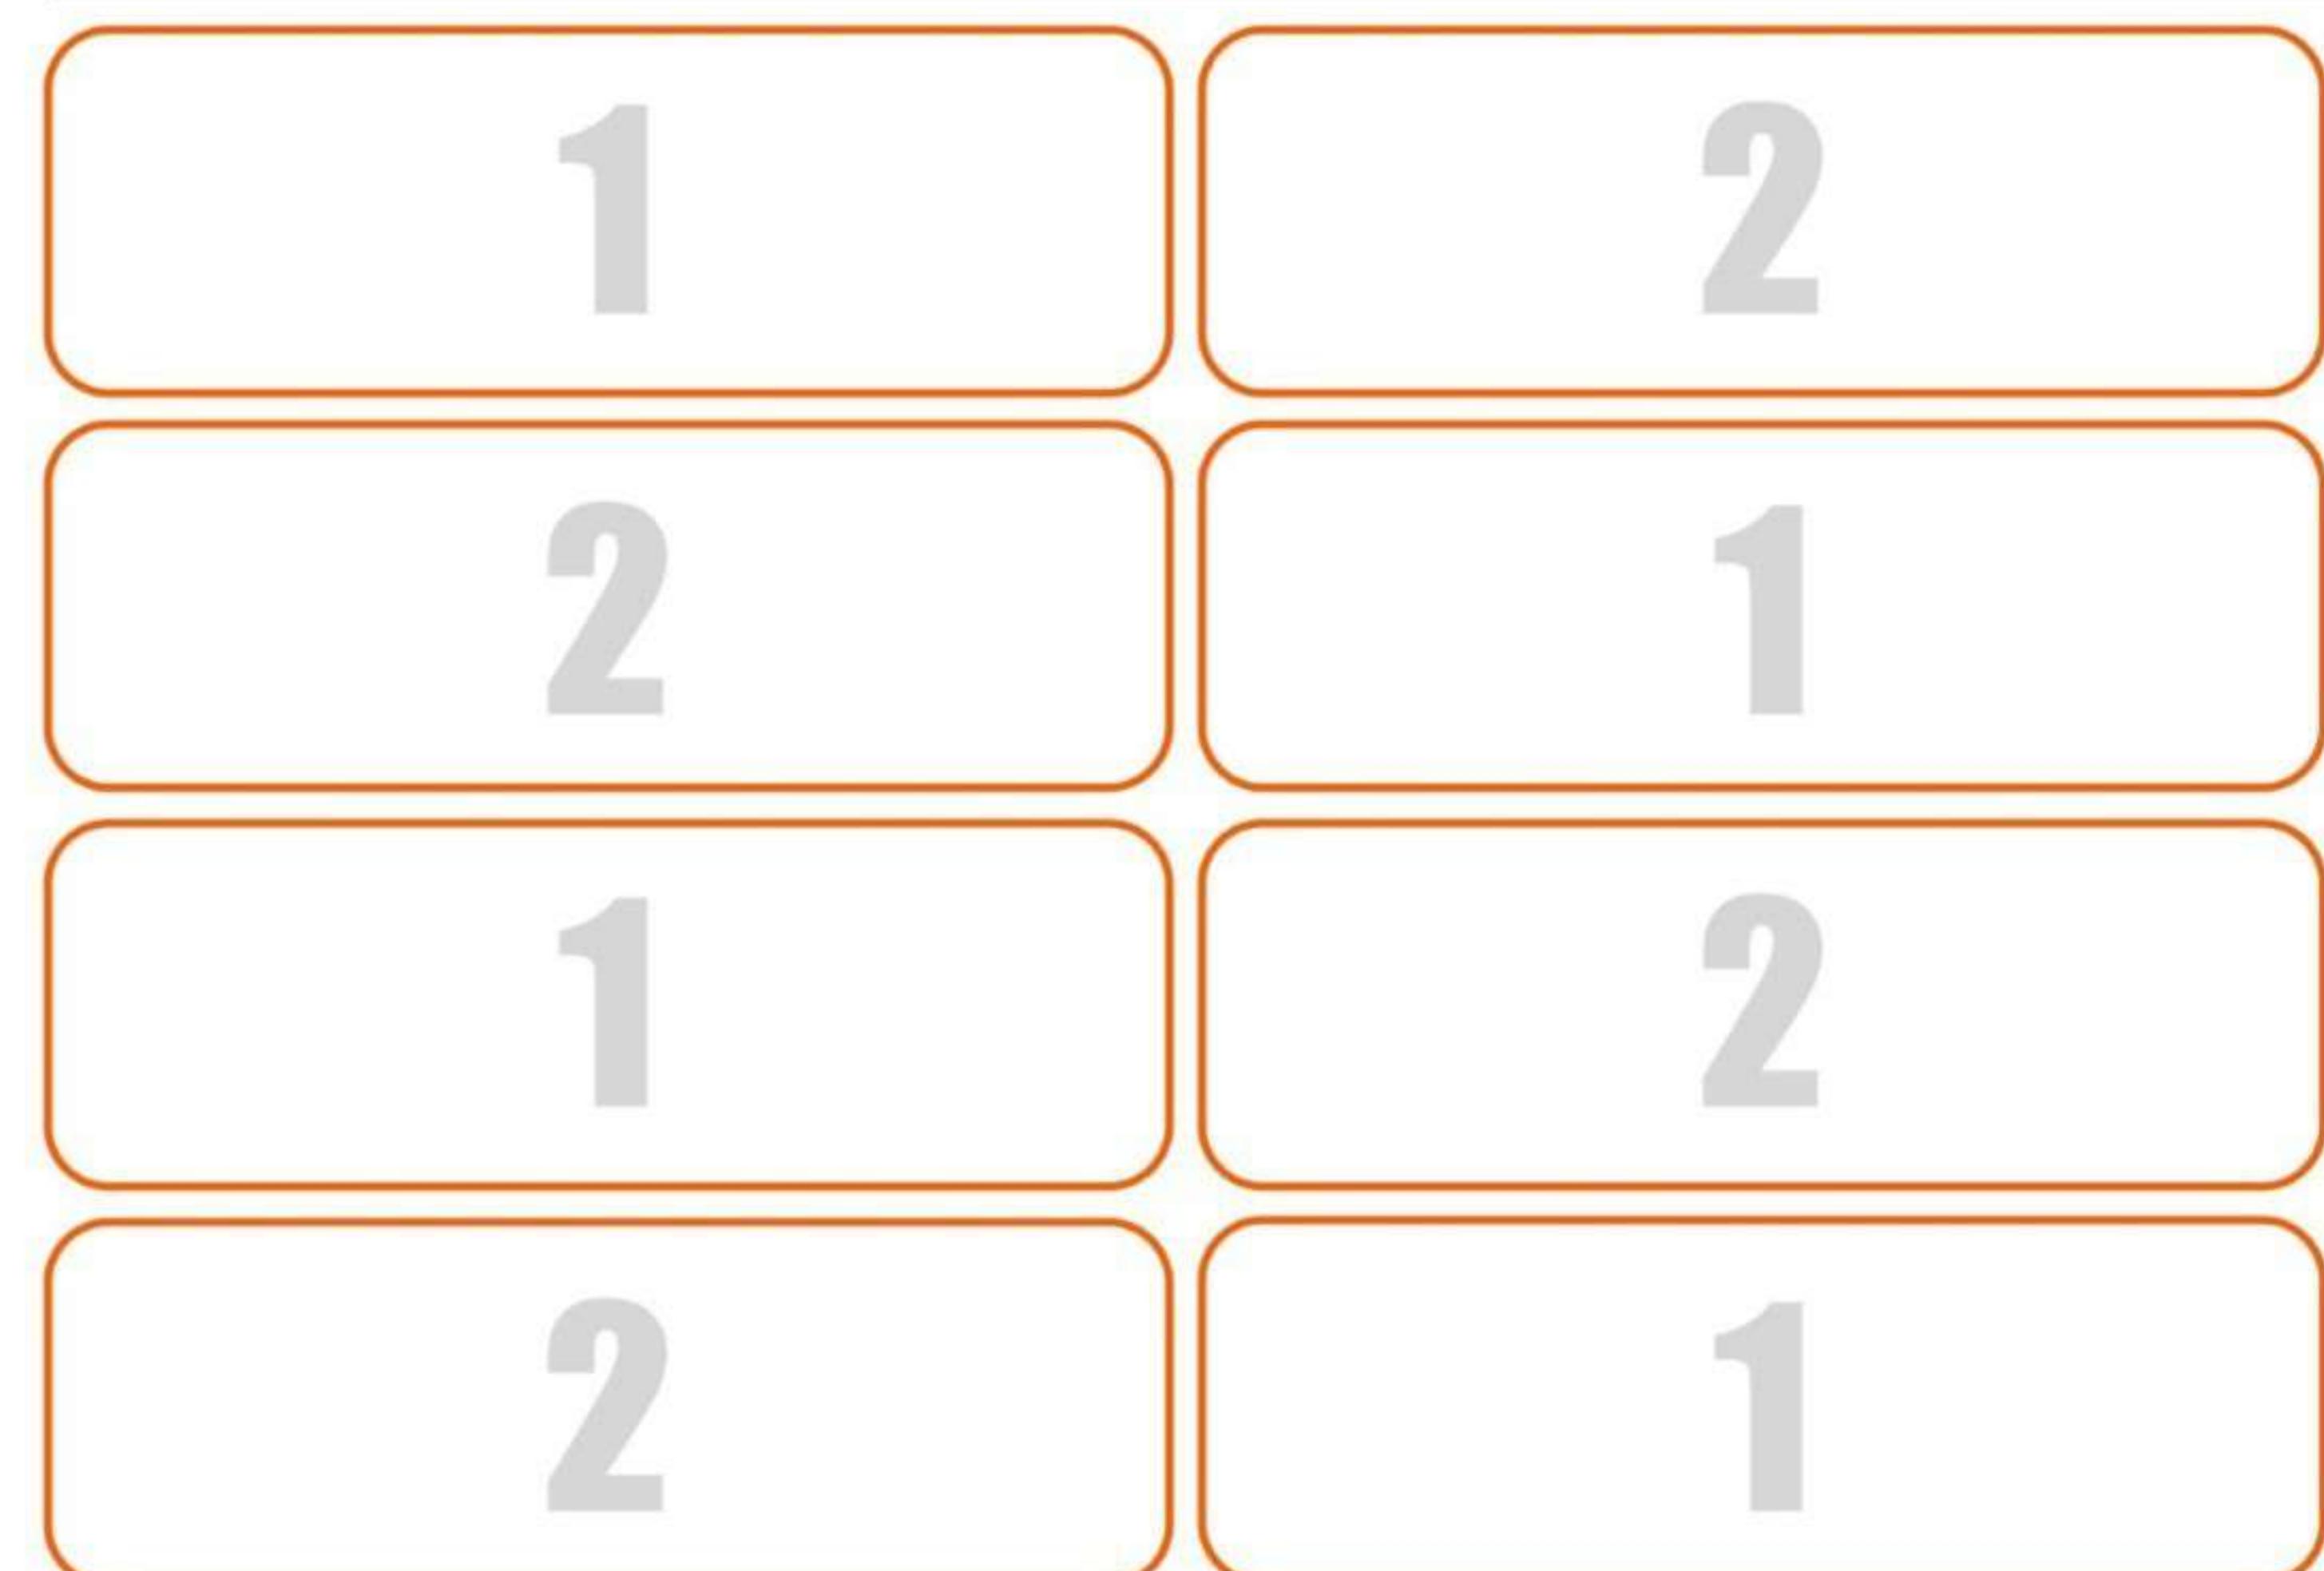


-2⁶³ to +2⁶³

Storing Data with Virtual Nodes

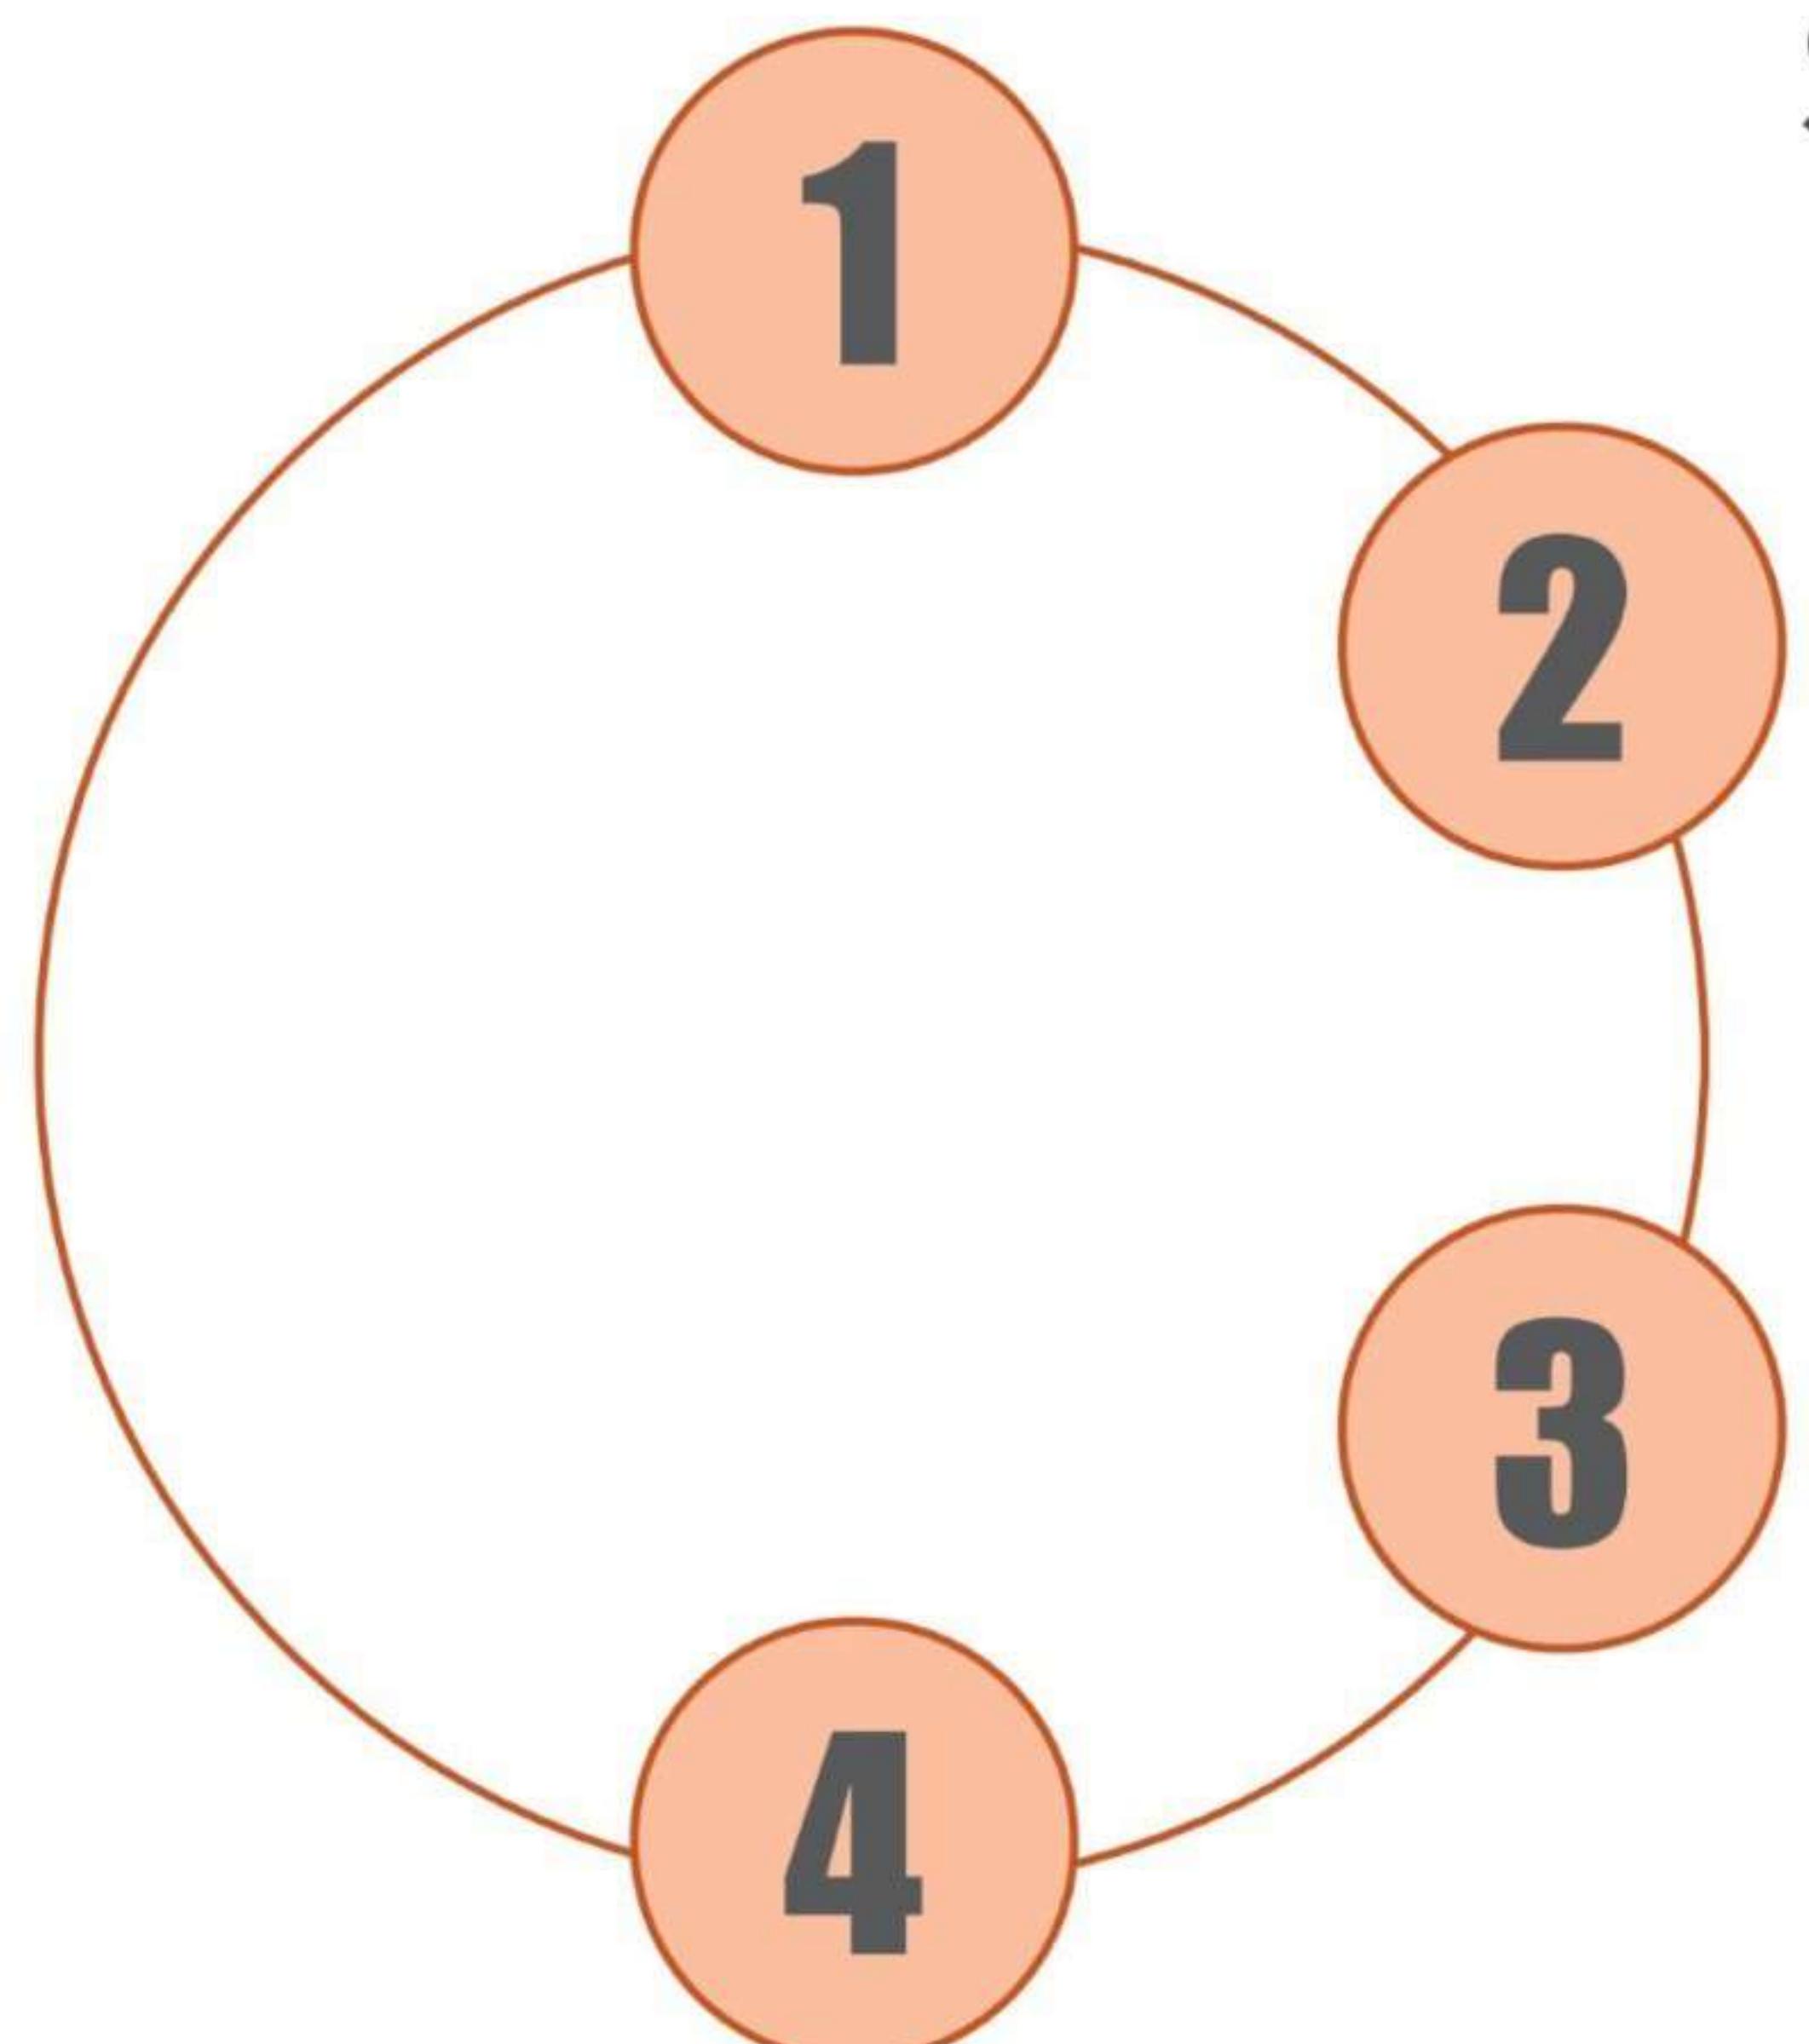


Range of all possible token values



-2⁶³ to +2⁶³

Storing Data with Virtual Nodes

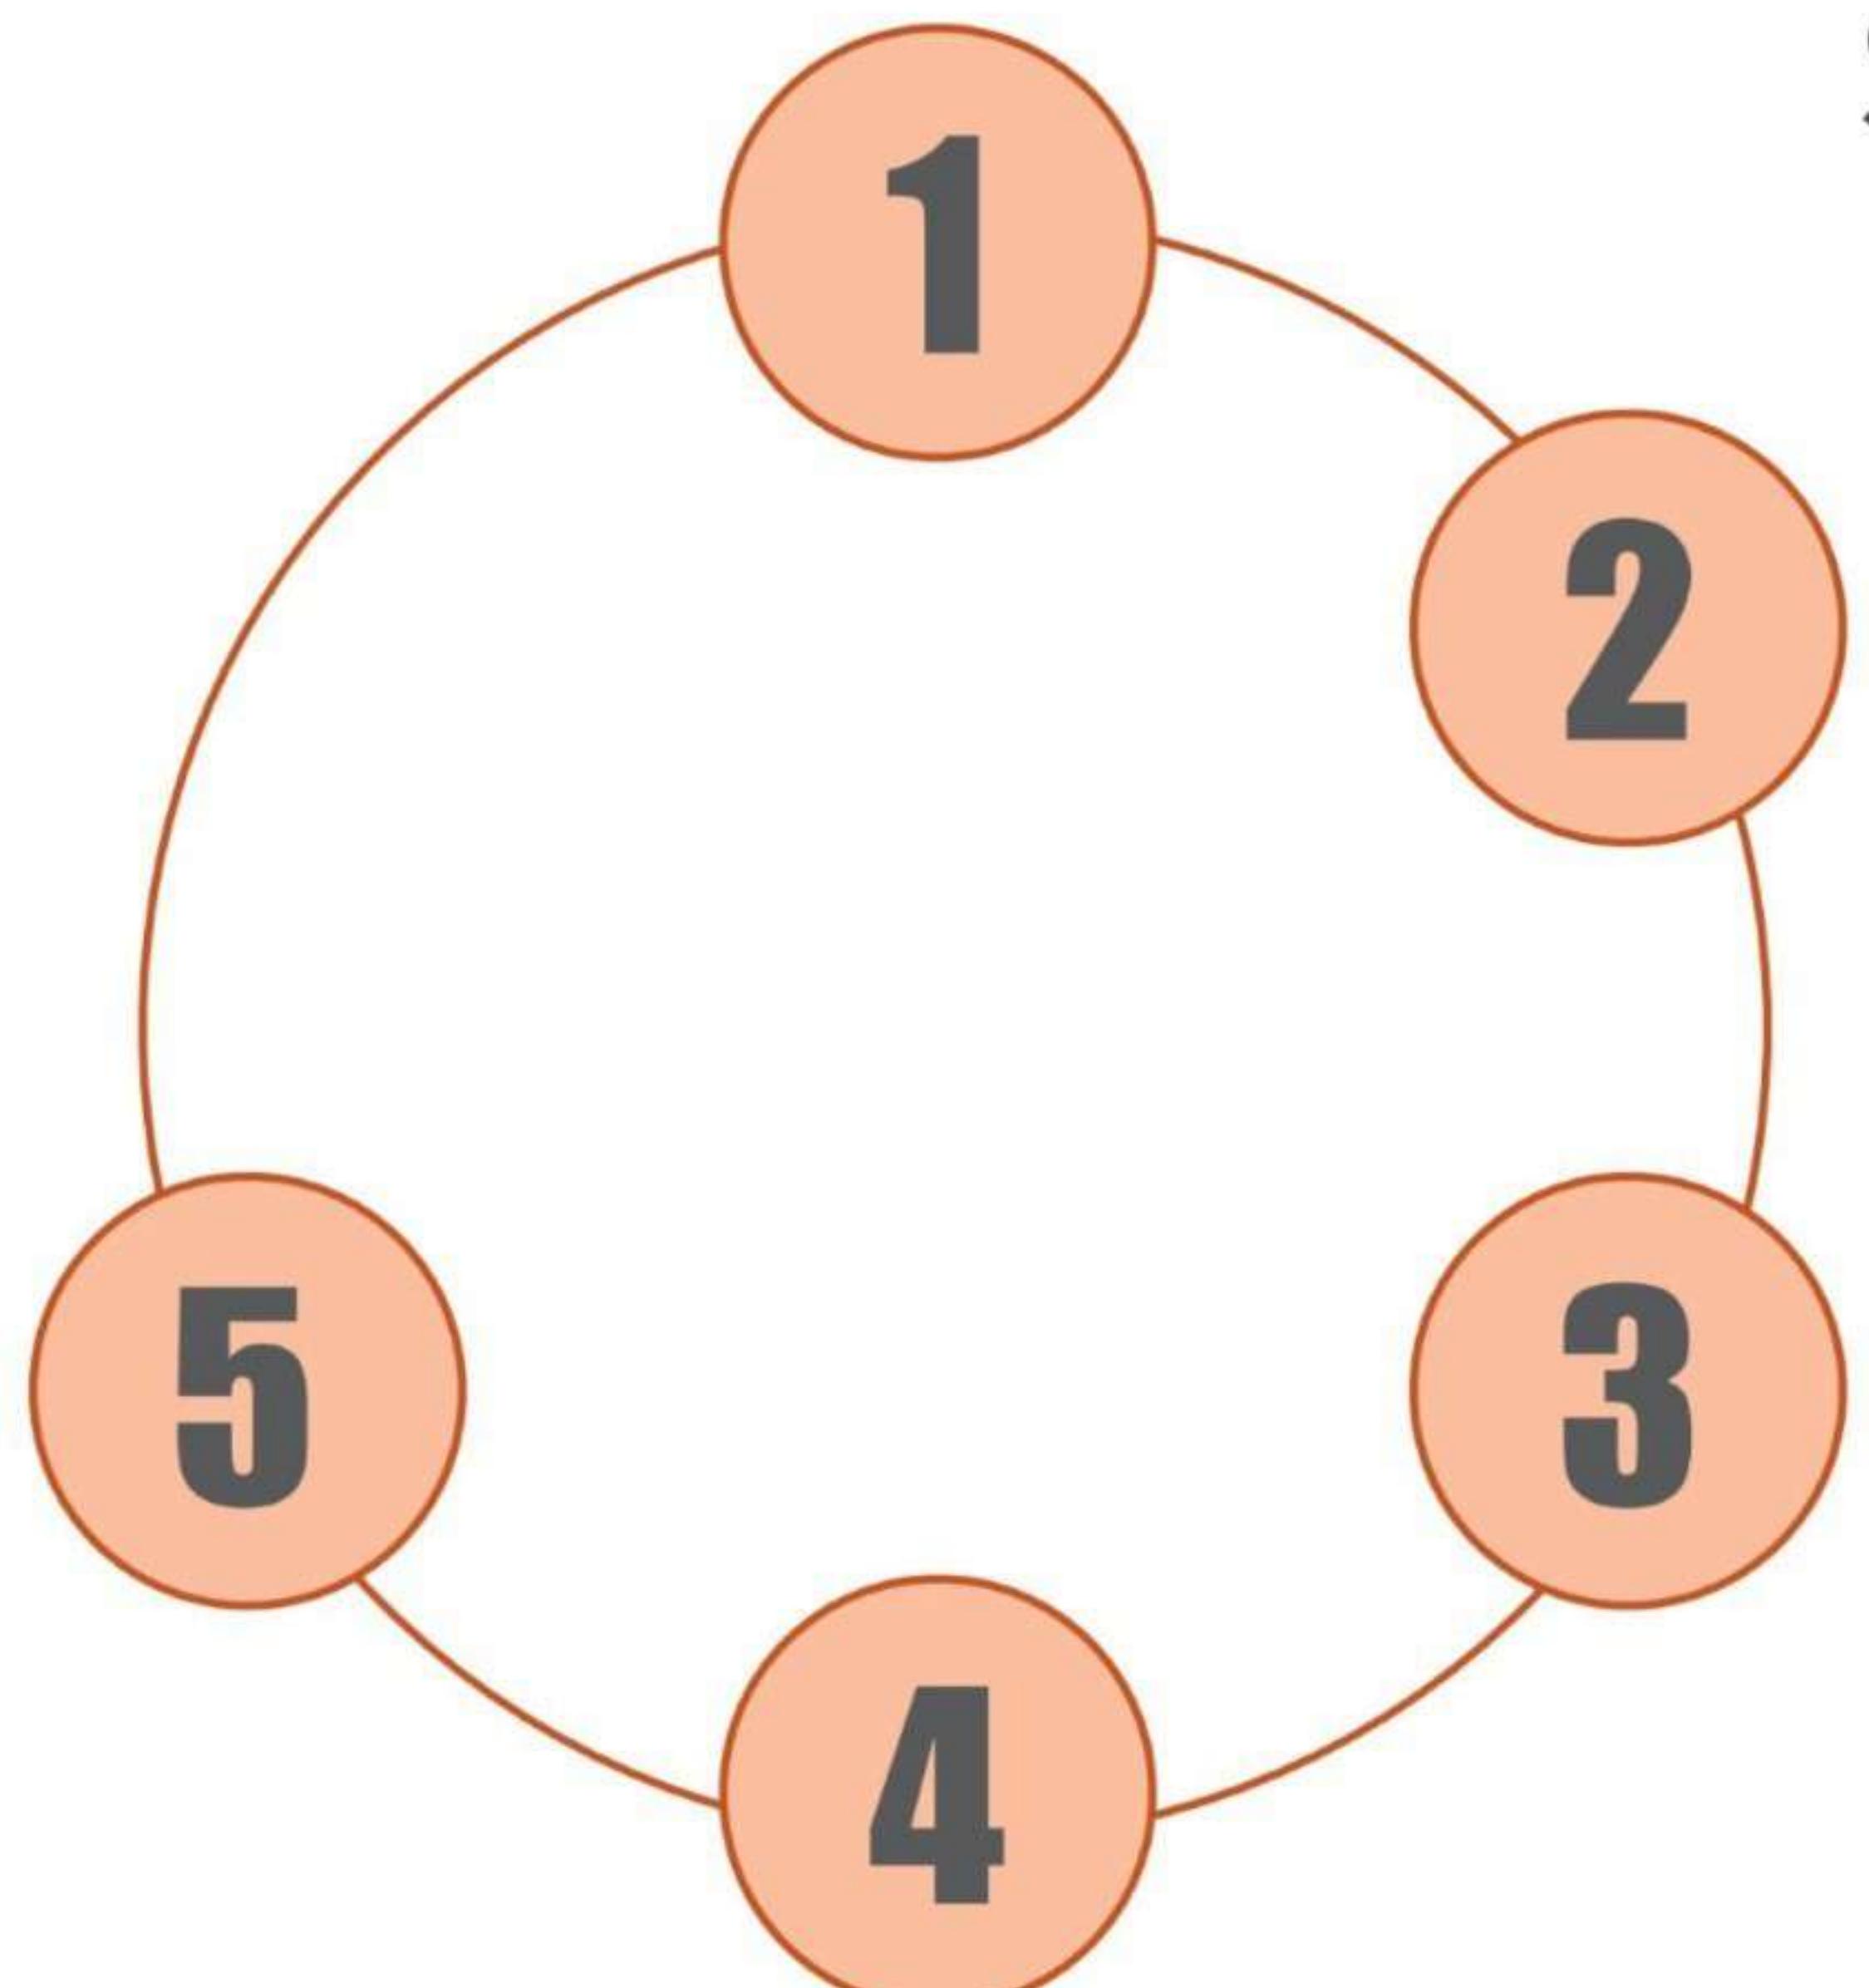


Range of all possible token values

1	2	3
2	3	1
1	2	3
3	2	1

- 2^{63} to $+2^{63}$

Storing Data with Virtual Nodes

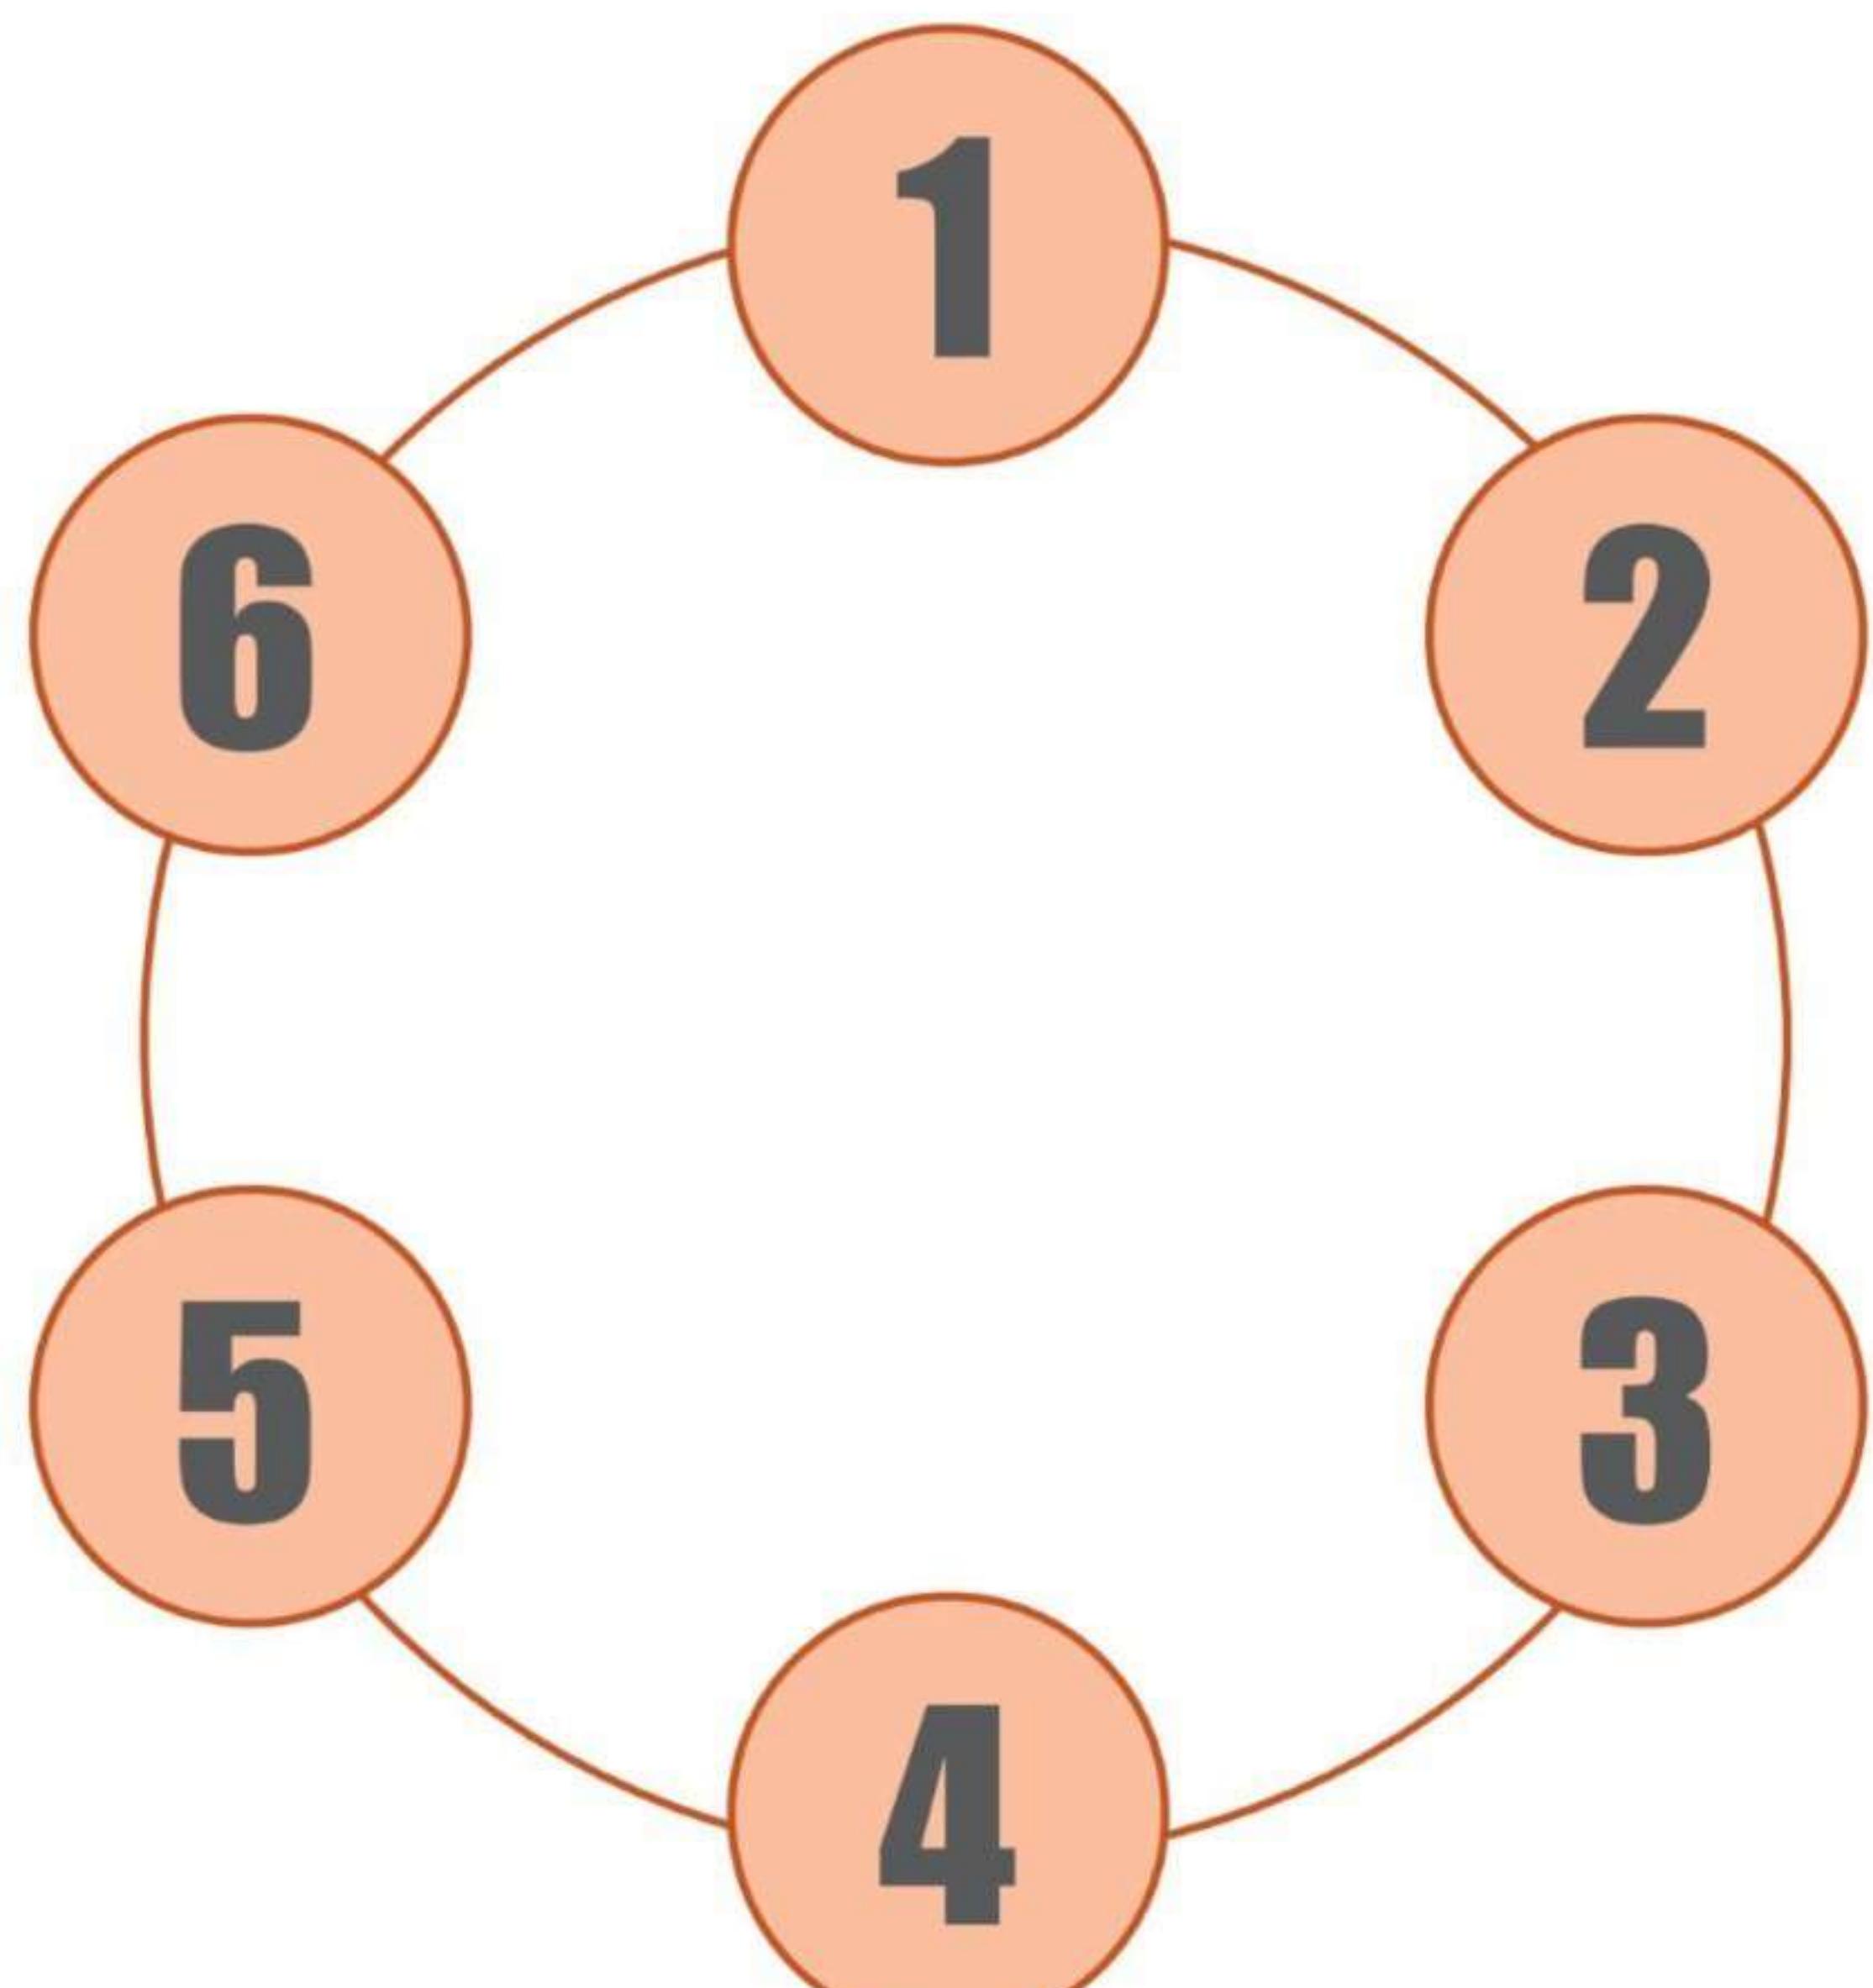


Range of all possible token values

1	2	3	4
2	4	3	1
4	1	2	3
3	2	4	1

- 2^{63} to $+2^{63}$

Storing Data with Virtual Nodes

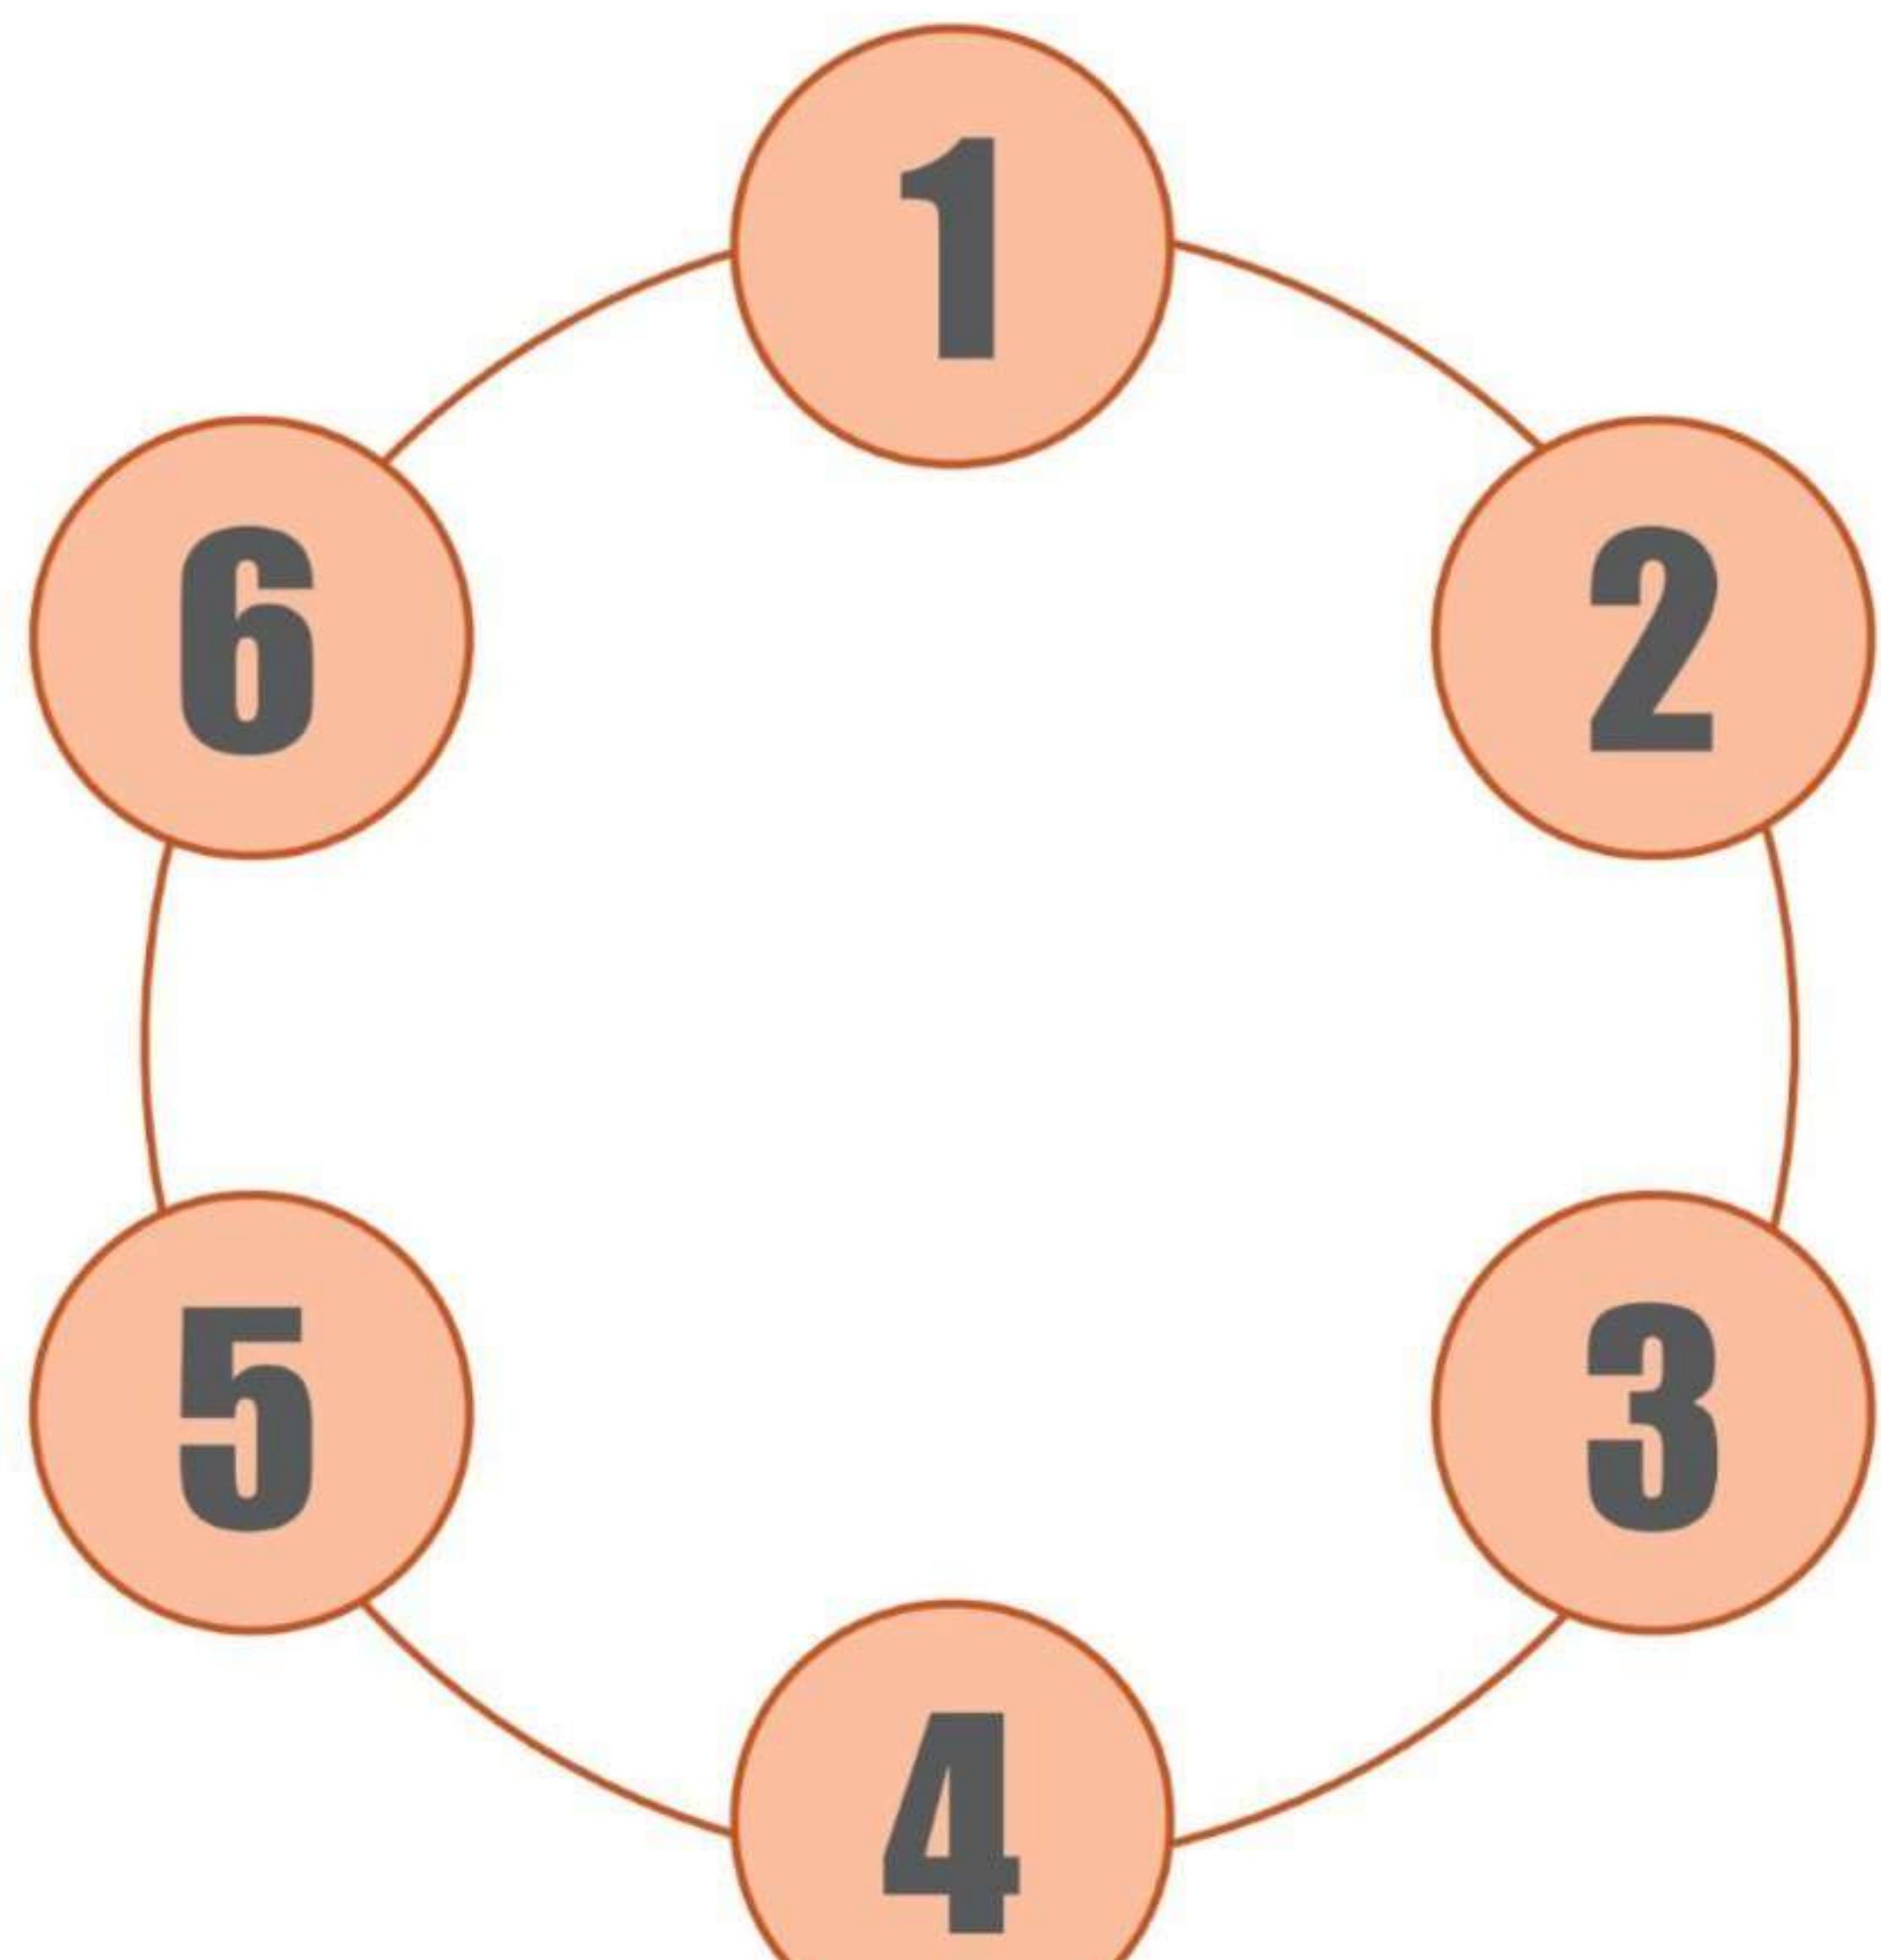


Range of all possible token values

1	2	3	5	4
5	2	4	3	1
4	1	2	3	5
3	5	2	4	1

- 2^{63} to $+2^{63}$

Storing Data with Virtual Nodes

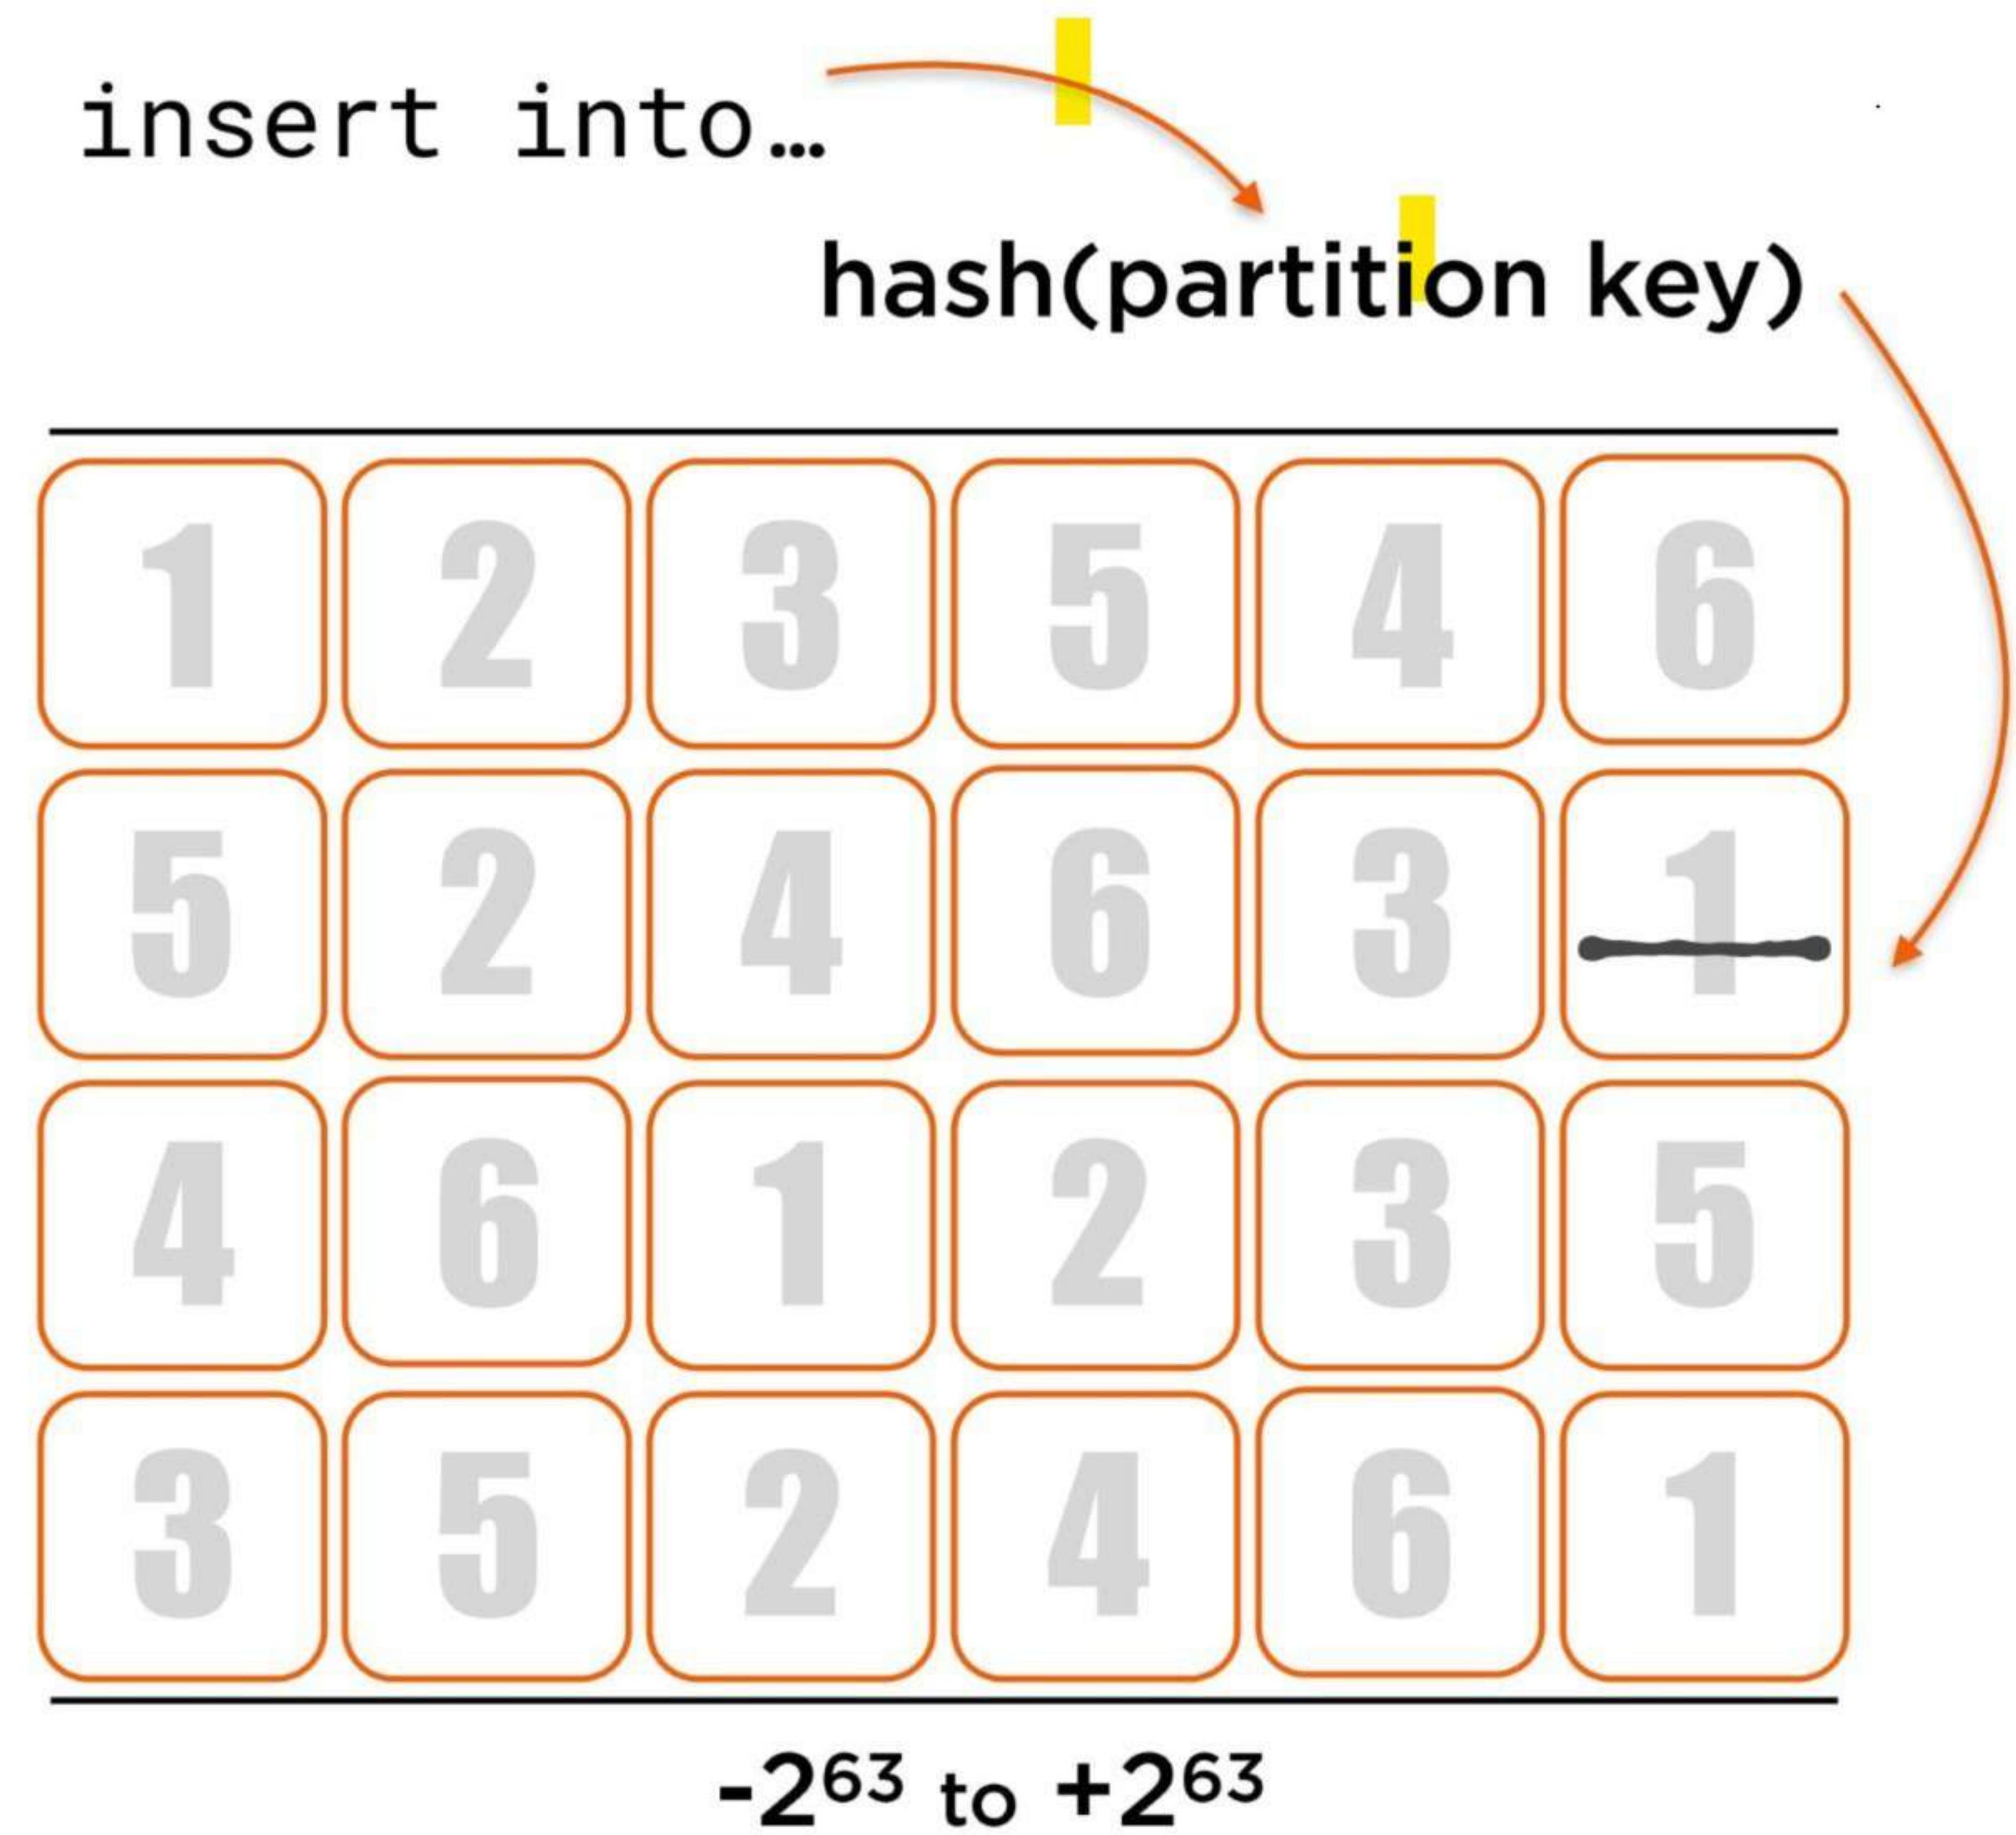
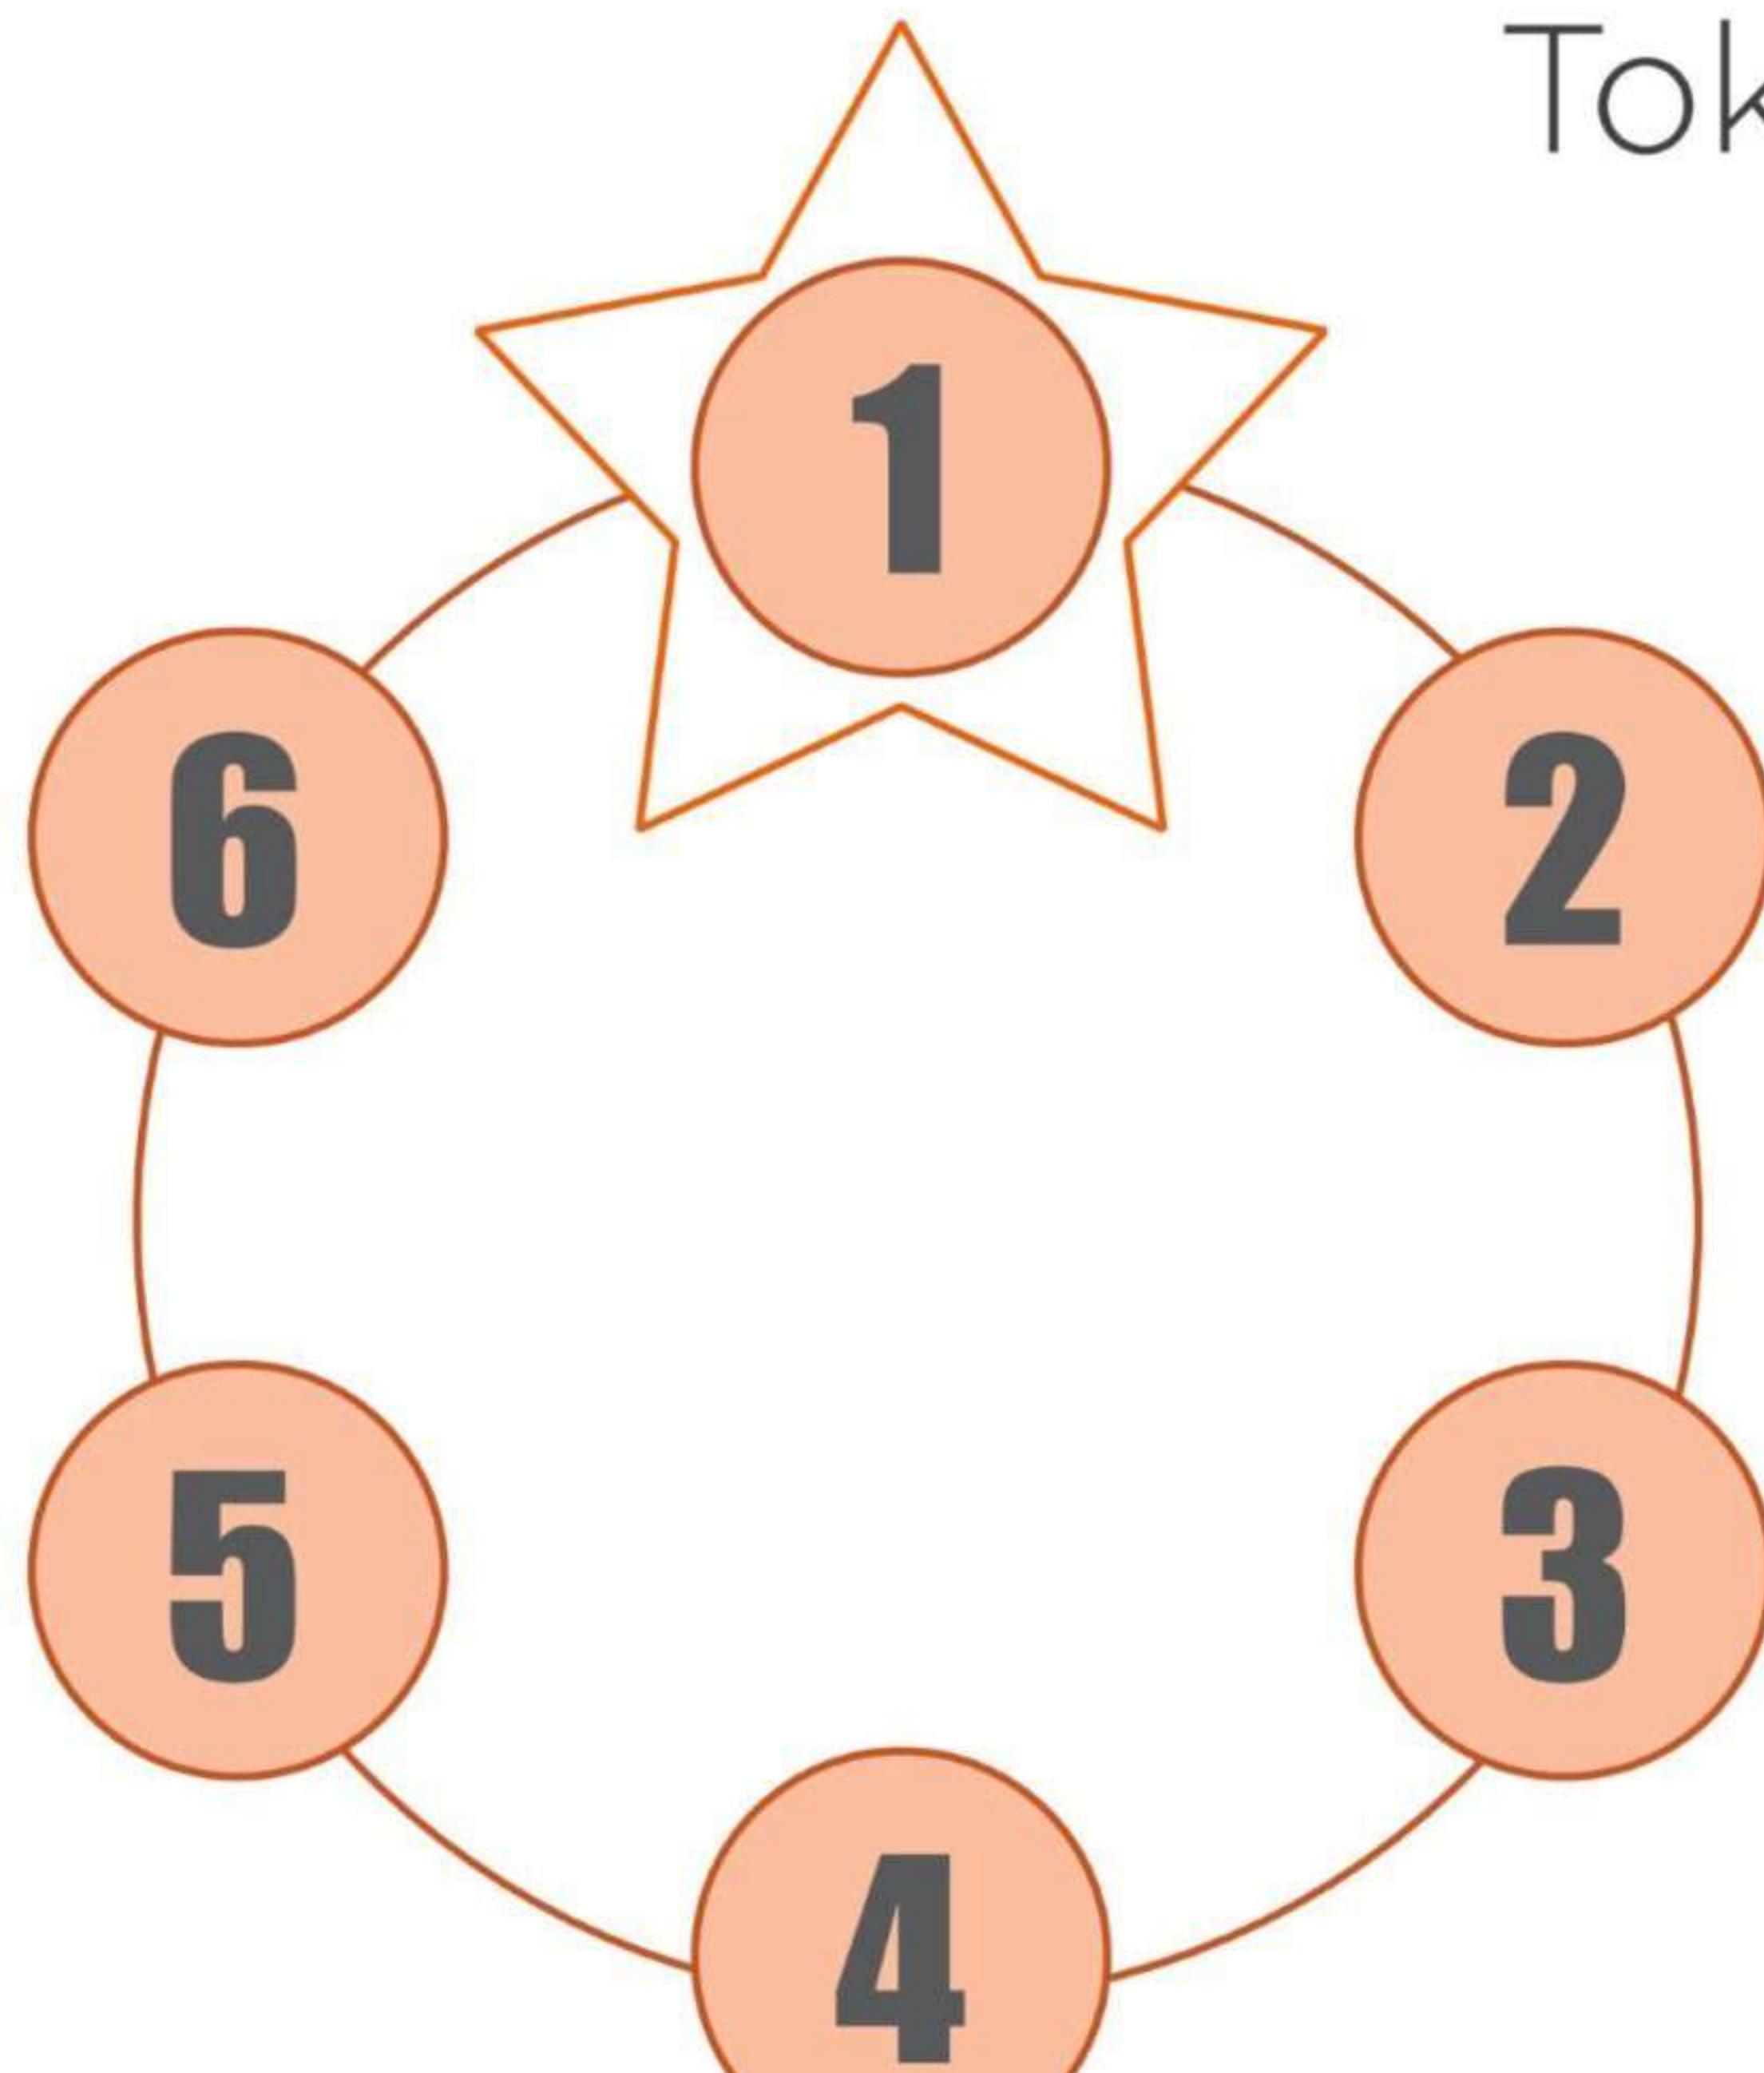


Range of all possible token values

1	2	3	5	4	6
5	2	4	6	3	1
4	6	1	2	3	5
3	5	2	4	6	1

- 2^{63} to $+2^{63}$

Token Allocation via the Partitioner



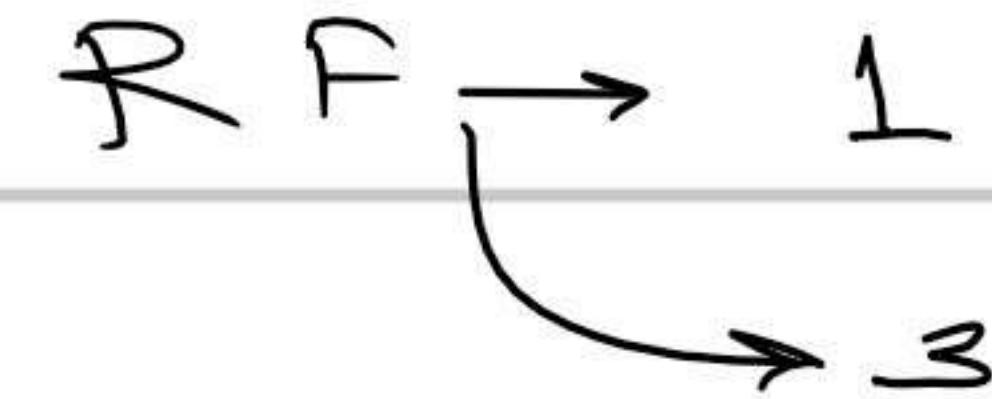
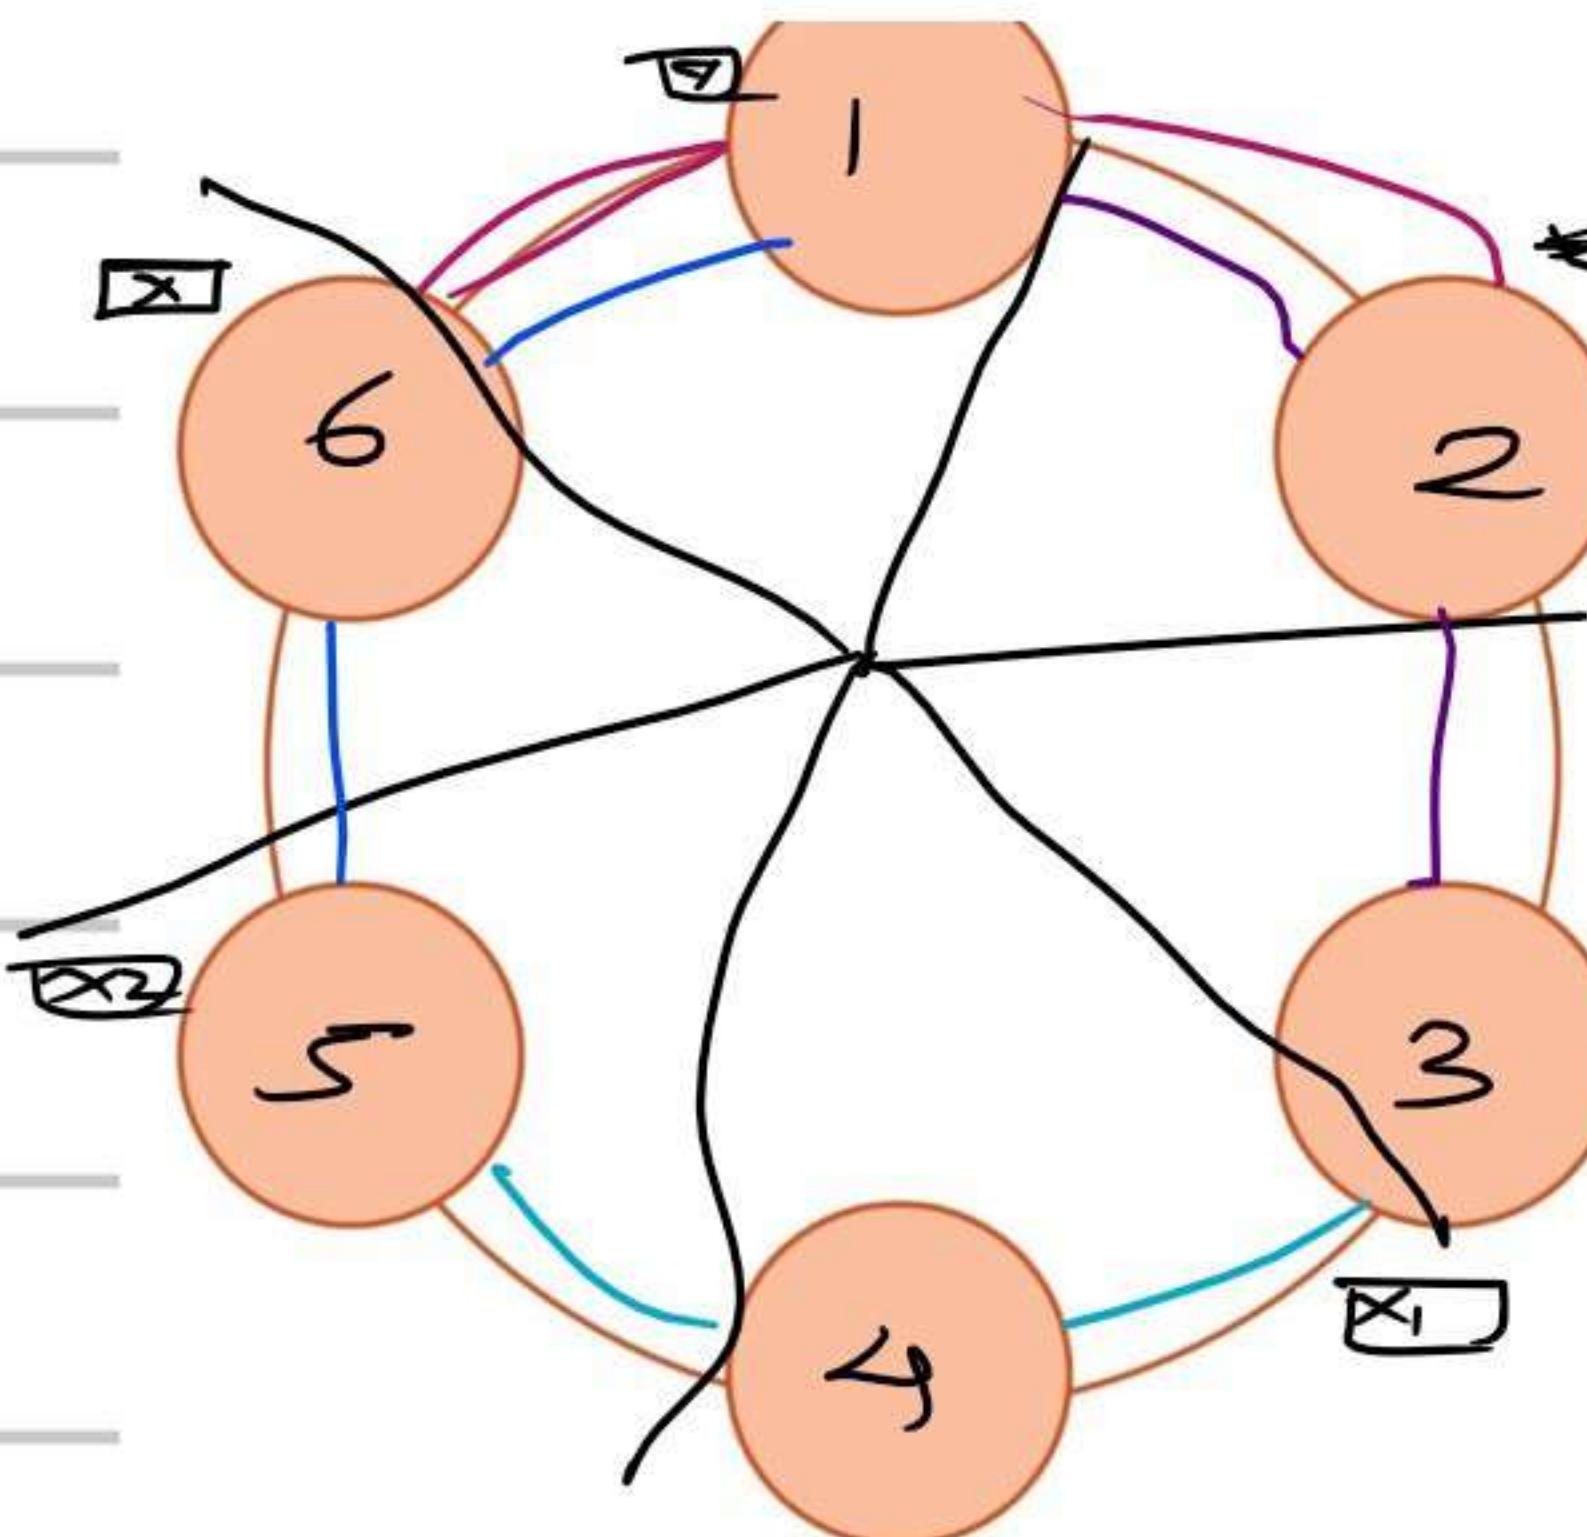
Replication Strategy → Simple

→ Once the token value is identified then the records are also stored in the next 2 adjacent nodes

$$1 N \rightarrow 256 GB - 8 DN$$

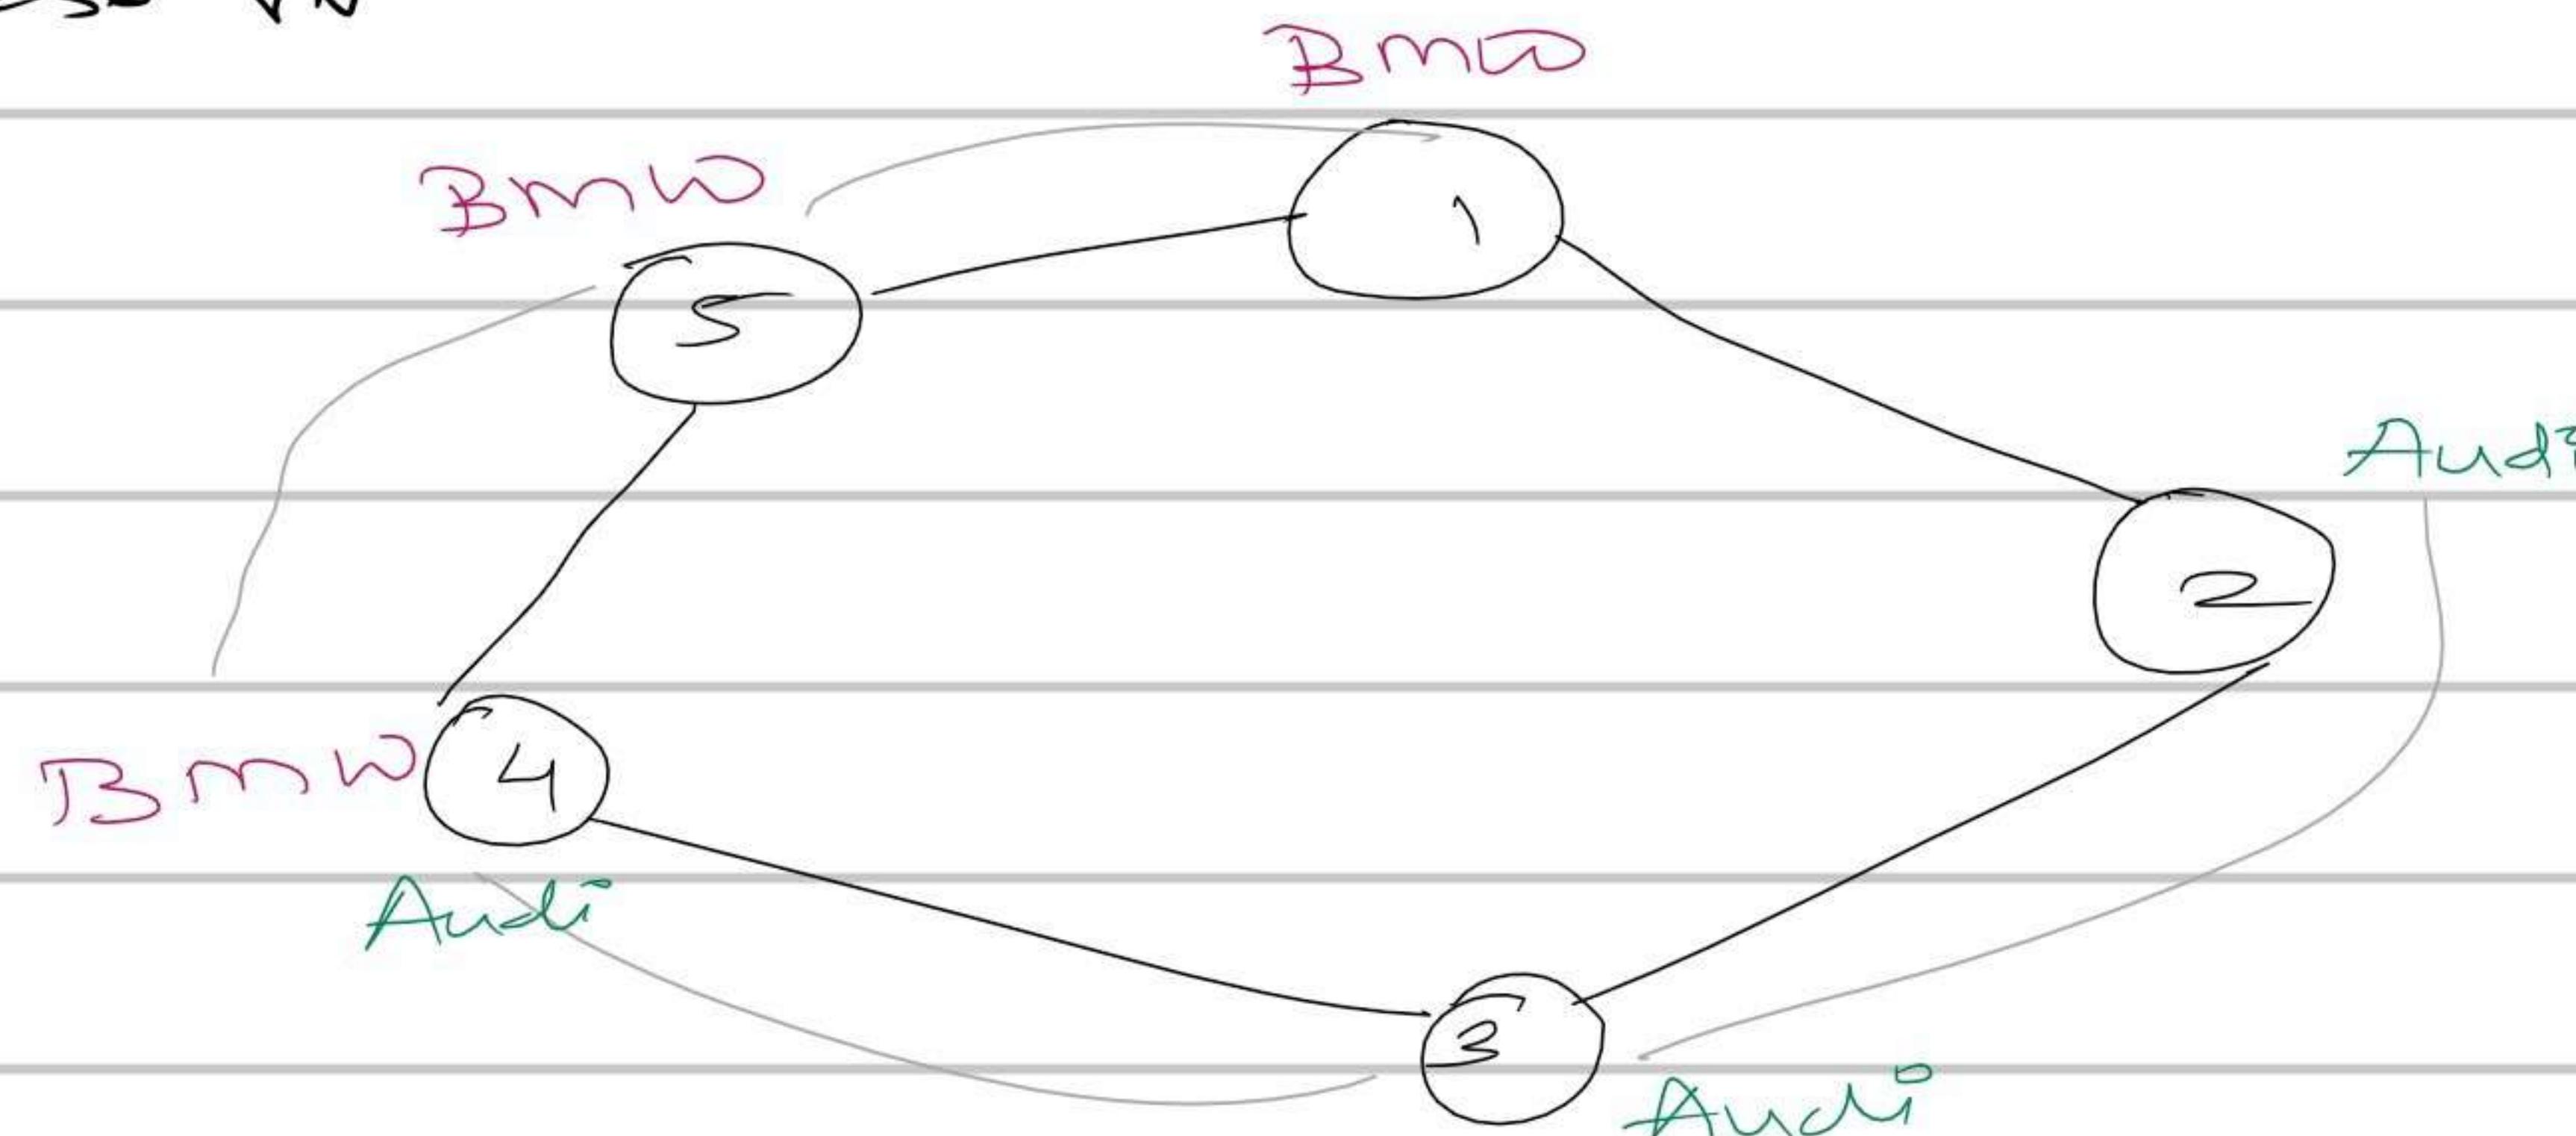
$$64 GB - 4 VM$$

$$1 N \rightarrow 256 VR$$

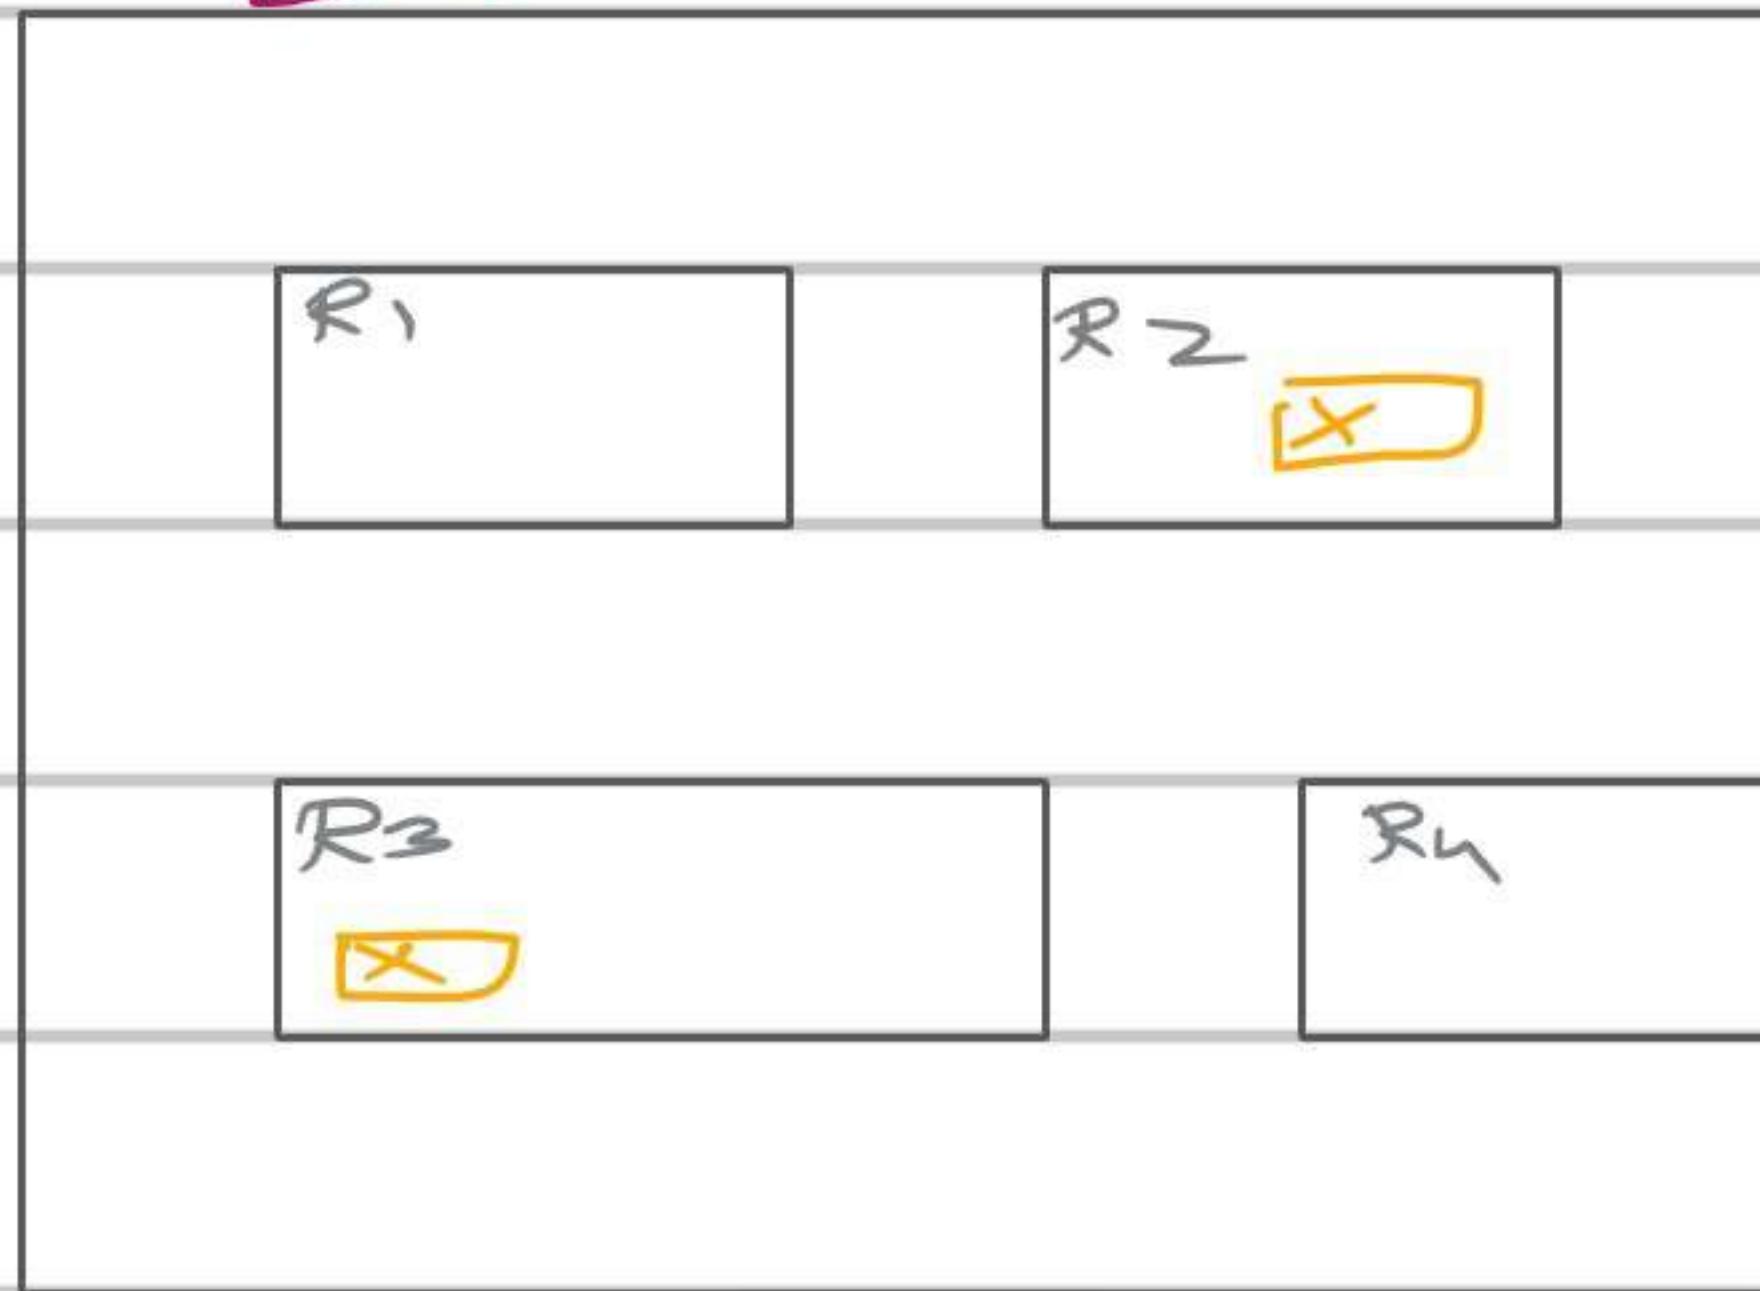


Cassandra 3 → 256
Cassandra 5 → 16

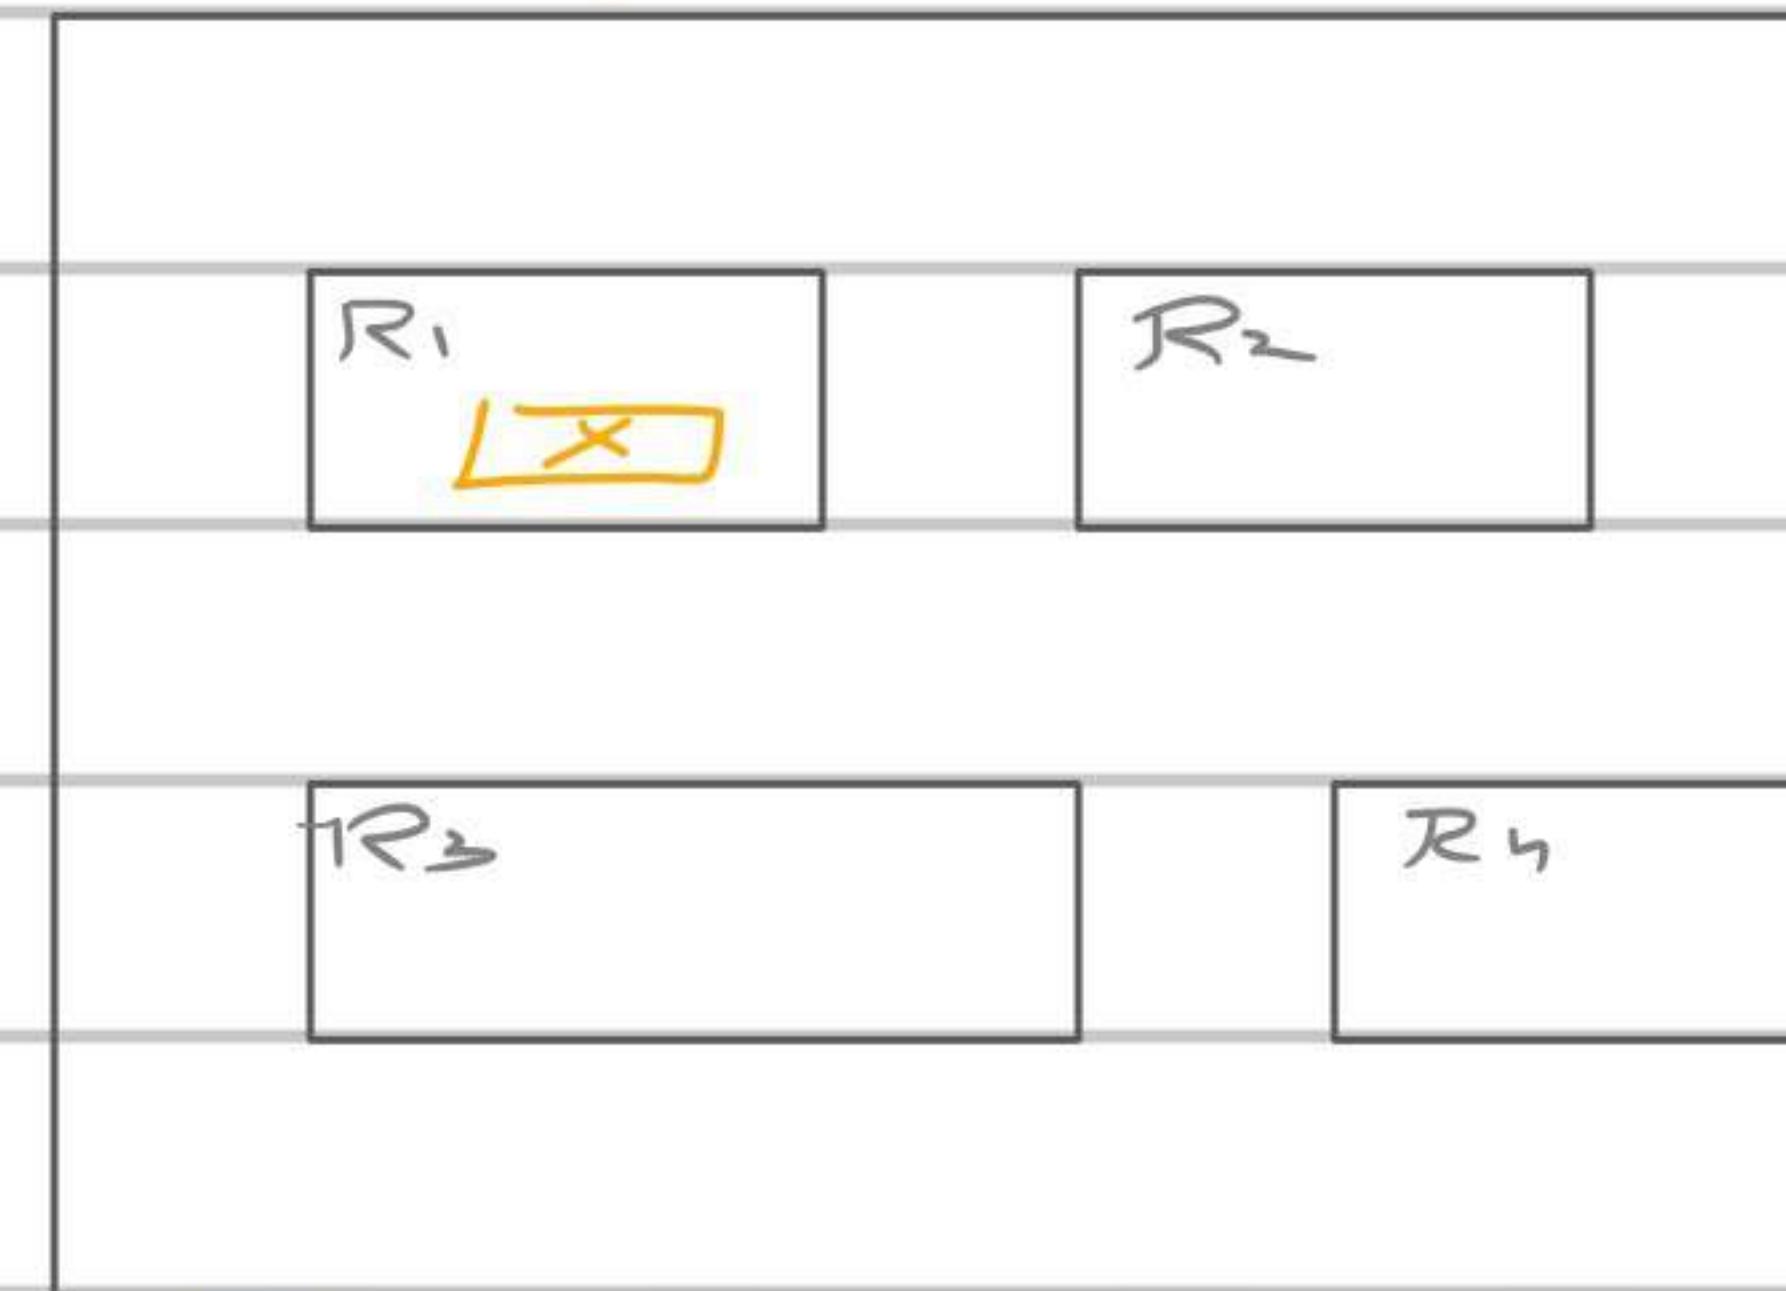
One DC \in
One Rack



EO



IN



Strategy - Network Topology.

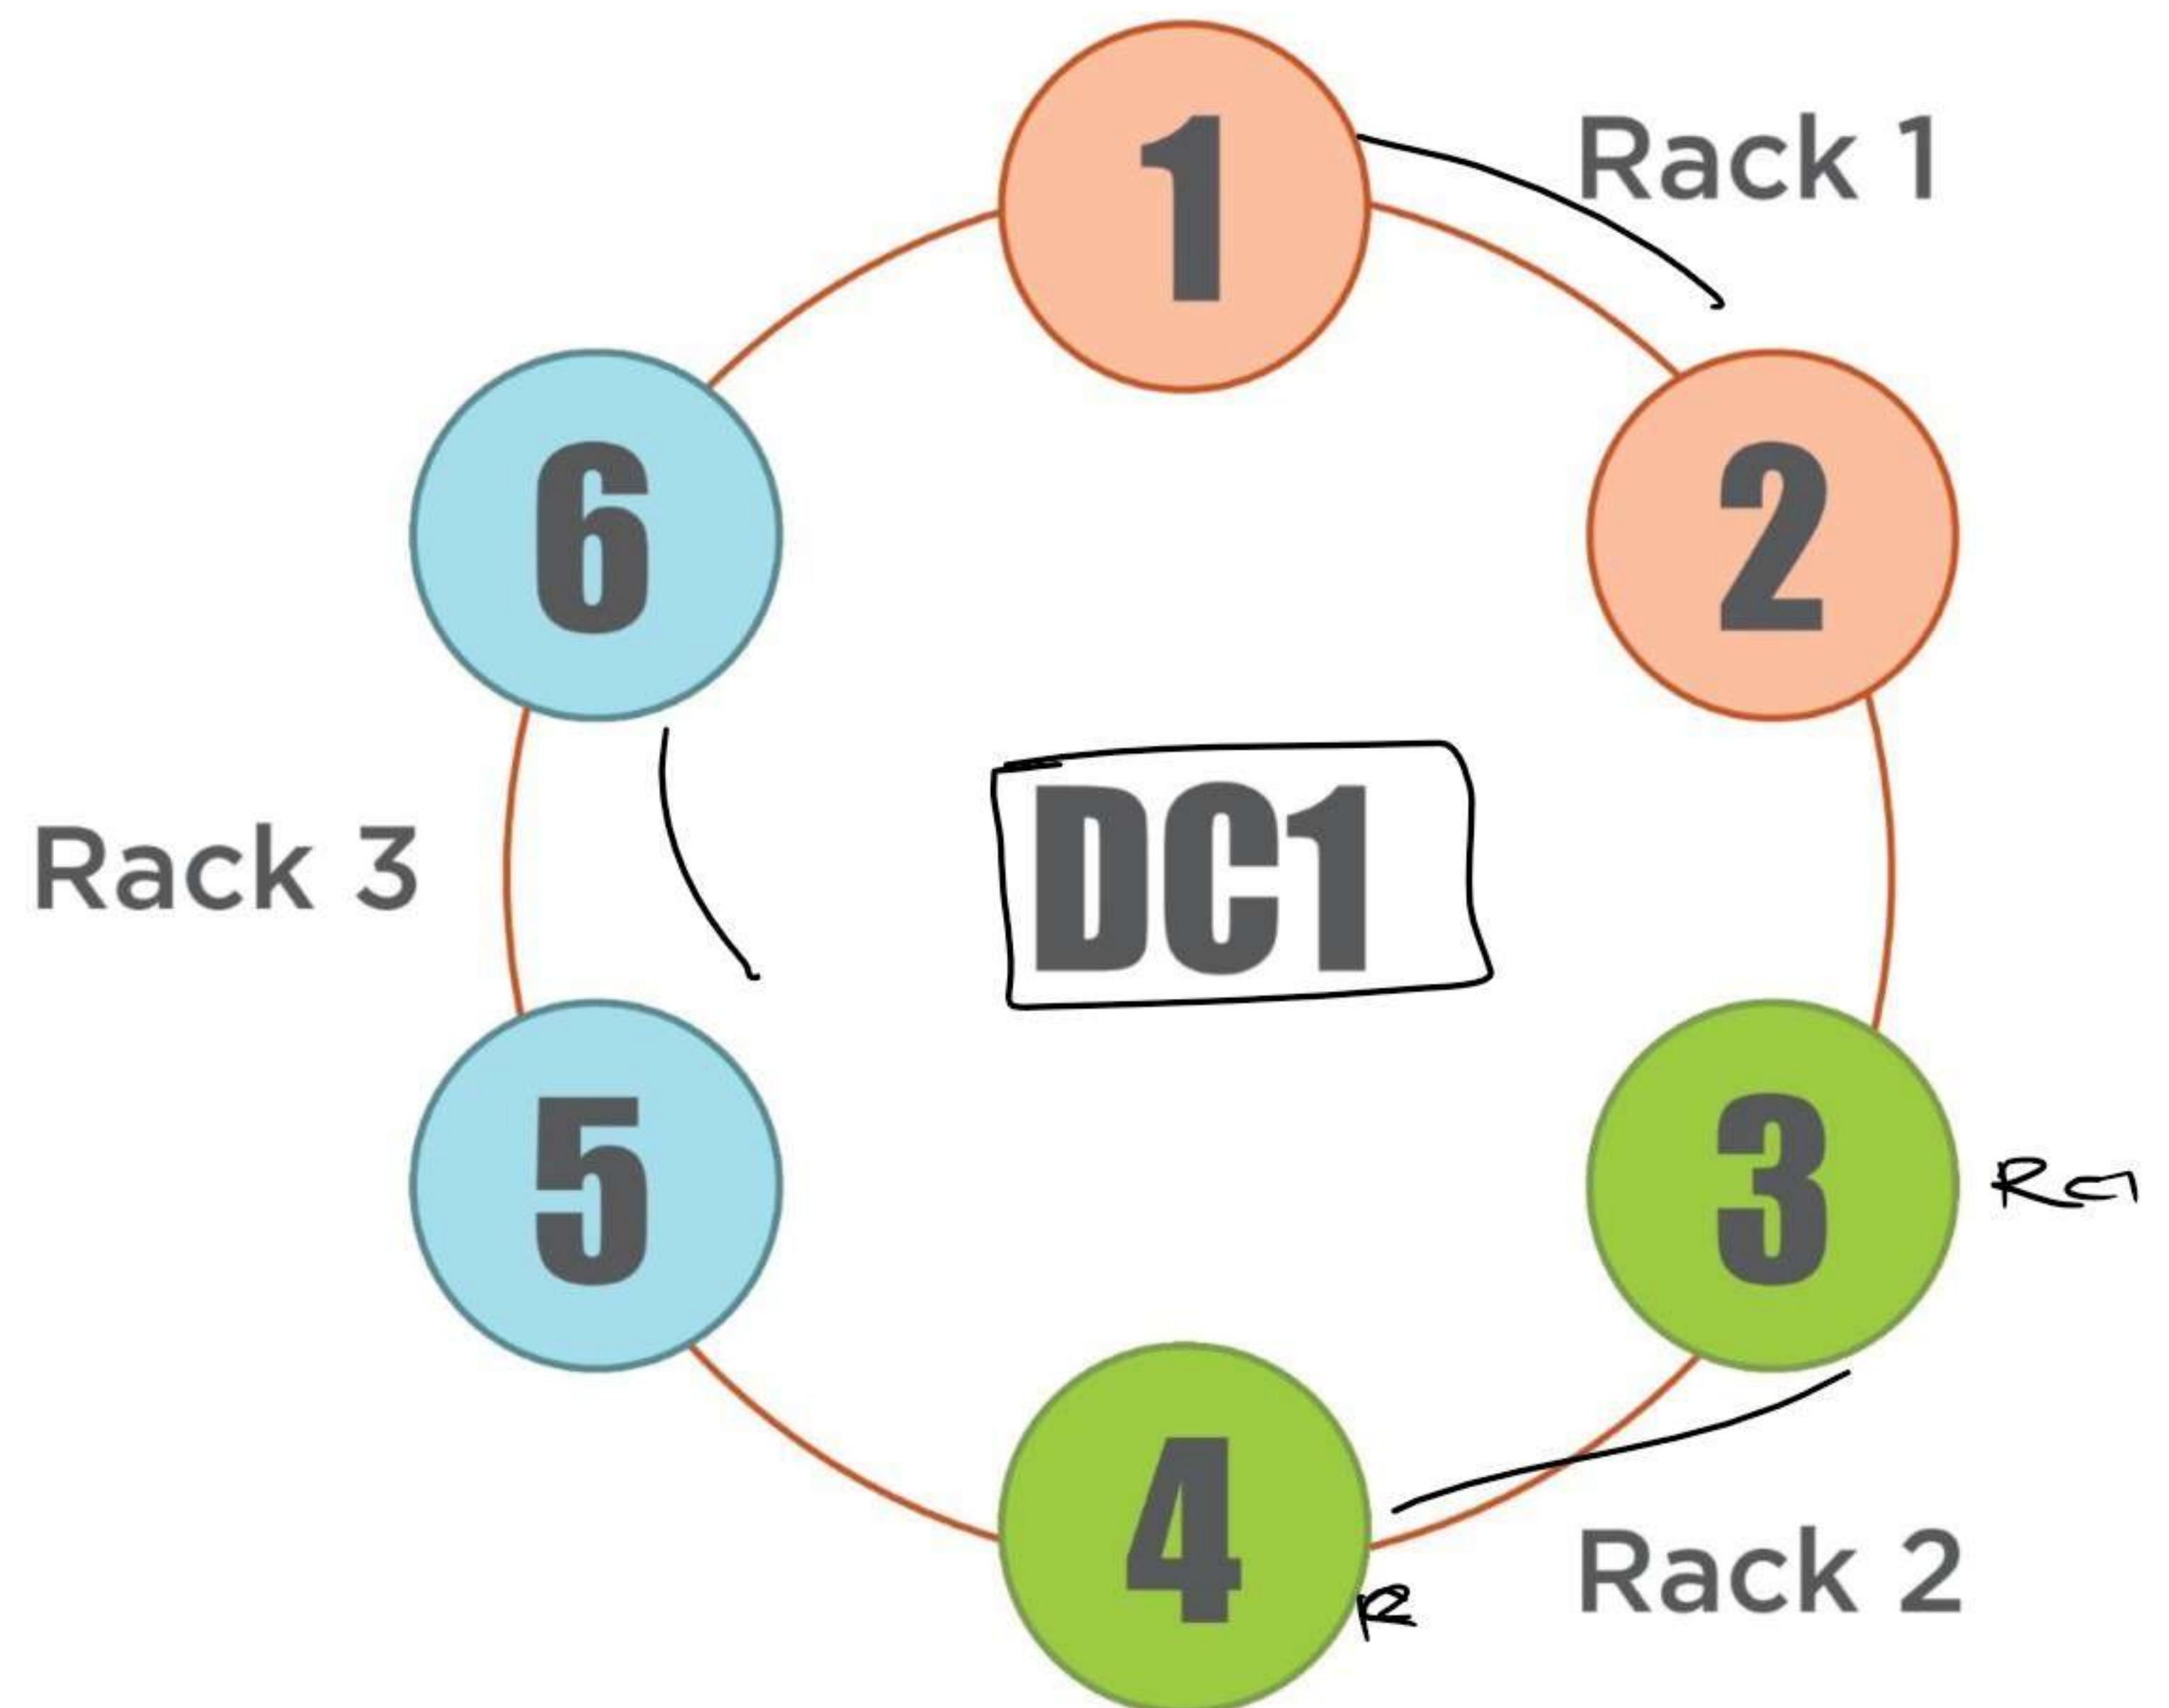
$RF \rightarrow 3$

First copy \Rightarrow rack which is finalized

Second copy on some other data
racks

Third one on any random data center

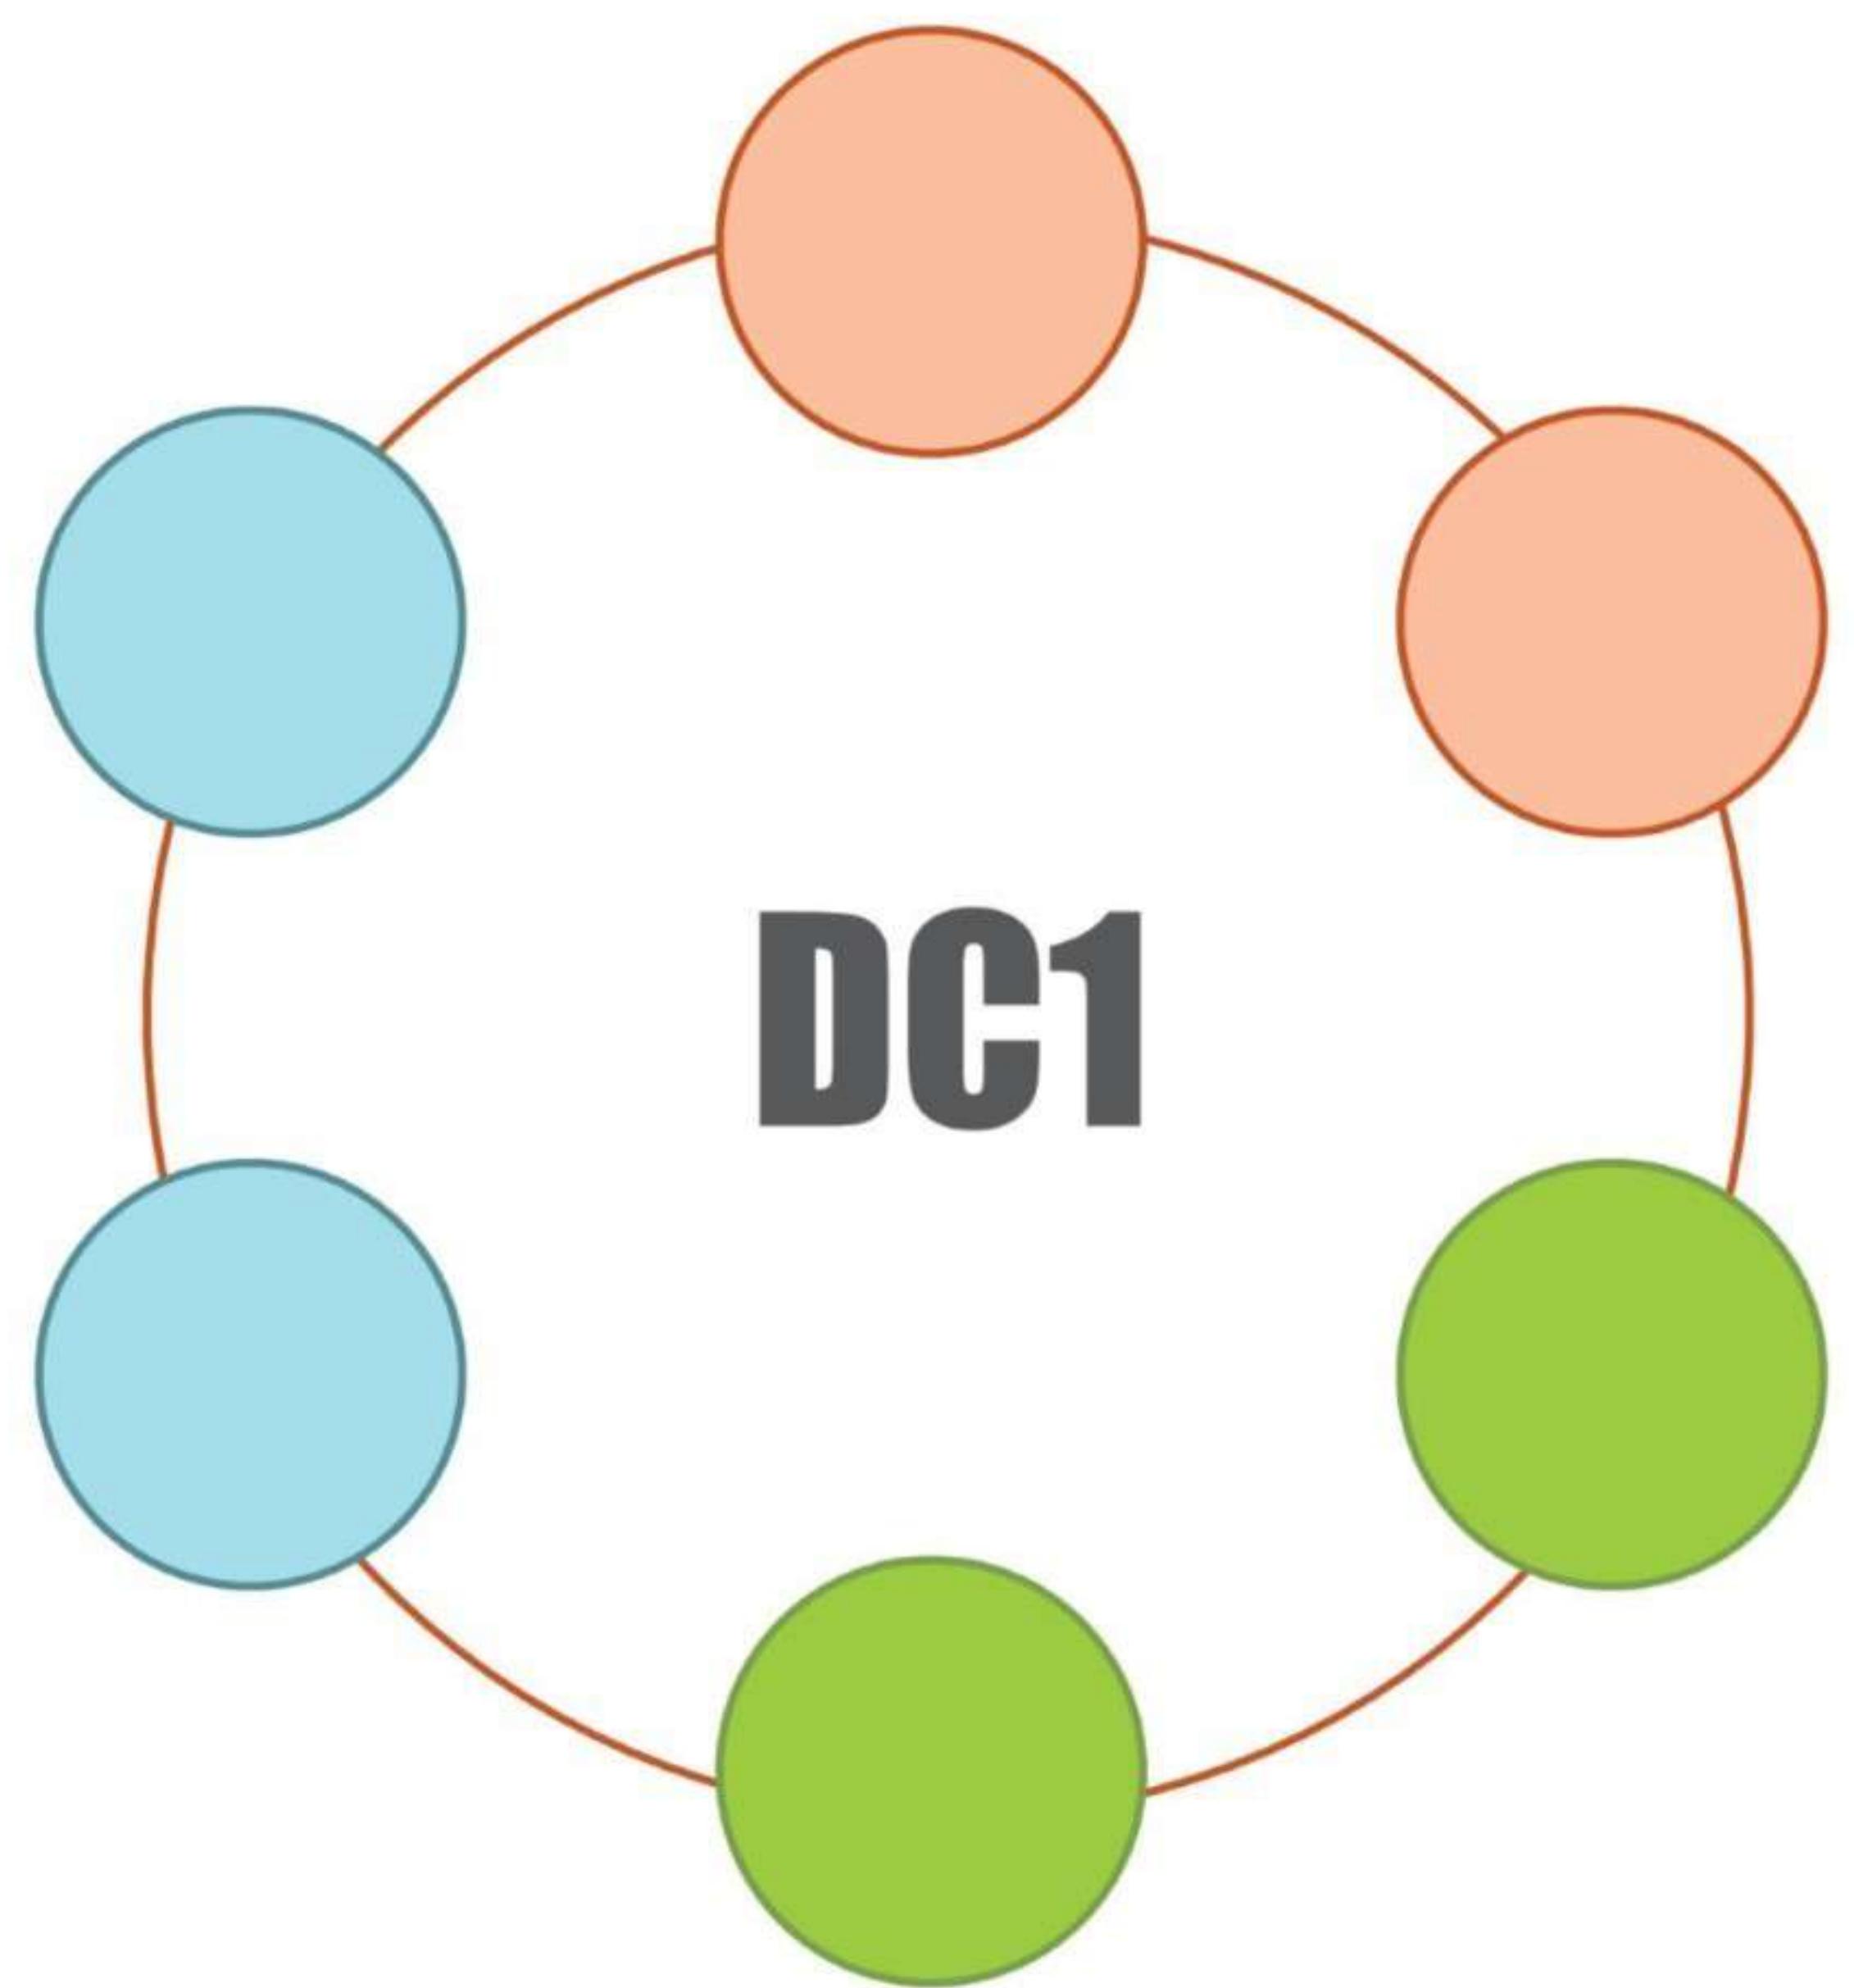
rack; which doesn't have it



GossipingPropertyFileSnitch

Snitch

SimpleSnitch

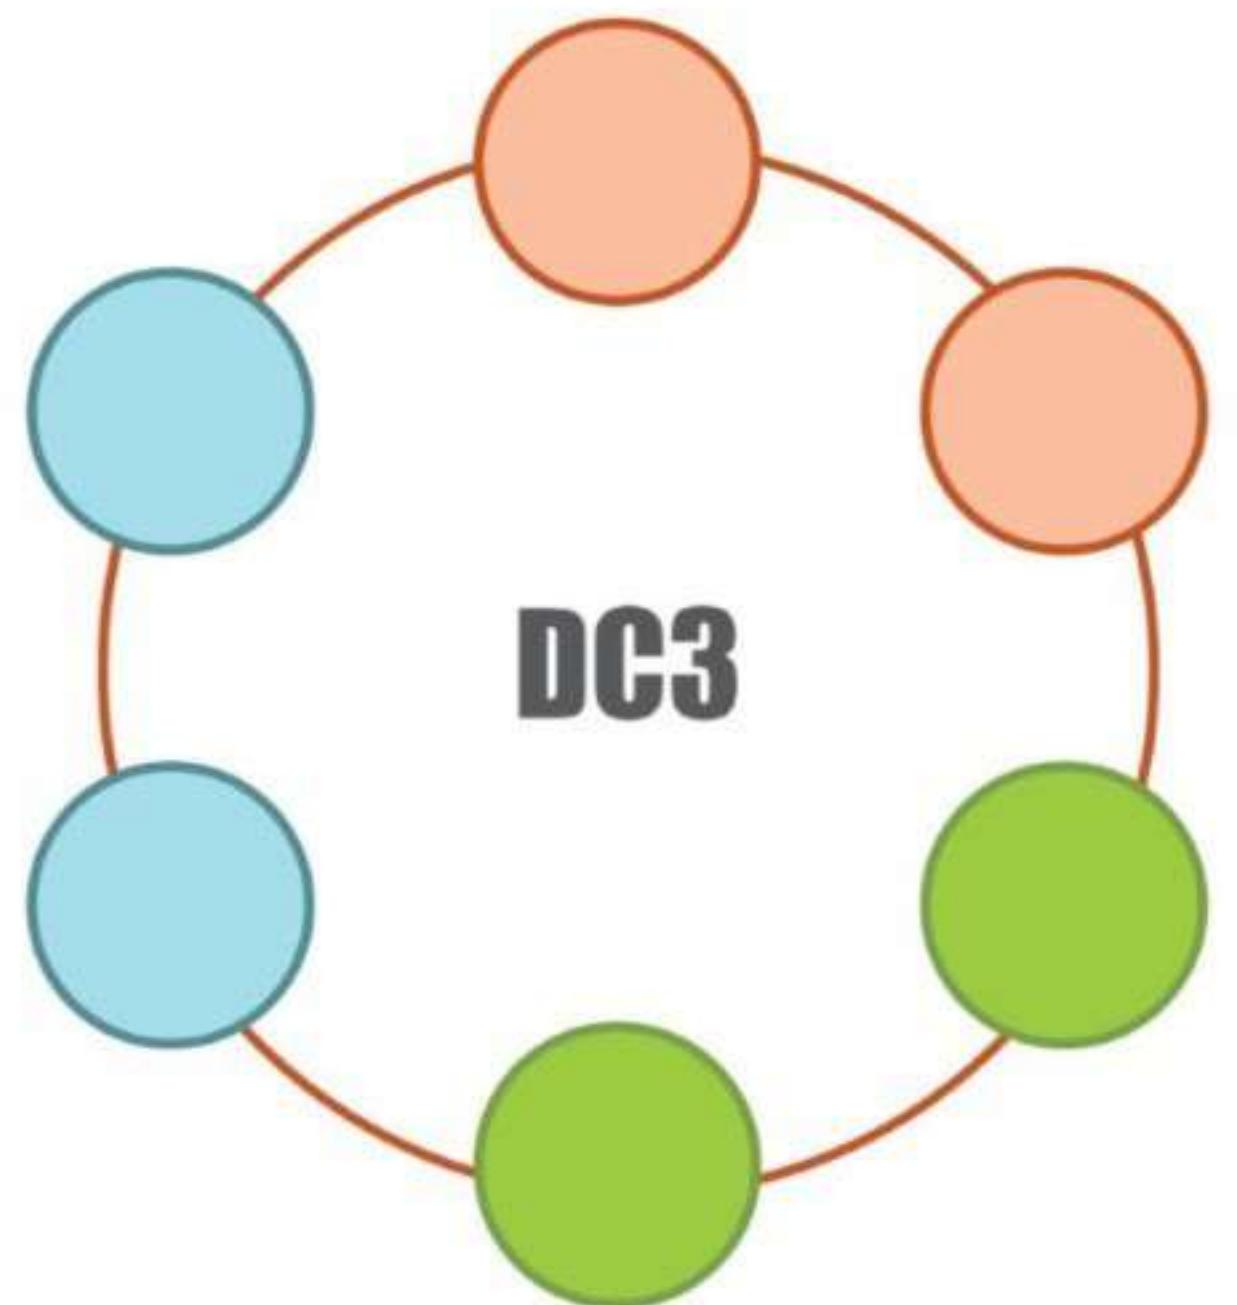
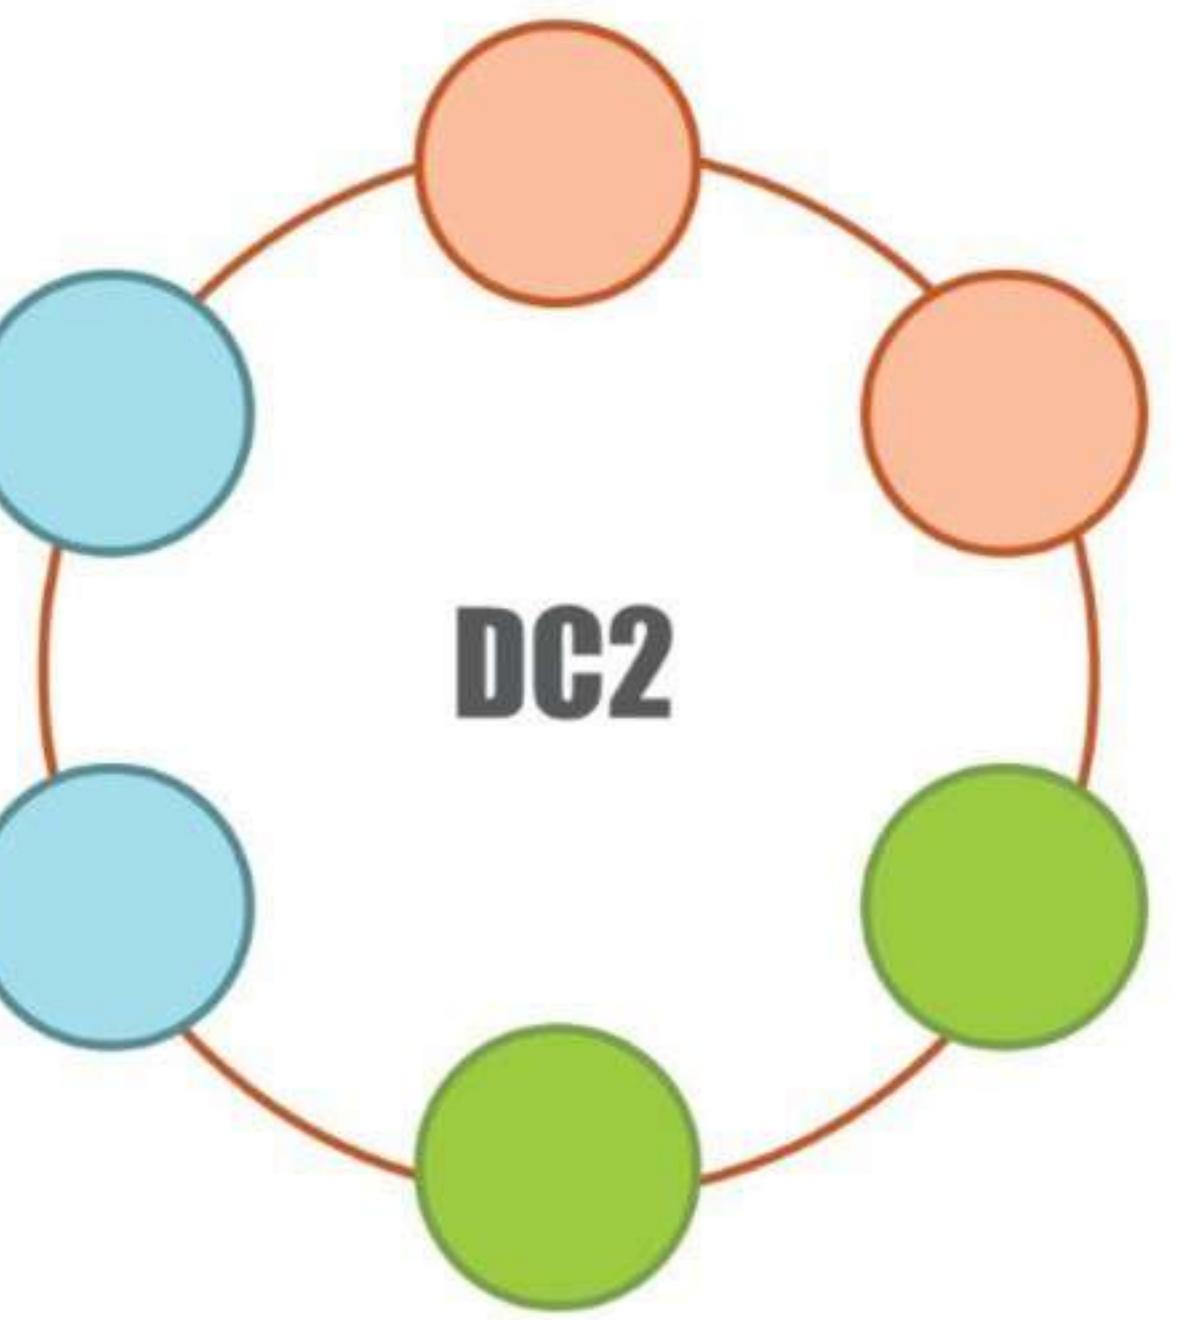
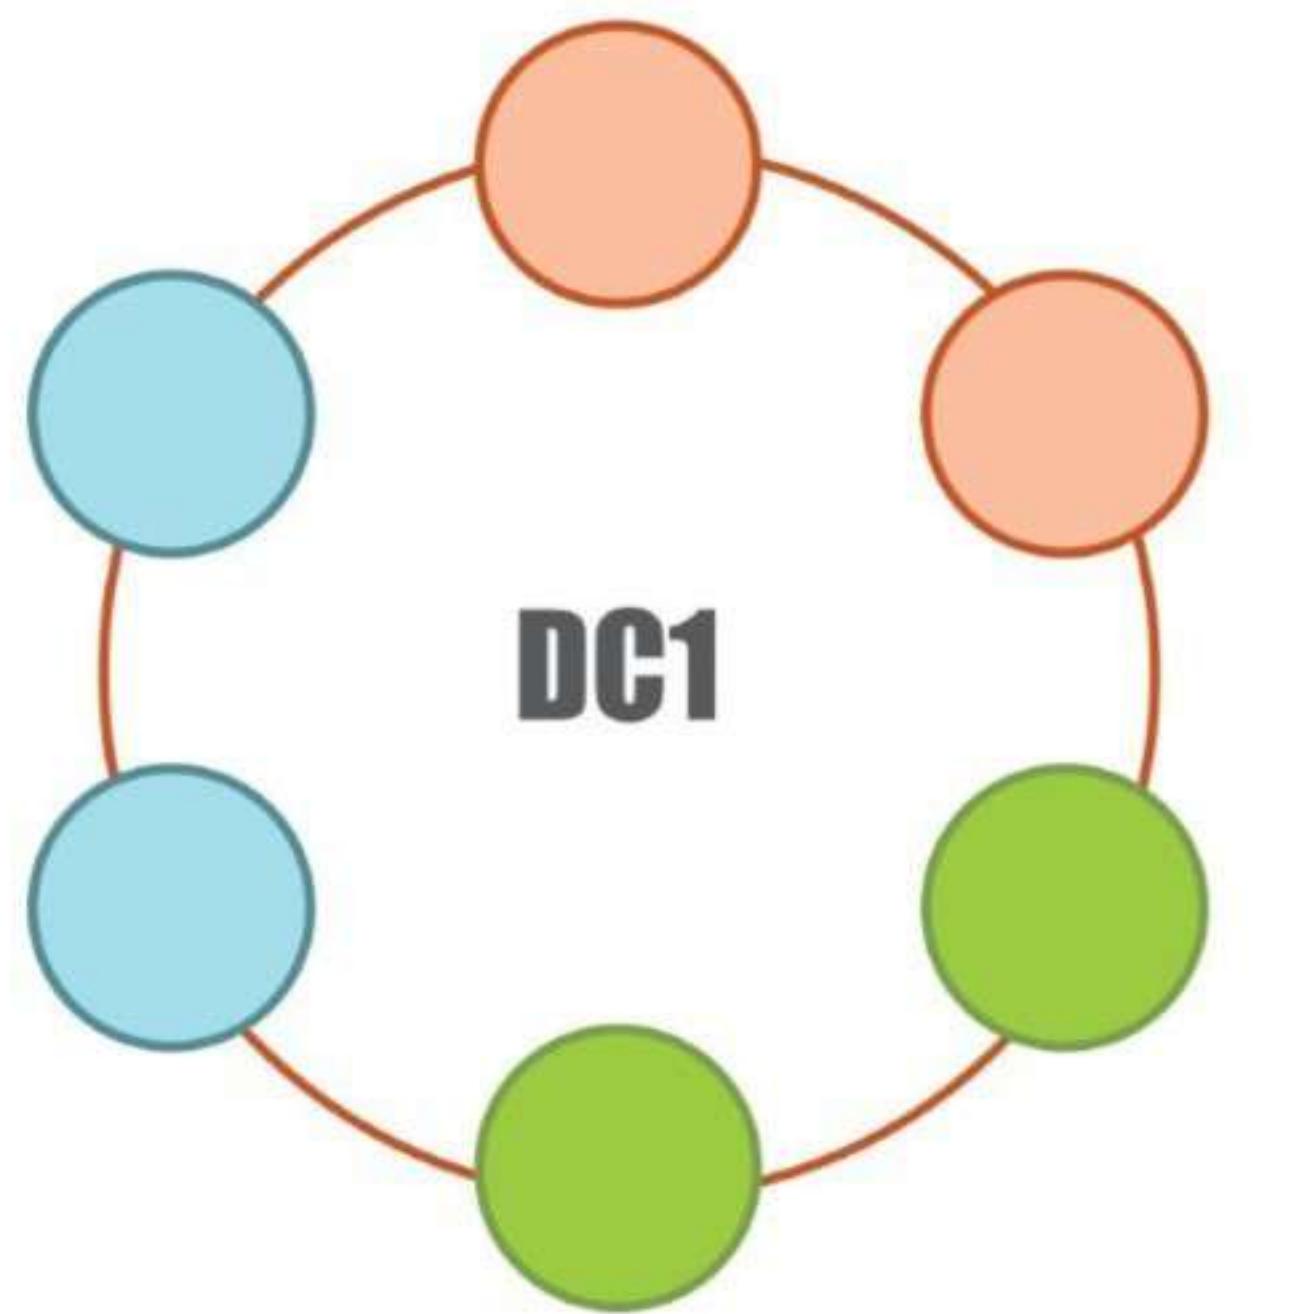


DC1

GossipingPropertyFileSnitch

SimpleSnitch

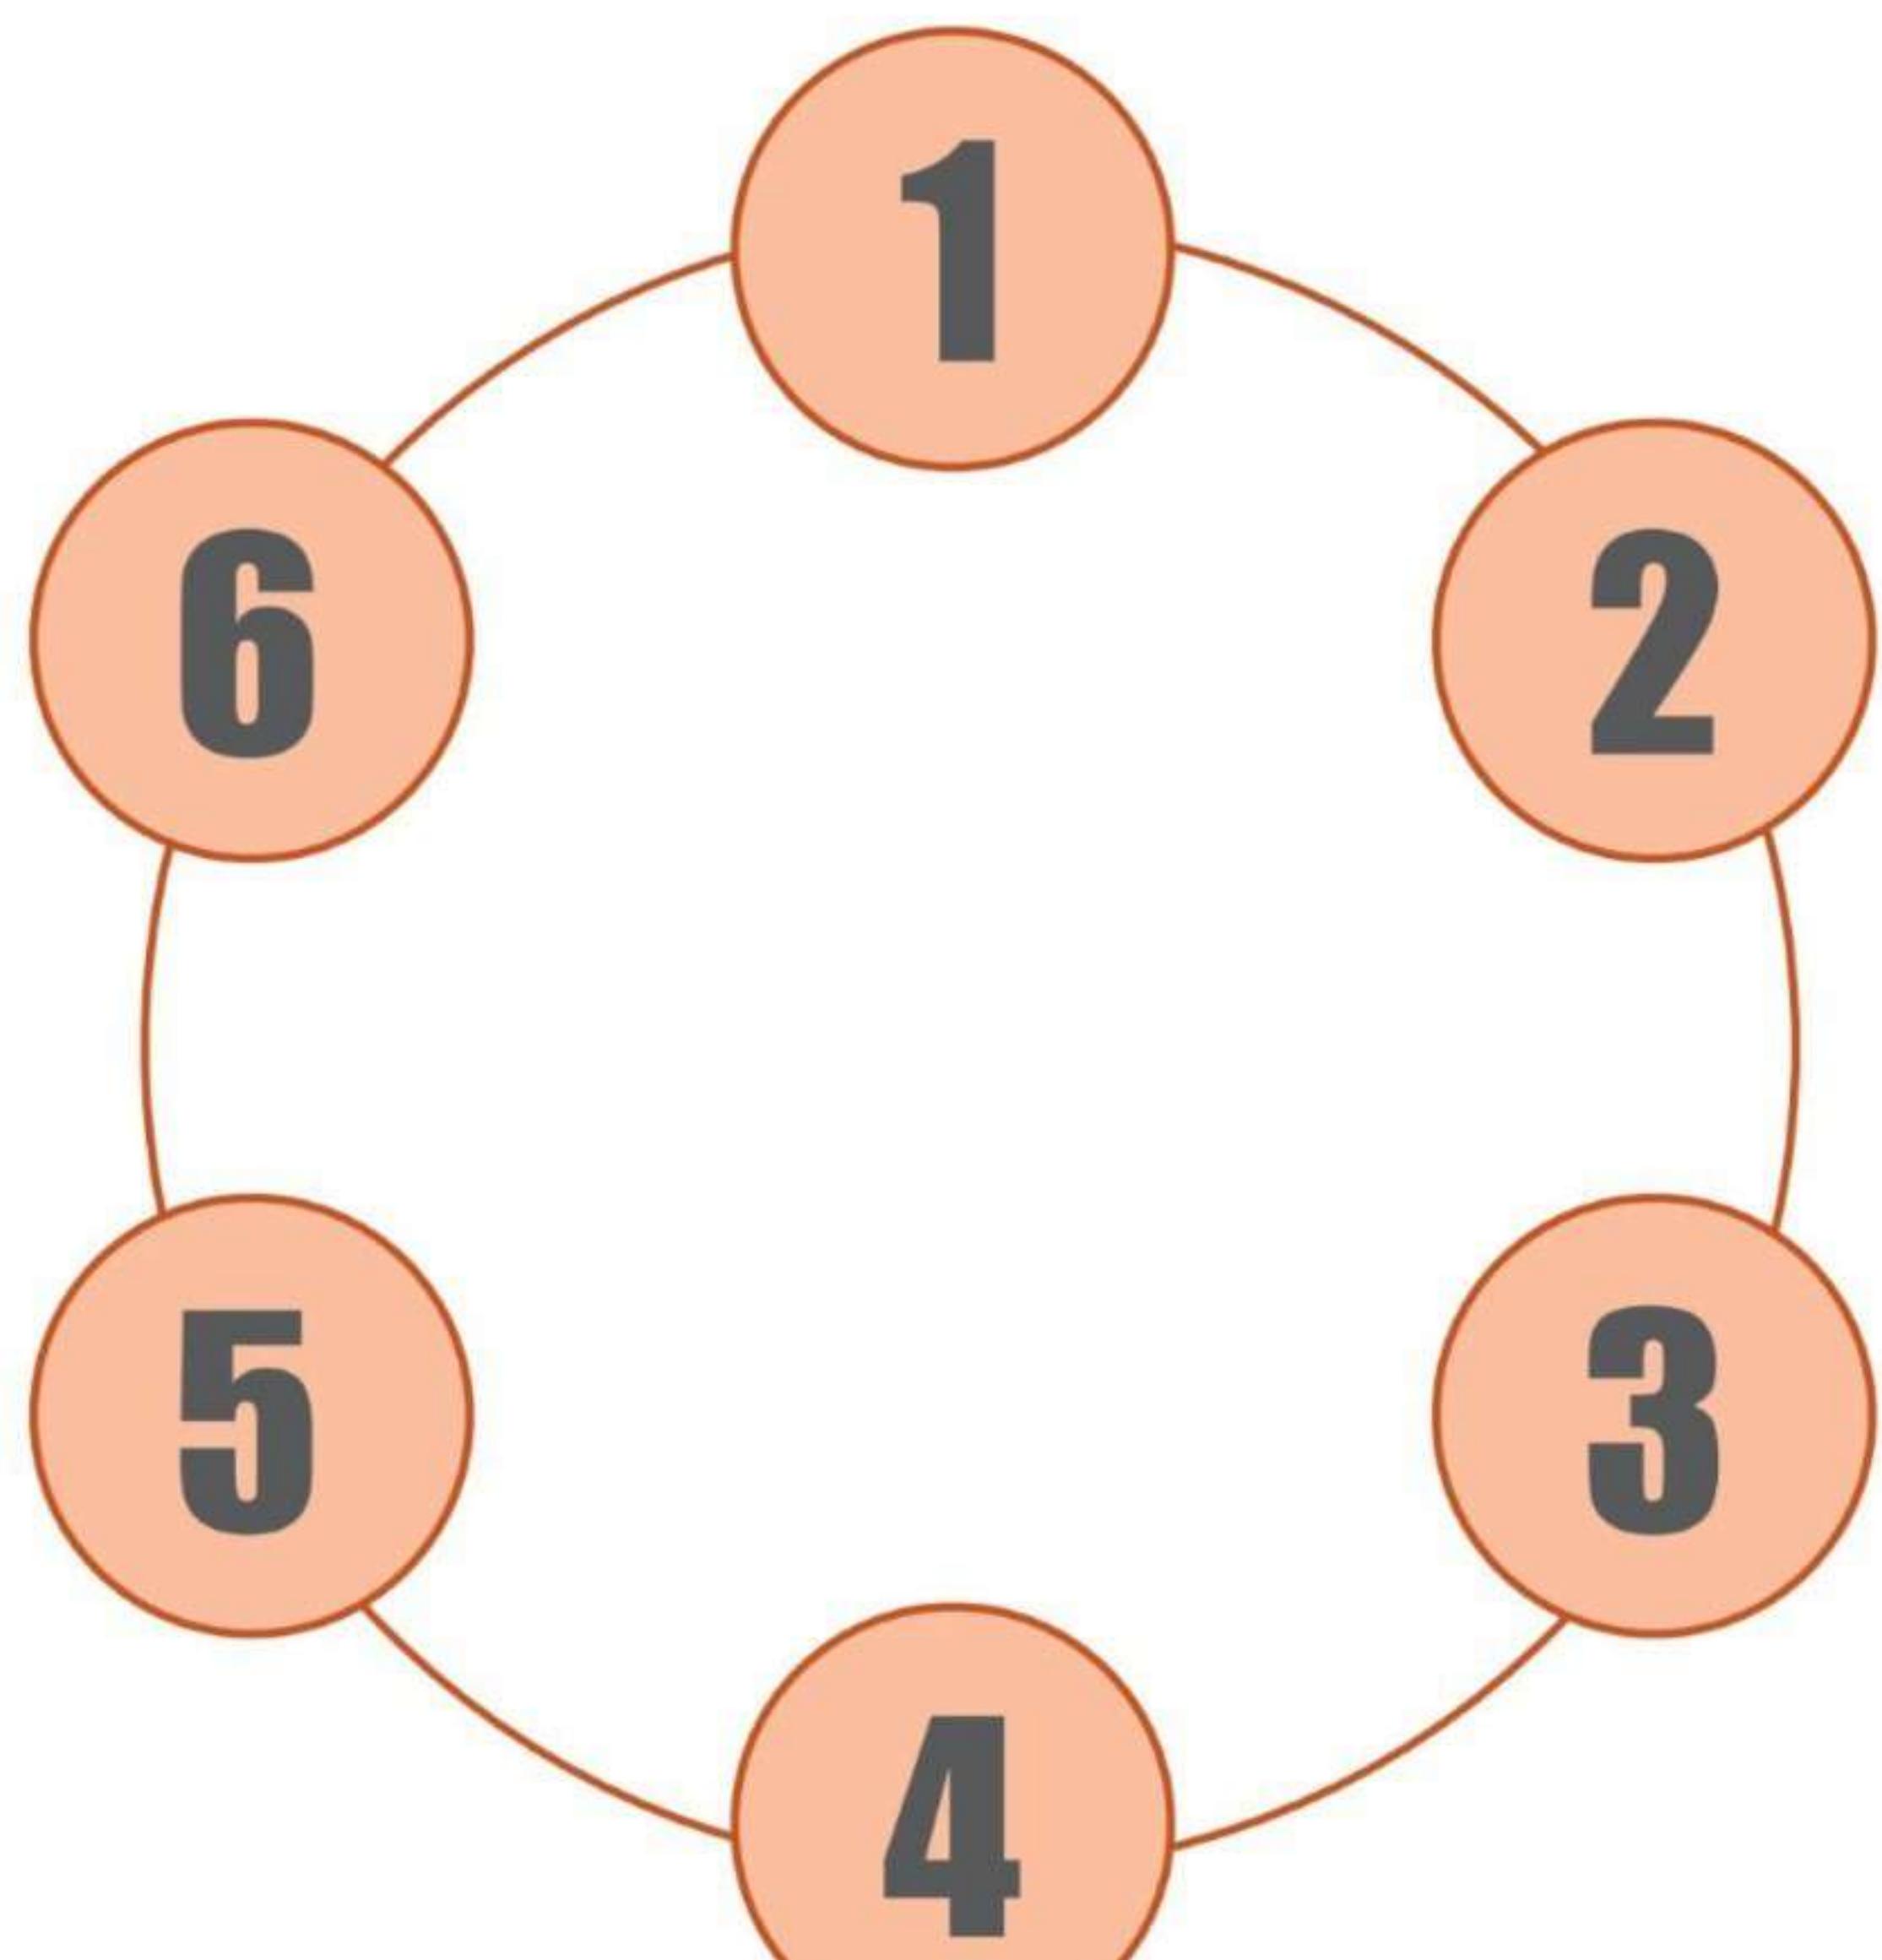
Snitch



Snitch
SimpleSnitch
GossipingPropertyFileSnitch
PropertyFileSnitch
EC2Snitch
EC2MultiRegionSnitch
RackInferringSnitch

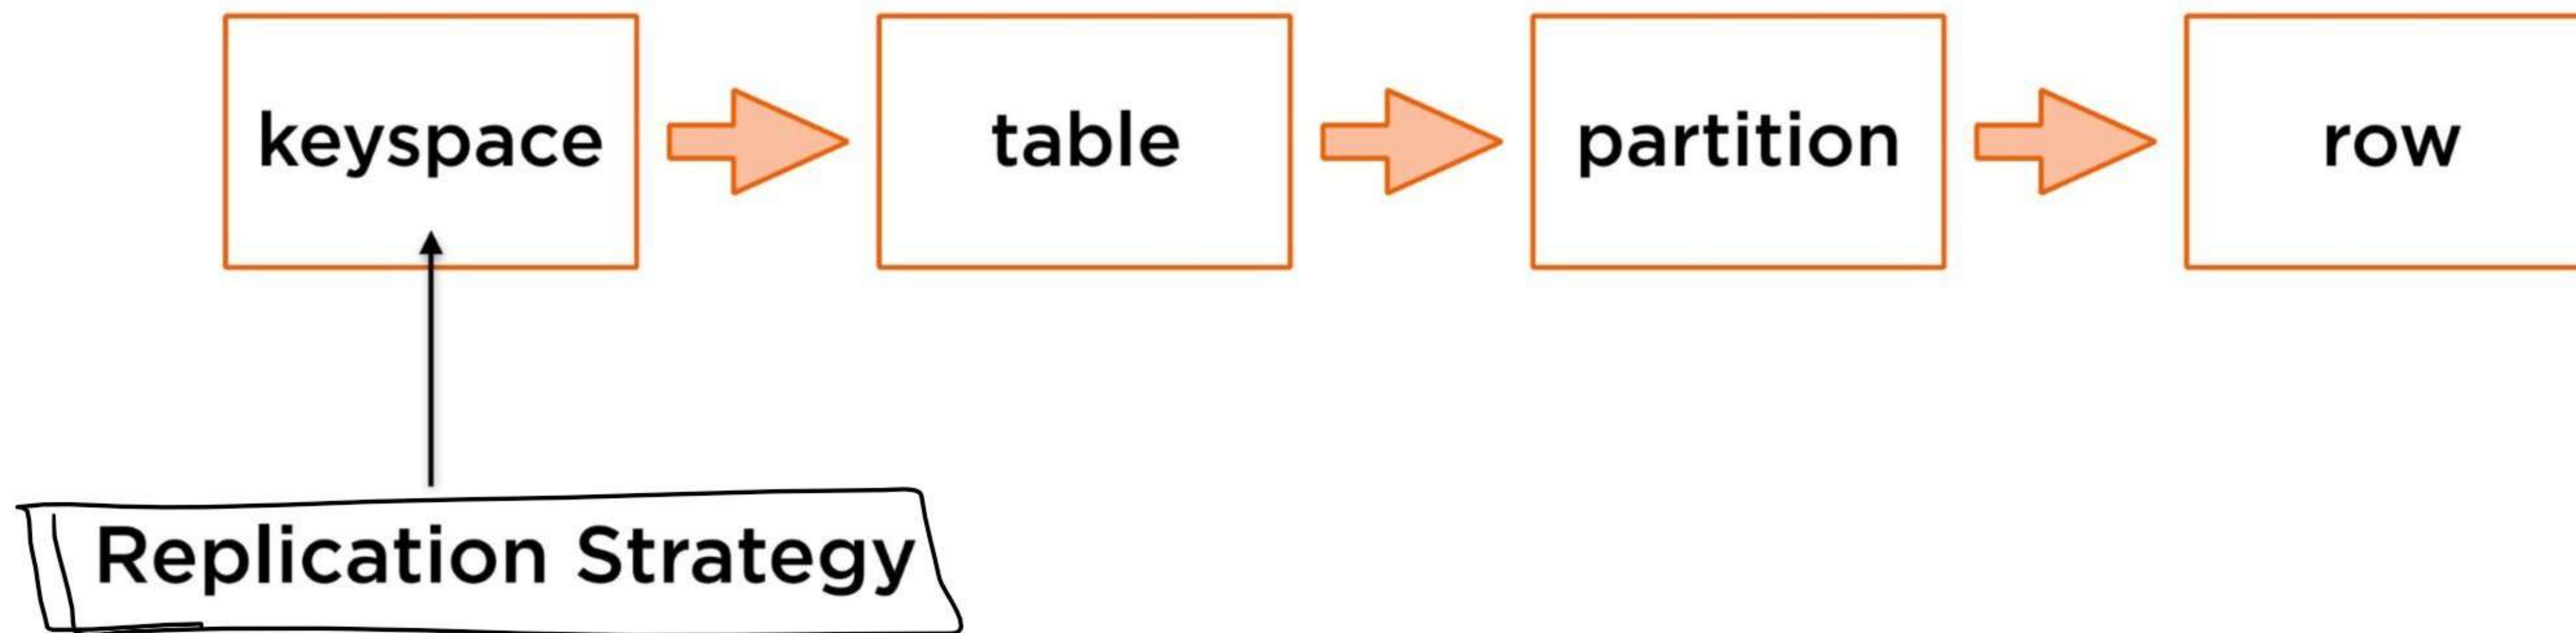
Replication Strategy and Consistency

Replication



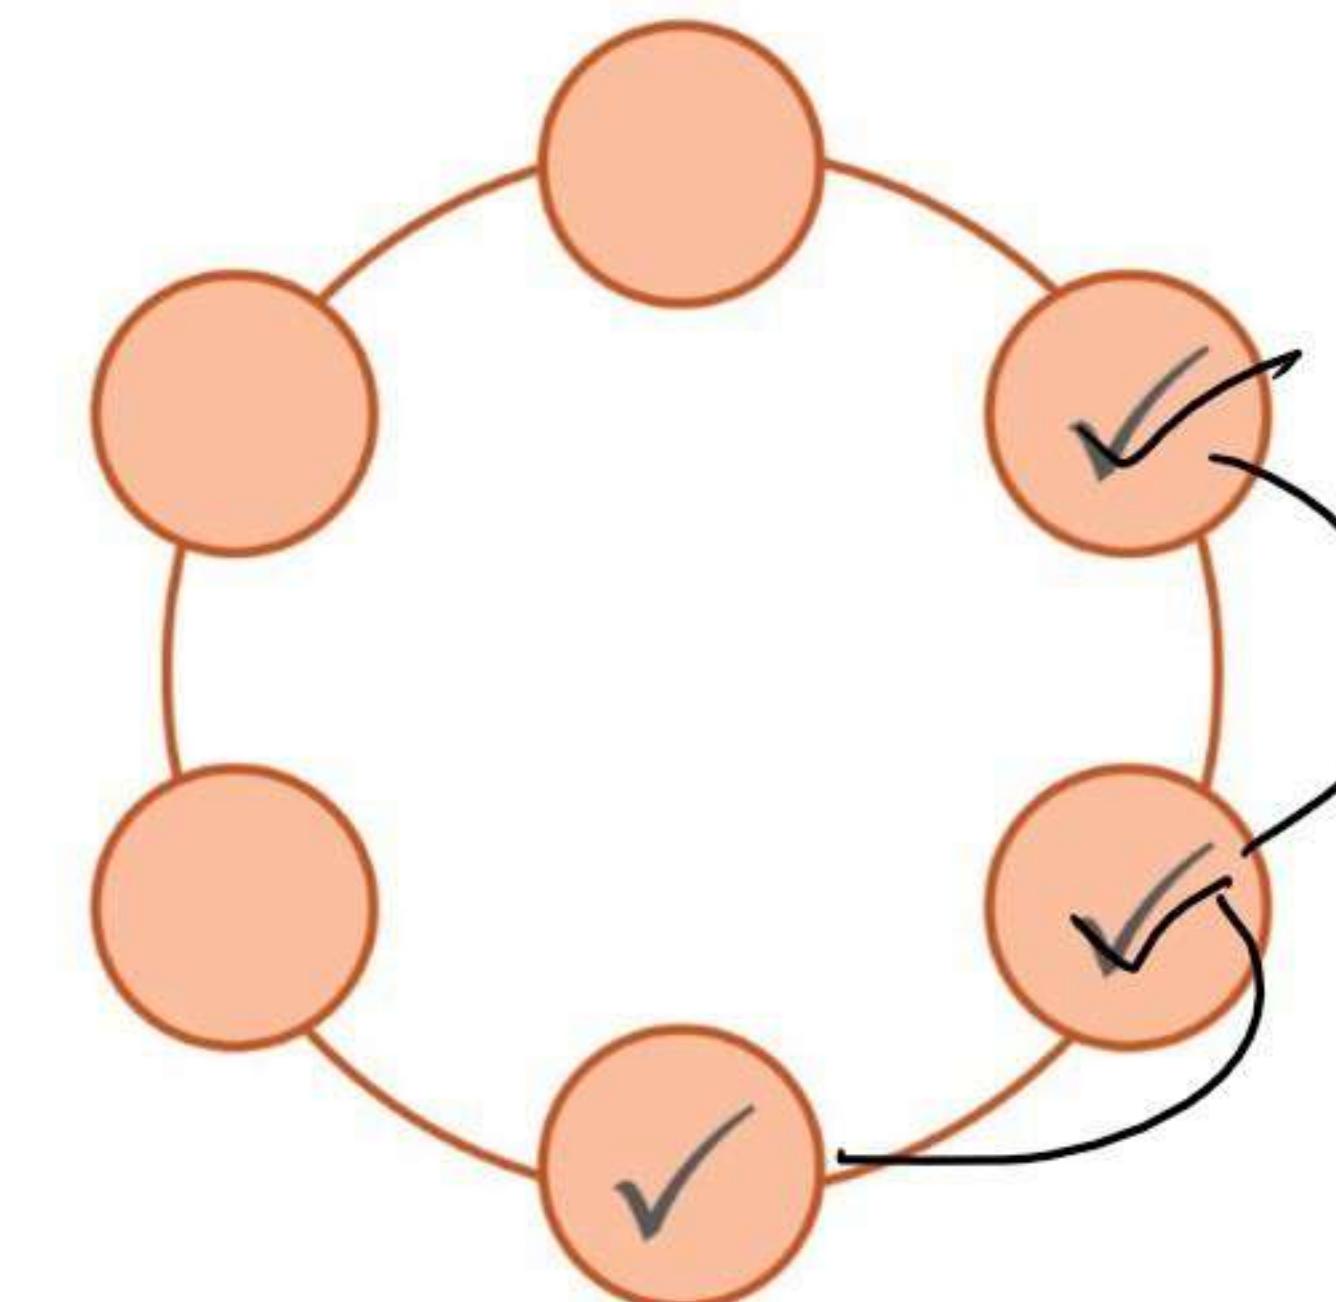
1	2	3	5	4	6
5	2	4	6	3	1
4	6	1	2	3	5
3	5	2	4	6	1

Cassandra Terminology



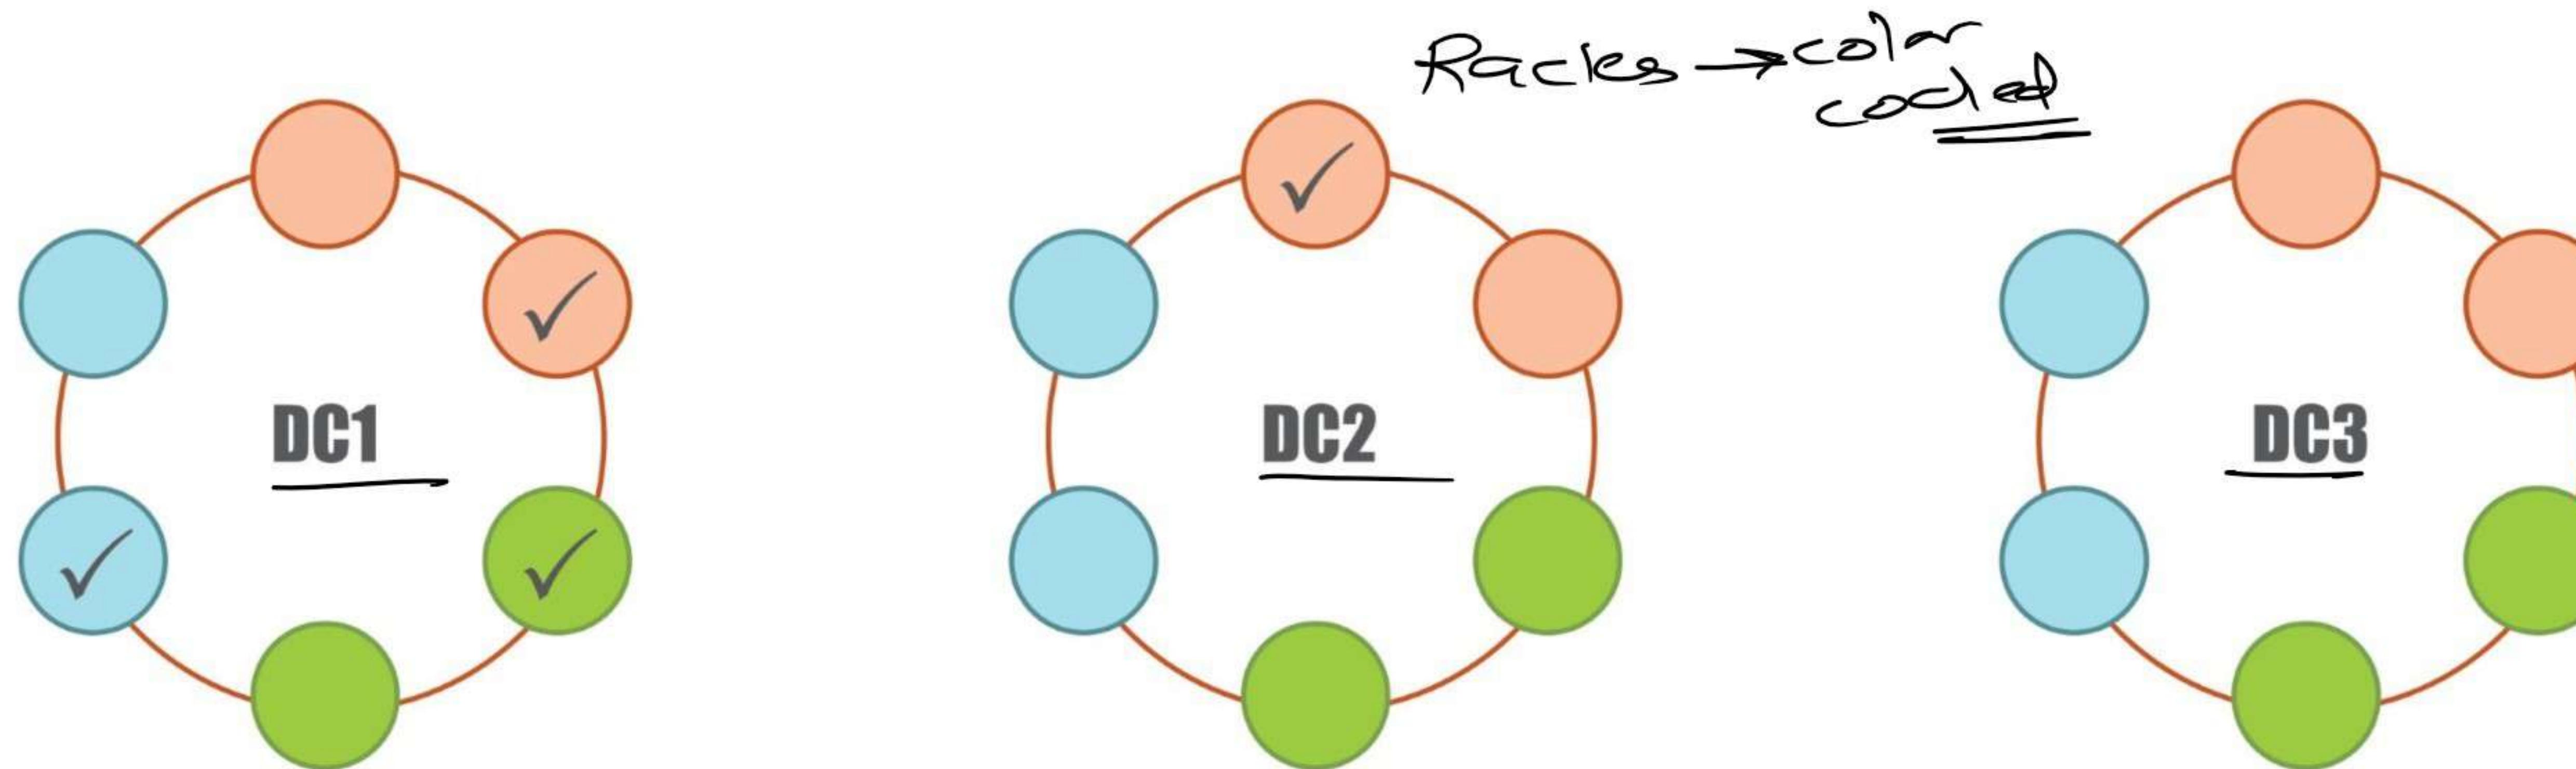
SimpleStrategy

```
create keyspace pluralsight with replication =  
  {'class': 'SimpleStrategy', 'replication_factor': 3},
```



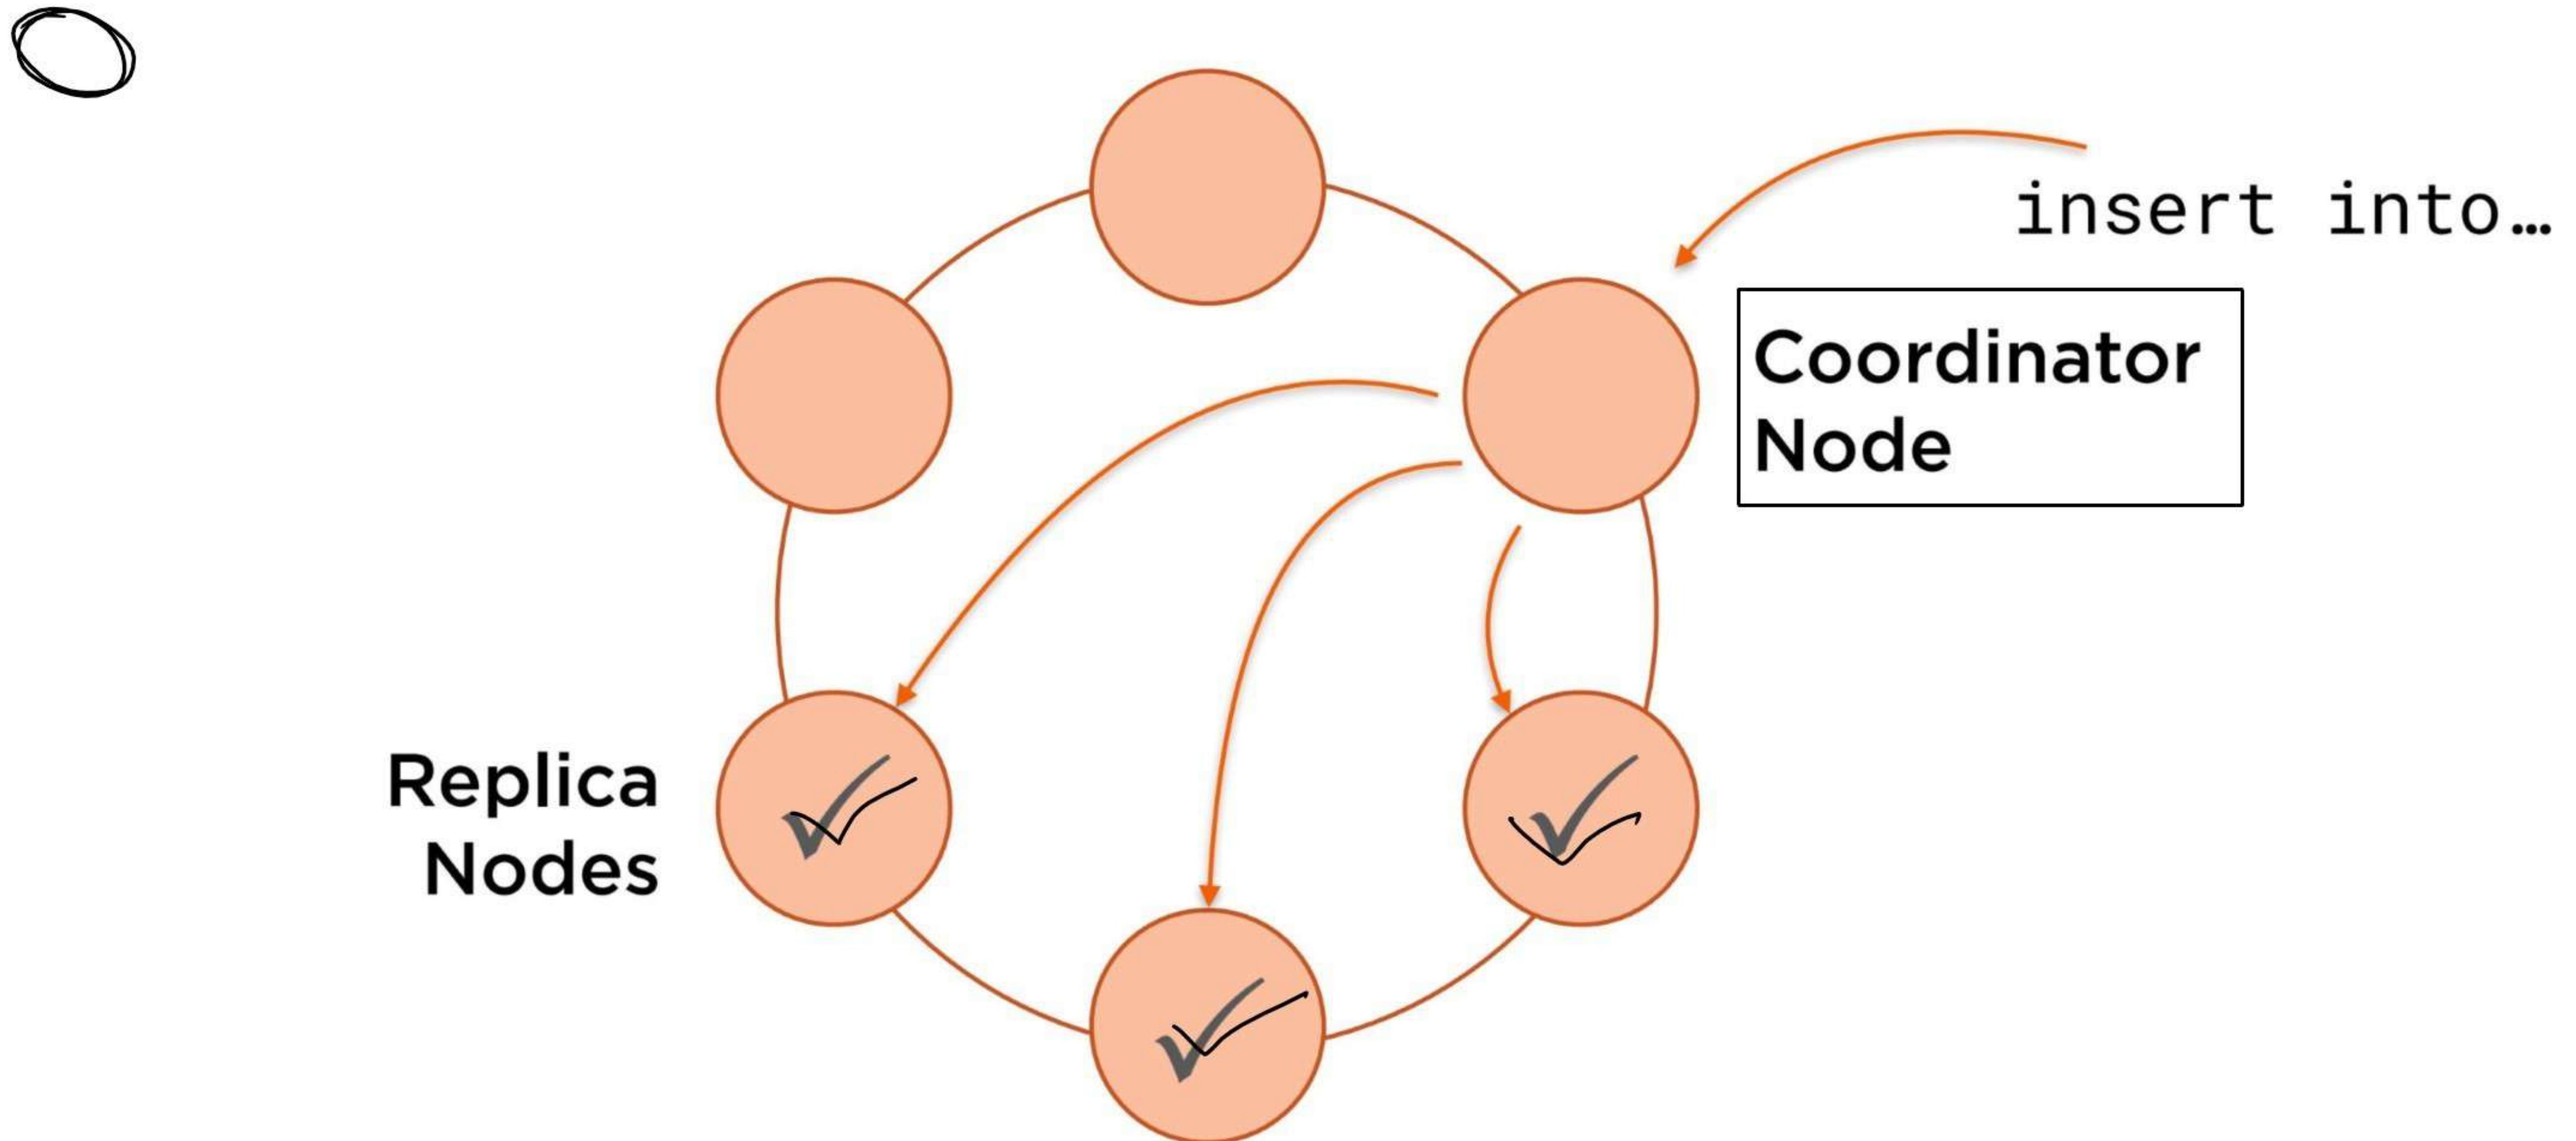
NetworkTopologyStrategy

```
create keyspace pluralsight with replication =  
  {'class': 'NetworkTopologyStrategy', 'DC1': 3, 'DC2': 1};
```

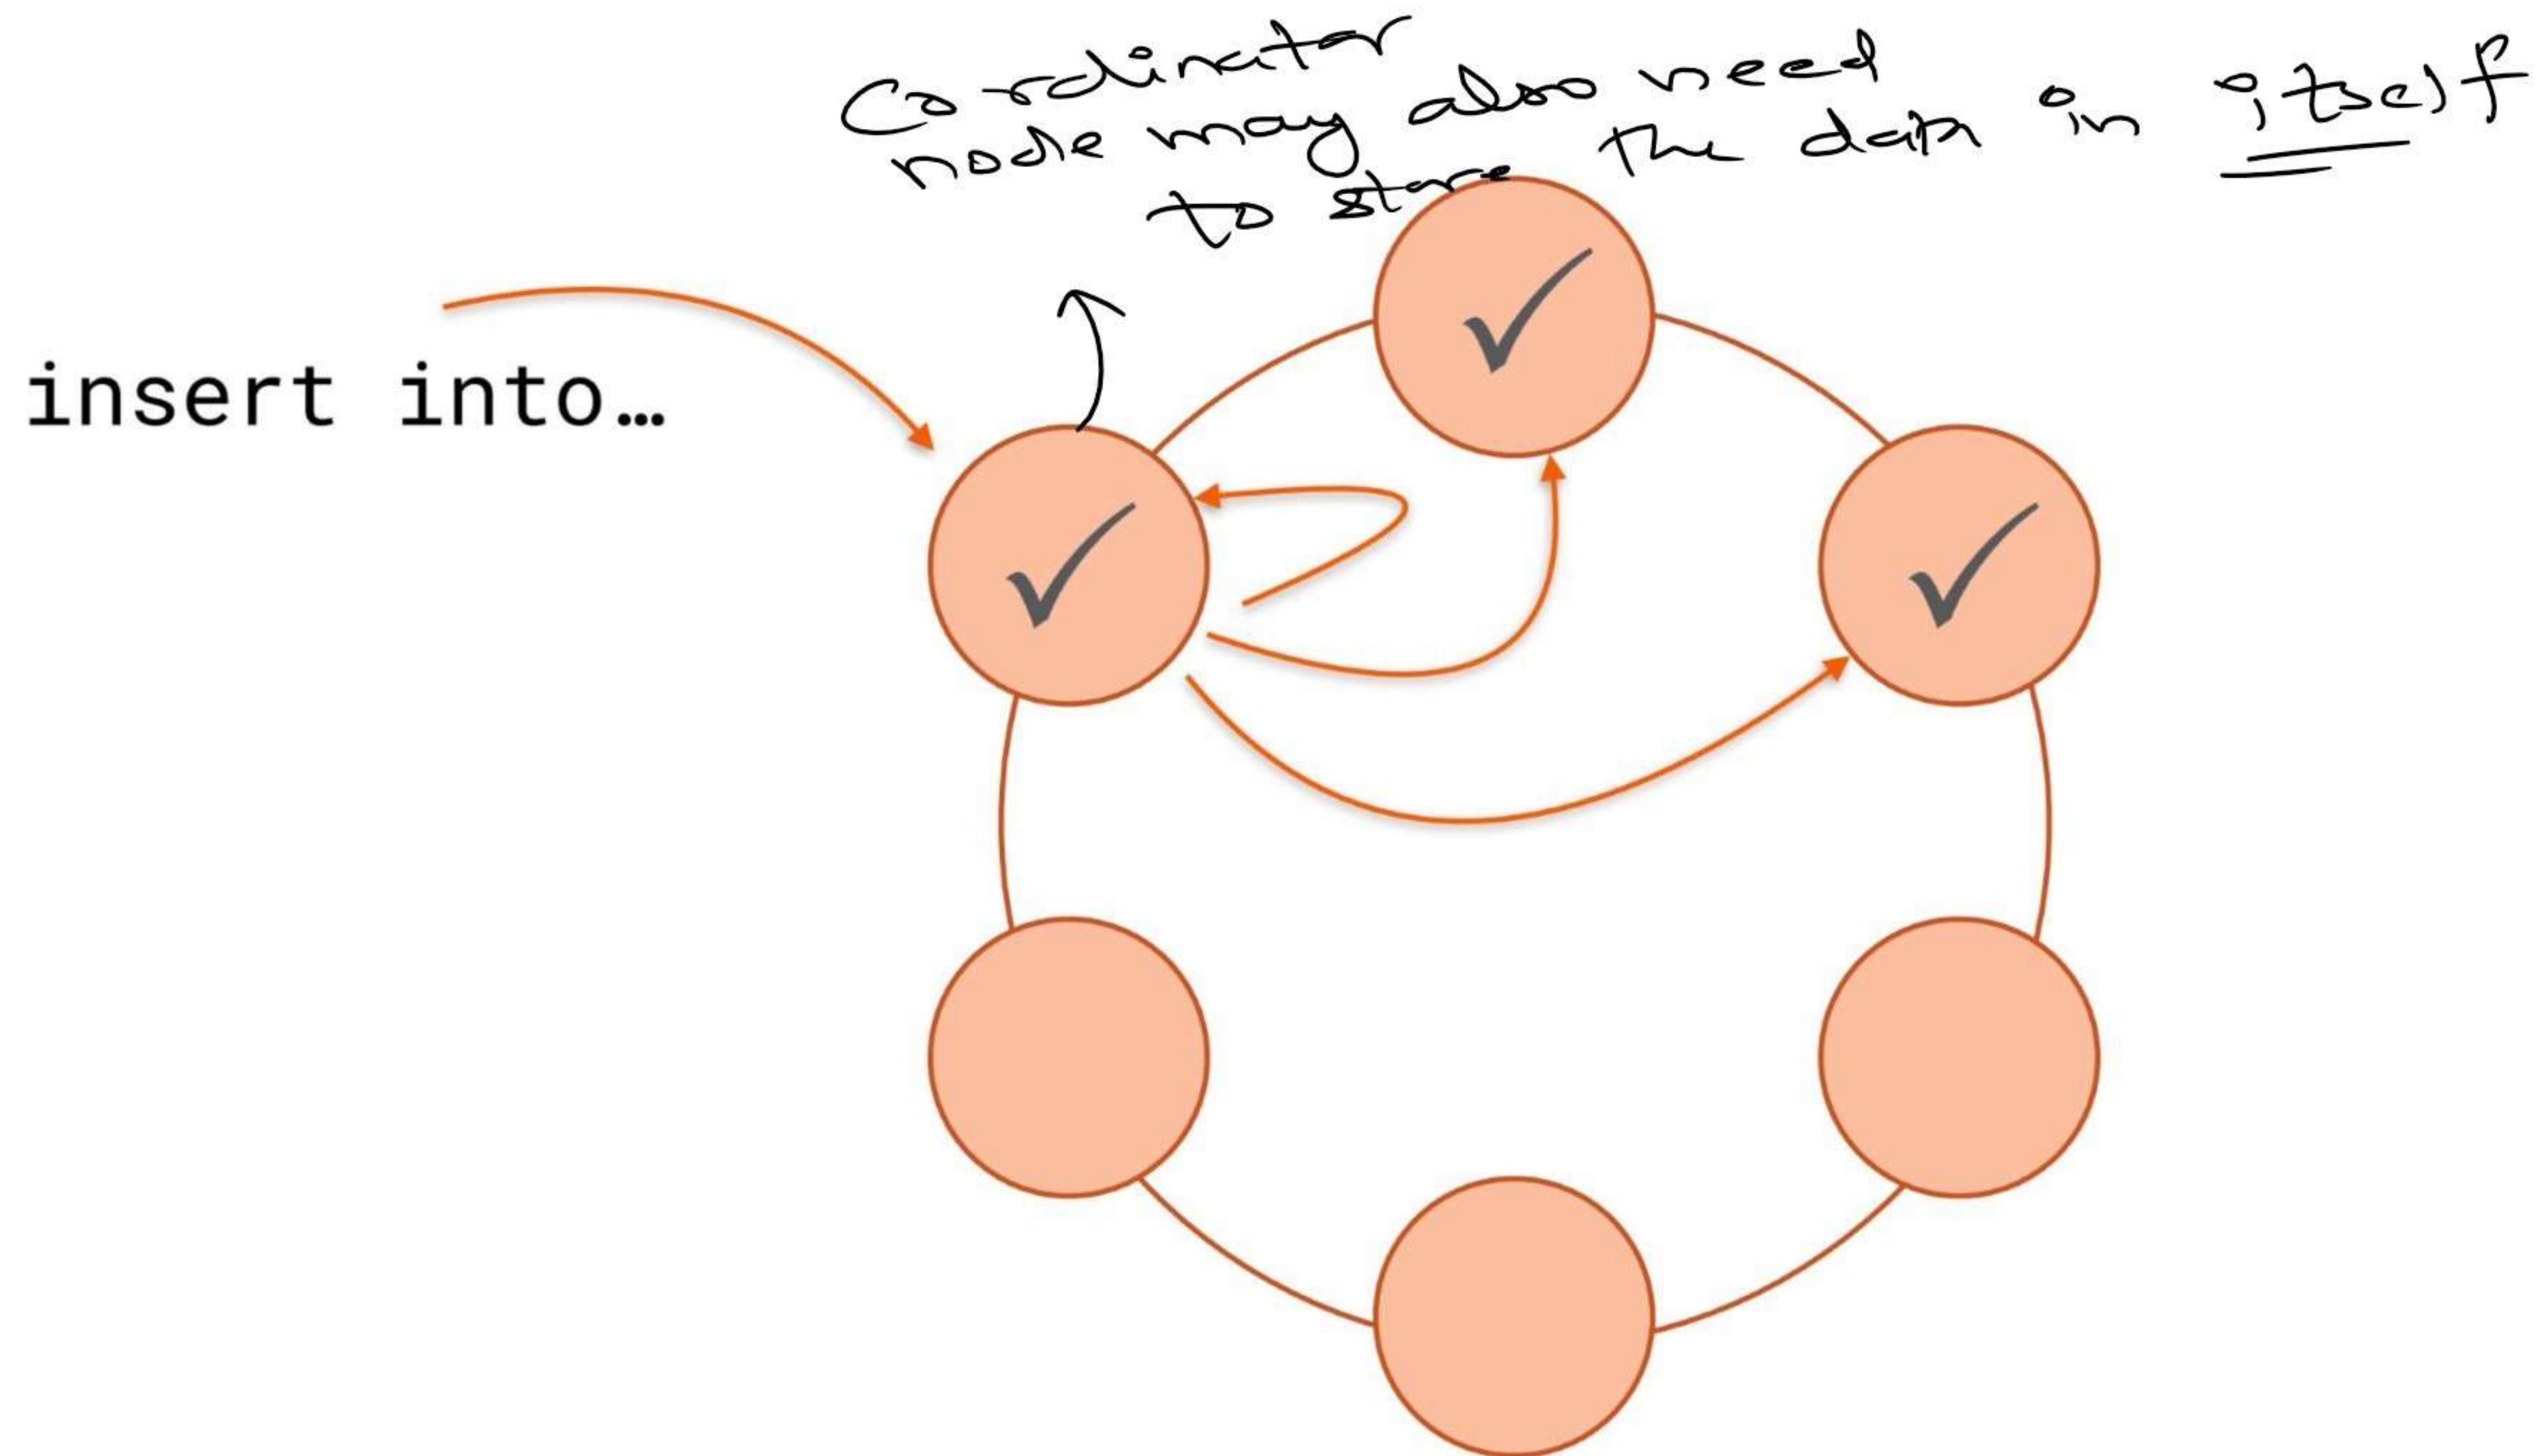


Consistency

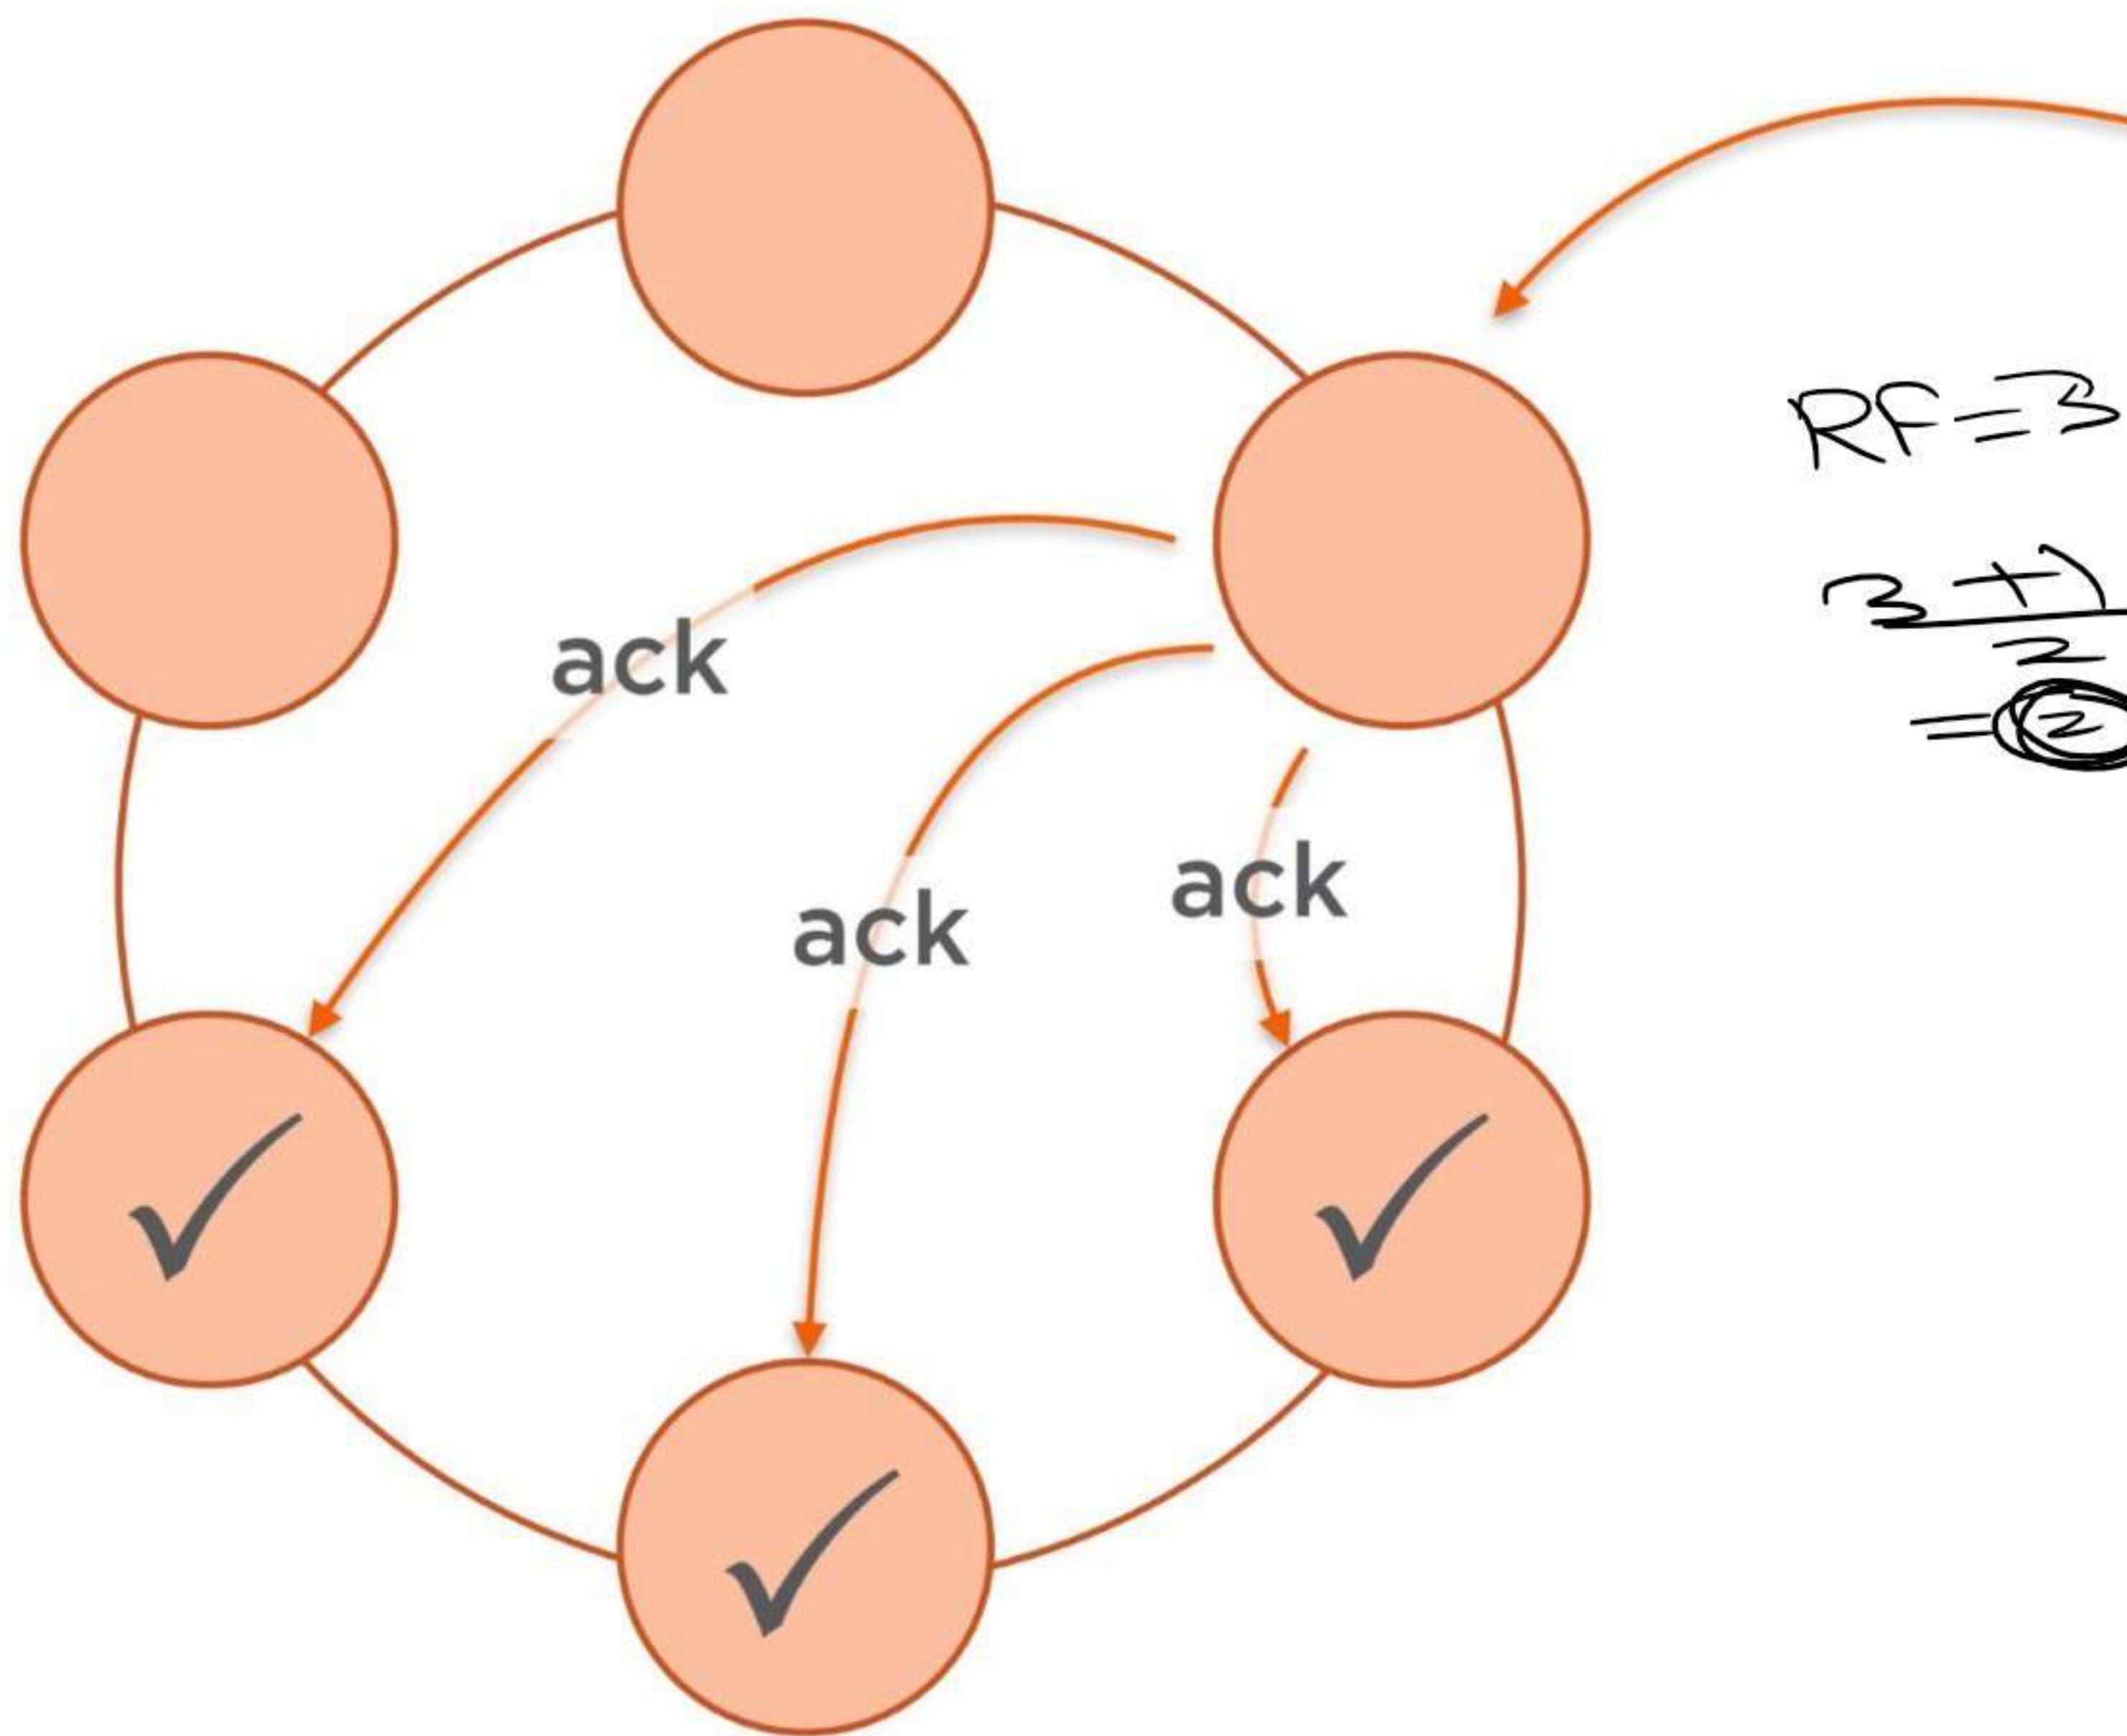
Reads and Writes in Cassandra



Reads and Writes in Cassandra



Tunable Consistency - Writes



writing

ONE

if we get ack
from one of the
node → write is
successful

TWO

THREE

QUORUM

ALL

ANY

Replication Factor + 1

2

even if we can
connect with
Coordinator node

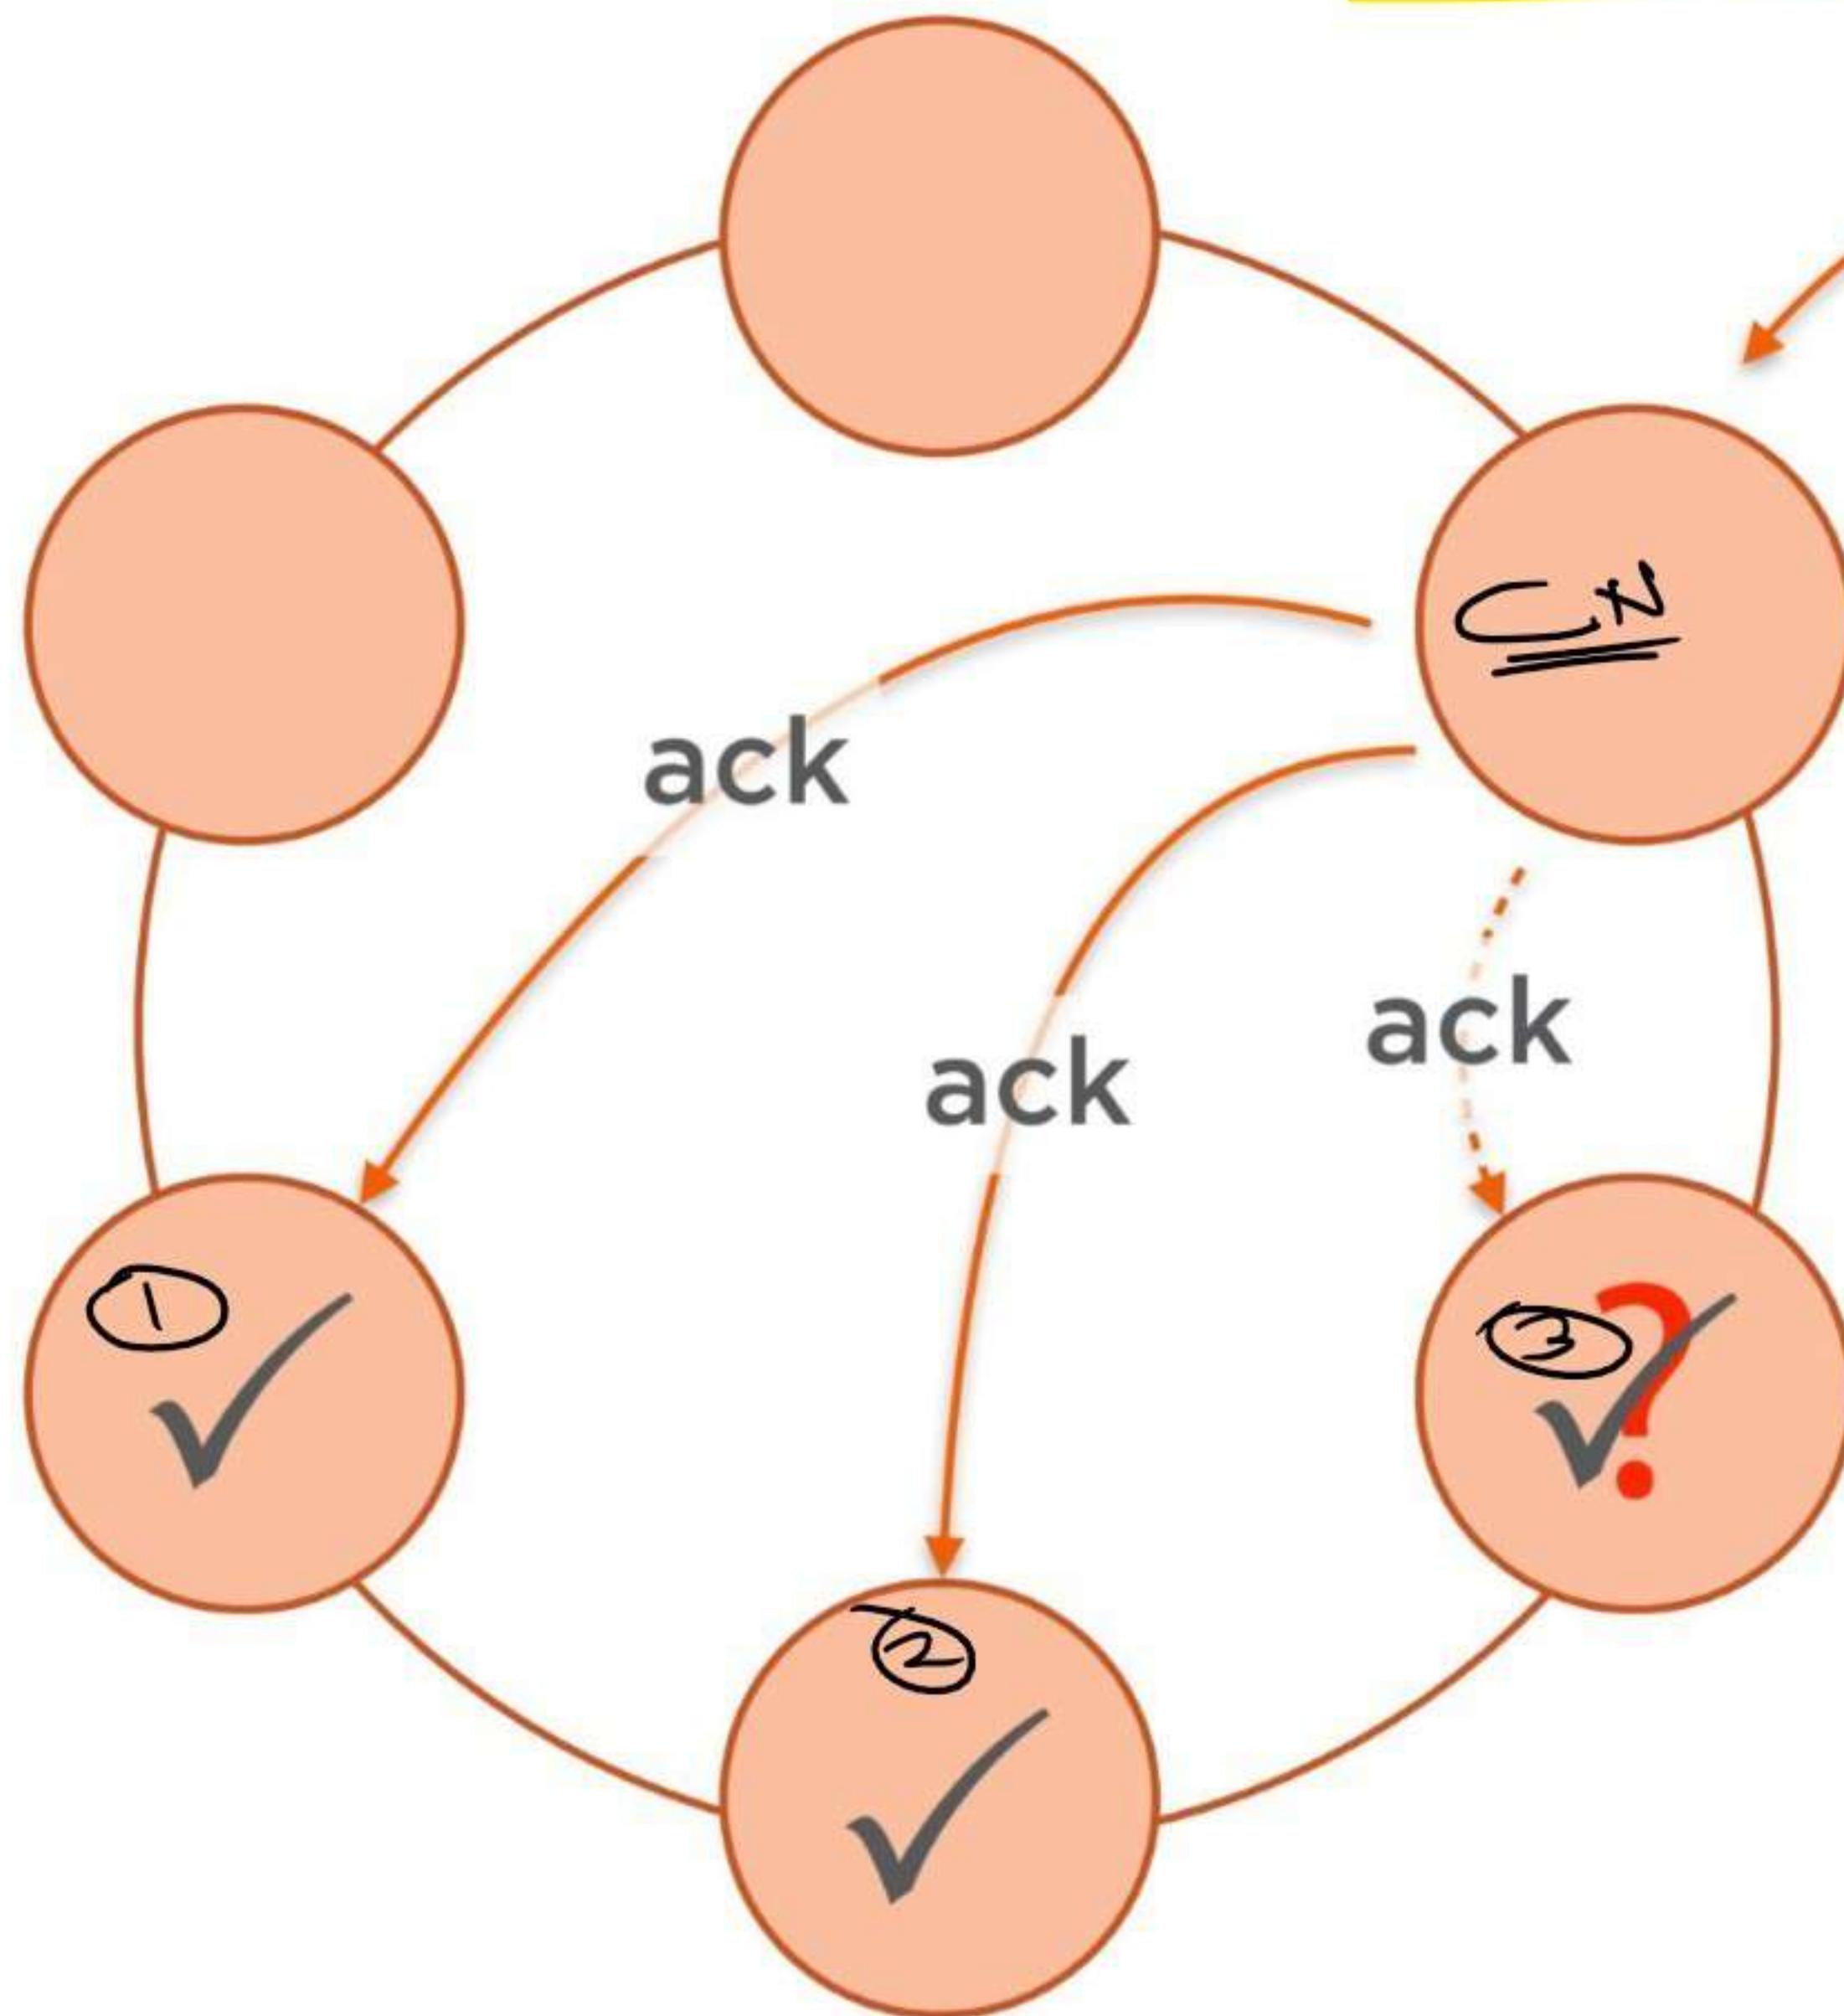
Hinted Handoff

important feature

RF < 3

ALL →

Successful write → ③



Solution

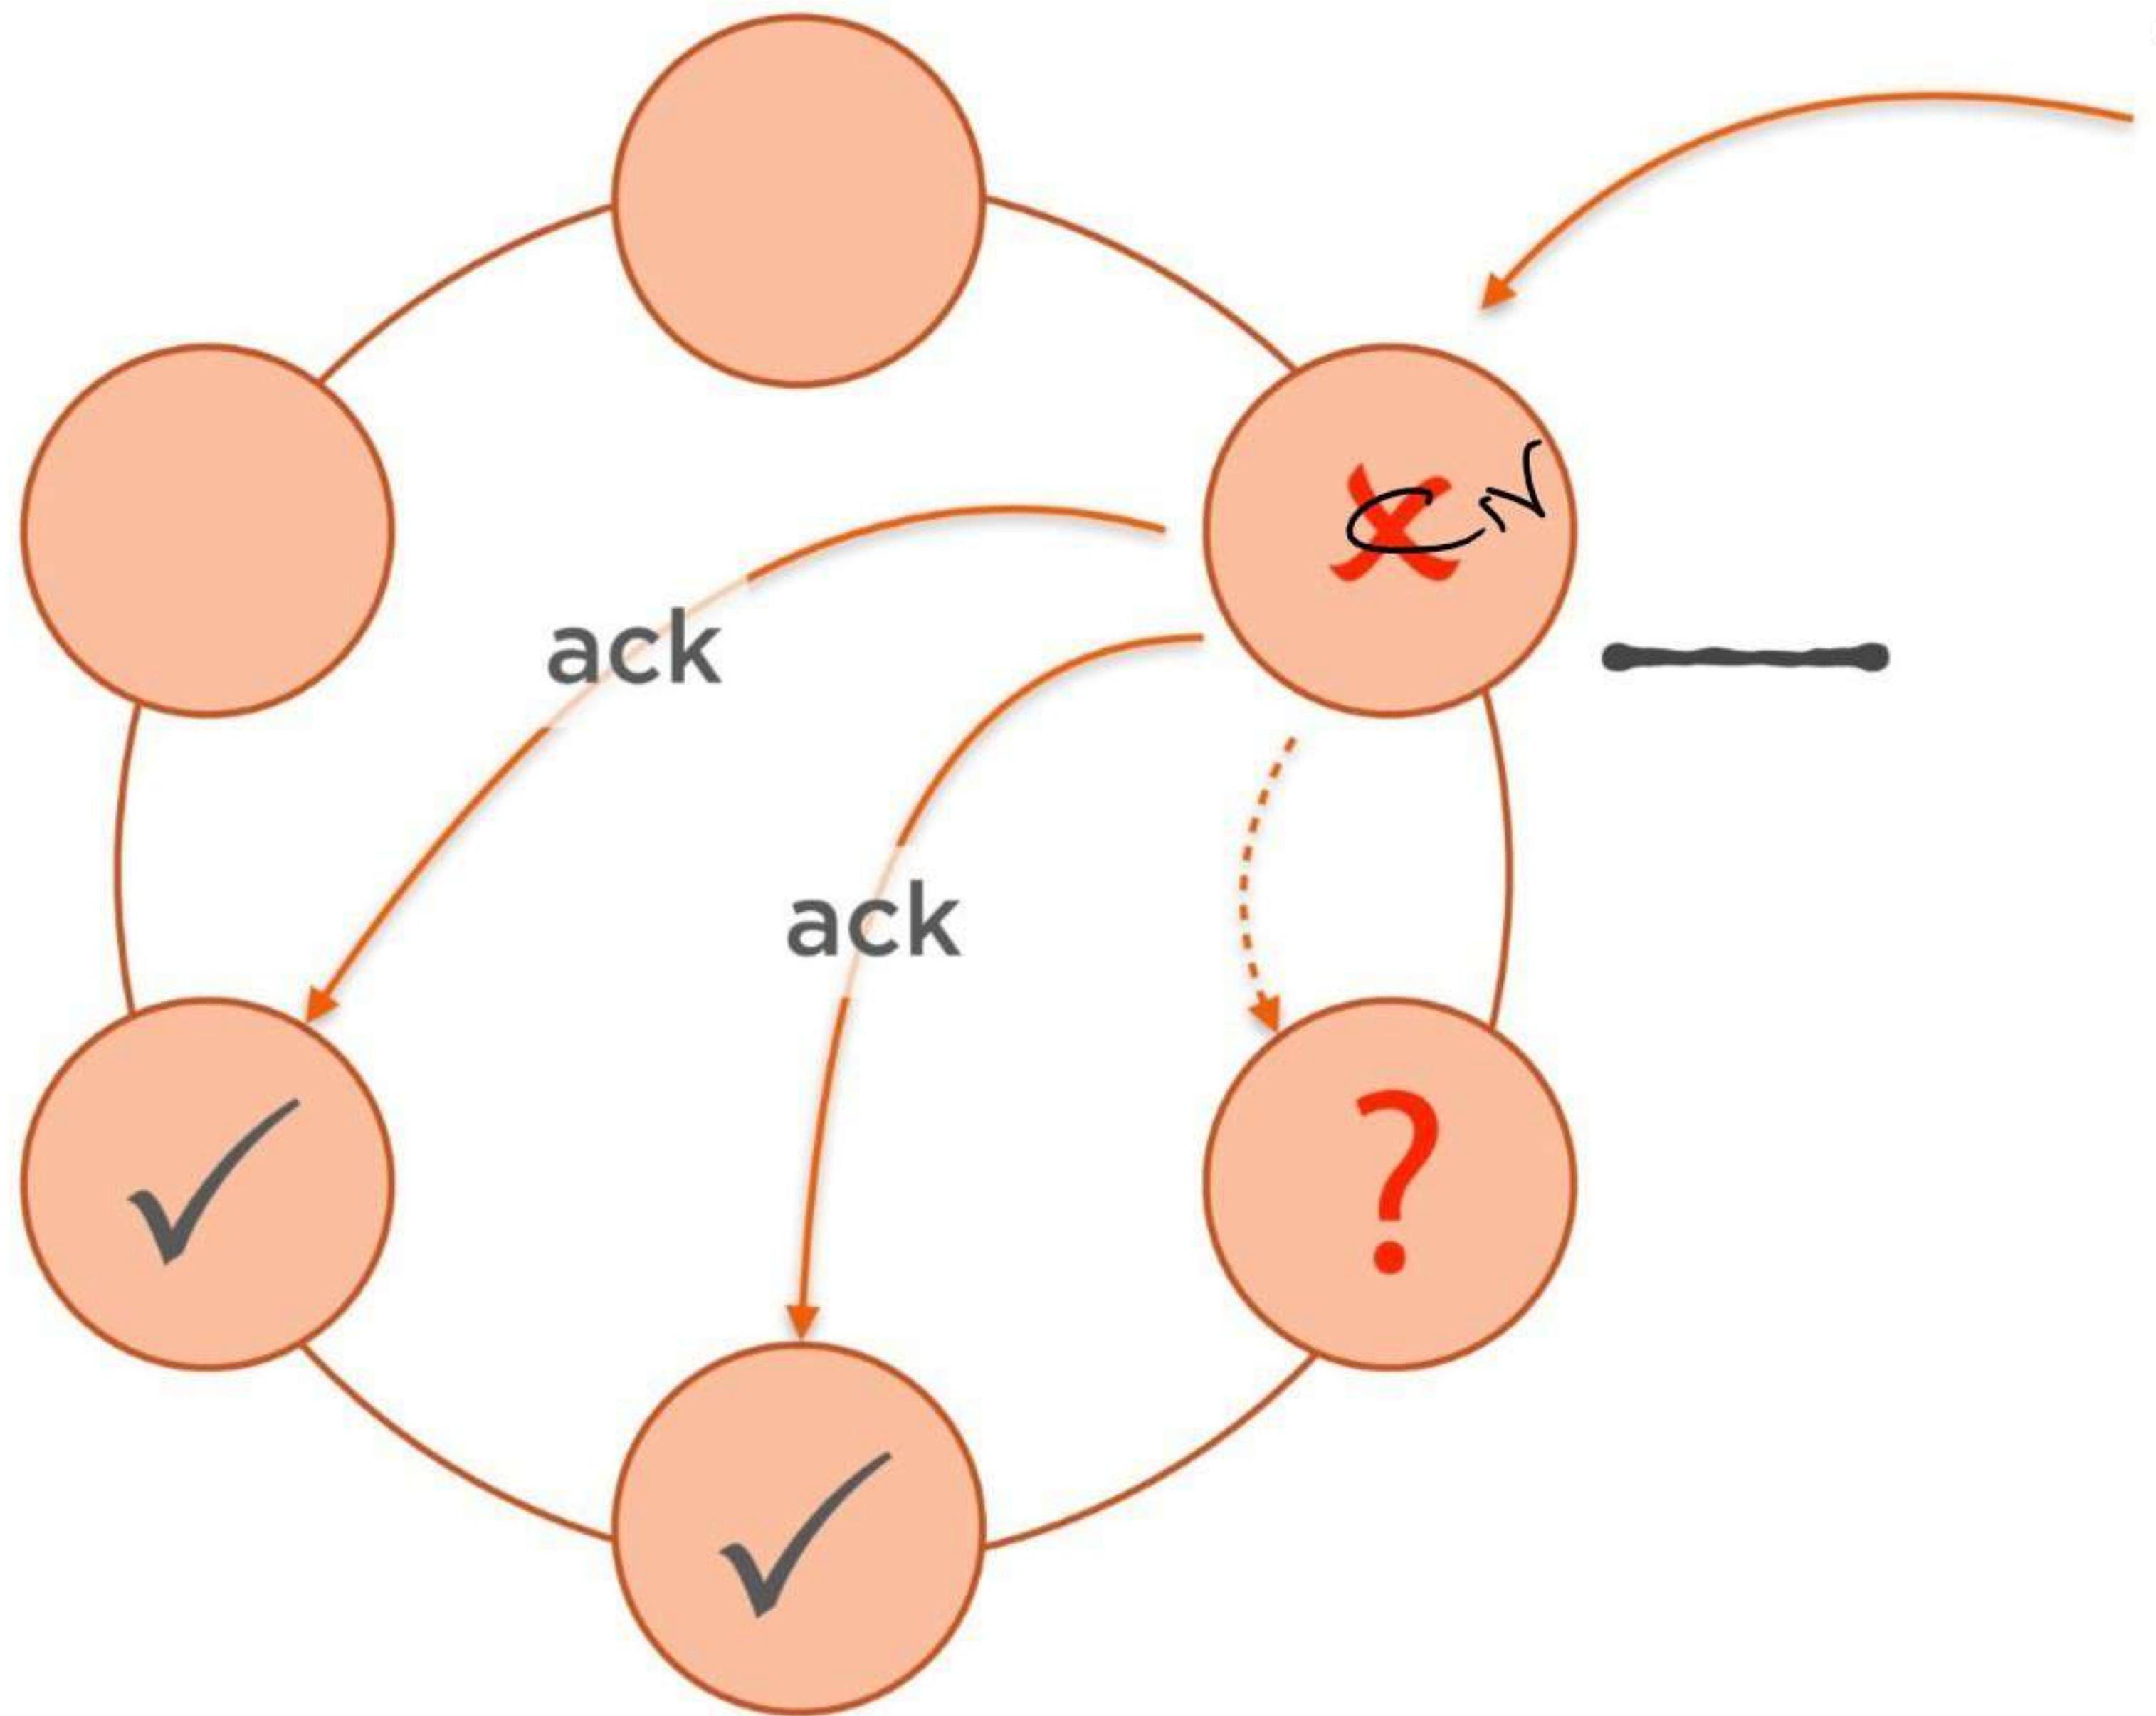
CN will store the data & keep it with itself until N3 comes back up & Then write to PT.

→ Node 3 had some network issue

Hinted Handoff

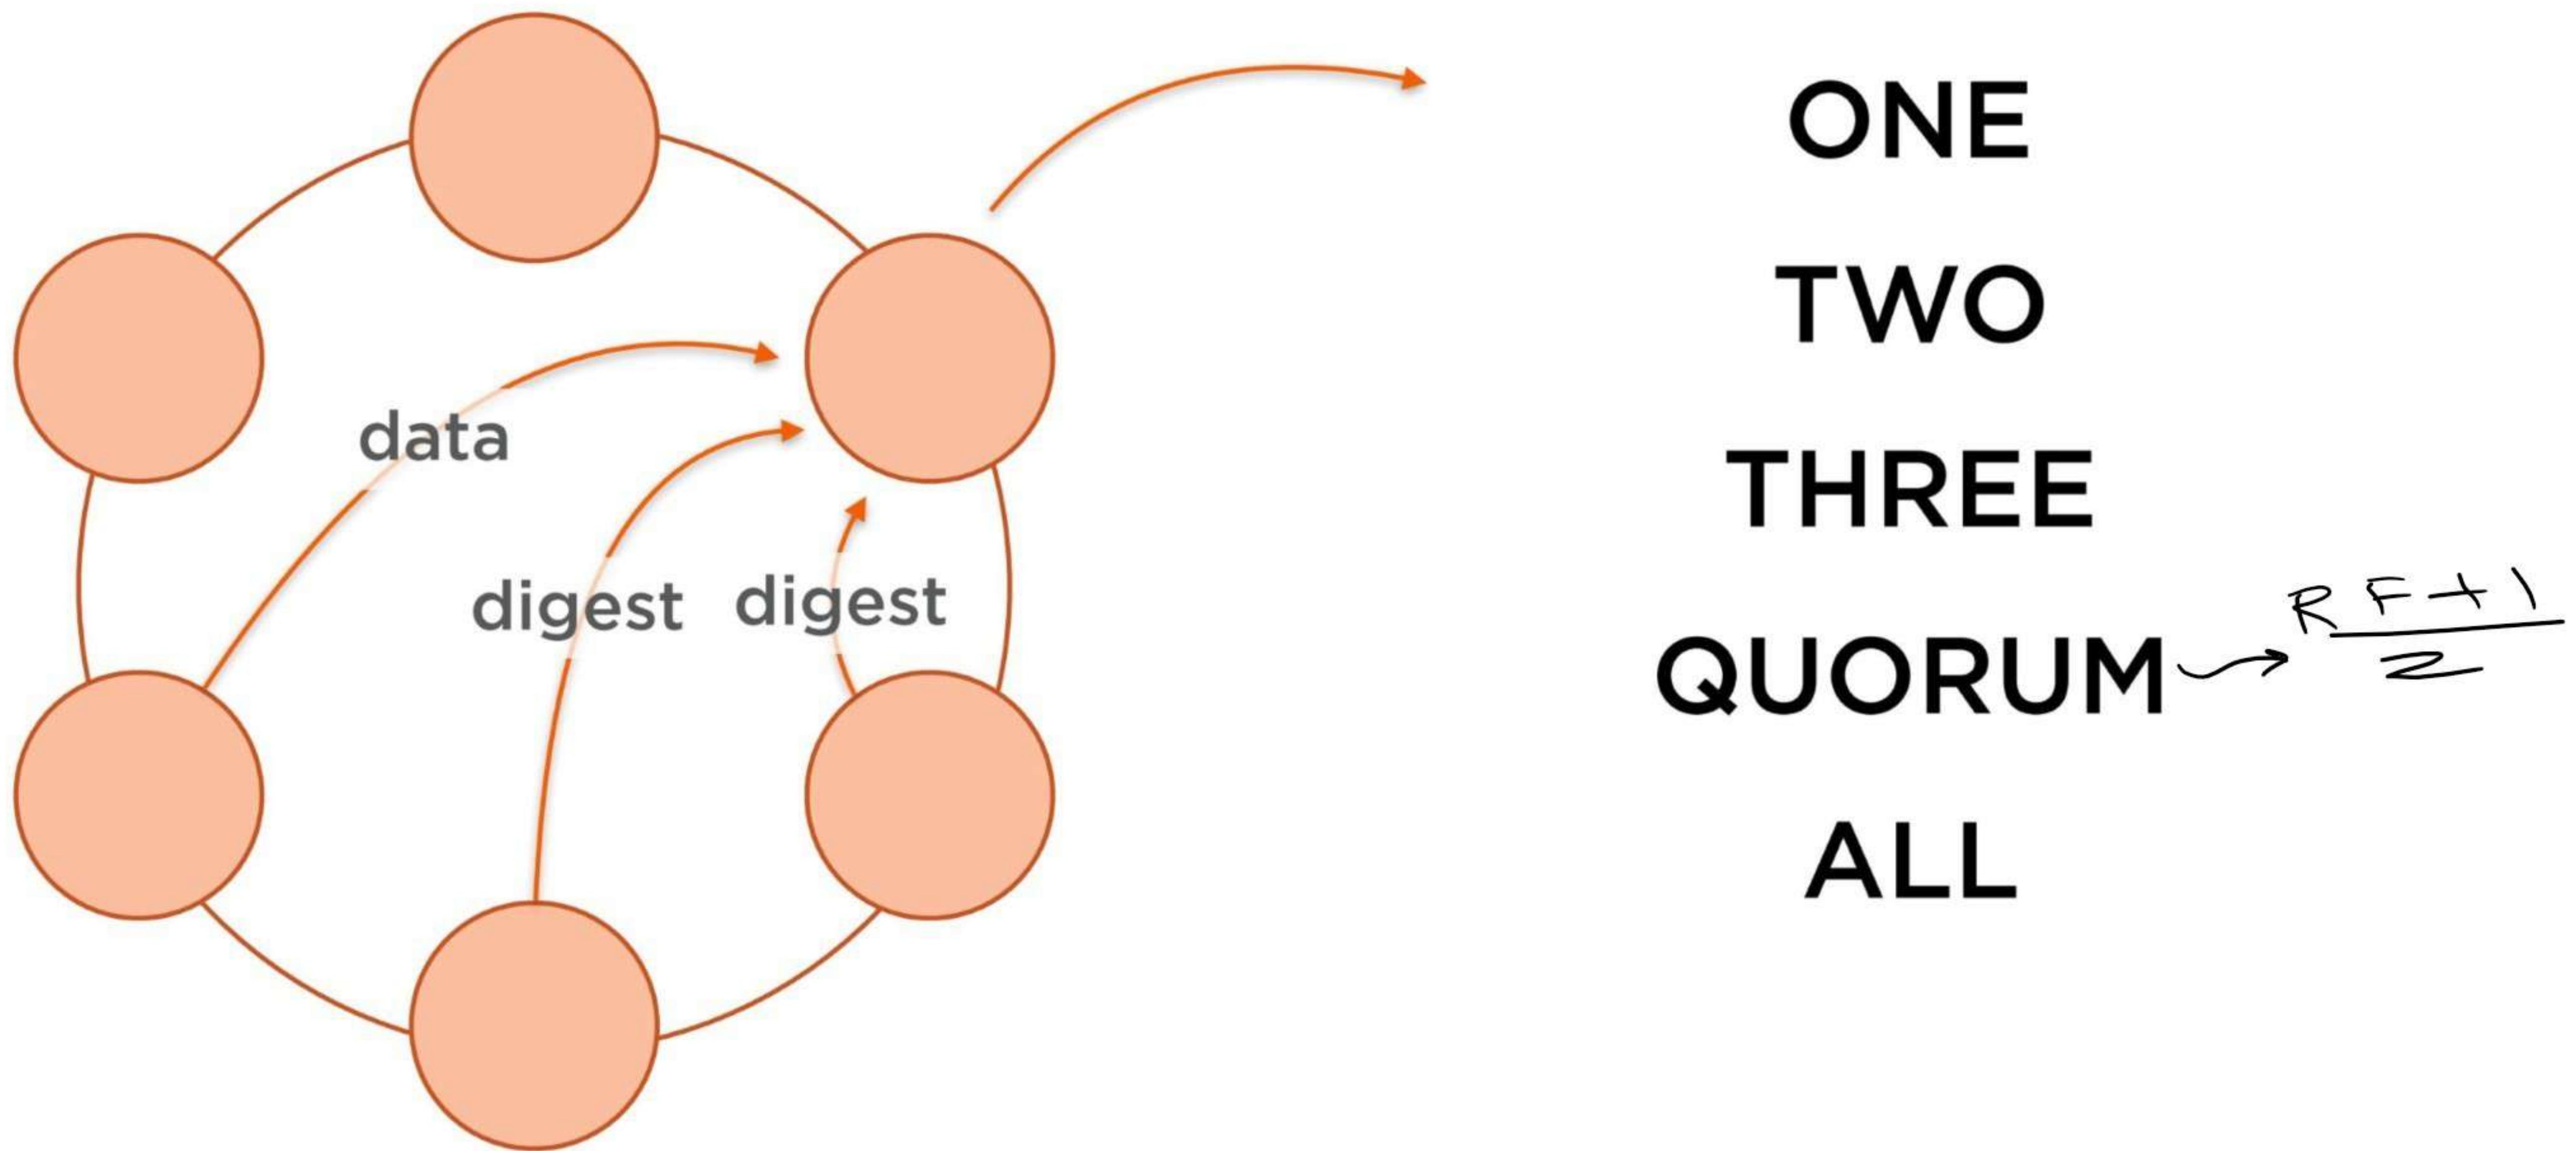
+ Read Repair

What if?



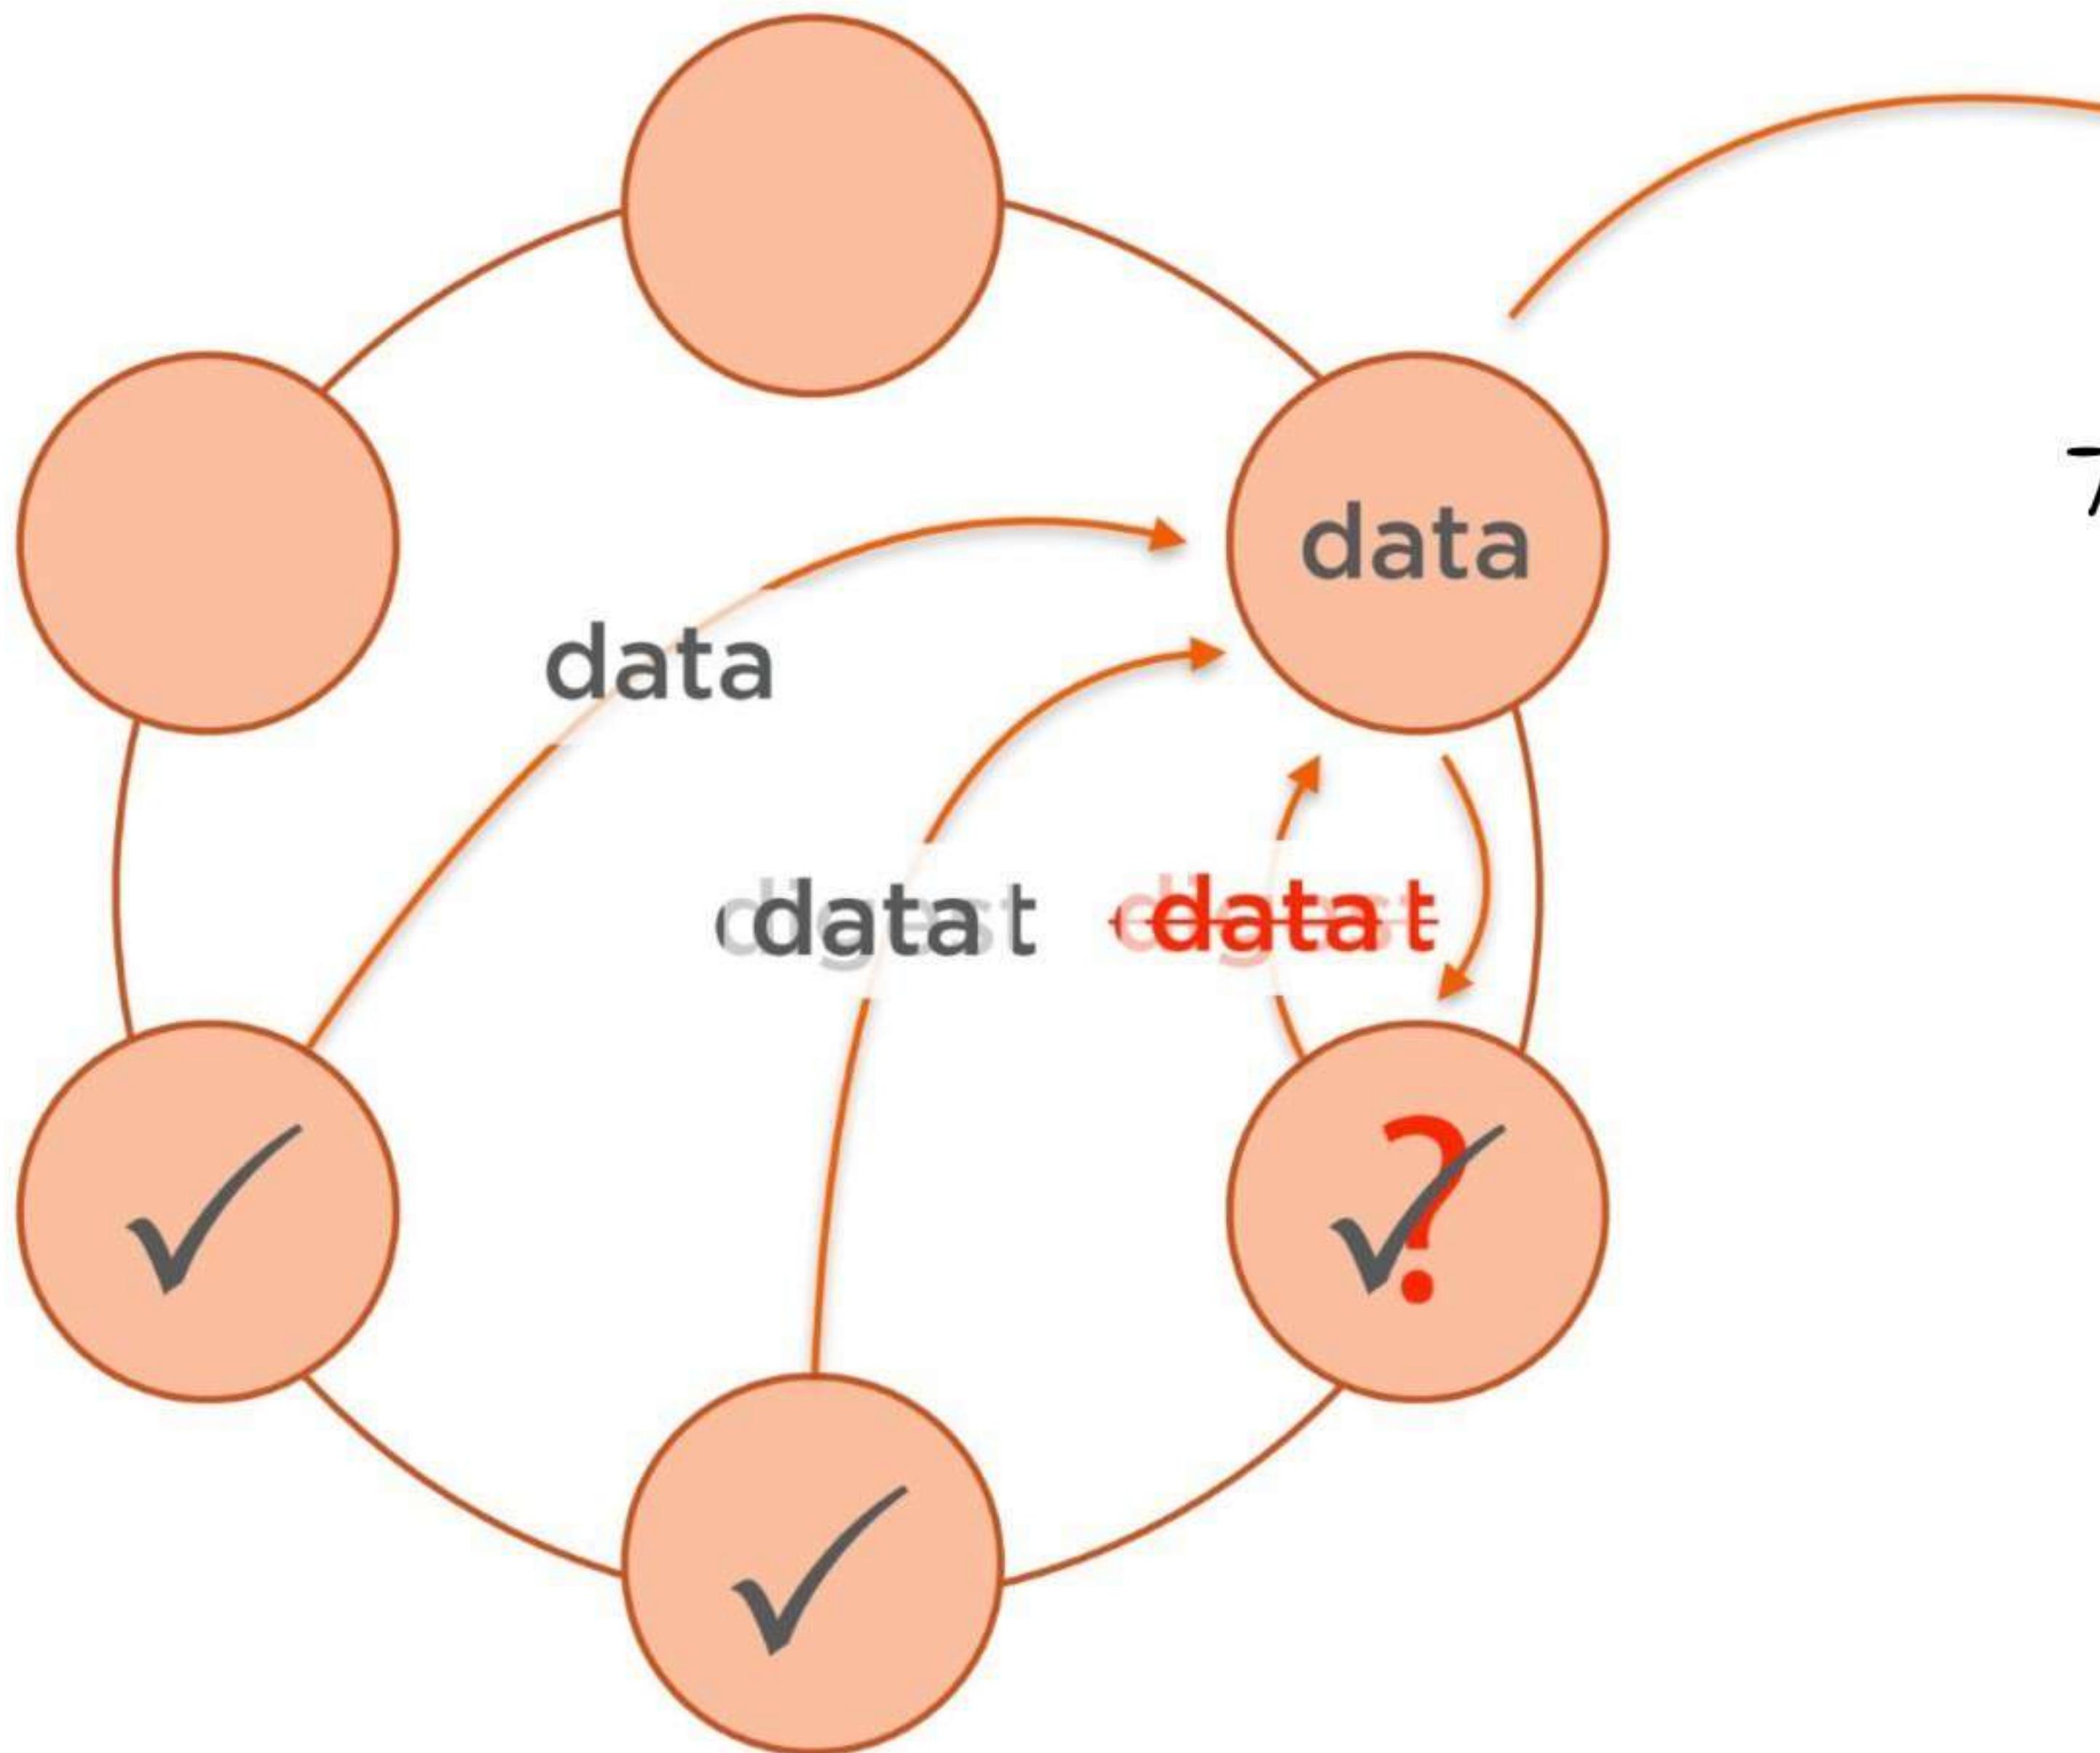
- N₃ went off while write
- CN was holding the data
- while coordinator node was holding the data & before N₃ came back
CN died.

Tunable Consistency - Reads



RF = 3
CL → ALL

Read Repair

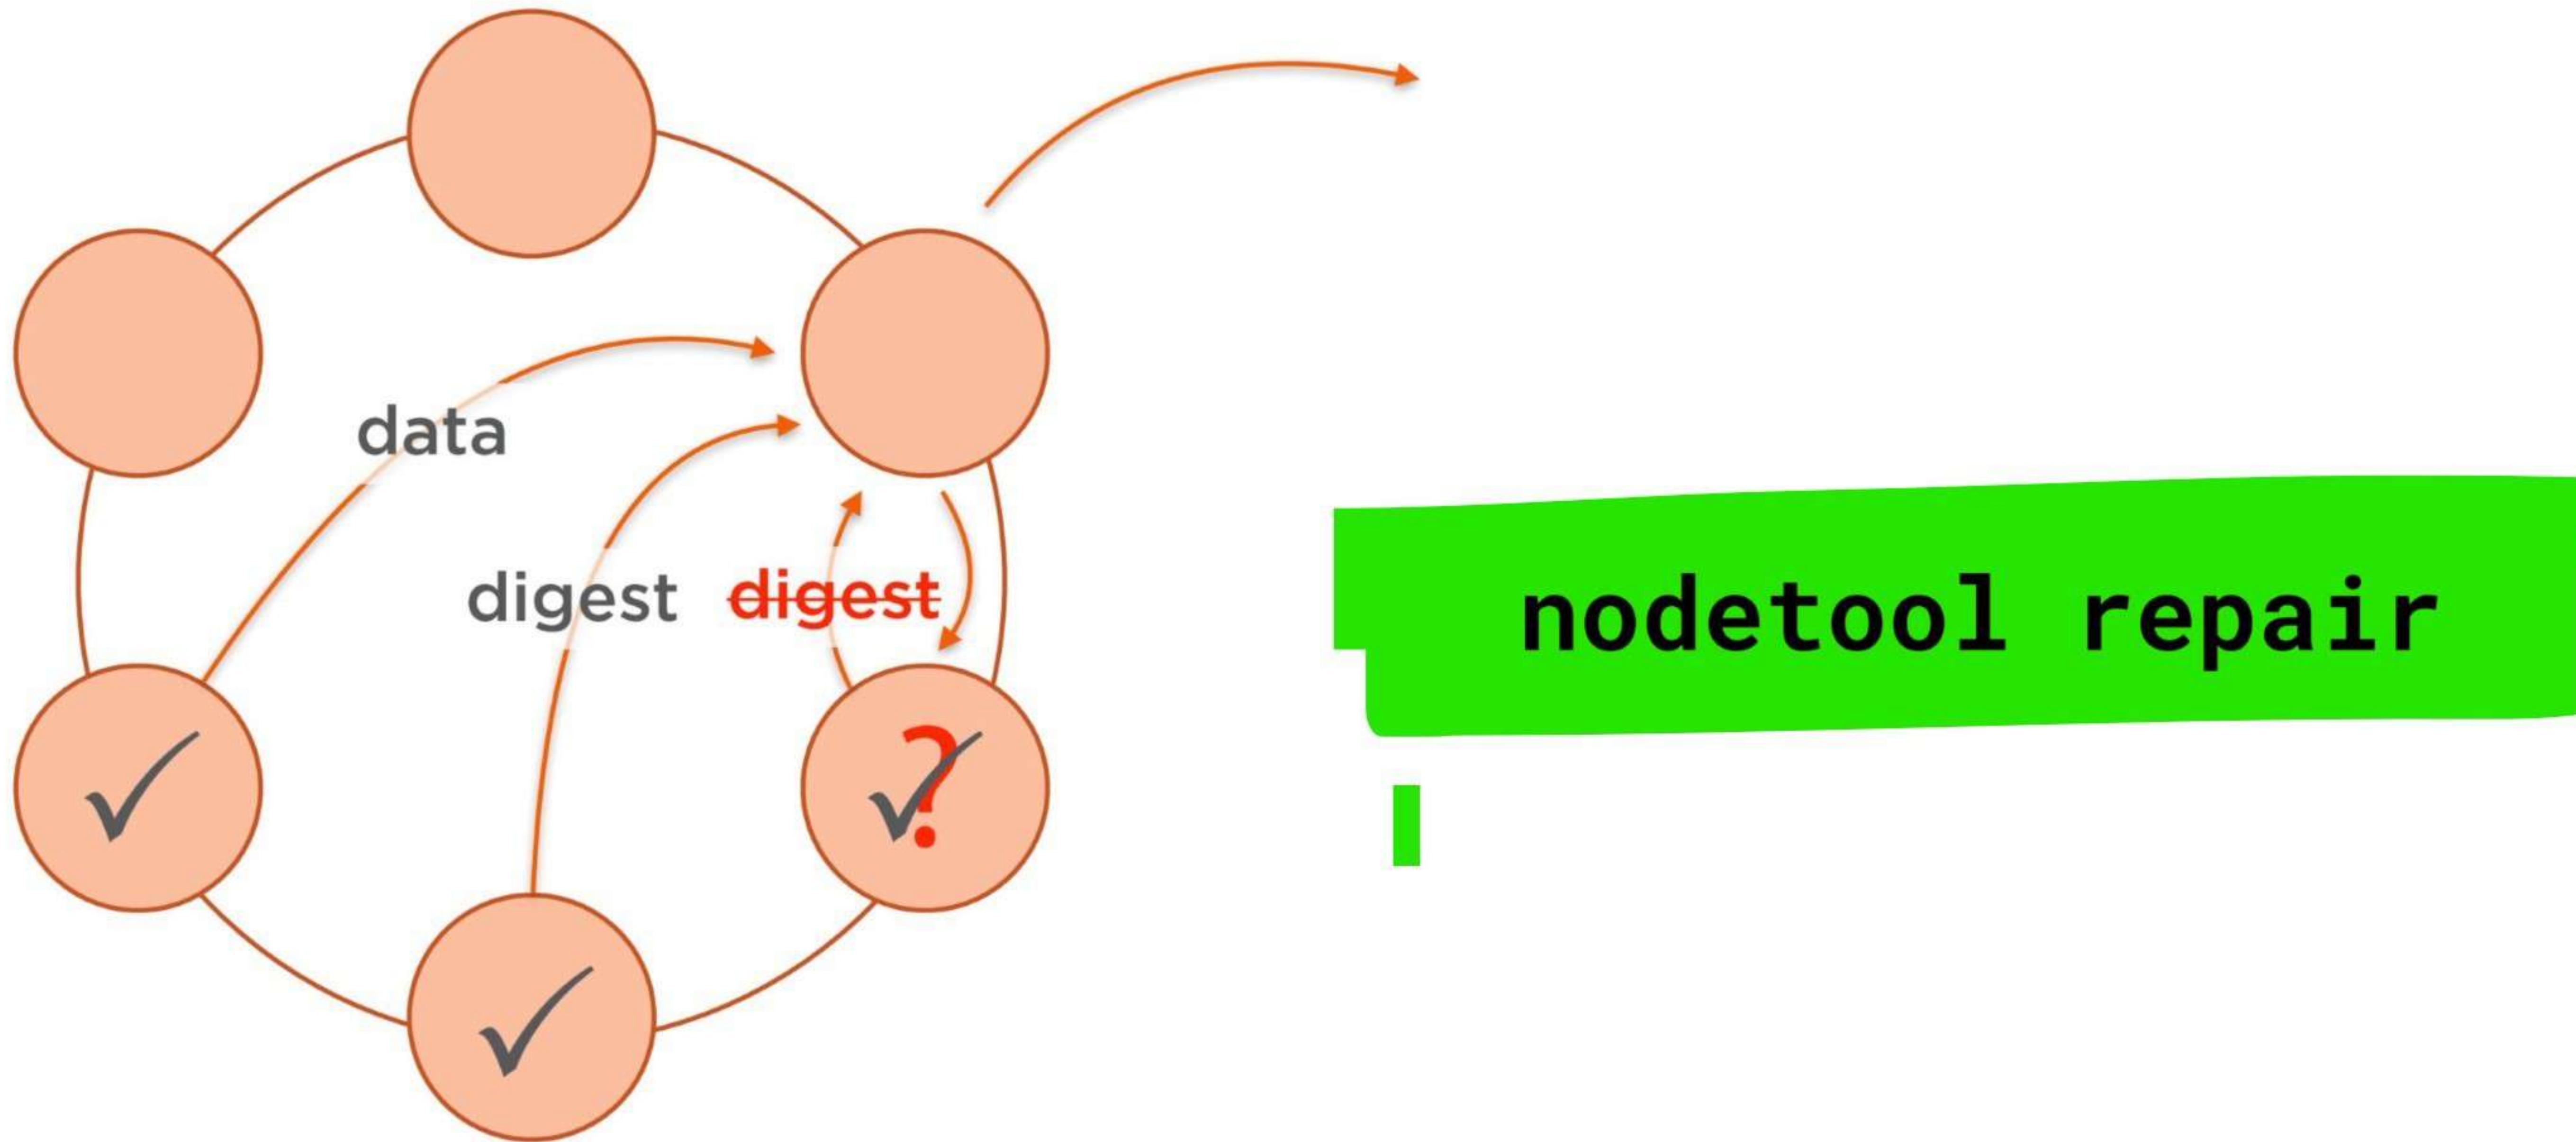


$N_1 \rightarrow Data$
 $N_2 \rightarrow digest \rightarrow data? \checkmark$
 $N_3 \rightarrow digest \underline{stale}$

Read Repair ↗

→ update the data in N_3
as a part of read process.

Read Repair

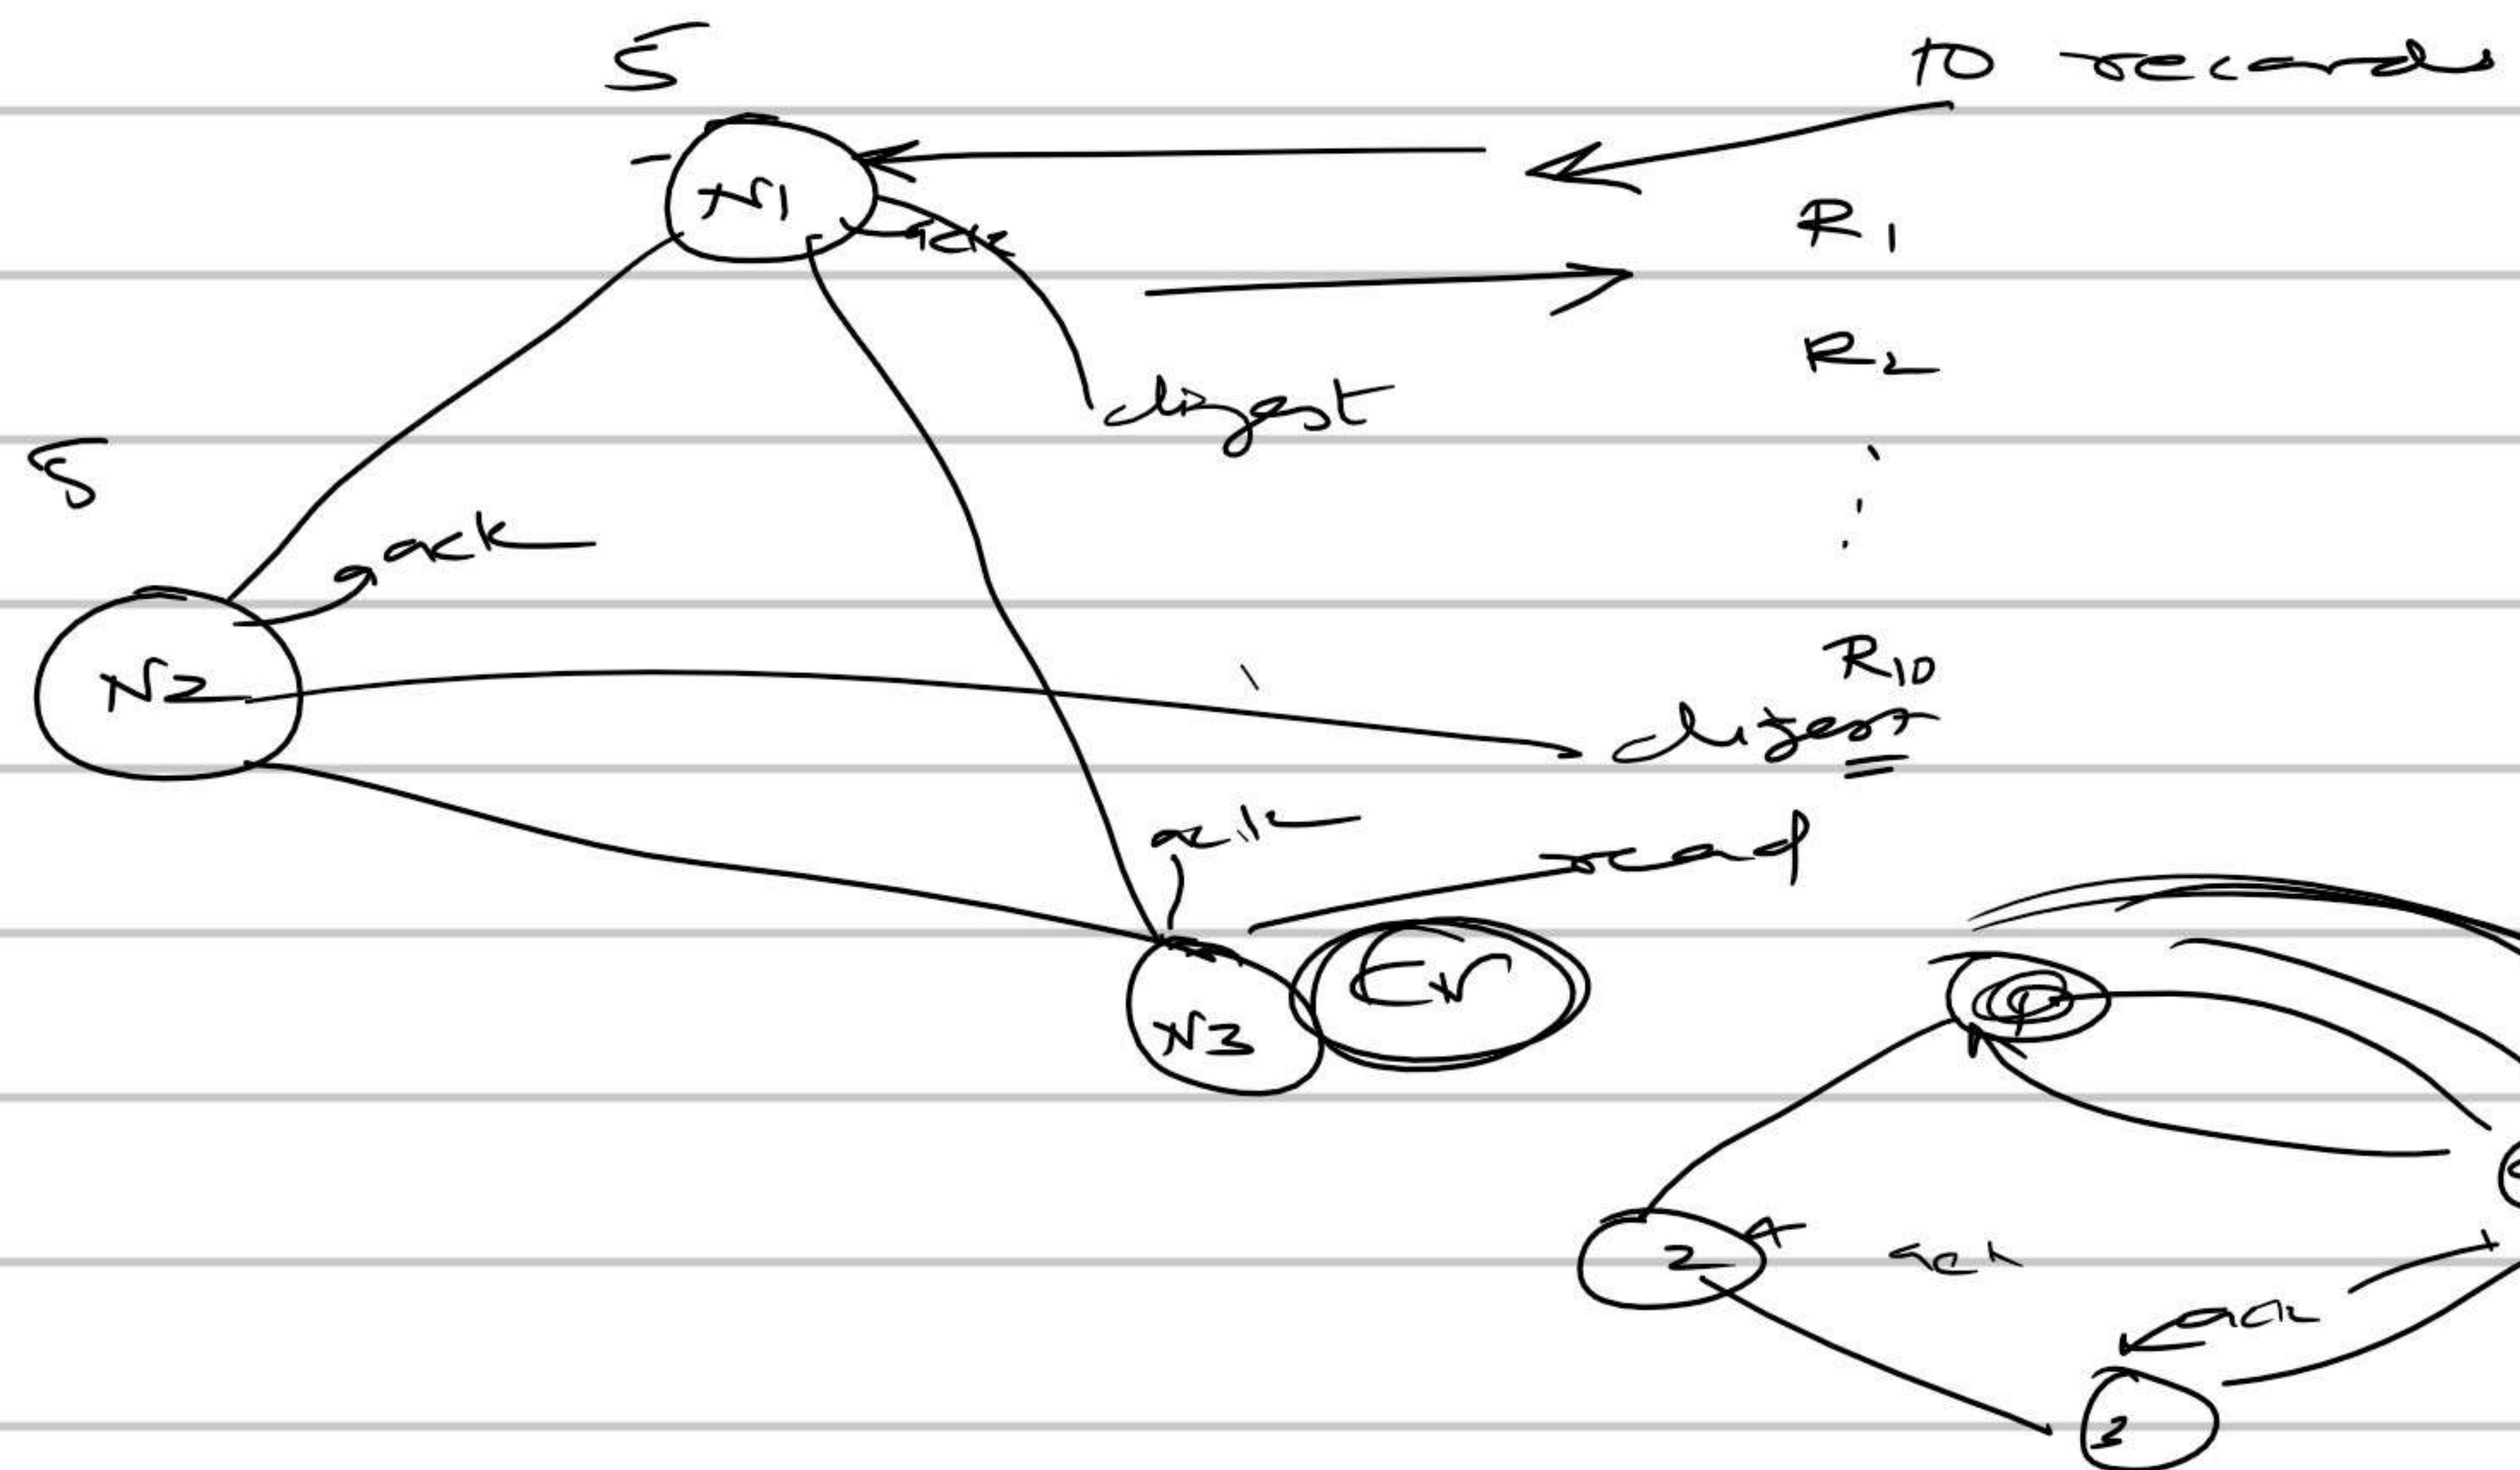


Achieving Strong Consistency

ensures that at least one replica is common between the write & read.

(Write Consistency + Read Consistency) > Replication Factor

$\frac{RF+1}{2}$	$= \frac{3+1}{2} = 2$	ONE	+	QUORUM	$2 > 3$	3	$1(W) + 2(R) = 3$	X
ONE	+	ONE	+	ALL	$3 > 4$	3		✓
ONE	+	ONE	+	ONE	$2 > 3$	3		X
QUORUM	+	QUORUM	+	QUORUM	$4 > 3$	3		✓

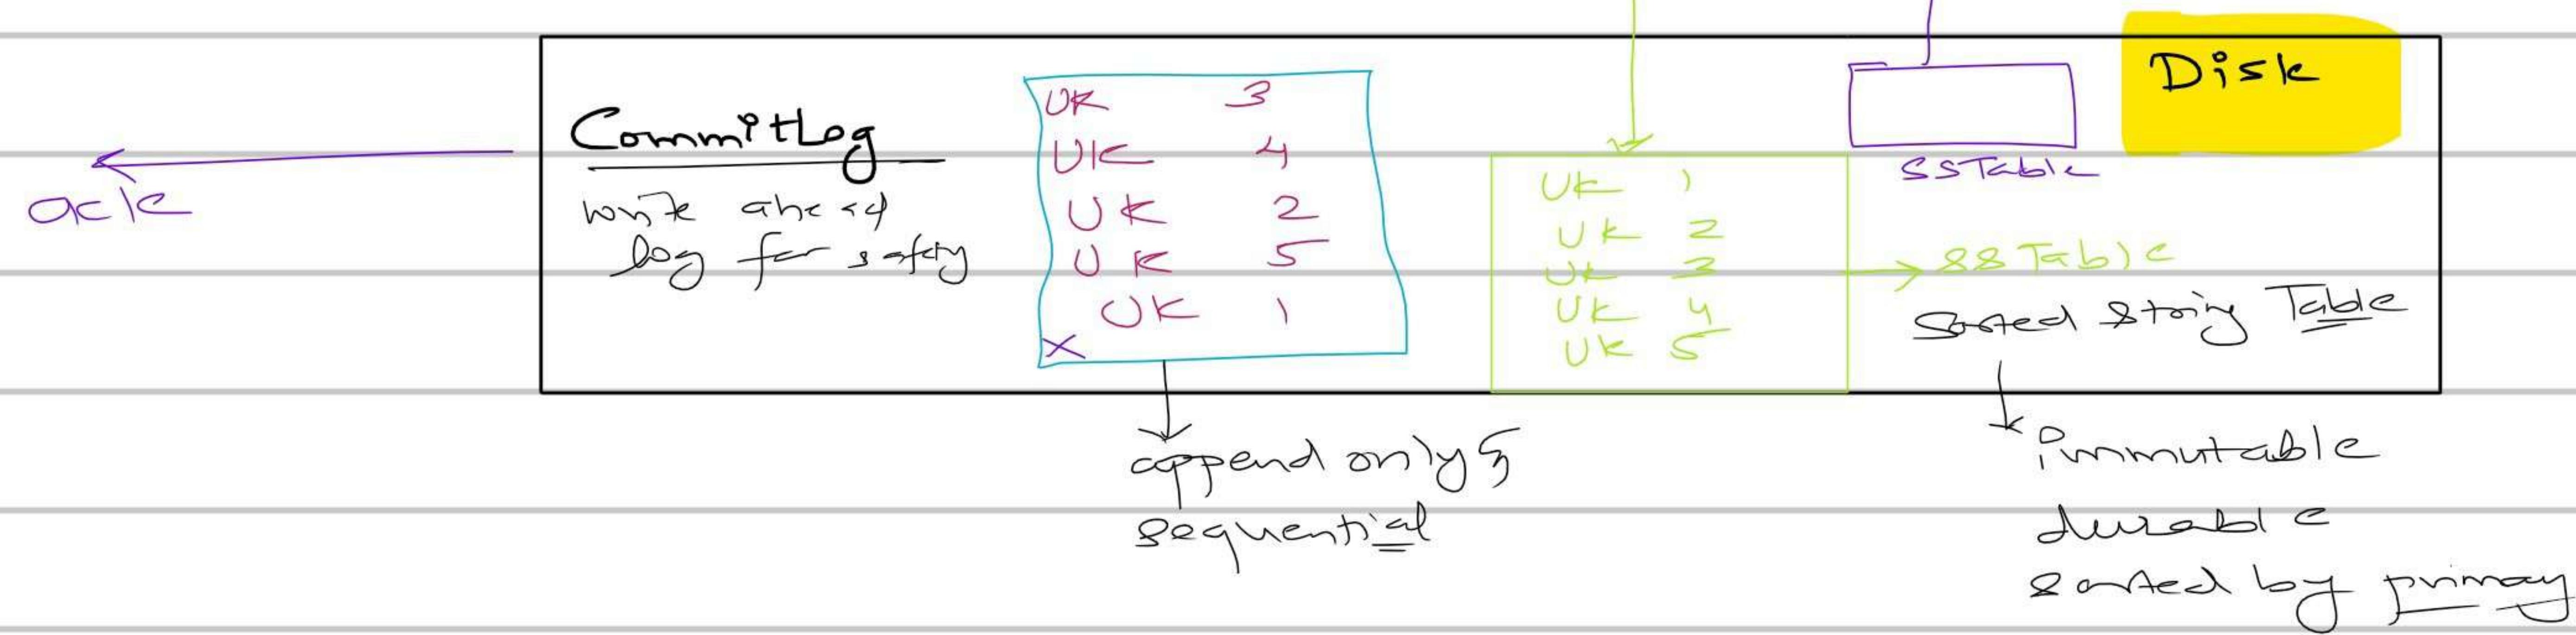
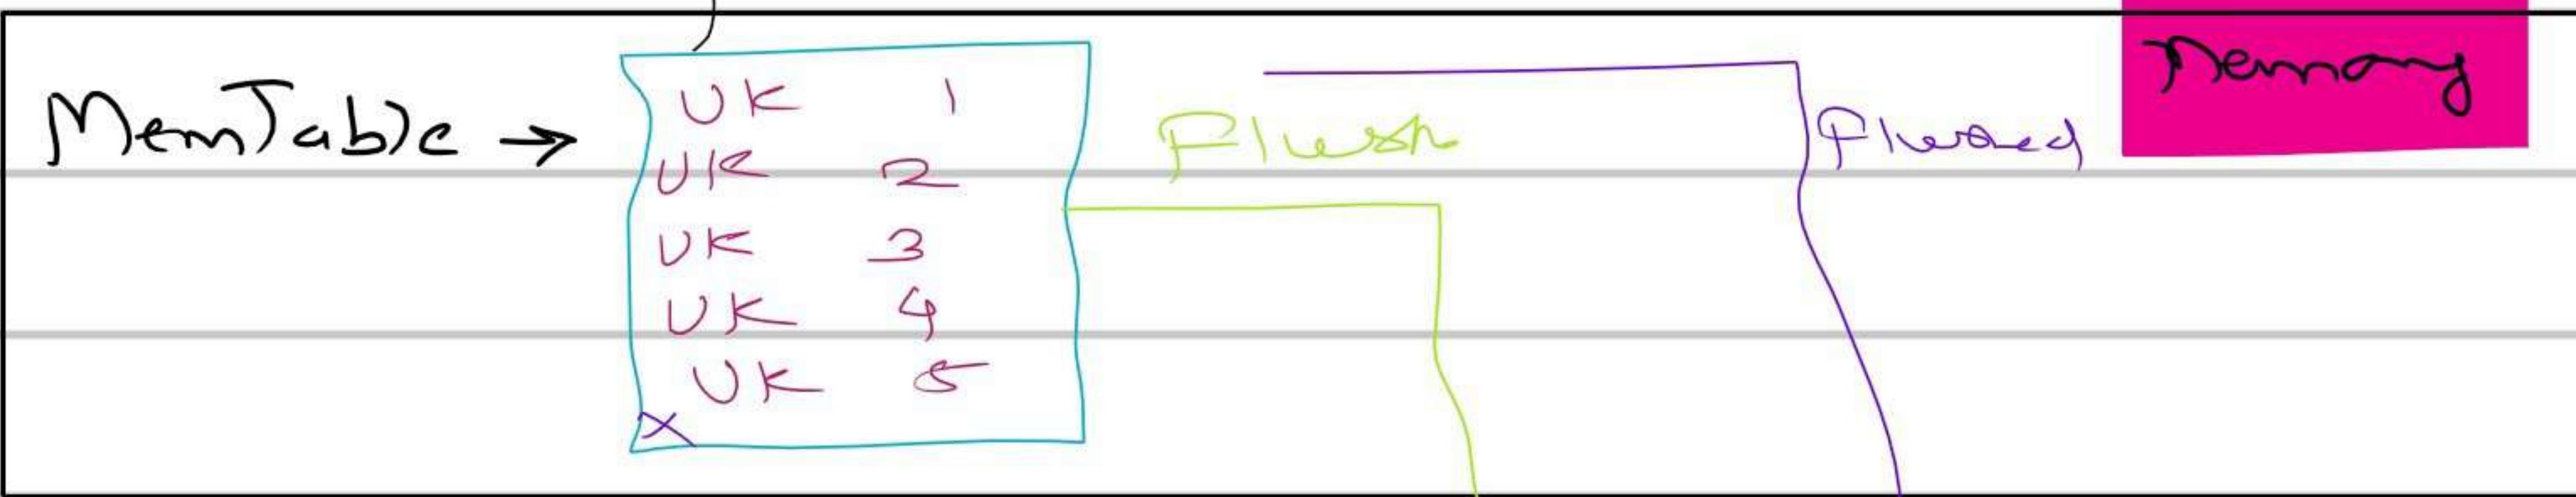


$$RF = \frac{3}{\pi}$$

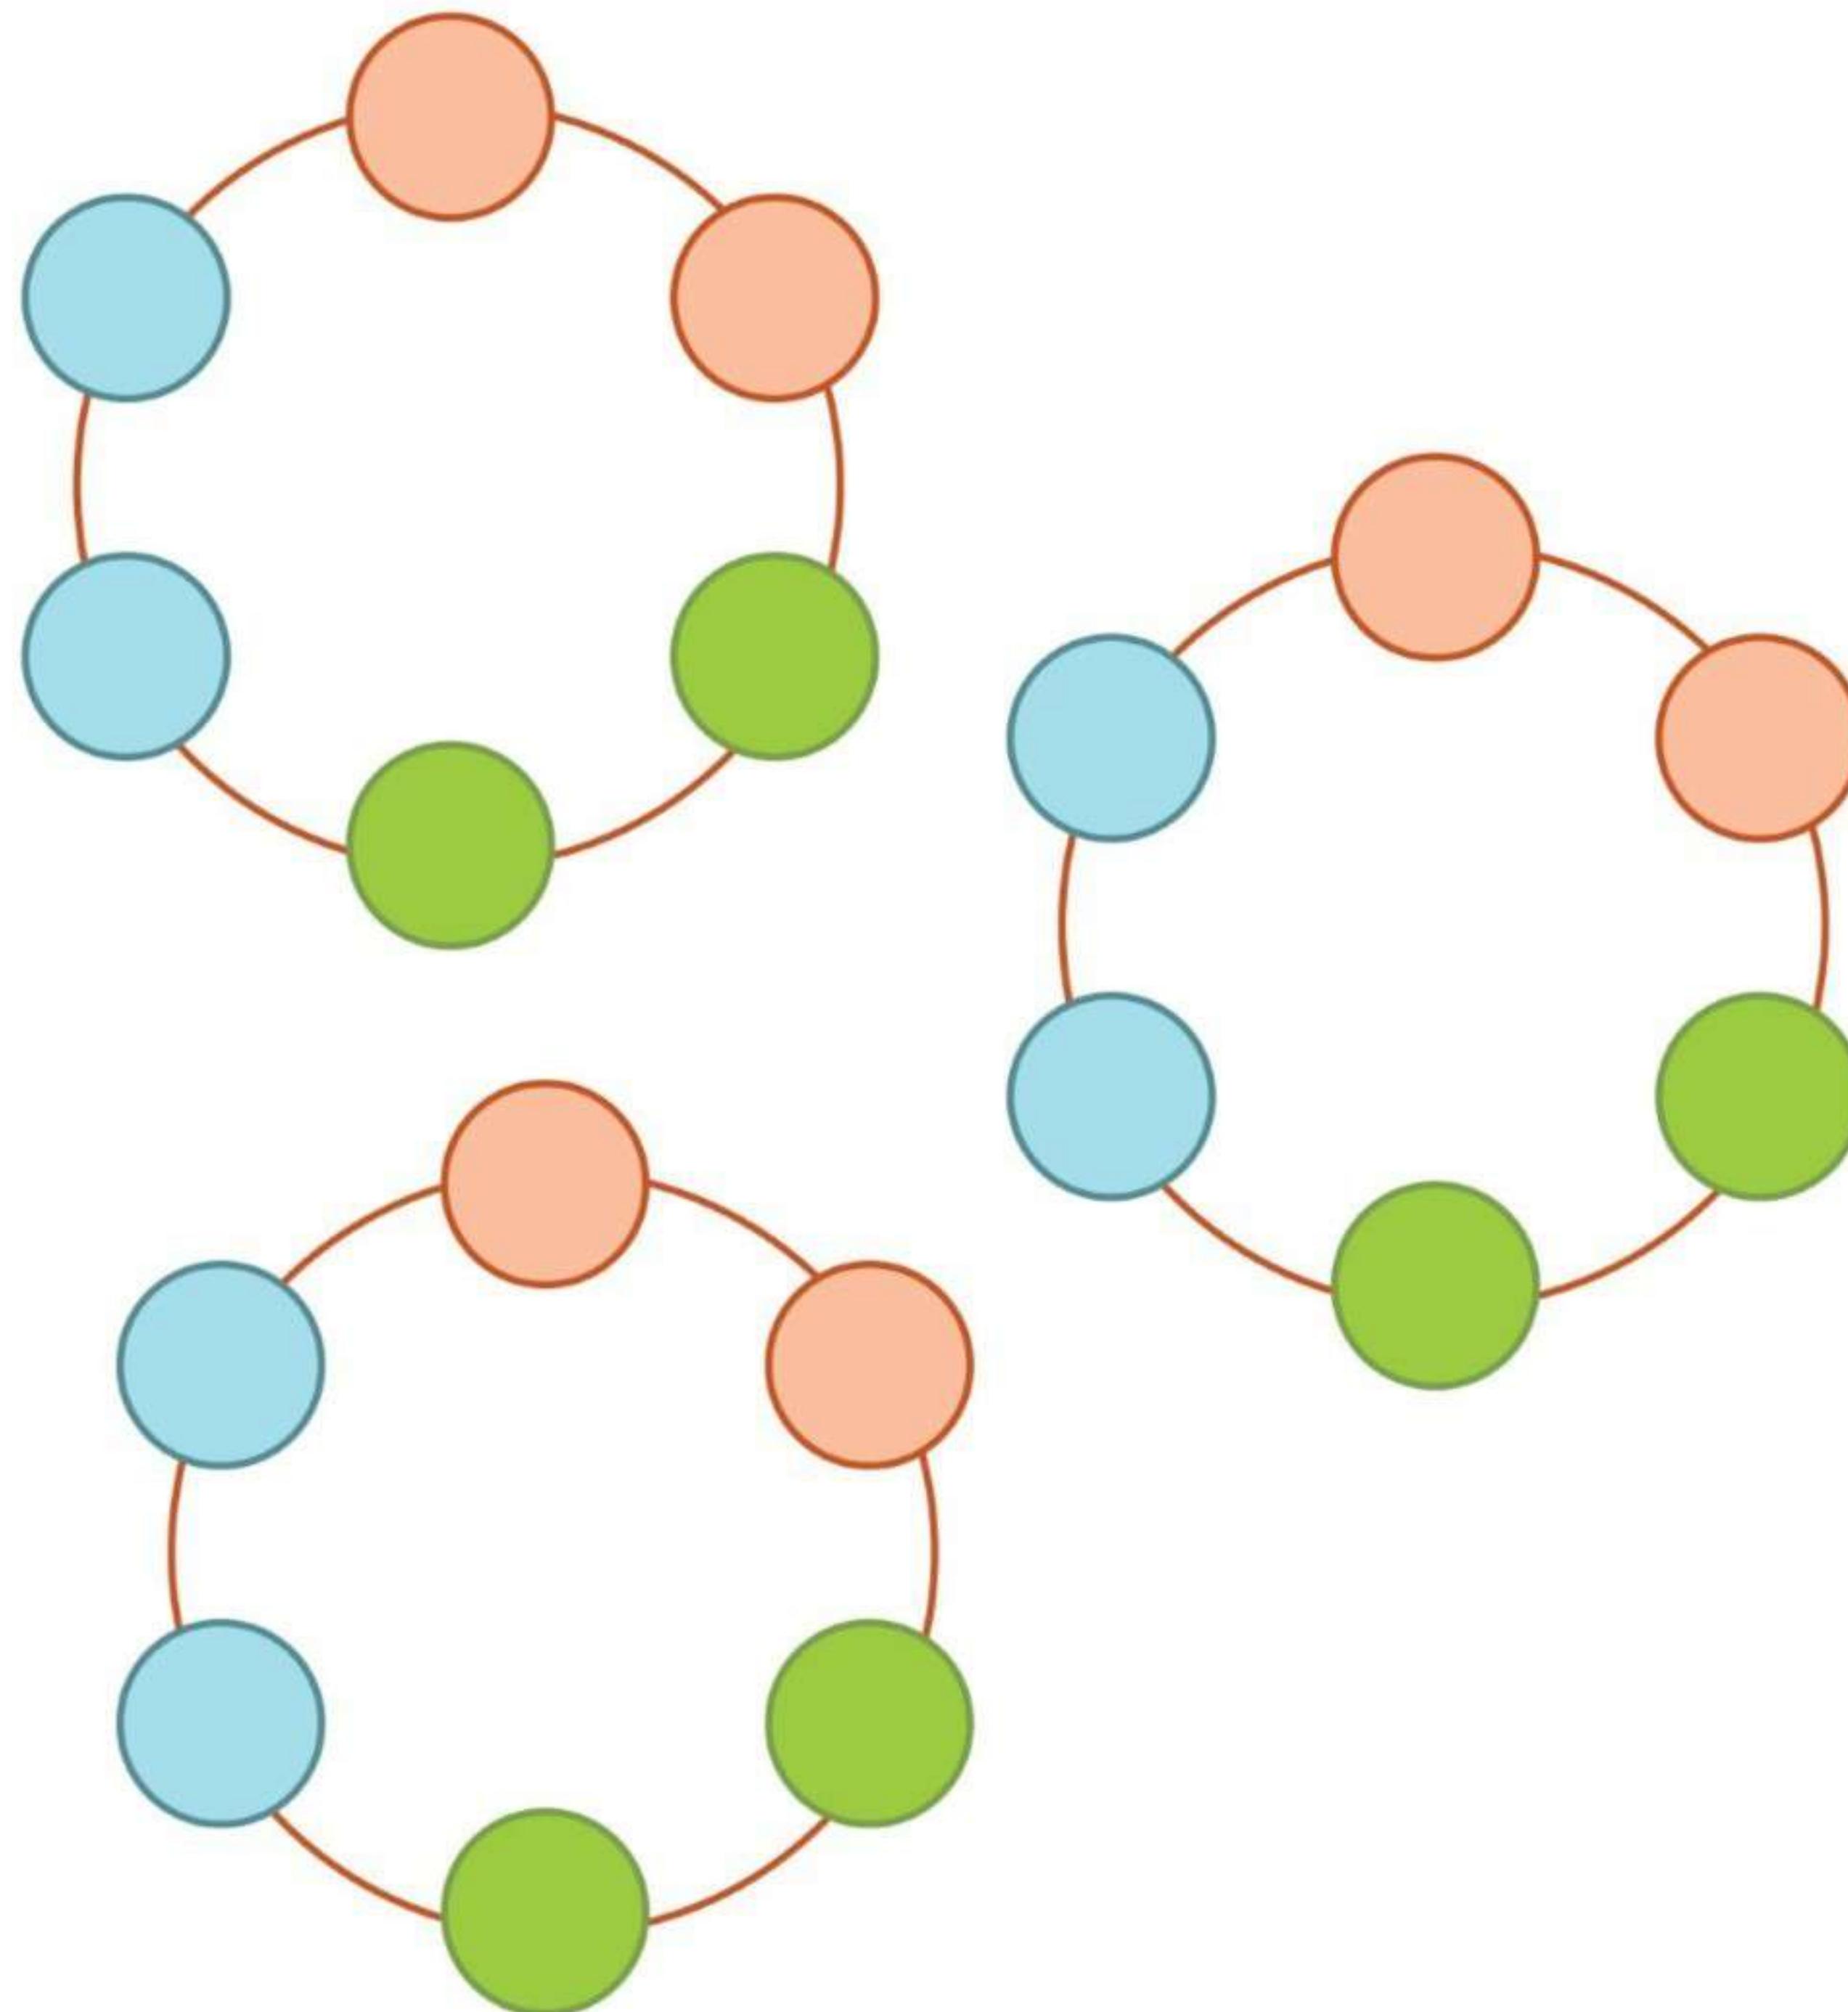
$$\underline{\underline{C_L = \alpha}} = 2 \quad \frac{3+)}{2} = 2$$

Cassandra Write Path

country	id
UK	3
UK	4
UK	2
UK	5
UK	1



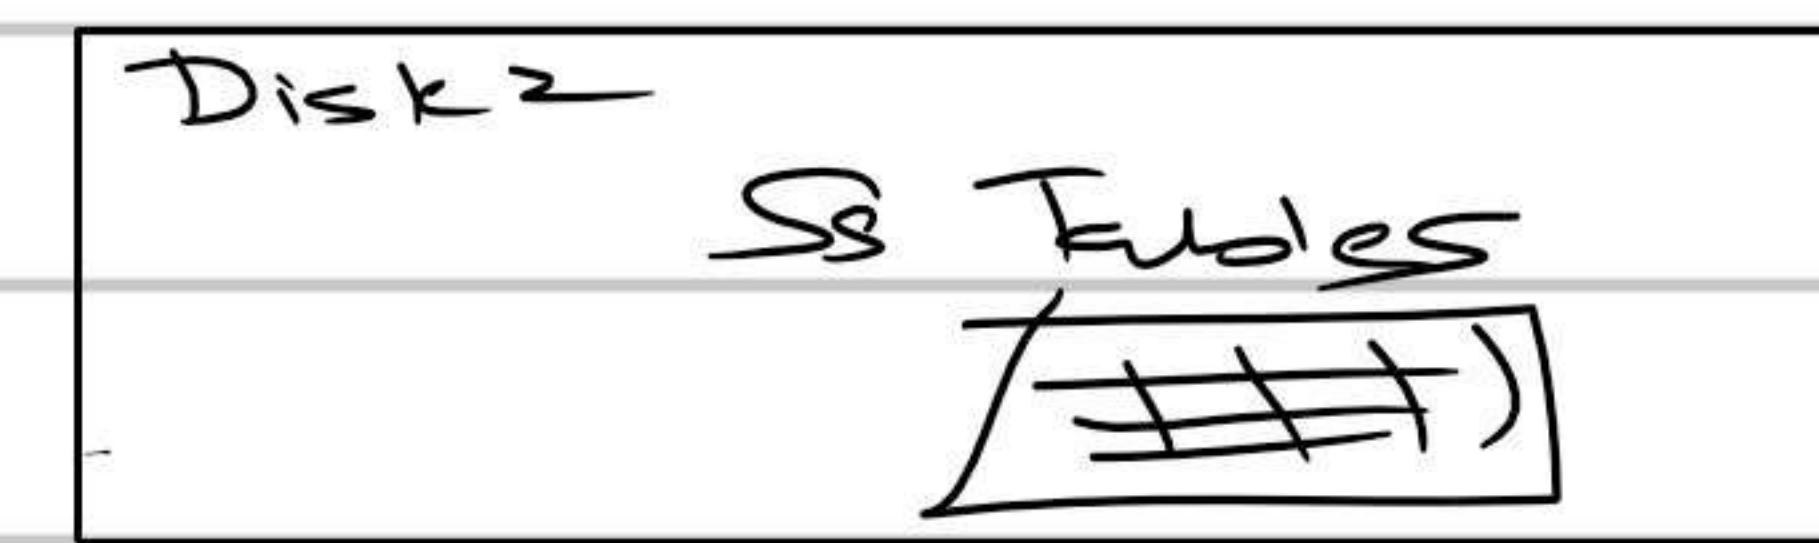
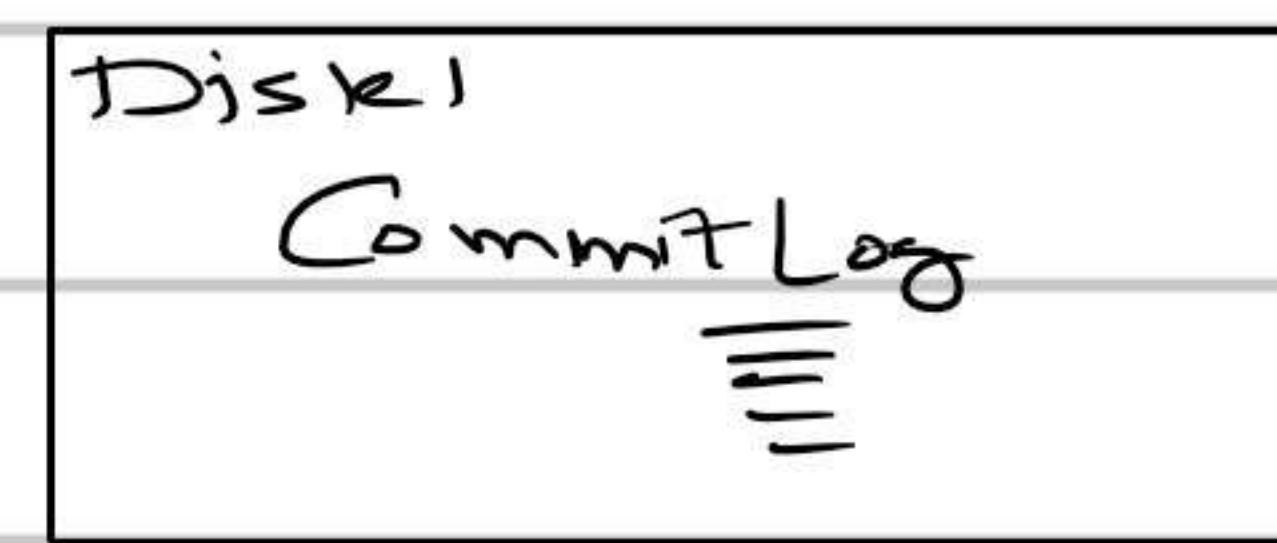
Consistency with Multiple Data Centers



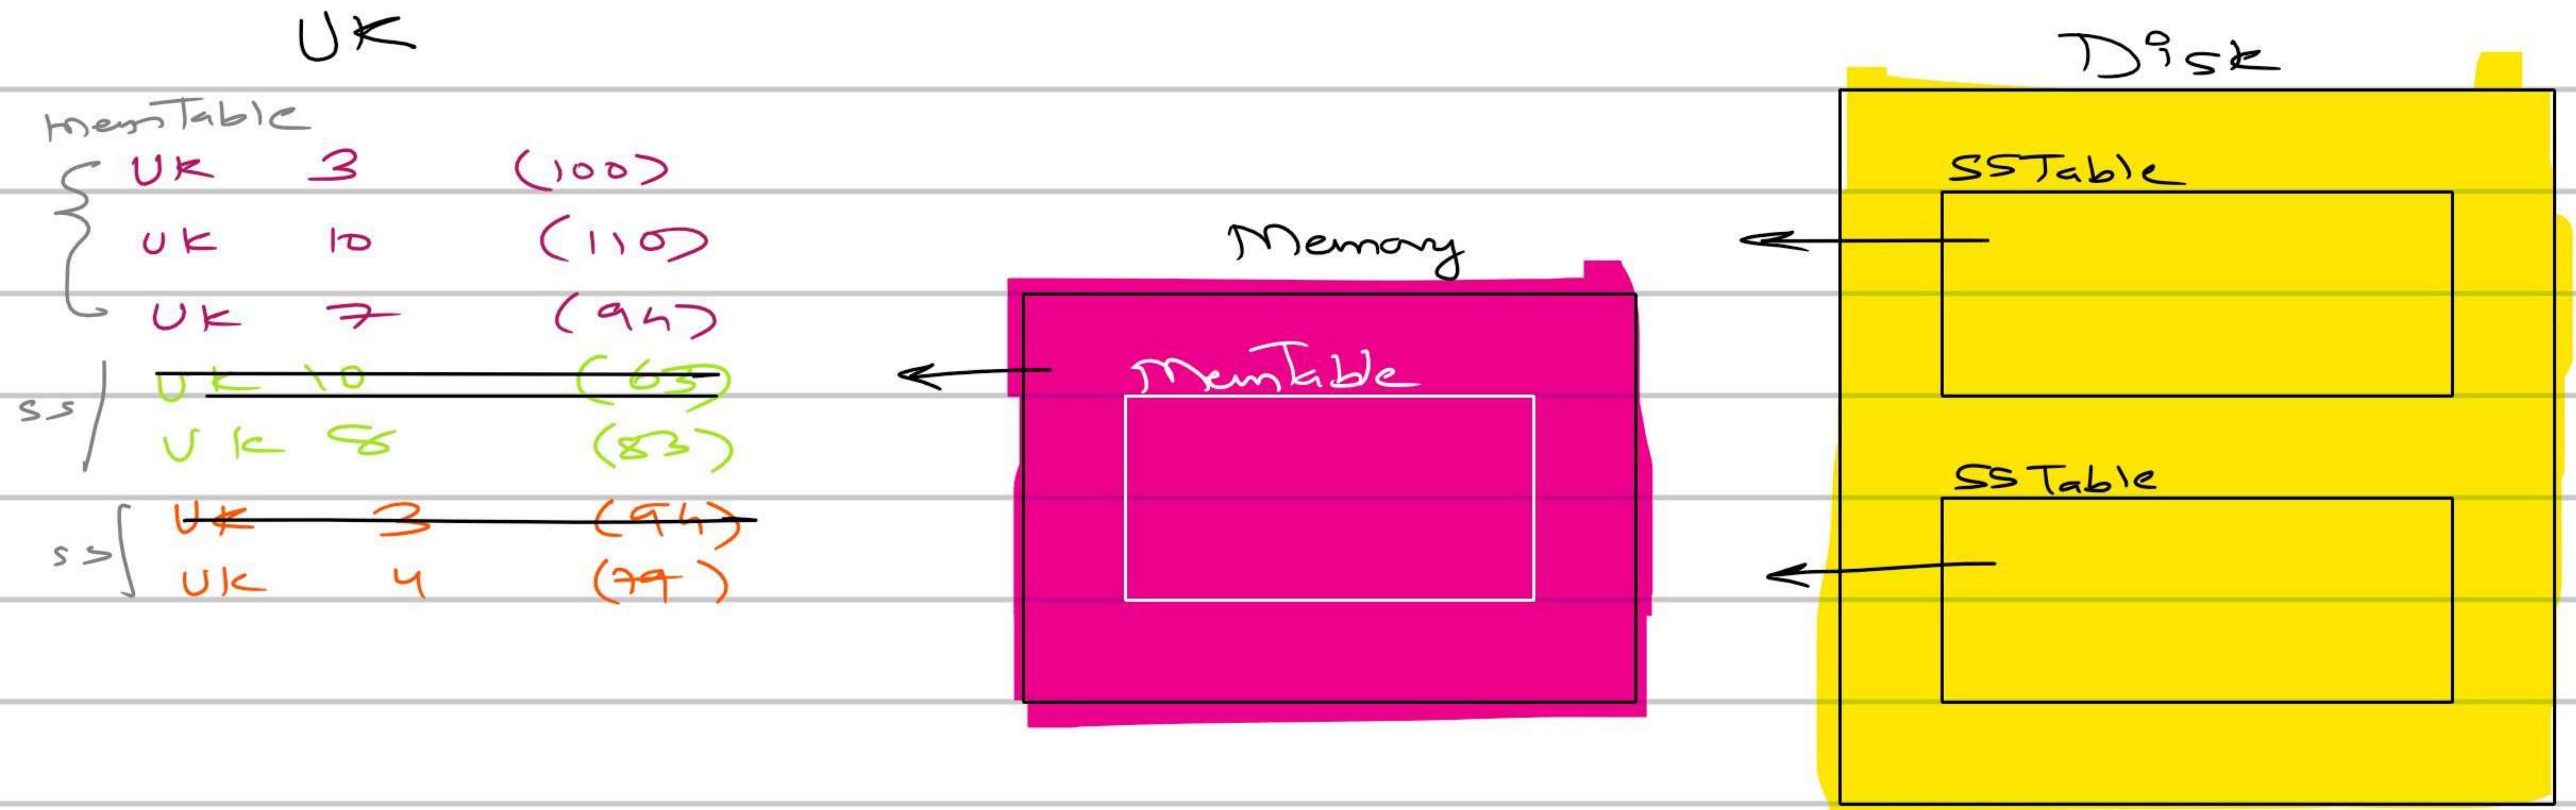
EACH_QUORUM
LOCAL_QUORUM
LOCAL_ONE

Cassandra Write Path

Spinning Disk

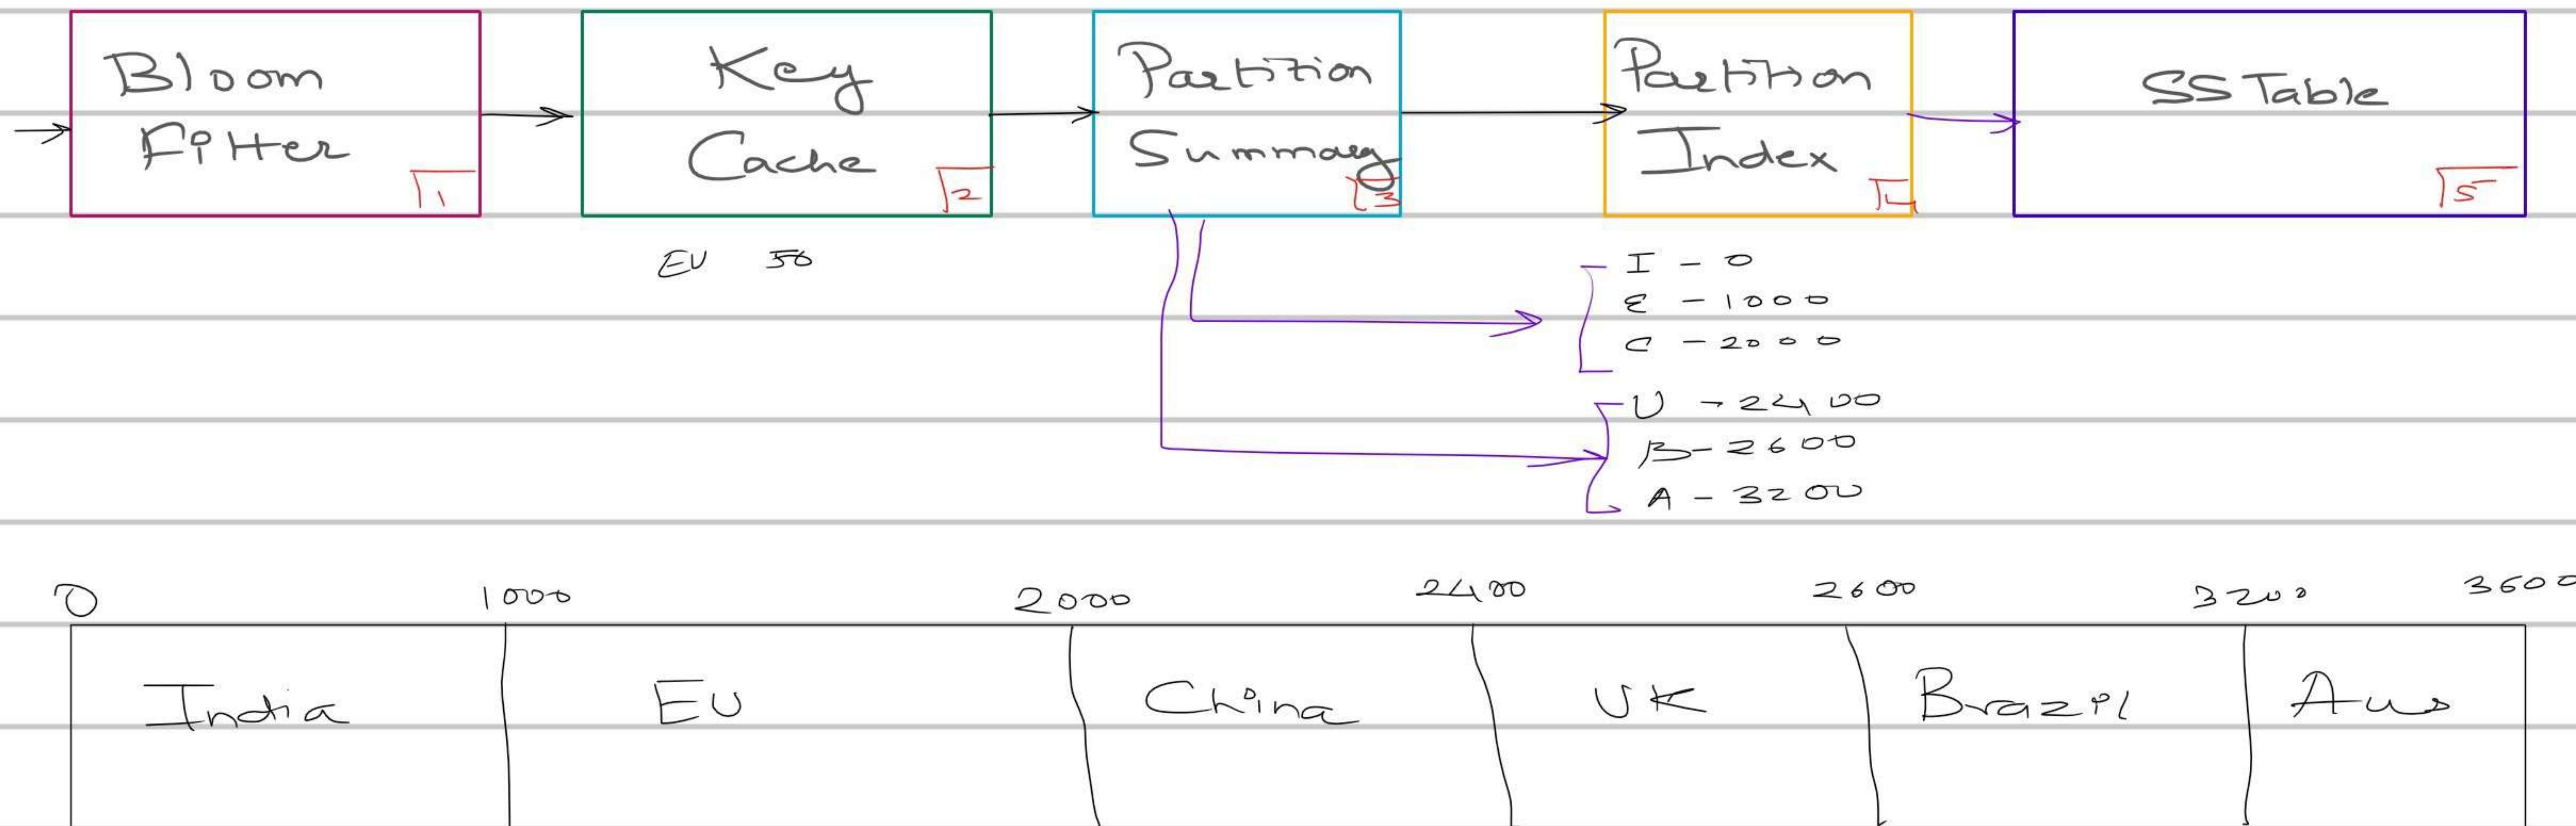


Cassandra Read Path



Latest Data is always in MemTable

UK ↗



Bloom Filter → quickly tells if data might be in SSTable or not (not 100% accurate but fast).

key Cache → If own key was accessed recently it will cache the exact disk location for it

Partition Index → notifies the start byte ^{location} of each partition

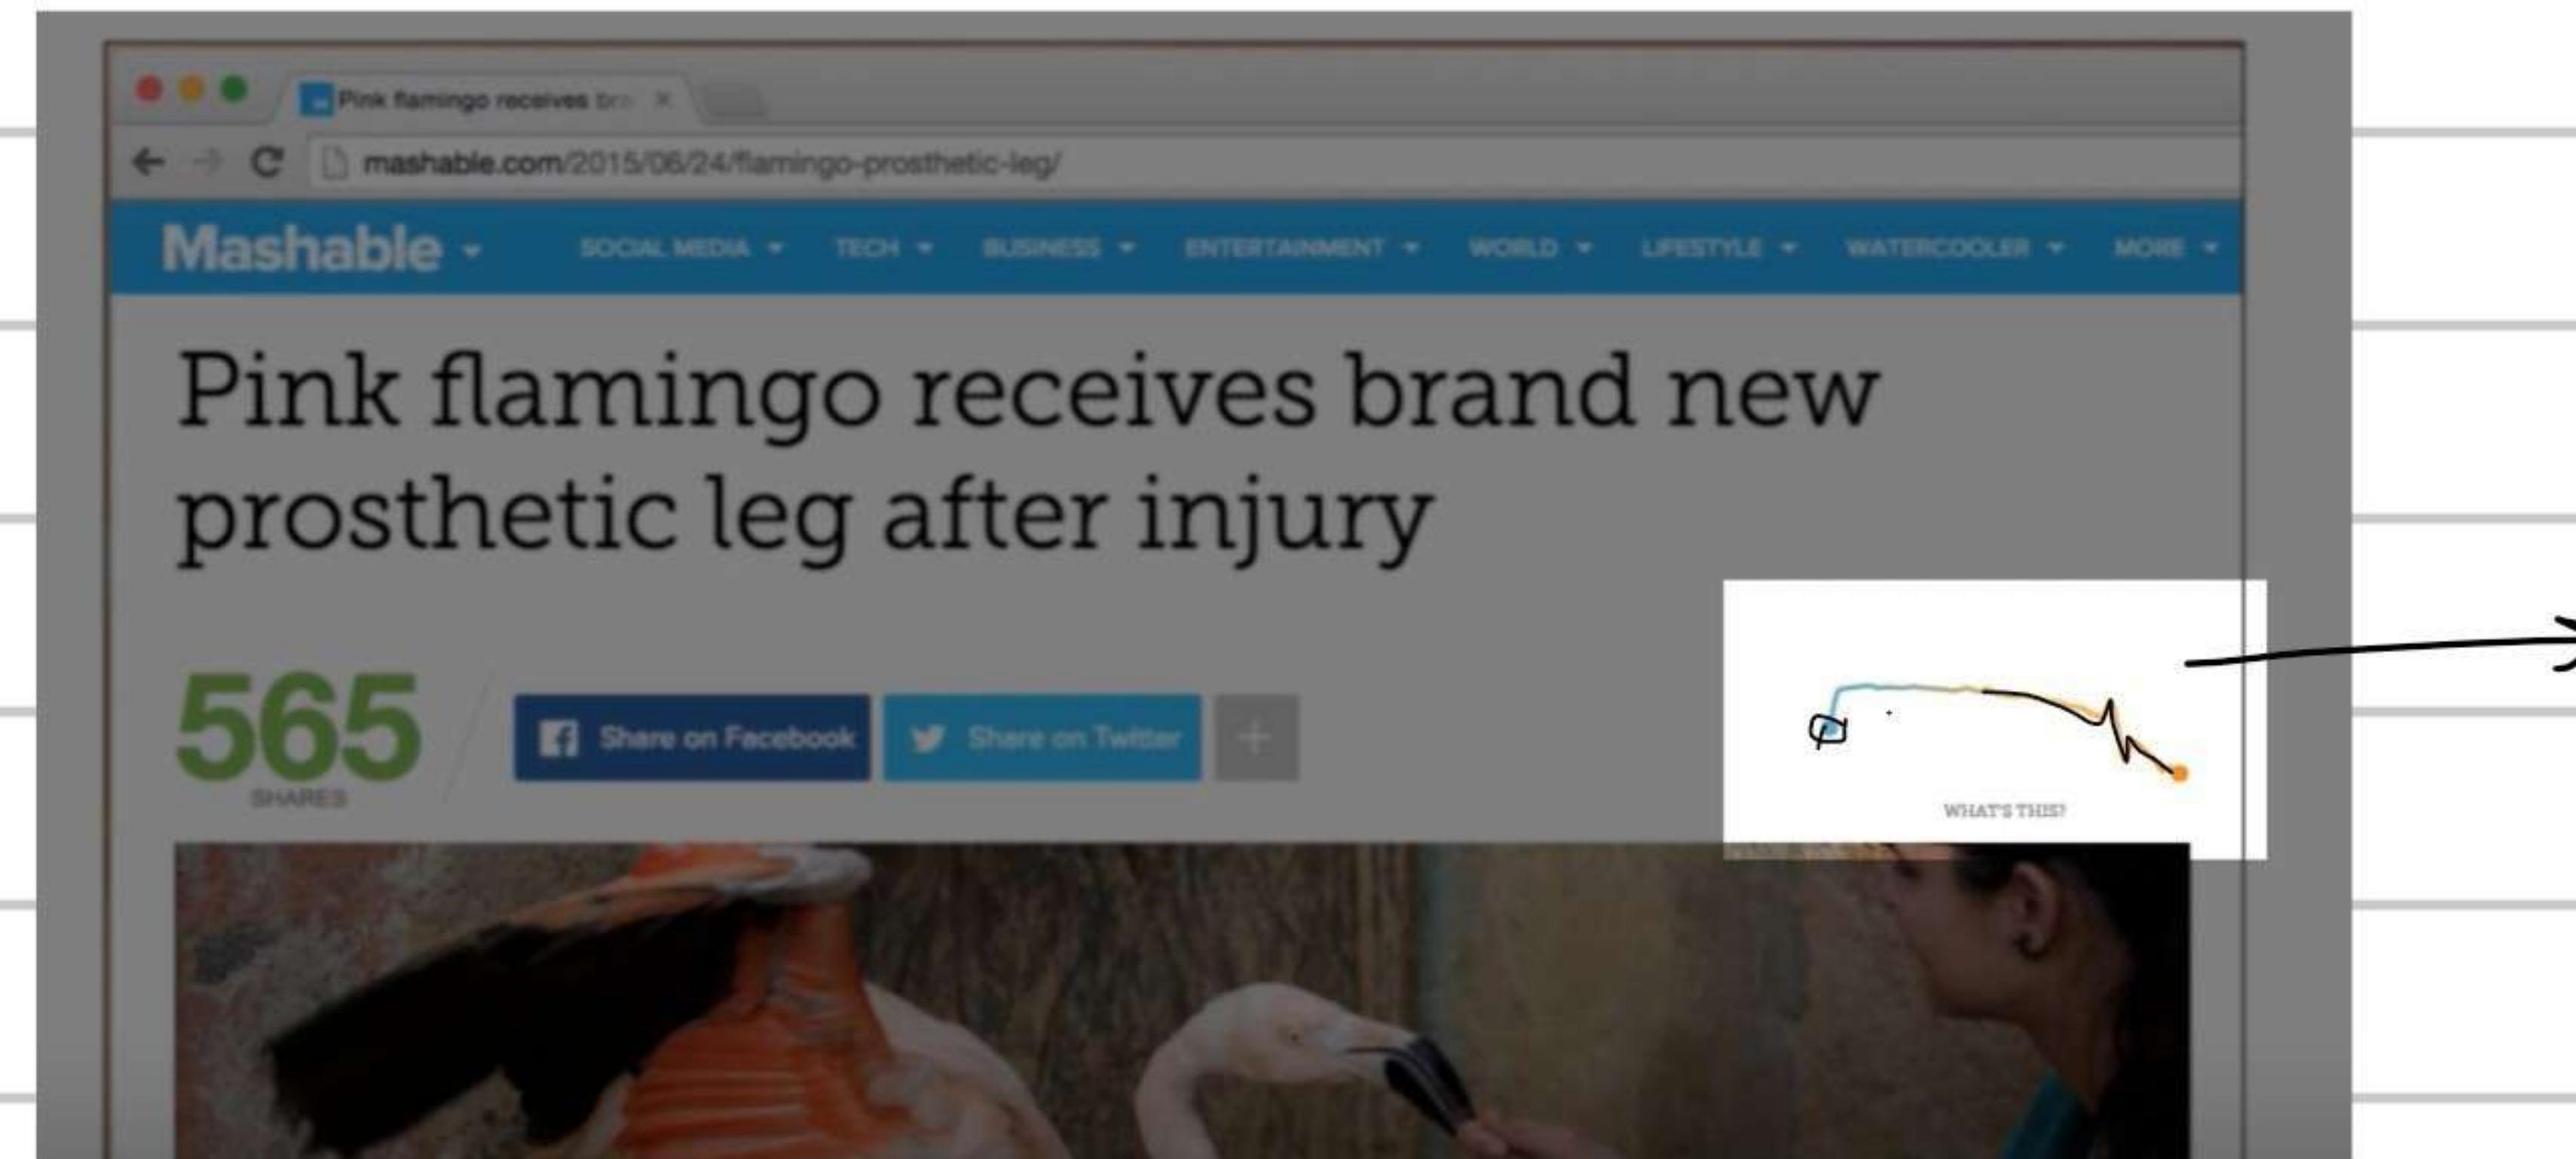
Partition summary → is like a partition index - for partition index

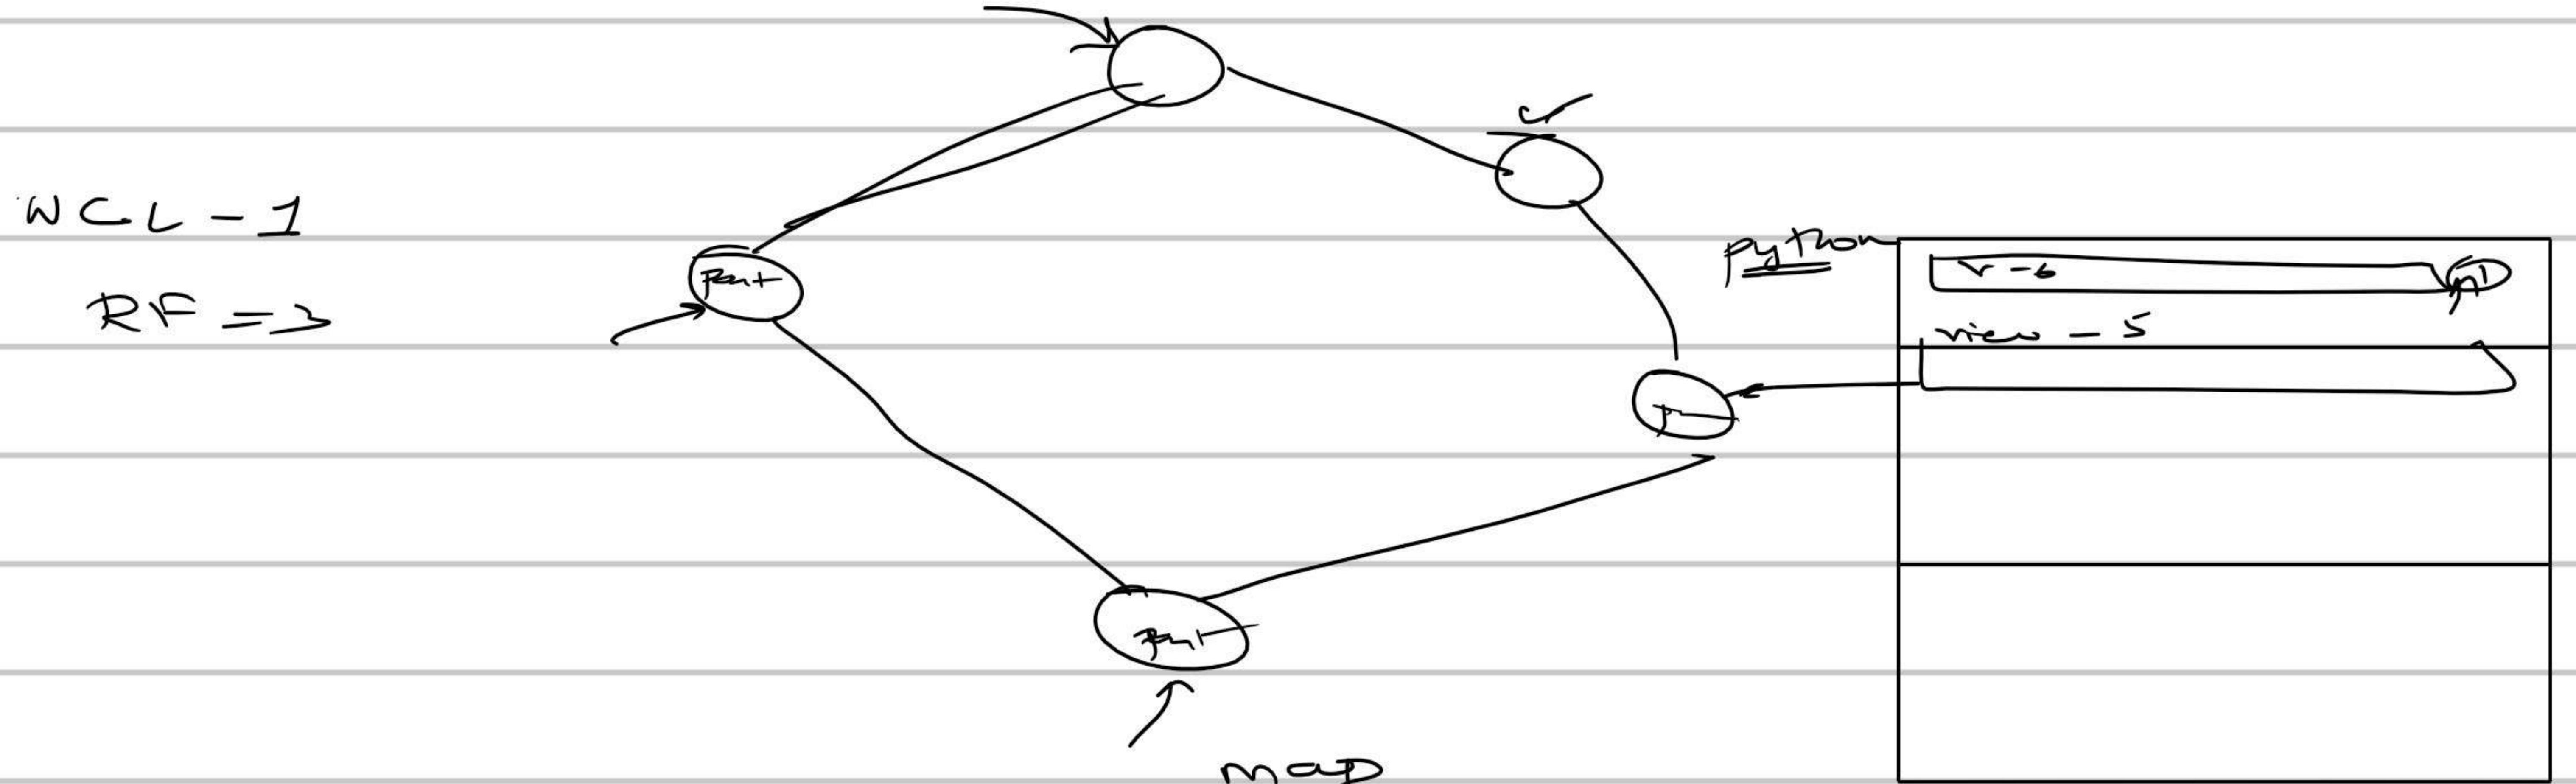
Time Series Data

Time UUID

45b94a50-12e5-11e5-9114-091830ac5256

- The number of 100 ns intervals since UUID epoch
- MAC address
- Clock sequence number to prevent duplicates





$\xrightarrow{\text{last_login}} \xrightarrow{\text{alter}}$ $\xrightarrow{\text{map}}$
 device identification
 $\text{tuple} < \text{timestamp}, \text{Pnet} >$

$\text{map} < \text{text}, \text{frozen tuple} < \text{timestamp}, \text{Pnet} >>>$

Partitions

A maximum of 2 billion
cells

[Row x Columns]

Must fit on single node

sensor

→ you cannot use ttl

→ data is coming in every 5 seconds

Restructure course-page-views in such a
way that P1 won't make partition
go out of memory